Generic programming in modern C++ for Image Processing

Michaël Roynard

February 1, 2022

Acknowledgement

Abstract

C++ is a multi-paradigm language that enables the initiated programmer to set up efficient image processing algorithms. This language strength comes from many aspects. C++ is high-level, so this enables developing powerful abstractions and mixing different programming styles to ease the development. At the same time, C++ is low-level and can fully take advantage of the hardware to deliver the best performance. It is also very portable and highly compatible which allows algorithms to be called from high-level, fast-prototyping languages such as Python or Matlab. One of the most fundamental aspects where C++ really shines is generic programming. Generic programming makes it possible to develop and reuse bricks of software on objects (images) of different natures (types) without performance loss. Nevertheless, conciliating genericity, efficiency, and simplicity at the same time is not trivial. Modern C++ (post-2011) has brought new features that made it simpler and more powerful. In this thesis, we first explore one particular C++20 aspect: the concepts, in order to build a concrete taxonomy of image related types and algorithms. Second, we explore another addition to C++20, ranges (and views), and we apply this design to image processing and image types in order to solve issues such as how hard it was to customize/tweak image processing algorithms. Finally, we explore possibilities regarding how we can offer a bridge between static (compile-time) generic C++ code and dynamic (run-time) Python code. We offer our own hybrid solution and benchmark its performances as well as discussing what can be done in the future with JIT technologies. Those three axes aim to address the issue of generic programming all the while remaining efficient and easy to use.

C++ est un langage de programmation multi-paradigme qui permet au développeur initié de mettre au point des algorithmes de traitement d'image. La force de langage se base sur plusieurs aspects. C++ est haut-niveau, cela signifie qu'il est possible de développer des abstractions puissantes mélangeant plusieurs styles de programmation pour faciliter le développement. En même temps, C++ est bas-niveau et peut pleinement tirer partie du matériel pour fournir un maximum de performances. Il est aussi portable et très compatible ce qui lui permet de se brancher à d'autres langages de haut niveau pour le prototypage rapide tel que Python ou Matlab. Un des aspects les plus fondamentaux où le C++ brille est la programmation générique. La programmation générique rend possible le développement et la réutilisation de brique logiciels comme des objets (images) de différentes natures (types) sans avoir de perte au niveau performance. Néanmoins, concilier la généricité, la performance et la simplicité d'utilisation tout en même temps n'est pas trivial. Le C++ moderne (post-2011) amène de nouvelles fonctionnalités qui le rendent plus simple et plus puissant. Dans cette thèse, nous explorons en premier un aspect particulier du C++20 : les concepts, dans le but de construire une taxonomie des types relatifs au traitement d'image. Deuxièmement, nous explorons une autre addition au C++20, les ranges (et les vues), et nous appliquons ce design au traitement d'image et aux types d'image dans le but résoudre les problèmes relatifs à la difficulté de customiser les algorithmes de traitement d'image. Enfin, nous explorons les possibilités concernant la façon dont il est possible de construire un pont entre du code C++ générique statique (compile-time) et du code Python dynamique (run-time). Nous fournissons une solution hybride et nous mesurons ses performances. Nous discutons aussi les pistes qui peuvent être explorées dans le futur, notamment celles qui concernent les technologies JIT. Ces trois axes visent à solutionner la problématique concernant la programmation générique tout en restant efficace et accessible.

Contents

List of Figures

List of Tables

Part I

Context

Chapter 1

Introduction

Outline

NOWADAYS *Computer Vision* and *Image Processing (IP)* are omnipresent in the day to day life of the people. It is present each time we pass by a CCTV camera, each time we go to the hospital do an MRI, each time we drive our car and pass in front of a speed camera and each time we use our computer, smartphone or tablet. We just cannot avoid it anymore. The systems using this technology are sometimes simple and, sometimes, more complex. Also the usage made of this technology has several different purposes: space observation, medical, quality of life improvement, surveillance, control, autonomous system, etc. Henceforth, Image processing has a wide range of research and despite having a mass of previous of work already contributed to, there are still a lot to explore.

Let us take the example of a modern smartphone application which provides facial recognition in order to recognize people whom are featuring inside a photo. To provide accurate result, this application will have to do a lot of different processing. Indeed, there are a lot of elements to handle. We can list (non exhaustively) the weather, the light exposition, the resolution, the orientation, the number of person, the localization of the person, the distinction between humans and objects/animals, etc. All of these is in order to finally recognizing the person(s) inside the photo. What the application does not tell you is the complexity of the image processing pipeline behind the scene that can not even be executed in its entirety on one's device (smartphone, tablet, . . .). Indeed, image processing is costly in computing ressources and would not meet the time requirement desired by the user if the entire pipeline was executed on the device. Furthermore, for the final part which is "recognize the person on the photo", one needs to feed the pre-processed photo to a neural network trained beforehand through deep learning techniques in order to give an accurate response. There exists technologies able to embed neural network into mobile phone such as MobileNets [112] but it is still limited. It can detect a human being inside a photo but not give the answer about who this human being is for instance. That is why, accurate neural network system usually are abstracted away in cloud technologies making them available only via Internet. When uploading his image, the user does not imagine the amount of technologies and computing power that will be used to find who is on the photo.

We now understand that in order to build applications that interact with photos or videos nowadays, we need to be able to do accurate, fast and scalable image processing on a multitude of devices (smartphone, tablet, . . .). In order to achieve this goal, image processing practitioners needs to have two kinds of tools at their disposal. One will be the prototyping environnement, a toolbox which allow the practitioner to develop, test and improve its application logic. The other is the production environnement which deploy the viable version of the application that was developed by the practitioner. Both environment may not have the same needs. On one hand, the prototyping environment usually requires to have a fast feedback loop for testing, an availability of state-of-the-art algorithms and existing software. This way the practitioner can

easily build upon them and be fast enough in order not to keep waiting for results when testing many prototypes. On the other hand, the production environment must be stable, resilient, fast and scalable.

When looking at standards in the industry nowadays, we notice that Python is the main choice for prototyping. Also, Python may not be enough so that a viable prototype can be pushed in production with minimal changes afterwards. We find it non-ideal that the practitioner cannot take advantages of many optimisation opportunities, both in term of algorithm efficiency and better hardware usage, when proceeding this way. It would be much more efficient to have basic low level building blocks that can be adapted to fit as mush use cases as possible. This way, the practitioner can easily build upon them when designing its application. We distinguishes two kind of use cases. The first one is about the multiplicity of types or algorithms the practitioner is facing. The second one is about the diversity of hardware the practitioner may want to run his program. The goal is to have building blocks that can be intelligent enough to take advantage of many optimization opportunities, with regard to both input data types/algorithms and target hardware. Then the practitioner would have a huge performance improvement, by default, without specifically tweaking its application. As such, the concept of genericity was introduced. It aims at providing a common ground about how an image should behave when passed to basic algorithms needed for complex applications. This way, in theory, one only needs to write the algorithm once for it to work with any given kind of image.

Different data types and algorithms

In Image Processing, there exists a multitude of image types whose characteristics can be vastly different from one another. This large specter is also resulting from the large domain of application of image processing. For instance, when considering photography we have 2D image whose values can vary from 8 bits grayscale to multiple band 32-bits color scheme storing informations about the non-visible specter of human eye. If we consider another domain of application, such as medical imaging, we now can consider sequence of images such as sequence of 3D image for an MRI for instance. More broadly there are two orthogonal constituent of an image: its topology (or structure) and its values. However, there are two more aspectes to consider here. Firstly, image processing provide plenty of algorithms that can or cannot operate over specific data types. There are also different kind of algorithms. Some will extract informations,(e.g. histogram) other will transform the image point-wise (e.g. thresholding), and some other will even combine several image to render a different kind of informations (e.g. background substraction). There are many simple algorithms and also many complex algorithms out there. Secondly, there are orbiting data around image types and algorithms that are also very diverse and necessary for their smooth operation. Indeed, a dilation algorithm will also need an additional information: the dilation disc. A thresholding algorithm may be given a threshold. A convolution filter requires a convolution matrix to operate. That is why, when considering both image types and algorithms, we need a 3D-chart (illustrated in fig. 1.1) to enumerate all possibilities, where one axis is the image topology, one axis is the color scheme and one axis enumerate the additional data that can be associated to an image.

Figure 1.1: Illustration of the specter of the multitude of possibilities in the image processing world.

Different user profiles and their use cases

The end user is a non programmer user who wants to occasionally use image processing software through UI-rich interface, such as Adobe Photoshop [141] or The GIMP [132]. Its skills are non-relevant as the end user is using the software to get work done even though he does not fully understand the underlying principles. For instance, the end user will want to correct the brightness of an image, of remove some impurities from a face or a building. The end user does not want to build an application but wants to save time. The needs of the end user mainly revolves around a clean and intuitive software UI as well as a well as support for mainstream image types and operation a photograph may need to do.

The practitioner is what we are called when we first approach the image processing area. A practitioner is the end user of image processing libraries. Its skills mainly revolve around applied mathematics for image processing, prototyping and algorithms. A practitioner aims at leveraging the features the libraries can offer to build his application. For instance, a practitioner can be a researcher in medical imaging, an engineer build a facial recognition application, a data scientist labeling its image set, etc. The needs of practitioner are mainly revolving around a fast feedback loop. The developing environment must be easily accessible and installable. This way a practitioner can judge quickly wether one library will answer his needs. The documentation of the library must be exhaustive and didactic with examples. When prototyping, the library must provide fast feedback loops, as in a python notebook for instance. Finally it must be easily integrated in a standard ecosystem such as being able to work with NumPy's array natively without imposing its own types. To sum up, practitioner's programmatic skills do not need to be high as his main goal is to focus on algorithms and mathematics formulas.

The contributer is an advanced user of a library who is very comfortable with its inner working, philosophy, aims, strengths and potential shortcomings. As such, he is able to add new specific features to library, fix some shortcomings or bugs. Usually a contributer is able to add a feature needed for a practitioner to finish his application. Furthermore he can then contribute back his features to the main project via pull requests if it is relevant. This way, a maintainer will assess the pull request and review it. The two main points of a contributer are his deep knowledge of a library and his ability to write code in the same language as the source code of it. Also, a contributer must have knowledge of coding best practices such as writing unit tests which are mandatory when adding a feature to an existing library. To facilitate contribution,

a library must provide clear guidelines about the way to contribute, be easy to bootstrap and compile without having heavy requirements on dependencies. The best case would be that the library is handled by standard packages managers such as system apt or python conan.

The maintainer is usually the creator, founder of the library or someone that took over the project when the founder stepped back. Also, when a library grows, it is not rare that regular contributers end up being maintainer as well to help the project. The maintainer is in charge of keeping alive the project by fulfilling several aspects: upgrade and release new features according to the user (practitioner) needs and the library philosophy. Also, a library may not evolve as fast as the user may want it because of lack of time from maintainers. A lot of open source projects are maintained by volunteers and lack of time is usually the main aspect slowing development progress. The maintainer is also in charge of reviewing all the contributers pull requests. He must check if they are relevant and completed enough, (for instance, presence of tests and documentation) to be integrated in the project. Indeed, merging a pull requests equals to accepting to take care of this code in the future too. It means that further upgrade, bug fix, refactoring of the project will consider this new code too. If the maintainer is not able to take care of this code then it should probably not be integrated in the project in the first place. Any project and library has its maintainers. A maintainer is someone very familiar with the inner working and architectural of the project. He is also someone that has some history in the project to understand why some decisions has been made, what choices has been made at some points and what the philosophy of the project is. It is important to be able to refuse a contribution that would go contrary to the philosophy of the project, even a very interesting one. Finally the profile of a maintainer is one of a developer that is used to the standard workflow in open source based on: forks, branches, merge/pull requests and continuous integration.

Different tools

Before stating the topic of the thesis, it is important to enumerate the different kind of tools the current market has to offer to know where we will be positioning ourself.

Graphic editors are what neophyte thinks about when they imagine what image processing is. Those are tools that allow a non expert user to apply a wide array of operation on an image from an intuitive GUI in a way the user does not have to understand the underlying logic behind each and every operation he is applying. Such tools are usually large complex software such as The GIMP [132] or Photoshop [141]. Their aim is to be useable by end users while supporting a large set of popular image format.

Command line utilities are binaries that perform one operation or more invocable from a console interface or from a shell script through a command line interface (CLI). This CLI usually offers several options to pass data and/or information to the programs in order to have an processing happening. The informations can be, for instance, the input image path, the ouput image name and the name of a mathematical morphology algorithm to apply. Usually command line utilities come as projects such as ImageMagick [148], GraphicsMagick [135] or MegaWave [76, 47].

Visual programming environment are software that allow the user to graphically and intuitively link one or several image processing operations while interactively displaying the result. The processing can easily be modified and the results are updated accordingly. Those software are usually aimed at engineer or researchers doing prototyping work not exclusive to image processing. Mathcad [130] is a good example of such a software.

Integrated environment are feature-rich platforms for scientists oriented toward prototyping. Those platforms provides a fully functional programming language and a graphical interface allowing the user to run commands and scripts as well as viewing results and data (image, matrices, etc.). The most well-known integrated environnement are Matlab [137], Scilab [139], Octave [144], Mathematica [140] and Jupyter [106] notebooks.

Package for dynamic language has known a surge in development these last few years and a multitude of libraries has been brought to dynamic languages this way. For instance, let us consider the python programming language. There are two main package provider: PyPi [149] and Conda [134]. Both allow to install packages to enable the user to program his prototypes in Python very quickly. In image processing, there are packages such as Scipi [35], NumPy [57], scikit-image [99], Pillow [142] as well as binding for OpenCV [22].

Programming libraries is the most common tool available out there. They are a collection of routines, functions and structures providing features through a documentation and binaries. They require the user to be proficient with a certain programming language and also to be able to integrate a library into his project. For image processing we have: IPP [49], ITK [90], Boost.GIL [52], Vigra [27], GrAL [51], DGTal [103], OpenCV [22], CImg [85], Video++ [94], Generic Graphic Library [26] Milena [65, 68] and Olena [161, 74, 77, 97].

Domain Specific Languages (DSL) are tools developed when a library developer deem he is unable to express the concepts and abstraction layers he wants to express through publishing a library. In this case, the barrier is often the programming language itself and so the developer does think that another layer of abstraction above the programming language would be a good thing. It leads to the genesis of a new programming language in some cases like Halide [91] and SYCL [123, 122] but can also be a case of having the current programming language be "upgraded" to include another subset of features that are not natively included. This is often the case in C++ where we have in-language DSL like Eigen [67], Blaze [78, 79], Blitz++ [30] or Armadillo [108]. They leverage a possibility of the C++ programming language (*expression templates* [16]) to achieve it.

Topic of this thesis

In the end, it is often known that there is a rule of three about genericity, efficiency and ease of use. The rule states that one can only have two of those items by sacrificing the third one. If one wants to be generic and efficient, then the naive solution will be very complex to use with lots of parameters. If one wants a solution to be generic and easy to use, then it will be not very efficient by default. If one wants a solution to be easy to use and efficient then it will not be very generic. To illustrate this rule, we can find examples among existing libraries. A notably generic and efficient library in C++ is Boost [143]: it is also notably known to be hard to use. Components such as Boost.Graph, Boost.Fusion or Boost.Spirit are hard to use. Also, a library which is generic and easy to use is the json parser written by Niels Lohmann [146] it strives to handle every use case while remaining very easy to integrate and to use in user code (syntax really close to native json in C++ code by providing DSL to parse C++ constructs into JSON). However, this has a cost and the parser is slower than Json parser optimized for speed such as simdjson [145] whose aim is to "parse gigabytes of JSON per second". Finally, there are plenty of example of user friendly and efficient code which is not generic. We can cite Scikit-image [99] and OpenCV [22] that are easy to use and efficient (lot of handwritten simd/gpu code) but not generic due to the design choices.

In this thesis, we chose to work on an image processing library though continuing the work on Pylene [117]. But only working at library level would restrict the usability of our work and

thus its impact. That is why we aim to reach prototyping users through providing a package that can be used in dynamic language such as Python without sacrificing efficiency. In particular, we aim to be useable in a jupyter notebook. It is a very important goal for us to reach a usability able to permeate into the educational side which is a strength of Python. In this library, we demonstrate how to achieve genericity and efficiency while remaining easy to use all at the same time. The scope of this library would be to specialize in mathematical morphology as well as providing very versatile image types. We leverage the modern C++ language and its many new features related to genericity and performance to break this rule in the image processing area. Finally, we attempt, through a static/dynamic bridge, to bring low level tools and concepts from the static world to the high level and dynamic prototyping world for a better diffusion and ease of use.

With this philosophy in mind, this manuscript aims at presenting our thesis work related to the C++ language applied to the Image Processing domain. It is organized as followed:

Genericity 2   presents a state-of-the-art overview about the notion of genericity. We explain its origin, how it has evolved (especially within the C++ language), what issues it is solving, what issues it is creating. We explain why image processing and genericity work well together. Finally, we tour around existing facilities that allows genericity (intrinsically restricted to compiled language) to exists in the dynamic world (with interpreted languages such as Python).

Images and Algorithms taxonomy 3   presents our first contribution which is a comprehensive work in the image processing area around the taxonomy of different images families as well of different algorithms families. This part explains, among others, the notion of concept and how it applies to the image processing domain. We explain how to extract a concept from existing code, how to leverage it to make code more efficient and readable. We finally offer our take about a collection of concepts related to image processing area.

Images Views 4   presents our second contribution which is a generalization of the concept of View (from the C++ language, the work on ranges [120]) to images. This allows the creation of lightweight, cheap-to-copy images. It also enables a much simpler way to design image processing pipeline by chaining operations directly in the code in an intuitive way. Ranges are the cement of news design to ease the use of image into algorithms which can further extend their generic behavior. Finally, we discuss the concept of lazy evaluation and the impacts of views on performances.

Static dynamic bridge 5   presents our third contribution which is a way to grant access to the generic facilities of a compiled language (such as C++) to a dynamic language (such as Python) to ease the gap between the prototyping phase and the production phase. Indeed, it is really not obvious to be able to conciliate generic code from C++ whose genericity is resolved at compilation-time (we call this the "static world"), and dynamic code from Python which rely on pre-compiled package binaries to achieve an efficient communication between the dynamic code and the library (we call this the "dynamic world"). We also cannot ask of the user to provide a compiler each time he wants to use our library from Python. In this part, we discuss what are the existing solutions that can be considered as well as their pros and cons. We then discuss how we designed an hybrid solution to make a bridge between the static world and the dynamic world: a static-dynamic bridge.

Chapter 2

Genericity

In natural language we say that something is generic when it can fit several purposes at once while being decently efficient. For instance, a computer is a generic tool that allows one to write documents, access emails, browse Internet, play video games, watch movies, read e-books etc. In programming, we will say that a tool is generic when it can fit several purposes. For instance, the gcc compiler can compile several programming languages (C, C++, Objective-C, Objective-C++, Fortran, Ada, D, Go, and BRIG (HSAIL)) as well as target several architectures (IA-32 (x86), x86-64, ARM, SPARC, etc.). Henceforth, we can say that gcc is a generic compiler. At this point it is important to note that even though a tool is deemed generic, there is a scope on what the tool can do and what the tool cannot do. A compiler despite supporting many languages and architectures, will not be able to make a phone call or a coffee. As such it is important to note that genericity is an aspect that qualifies something. We will now study the generic aspect related to libraries and programming languages.

This thesis voluntary leaves out the generic aspect related to the target architecture. Indeed, being able to write and/or generate code that is able to run on a large array of different hardware architecture is a field of research on its own and is not the main focus of this thesis. It is also known as *heterogeneous computing*. This field saw the birth of its own standards (SYCL [123, 122]) and libraries solving different problems, such as Halide [91], which provides its own DSL (Domain Specific Language) to write code that will run on GPUs. In pure C++ there exists several high performance math library for linear algebra, dense and sparse arithmetic which are optimized to produced very optimized code (vectorized instruction support, parallel execution etc.). The most popular libraries are Eigen [67], Blaze [79, 80, 78], Blitz++ [30] and Armadillo [108, 109, 121] leveraging *expression templates* [16] to achieve their goal. Also, we note that Eigen is compatible with GPU source code [151] and can be used inside Cuda kernels. A Cuda extension for Blaze was released recently [128] and allow its use in GPU code as well. Armadillo uses BLAS [36] as underlying linear algebra routines which enables one to link against the GPU-accelerated NVBLAS (nvidia) [152] or ACML-GPU (AMD) [87] as drop-in replacement for BLAS to offload the work on GPU. All those libraries have set performance as their main goal. They try to provide generic ways to solve issues related to parallelism and/or vectorization while making use of expression templates for lazy computing (which will be seen in section 4.3.3). They do not aim to be able to handle as many input types as possible, however, the lazy-computing techniques is used to generate new types on-the-fly. Henceforth, those libraries still need to embed generic facilities to handle their own internal set of types. This thesis addresses genericity at the input level rather than the target architecture level, henceforth, we will not address this topic here.

Rework style

History of genericity takes its root in Aug. 1978, year when Backus publishes his paper about functional programming [2]. Backus thinks that there exists five computation forms with which

one can build up all the rest of the computational infrastructure. Every piece of software, in theory, is built from those five functional forms. Furthermore, the initial work of Backus does not use the possibility offered by mutations in his five computational forms. These forms will lead to the birth of the functional programming paradigm (notably famous for its value immutability). Stepanov, a mathematician, thinks that those forms are theorems. He also thinks that there is an infinite number of theorem (as in mathematics) and that reducing their number to five for software programming is taking things a little too far. He publishes in 1987 [6] that one cannot ignore the mutability of states if one wants to achieve maximum efficiency. Stepanov reasons about software programming by drawing a parallel with algebraic structures. Indeed, let us consider the classical parallel computation model. In this model, it is needed to be able to reorder computation in order to have a reduction that works. *Reordering computation* can be reworded as *associative property* of an algebraic structure: the monoïd which is a triplet consisting of a data structure, an associative binary operation and a neutral element. Stepanov thinks that we extract those data structures and those laws/properties from software program the same way as we discover theorems and axioms in mathematics. Software is then defined on top of algebraic structures and it's the software programmer's job to discover the data structures and laws that compose them.

This reflection leads to the publication of [7] in which the term *Generic Programming* first appears. "By generic programming, we mean the definition of algorithms and data structures at an abstract or generic level, thereby accomplishing many related programming tasks simultaneously. The central notion is that of generic algorithms, which are parametrized procedural schemata that are completely independent of the underlying data representation and are derived from concrete, efficient algorithms." This articles lead to the genesis of the book *The Ada Generic Library linear list processing packages* [8] in which was published the first work about a generic library of algorithms and data structures. Then Stepanov and Lee wrote the first version of the Standard Template Library [15] in 1995. This standard template library was then incorporated alongside the C++ language for the release of the first ISO standard of the language in 1998 [19]. That same year is published [23]. This is the first place where the term *concept* appears as "a set of axioms satisfied by a data type and a set of operations on it." This term is designed to include the complexity of an operation as part of an axiom in software programming. This term is introduced to replace the previously used mathematical terms that could not include this notion of complexity. It is also the first place where the notion of *regular* type appears: "Since we wish to extend semantics as well as syntax from built-in types to user types, we introduce the idea of regular type, which matches the built-in type semantics, thereby making our user-defined types behave like built-in types as well." Efforts were made by Gregor et al. in [56, 55] in 2006 to introduce them into C++ but it ultimately failed, and the feature was pulled off of the C++ standard.

Stepanov presented *Generic programming* to Backus before publishing [66] and Backus "always knew that at some points he needed to figure out mutation into functional programming and one view of generic programming is that generic programming is functional programming with a well-defined way to handle mutation." Unfortunately Backus passed away before being able to write the forewords of *Elements of Programming* book. The term *generic programming* never appears in the book because Stepanov thought he lost control over it. It has evolved in practice into metaprogramming, effectively associated to C++ template metaprogramming instead of being associated with the underlying mathematics, algebraic structures, data structures and algorithms. This book introduces the *require* clause on algorithms in order to achieve constrained genericity.

A workshop was help and its summary was published in 2012 [84], referred to as the *Palo Alto* report, to summarize what design the committee wanted for concepts and what problem it would solve. *Elements of Programming* argues that just having constrained template was already incredibly useful, and the STL could be described in terms of *require* clauses. This

subset becomes known as "concepts light" and was enriched to becomes later what would be standardized in C++20.

Stepanov then published *From mathematics to generic programming* [98] that traces the history of algorithms and ties the history of mathematics with the history of generic programming. Indeed, Stepanov already gave a lecture in 2003 [44] where he traces in particular the history of the algorithm of gcd/gcm for 2500 years and explain how successive generation of mathematicians improved on by always looking for more generic ways to solve the same problem. In essence, the book presents generic programming as an extension of the evolution of mathematical algorithms over time. In this book, Stepanov is also reclaims the term of *generic programming* and differentiates it from template metaprogramming once and for all.

In this book Stepanov states that "Generic programming is about abstracting and classifying algorithms and data structures. It gets its inspiration from Knuth [96] and not from type theory. Its goal is the incremental construction of systematic catalogs of useful, efficient and abstract algorithms and data structures. Such an undertaking is still a dream." Stepanov thought of STL as a starting point to a very large library of data structures and algorithms, written in a generic form, that would all work together nicely. STL is intended to be an example of how industry should move forward and work on building a large library of this form. Hopefully, this is the direction the C++ standard committee is going toward.

Genericity within libraries   is described by the cardinality of how many use-cases it can handle. Libraries always provides their own data structures, to represent and to give a meaning to the data the user wants to process, as well as algorithms to process those data and provide different type of results. A library will be then labeled as *generic* [12] when (i) its data structure allows the user to express himself fully with no limitation and when (ii) its algorithm bank is large enough to do anything the user would want to do with its data. In reality such a library does not exist and there are always limitations. Studying those limitations and what reason motivates them is the key to understand how to surpass them in the future, by developing new hardware and/or software support for new feature enabling more genericity.

Genericity within programming language   is described by the ability of the language to execute a code statement over a large amount of data structures [23], be they native (char, int, ...) or user defined. It is nowadays primordial for a programming language to be able to do so. Indeed, in a world where Information Technologies are everywhere, the amount of code written by software developers is staggering. And with it so is the amount of bugs and security vulnerabilities. Being able to natively have a programming language that enables to do *more* by writing *less* mathematically results in a reduced development and maintenance cost. Programming languages offers many ways to achieve genericity which is dependent of the language intrinsic specificities: compiled or interpreted, native or emulated, etc.

Before delving into the specifics of what genericity implies for libraries and programming languages, let us introduce some vocabulary for the sake of comprehension. First is the notion of *type*. A *type* (or *data type*) is an attribute of data which tells the compiler or interpreter how the programmer intends to use the data. Most programming language support basic data types (also called primitive types) such as integer numbers, floating point numbers, boolean and characters (ASCII, Unicode, etc.). This data attribute defines the operation that can be done on the data, the meaning of the data and the size of the data in memory (the data can then be stored on the heap, stack, etc.). A data type provides a set of values from which an expression (i.e. variable, function, etc.) may take its values. Among programming language, we can distinguish those who are dynamically types and those that are statically typed. Statically typed languages are those whose variables are declared holding a specific type. This variable cannot hold data from another type in the scope it is declared. Statically typed programming languages are Ada, C, C++, Java, Rust, Go, Scala. Dynamically types languages are those whose variables can be

reassigned with a value of different type from the one it was initially declared to hold. The variable type is then dynamically changed to fit the new value it is holding. Dynamically typed programming languages are PHP, Python, JavaScript, Perl.

The consequence of being able to tell which type a variable is holding at all time (statically-typed language) is two-fold. For the developer, it is easier to reason about code and to spot bugs. For the compiler, it is possible to generate optimized binary code specific to this data type (vectorization, etc.). The consequence of being able to morph the types a variable can hold at run-time is mainly to serve prototyping purpose. When tweaking a Jupyter notebook, it is much appreciated not to be limited to a single type for each variable to be able to iterate on the prototype much faster.

In image processing, an image Im is defined on a domain $\mathcal{D}$ (which contains points) by the relation $\forall x \in \mathcal{D}, y = Im(x)$ where y is the value of the image Im for the point x. This definition always translates into a complex data structure when transposed into a programming language. This data structure must be aware of the data buffer containing the image data as well as information about the size and dimensions of the image. Furthermore, to add to the difficulty, the information needed to define precisely the data structure is not always known when writing the source code. Indeed, a very simple use-case consists in reading an image from a file to load it in memory. The file can contain an image of varying data type and the program should still work properly. There are multiple approach to solve this issue and we will address them in the following section 2.1 and section 2.2.

2.1   Genericity within libraries

Projecting the notion of genericity to Image Processing, we can deduce that we need two important aspects in order to be generic. First, we need to decorelate the data structures and its topology and underlying data from the algorithms. Indeed, we want our algorithms to support as much data structures as possible. Second, many algorithms share the same computational shape and can be factorized together.

image 2D   graph   mesh

input:

output:
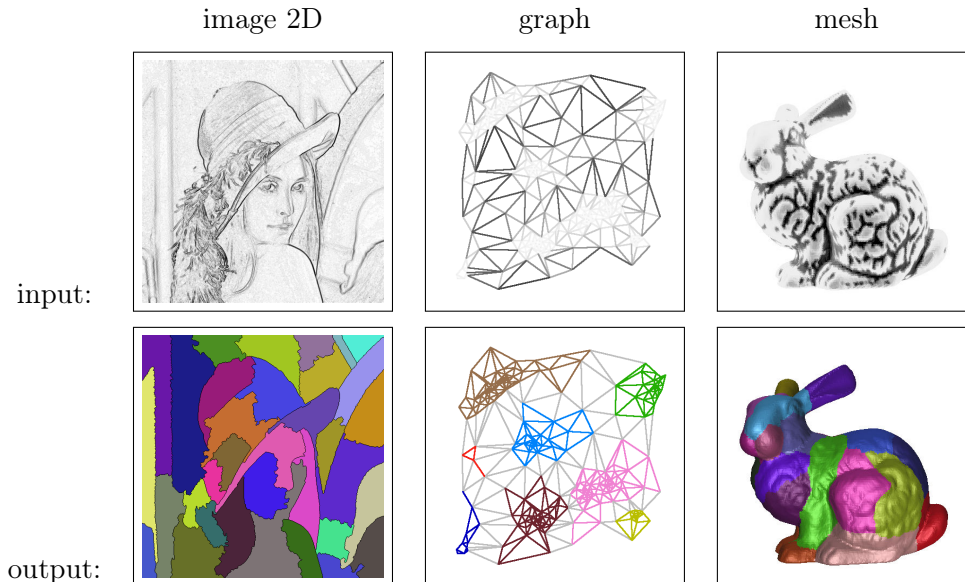
The same code run on all these input.

Figure 2.1: Watershed algorithm applied to three different image types.

Genericity can have two different meanings depending on the people you ask. For instance, some will argue that genericity is high level and qualifies a tool which is "generic enough" to

handle all of his use-cases. Others will argue that genericity is about how the code is written: generic enough to handle all the use cases possibles. Neither is wrong. However, for the sake of comprehension we will use different words for each of these cases. A tool generic enough to handle a lot of use-case will be called *versatile*. Finally, for a tool whose aim is to provide a programming framework to handle code of any use-case we will use *generic*. In this thesis, genericity will be about code. The figure 2.1 illustrates this result of the same generic watershed implementation applied on an image 2D, a graph as well as a mesh.

In image processing, there are three main axes around which genericity is applying. The first axis is about the data type: grey level or RGB color (8-bits, 10-bits), decimal (double) and so on. The second axis, is about the structure of the image: a contiguous buffer (2D or 3D), a graph, a look-up table and so on. Finally, the third axis is about additional data that can be fed to image processing algorithms: structuring element (disc, ball, square, cube), labels (classification), maps, border information and so on. In the end, an image is just a point within this space of possibilities, illustrated in 2.2. Nowadays, it is not reasonable to have specific code for every existing possibility within this space. It is all the more true when one wants efficiency.

Figure 2.2: The space of possible implementation of the *dilation(image, se)* routine. The image axis shown in (a) is in-fact multidimensional and should be considered 2D as in (b).

Genericity is not new and was first introduced in 1988 by Musser et al. [7]. The main point is to dissociate data structures and algorithms. The more your data structures and algorithms are tied together, the less you will be generic and will fail to handle multiple data structures in the same algorithm. Further work has been made about genericity in [12, 23]. Those works highlight the notion of abstraction to be able to turn an algorithm tied to a data structure into a generic algorithm. Notably in [66], Stepanov digs further and introduce the notion of *Concepts*, which are static requirements about the behavior of a type, by showing how to design a generic library and its algorithms. He highlights the importance of having the algorithms driving the behavior requirements, and not the opposite. These works are very suitable to be applied in the area of Image processing where we typically have a lot of algorithms (also called operators) that are required to work on a lot of different data structures (also called image types).

The authors explain in [131] how to capitalize on those works to turn a data-structure-specific image processing algorithm into a generic algorithm. We also explain how *concepts* can ease the implementation of generic algorithms. This approach is implemented in a library [117] which allows us to provide a proof of concept over the feasibility of having generic image processing operators running on multiple image types with near-native performances. Let us first explain briefly how we achieved this.

2.1.1 Different approaches to get genericity

First, let us consider the morphological *dilation* that takes two inputs: an image and a flat structuring element (SE). The set of some possible inputs is depicted in 2.2. Without genericity, with s the number of image type, v the number of value type and k the number of structuring elements, one would have to write $s * v * k$ different *dilation* routine.

There are several ways to reach a high level of genericity. First there are the *code duplication* approach as well as the *generalization* approach. Finally, there is a way that consists in using expert, domain specific tools specifically engineered for this purpose and build upon them: those tools usually make heavy usage of *inclusion & parametric polymorphism*, also known as template metaprogramming in C++, to provide the basic bricks the user needs to build upon.

Code duplication approach   consists in writing and optimizing the algorithm for a particular type in mind. Then, each time a new type is introduced, all the algorithm must be rewritten for this specific type. Additionally, each time a new algorithm is introduced, it must support all the existing types and thus be written multiple times. This approach does not scale well when the complexity of algorithms grows, and the number of data types increases. Neither it does allow the implementer to easily make use of optimization opportunities that can be offered by different data types having a common property. This translates into heavy switch/case statement in the code as show in 2.3 that illustrate how the *fill* algorithm needs to dispatch according to the input data type.

```
1    // image types parametrized by their
2    // underlying value type
3    template <ValueType V> struct image2d<V> { /* ... */ };
4    template <ValueType V> struct image_lut<V> {/* ... */};
5    // ...
6    void fill(any_image img, any_value v)
7    {
8      switch((img.structure_kind, img.value_kind))
9      {
10     case (BUFFER2D, UINT8):
11       fill_img2d_uint8( (image2d<uint8>) img,
12                         (uint8) any_value );
13     // ...
14     case (LUT, RGB8):
15       fill_lut_rgb8( (image_lut<rgb8>) img,
16                      (rgb8) any_value );
17     }
18   }
```

Figure 2.3: Fill algorithm skeleton with a switch/case dispatcher to ensure exhaustivity.

In addition, it is important to note that the exhaustivity aspect is only illustrated regarding the data structure types here. Indeed, the data structures are all already generic for their underlying data type (named *ValueType* in the code). When one write `image2d<uint8>` (l.10), it means *2D-image whose pixels' have a single channel 8-bits value*. This approach enables one to write an algorithm at maximum efficiency for a particular data type, however one can easily miss optimization opportunities if not knowledgeable enough too. This approach is best for early prototypes and trying to find common behaviors pattern among algorithms, or common properties across different data types. No IP library has chosen this approach due to the obvious maintenance issue induced.

Generalization approach   consists in finding a common denominator to all the image types. Once designed, this common denominator, also called supertype, can store informations about all the supported image types by the library. This supertype allows the library developer to write all the algorithms only once: for the supertype. The processing pipeline will then consist in three steps. First convert the input image type into the supertype, second process the supertype into the algorithm pipeline requested by the user, finally convert back the resulting image into the specific image type the user is expecting. This approach offers the advantage of being maintainable. Adding a new image type is just a matter of providing the two conversions facilities: to and from the supertype. Adding an algorithm is also just a matter of writing it once for the supertype. This mechanism is shown in 2.4. However, one must keep in mind that the

conversion can be costly. Also, processing the supertype may induce a significant performance trade-off while processing the original type would be much faster. Furthermore, it is not always possible to find this common denominator when enumerating through some esoteric data types. Finally, the provided interface (from the supertype) may allow the image to be used incorrectly, such as a $2D$ image being processed into video $(3D + t)$ algorithm. Widely use libraries such as OpenCV [22], scikit-image [99] use this technique to handle as many image types as possible. Another library making use of this generalization technique in its implementation is CImg [85]. CImg generalize its data type to a 4D image type templated by its underlying data type.

```cpp
struct image4D { // generalized supertype
  // generalized underlying value-type
  // every value is converted to this one
  using value_type = std::array<double, 4>;
  /* ... */
};
// specific types w/ conversion routines
struct image2D { image4D to(); void from(image4D); };
struct image3D { image4D to(); void from(image4D); };
// ...
void fill(image4D img, const std::array<double, 4>& v) {
  for(auto p : img.pixels())
    p.val() = v;
}
```

Figure 2.4: Fill algorithm for a generalized supertype.

Inclusion & Parametric polymorphism approach consist in extracting behavior patterns from algorithms to group them into logical brick called *concepts* (for static parametric polymorphism), or *interface* (for dynamic inclusion polymorphism). Each algorithm will require a set of behavior pattern that the inputs need to satisfy. In C++, it can be done either by using inclusion polymorphism, or by using parametric polymorphism, as shown in 2.5. In [131], we leverage a new C++20 feature (the concept) to show how it is possible to turn an algorithm, specific to an image type, into a more abstract, generic one that does not induce any performance loss. This approach, especially applied to image processing, will be seen more in-depth in chapter 3.

(c)   (d)
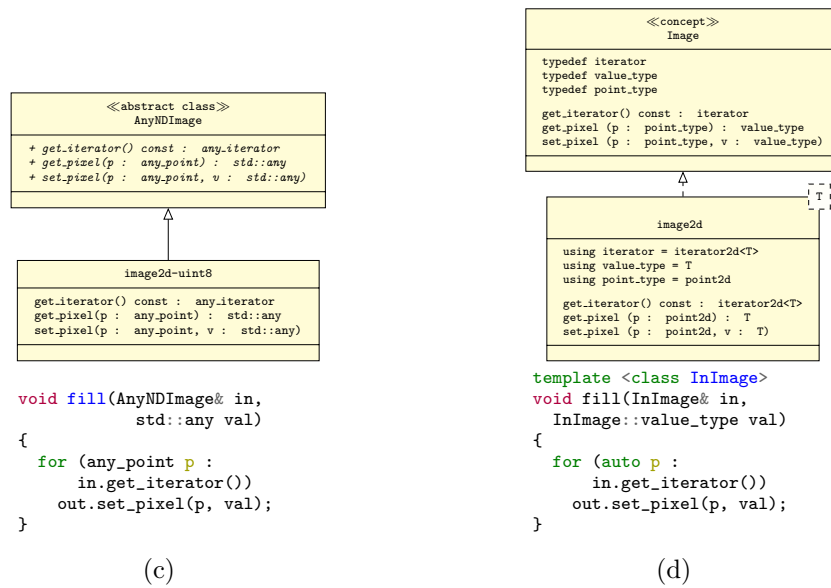
Figure 2.5: Dynamic, object-oriented polymorphism (a) vs. static, parametric polymorphism (b).

Multiple libraries exist and leverage this approach to try to achieve a high genericity degree as well as high performance by offering varied abstract facilities over image types and underlying data types. Those are ITK [90], Boost.GIL [52], Vigra [27], GrAL [51], DGTal [103], Milena [65, 68], Olena [161, 74, 77, 97] and Pylena [117]. Most of them have been written in complex C++ whose details remain visible from the user standpoint and thus are often difficult and complex to handle. It is also harder to debug because errors in highly templated code shows up very deep in compiler error trace.

The table comparing all the pros. and cons. from the aforementioned approaches is presented in table 2.1. We can see in this table that Generic Programming in C++20 check all the boxes that we are interested in.

Table 2.1: Genericity approaches: pros. & cons.

Paradigm	TC[1]	CS[2]	E[3]	One IA[4]	EA[5]
Code Duplication	✓	✗	✓	✗	✗
Code Generalization	✗	≈	≈	✓	✗
Object-Orientation	≈	✓	✗	✓	✓
Generic Programming:					
with C++11	✓	≈	✓	✓	≈
with C++17	✓	✓	✓	✓	≈
with C++20	✓	✓	✓	✓	✓

[1] TC: type checking.
[2] CS: code simplicity.
[3] E: efficiency.
[4] One IA: one implementation per algorithm.
[4] EA: explicit abstractions / constrained genericity.

2.1.2   Unjustified limitations

Image processing community operates mostly with either Python or Matlab [38]. As such this subsection will focus on those two technologies. Python offers access to two major libraries for image processing: OpenCV and scikit-image. Matlab has built-in support as well as toolboxes for more advanced features. When we intersect scikit-image and Matlab, we can notice that both are very similar both in feature and interface. As such, it is possible to regroup them both here for the sake of comprehension. As stated above, when considering a generic library, one must consider the three axes: underlying data type, domain structure and additional data. Let us compare how the mentioned library behave along those axes with a simple algorithm such as the morphological dilation.

Limitations regarding feasibility

Data type   Dilating a grayscale or a binary image works fine as intended with all the libraries. However, there is no trivial way of dilating a RGB colored image [59, 50] as this operation is not defined for colored images. Indeed, the algorithm is able to work if a supremum function is provided. Such functions have multiple possible implementation and selecting the correct one is not trivial. However, provided a supremum function, the dilation algorithm should normally be performed. Despite that fact, scikit-image does not allow one to dilate a colored RGB image and raises an error: it is required to convert into a grayscale the image beforehand.

OpenCV arbitrarily decides that the colored dilation consists in dilating each channel of the colored image separately from one another. It is effectively selecting a partial marginal order relation under the hood. This arbitrary choice may cause false colors to appear in the resulting image (which most of the time is not what the user wants). Furthermore, it is not possible to

provide a supremum function to the dilation algorithm to customize the behavior which is a server drawback.

Domain structure   To perform a dilation, it is required to have a structuring element whose shape matches the structure of the domain of the image. For instance, dilating a 2D-image requires the use of a structuring element whose shape may be a disc or a rectangle. To dilate a 3D-image, one would need to use a structuring element whose shape is of a ball or a cube. Scikit-image supports 3D-images as well as structuring element whose shape are compatible (ball, rectangle and octahedron). This naturally leads to having a support for the dilation of 3D-images. On the other hand, OpenCV does support 3D-images whereas its dilation algorithm cannot handle them. The algorithm exits with an error. Worse, when passing a wrong structuring element (a rectangle) to the dilation algorithm alongside the 3D-image, the algorithm works and produce a result which is false: it is similar to the application of the 2D-structuring element on each 2D-slice of the 3D-image.

Limitations regarding optimizations

Each library has its own strategies to optimize its routines when implementing them.

Scikit-image   for instance, will check whether the structuring element is separable (only for rectangle shapes) so that it can dispatch on an optimized multi-pass 1D routine for each part separated which linearize the execution time and greatly improve performances for large structuring element.

Also, Scikit-image relies on SciPy internals which does not abstract the underlying data type for the algorithm implementer. As such, each algorithm must provide a switch/case dispatch for every supported type (floating points, 8-bit channel, 16-bit channel, RGB, etc.), and it must provide it in the middle of the algorithm implementation. If one type is not natively supported; an error occurs and the program halts. Henceforth, handling a new supported data type will requires to review every single written algorithm.

On the other hand, SciPy provides an abstraction layer over the dimensional aspect of the image by providing a tool named point iterator. This tool allows one to iterate over every point of the image, without being aware of the number of its dimension, and make the translation from the abstract iterator to the actual offset in the data buffer of the image. The implementer can then only worry about handling the underlying data type to provide a generic algorithm. This approach, sadly, is fully dynamic (that is, runtime) and does not allow the compiler to provide native optimization such as vectorization out of the box.

OpenCV & Matlab   In OpenCV as well as in Matlab, the choice was made to systematically attempt to decompose big rectangular structuring elements into smaller $3*3$ structuring elements. This is not as effective as using multi-pass 1D algorithm but still allows for relatively stable performances.

Also, OpenCV let the implementer handle the cases he wants to support by himself. For instance, the dilation algorithm is written with a dispatch on the data type before the actual call to the algorithm. This enables compiler optimizations such as vectorization because all the required information is known at the right time. It also enables offloading the computation into GPU kernels when feasible. However, the downside is that few algorithms are written in a way to handle multidimensional images. Most are written to only handle specific subsets. As such, conversion from one subset to another may be unavoidable when writing an algorithm pipeline for a more complex application. For instance, it is currently not possible to dilate a 3D image with a 3D ball (as stated above).

Another point to note with OpenCV is the requirement to do temporary copies (to extract data or to have working copy) when writing an algorithm. For instance, it is currently not

possible to write a dilation algorithm operating only on the green channel of an RGB image. One must first extract the green channel into a single channel temporary image, blur that image, to finally put the result back into the original image. Generally, in-place computation is poorly handled in OpenCV.

Performance discussion   When comparing performances of the simple dilation between Matlab and OpenCV, which is done in [83], shows that Matlab is very oriented toward prototyping and not toward production. The performance gap between the two libraries shows that performances may not a major concern for MatLab in this case. Opposite to this, OpenCV and scikit-image both have a C/C++ core to provide fast basic algorithms such as the dilation and erosion mathematical morphology.

As such, when comparing the performances of OpenCV, Scikit-image and Pylene in fig. 2.6, we can notice some interesting facts. Both scikit-image and Pylene have a very stable execution time even though the size of the structuring element grows by power of two. This corroborates the fact that the author did see code taking advantage of the structuring element's properties, such as the decomposability/separability. OpenCV has very good performances for a square because it has specific handwritten code for both vectorization and GPU offloading when possible: even if OpenCV decomposed its square into smaller sub square (and not periodic lines), it remains steady fast.

Figure 2.6: Benchmark: dilation of a 2D image (3128x3128 $\approx$10Mpix) with a 2D square and a 2D disc.

In the case of a structuring element shaped as a disc (also in fig. 2.6), we can observe that the execution time raises exponentially for both Scikit-image and OpenCV whereas Pylene remains regular and steady fast. These results show that Pylene's attempt to decompose each structuring element into periodic lines when possible may be slightly slower for smaller structuring elements whereas it is much more regular and faster when the structuring element start to be of a certain size.

Limitations regarding static types in a dynamic world

Being statically typed (at language level) has its advantages (performance, optimization) but also has sever drawbacks, especially when interfacing with dynamic languages such as Python. Indeed, being statically typed in a dynamic world disqualify the library from being able to select, for instance, a new, custom structuring element, when performing a dilation. Indeed, all

the supported structuring elements must be listed in advance and code must be compiled for each supported type. This induces a strong inertia when the practitioner wants to try out new things not yet supported by the library. This also defeats the initial purpose of being generic. Indeed, one would expect a generic library to work with any type and not just a specified set of pre-compiled types. Being able to break through this limit is addressed more in-depth in chapter 5.

2.1.3   Summary

To conclude, a generic algorithm will not be faster after it is first written, but will provide acceptable performances for most cases. However, a generic algorithm provide opportunities for the implementer to take advantage of some properties from input types in order to be faster. The main advantage is that once those optimisation opportunities are written once, they are available for every input types (that match the property). Genericity enables code dispatching very easily based on type property which may induce almost no runtime overhead, depending on which generic strategy was chosen in the library.

2.2   Genericity within programming language

Genericity is a more than 45 years old notion. It was first introduced alongside the CLU language in 1974 by Barbara Liskov and her students [11]. The language offered many features such as data encapsulation, iterators and especially *parametrized modules*. A module in CLU is represented by a *cluster*. A *cluster* is a programming unit grouping a data structure and its related operations. In modern programming, we would call it a class where the data structure correspond to the member variables and the operations correspond to the member functions (or methods). In CLU, the clusters can be parametrized which is a way to introduce the notion of parametric polymorphism. Indeed, clusters offered the ability to define a generic data structure and its functions whose behavior will not change whatever type the cluster is parametrized with. At first, type-safety was enforced at run-time in CLU. Later, *where clauses* were introduced to specify specific requirements over cluster type parameters. This had the consequence of allowing only the operations required in the where clause to be used in the cluster implementation code. This enables proper compile-type checks, type-safety and the compilation into a simple native and optimized code by the compiler. The following code illustrate how to create a cluster (or class) named *vector*, declaring a set of member functions and requiring a specific behavior (operator = and <) from its underlying stored elements:

```
vector = cluster [T: type] is
  create, size, contains, sort, remove, push_back
where
  T has equal: prototype(T, T) return (bool)
  T has less: prototype(T, T) return (bool)
```

When implementing the member functions declared for `vector`, the only valid operations that can be performed on a value of type `T` are `equal` (= for comparison) and `less` (< for ordering). In CLU the actual instantiation of the parameter (here `T`) is done at run-time. Indeed, each type is represented by a descriptor, even a parametrized type parameter. At run-time, a concrete type is used to instantiate the cluster. To achieve this, the placeholder type descriptor is replaced by the one of the actual concrete type at this moment. The instantiation can then be considered dynamic which differs from the C++ compilation model where the template instantiation is considered static (i.e. done at compile-time). The pros and cons of this approach are discussed in [1] which essentially are the resulting binary code size due to combinatorial explosion combined with optimized code generation versus small, flexible and slower binary.

Later on came the Ada programming language whose conception work started in 1977 and first version was released in 1980. It was first standardized in 1983 [3] in the United States by

the American National Standards Institute (ANSI) before being Internationally standardized by
the International Organization for Standardization (ISO) in 1987 [5]. From this point on, the
Ada language released a new version in 1995 [14, 33], then the standard committee published
an Amendment in 2007 [60] which is often referred as Ada 2005. Finally, a new version of the
language was published in 2012 [81, 105]. What interest us in the Ada language is the fact that
the language features *generics* since it was first designed in 1977-1980 which *ten years* before
Musser and Stepanov published their first work about genericity in 1988 [7]. Indeed, it took
around 20 years for the fist Standard Template Library (STL) will then be standardized in the
C++ programming language in C++98 [19]. It then took almost 10 more years for an STL to
land into the Ada programming language in the 2007 Amendment for Ada 2005.

In the Ada programming language, it is possible to mark a package or a procedure as generic
with the `generic` keyword. The developer then lists the parametrized parameter the package
and/or the function requires to be implemented. The following code demonstrates how this fea-
ture work by implementing a function replacing the first argument by a new value and returning
its previous value (also known as exchange).

```
generic
  type T is private;
function exchange (x : in out T; v : in T) return T is
  tmp : T;
begin
  tmp := x;
  x    := v;
  return tmp;
end exchange;
```

In Ada, generic packages and routines (procedures, functions) must be instantiated explicitly:
the compiler cannot infer the parametrized type at compile time from the context of use (whereas
a C++ compiler can). Thus, the following three lines must be written explicitly for the compiler
to generate the binary code of the above function:

```
function int_exchange is new exchange (Integer);
function float_exchange is new exchange (Float);
function str_exchange is new exchange (String);
```

In Ada, this model enables the possibility of sharing a generic across several compilation units
since its compilation is independent of its use whereas in C++, until C++20 the sharing model
consists in copy-pasting the whole source code (and transitive recursive dependencies) each time
a generic (a template) code is compiled. C++20 (standardized in 2020, 42 years after C++98)
bring a solution to this issue by standardizing C++ modules. However, this feature is out of
scope of this study and will not be discussed in this thesis.

Also, Ada support syntactic constraints on parameters similar to the CLU language. It
is translated into a `with` clause listing the constraints on the parameter(s). The following code
shows how to implement such a constrained generic function by using the mathematical maximum
operator as an example:

```
generic
  type T is private;
  with function "<" (x, y: T) return Boolean is <>;
function maximum (x, y : in T) return T is
  tmp : T;
begin
  if x < y then
    return x;
  else
    return y;
  end if;
end maximum;
```

Let us not forget to instantiate the constrained generic function for integers:

```
function int_maximum is new maximum (Integer);
```

This idea of constrained generic is 30 years old and has only made his way very recently (2020) in the C++ programming language under the feature name of *concepts*.

Now that we have introduced where generic programming is coming from, let us focus on the C++ programming language and how it has made his way into the language along the last 30 years, as well as along the years to come.

2.2.1 Genericity in pre C++11

Before C++11 [73] came out the genericity facilities offered by the C++ programming language were already Turing-complete [45]. However, it was lacking a certain amount of features the language now have that made writing generic code a real challenge at that time. For instance, when writing code with a variable number of types (nowadays designated as variadic templates) one had to write the generic code for each and every number of type supported. This meant that to implement `std::tuple`, one had to copy the implementation for every number of type supported by `std::tuple`. This limitation defeated the very first principle and motivation of generic programming which is to write *less* Code. To compensate, library implementers used tricks with macro not to have to rewrite code which made the initial code even harder to understand for outsiders.

Functions on types (or traits)

The very first feature every developer writing metaprogramming C++ code has used is called *type-traits*. Those *traits* are a way to mutate a type depending on the way a template declared on a data structure is resolved. Indeed, in C++ the instantiation of a template depends on the context, which means the compiler is required to build a set of possible way to resolve a template instantiation and in order to determine the best match to resolve the instantiation. This mechanism is well known and documented in the C++ standard, it can then be used (and abused) to do a great number of things. Among all the traits that exist (a lot of them have been standardized in C++11, 2011, in the header `<type_traits>`), some in particular are used in every codebase: `remove_const`, `remove_volatile` and `remove_reference`. Let us see how they are implemented:

```
template<class T> struct remove_const                  { using type = T; };
template<class T> struct remove_const<const T>         { using type = T; };

template<class T> struct remove_volatile               { using type = T; };
template<class T> struct remove_volatile<volatile T>   { using type = T; };

template<class T> struct remove_reference              { using type = T; };
template<class T> struct remove_reference<T &>         { using type = T; };
template<class T> struct remove_reference<T &&>        { using type = T; };
```

For `remove_const`, first is defined the structure whose underlying alias `type` points to the passed template parameter `T`. Then we define a template specialization whose matching parameters are all `T` parameters that are `const`. The defined underlying alias `type` for this specialization then is `T` without the qualifying `const`. This way, there are two possibilities when calling this `trait` (or metafunction): either the passed parameter is not `const` which means it will be forwarded as-is to the underlying alias `type` or the passed parameter is `const` which means the underlying alias `type` will be defined by dropping the `const-qualifier` off of the passed parameter type. For instance:

```
using T1 = remove_const<double>::type // T1 is double
using T2 = remove_const<const double>::type // T2 is double too, const-qualifier is dropped
```

This language construct is very useful when developing generic libraries because it enables performing "functions" on types, and even chain them. It is also possible to perform checks to extract information about those types. We can easily write a `is_const` metafunction if we need it.

In image processing in particular, the usage of traits in the generic library Milena [65, 68] was very useful to achieve standard ways to compute very useful types from other complex type. From the image processing definition of an image type, we can already see a number of traits that a generic library would want to provide. Indeed, it would be useful to be able to extract the type of the domain (box2d, box3d, etc.), the type of the point (point2d, point3d, ...), the underlying value type (uint8, rgb8, etc.) and so on. We can also already see emerging consistency issues between those types. Indeed, a box3d domain would not accept access via a point2D. It would instead require a point3d. All those issues will be addressed later in the chapter 3. However, we can already give here minimal working example as to how type traits are especially useful in a generic image processing library. First let us define some minimal data structures:

```cpp
struct point2d {int x, y; };
string rgb8 { uint8_t r, g, b; };
struct box2d {
  using value_type = point2D;
  // ...
};
struct image2d {
  using domain_type = box2d;
  using point_type  = point2d;
  using value_type  = rgb8;
  // ...
};
```

Now we would want to implement `traits` to extract type information from those structures. Here is how we do it in a generic library:

```cpp
template <class T> struct domain_value_type { using type = typename T::value_type;  };
template <class T> struct image_point_type  { using type = typename T::point_type;  };
template <class T> struct image_value_type  { using type = typename T::value_type;  };
template <class T> struct image_domain_type { using type = typename T::domain_type; };
```

These traits extract information about types in a generic way and can be used in any algorithm taking an image as a template parameter. For instance, here is how an image processing algorithm (trivially extracting the max value) would b e written:

```cpp
template <class Ima>
typename image_value_type<Ima>::type // usage of trait
min_value(const Ima& ima) {
  using value_t = typename image_value_type<Ima>::type; // usage of trait
  auto min = std::numeric_limits<value_t>::max();
  for(auto v : ima.values()) {
    auto min = std::min(min, v);
  }
  return min;
}
```

SFINAE: Substitution-Failure-Is-Not-An-Error

Additionally, another feature related to metaprogramming allowed the developer to design generic libraries: the SFINAE [39] (substitution-failure-is-not-an-error) technique that leads to the popularization of the usage of the `std::enable_if` metaprogramming construct. The SFINAE technique relies on a feature of the C++ programming language. Indeed, when standardizing how the compiler should resolve and select function overloads, in a templated context, the standard committee chose to have the following behavior: "when substituting the explicitly specified or deduced type for the template parameter fails, the specialization (function overload candidate) is discarded from the overload set (of matching functions) instead of causing an error".

This feature allows writing code that seem to be ill-formed, for instance in a function, trying to access a class member type, variable or function that does not exist should be ill-formed. However, because it happens in a templated context during the instantiation resolution, when the compiler tries to instantiate a function template with a parametrized type, the compiler will just discard the function from the overload resolution set at call-site instead of throwing a hard error. An error can occur only when the compiler tried all the overload it knows in the overload set and still could not find an overload that was not ill-formed. If this happens, the compiler will then proceed to list all the overloads it tried, to list all the template substitution it tried and finally to list why it failed. This mechanism is the very reason of the unpopularity of this technique because it leads to situation where the compiler can output *several Mos* of error message for one single file. Error messages become incomprehensible very fast and programs are hard to debug because everything happens at compile-time. But still, it was the only technique we had to perform any kind of detections on types at compile time to require some any given constraints on them.

For instance, here is some real-world example extracted from generic image processing code that allows implementing the *fill* algorithm in two very different ways depending on how behaves the input image type. First we need to write the detector which is a structure whose templated context will be ill-formed during template instantiation. This detector will inherit either `std::true_type` or `false_type` depending on whether the detection is successful or not:

```
// Step 1 write detector
template <class Ima, class = void>
struct is_image_with_lut : std::false_type {};

template <class Ima>
struct is_image_with_lut<Ima,
   typename Ima::lut_type // constraint over the existence of the lut_type field
> : std::true_type {};
```

Now let us introduce a new image type for the sake of this example:

```
struct image2d_lut : image2d {
  using lut_type = std::array<uint8_t, 256>;
  using value_type = uint8_t:
  lut_type& get_lut();
};
template <class Ima>
struct image_lut_type {
  using type = typename Ima::lut_type;
}
```

The next step is to implement the `enable_if` facility. This is included in the C++ STL starting from C++11 and onward:

```
template<bool B, class = void>
struct enable_if {};

template<class T>
struct enable_if<true, T> {
  using type = T;
};
```

Now we are all set to use all those construct to dispatch our algorithms from call-site depending on our input image type:

```
// Overload #1 : with lut
template <class Ima,
void fill(Ima& ima, typename image_value_type<Ima>::type val,
          typename enable_if<is_image_with_lut<Ima>::type::value, void*>::type = 0)
{ // Image with lut
  using lut_t = typename image_lut_type<Ima>::type;
  using value_t = typename image_value_type<Ima>::type;
```

```
  lut_t& lut = ima.get_lut();
  for(value_t& v : lut)
    v = val;
}


// Overload #2 : without lut
template <class Ima,
void fill(Ima& ima, typename image_value_type<Ima>::type val,
          typename enable_if<not is_image_with_lut<Ima>::type::value, void*>::type = 0)
{ // Image without lut
  using value_t = typename image_value_type<Ima>::type;
  for(value_t& v : ima.values())
    v = val;
}
```

Finally, let us give some call-site example to finish illustrating our point:

```
image2d_lut ima1;
image2d ima2;

fill(ima1, 0);   // will dispatch over overload #1
fill(ima2, 255); // will dispatch over overload #2
```

Here, no hard error occur. Both overload are dispatched according to the constraint built with the SFINAE construct. Now we can talk about *constrained genericity* in C++. If we had an algorithm returning a value instead of performing an in-place computation, we would have been able to write the code with the SFINAE construct in the return-type instead of having it in the templates parameters. However that is not the case here.

CRTP: Curiously Recurring Template Pattern

Another features that precede C++11 and was available in C++98 [19] and C++03 [43] is the curiously recurring template pattern (CRTP) introduced in 1996 by Coplien [18]. This programming technique allows a base class (in its specific code) to be aware of its derived class at compile type. We made extensive use of this pattern in the past to design a new paradigm, SCOOP [42, 54, 63, 74] which combined multiple inheritance (via CRTP) and concept checking (via Boost.Concept [29, 58]) to implement a solution to provide a library with constrained genericity in C++ for the image processing area. The Scoop paradigm relied on the fact that multiple (especially diamond) inheritance did not pose much of an issue as long as there was no member variable involved. This way, one could consider a class hierarchy as a hierarchy of constraint indeed. Using CRTP, it was possible to find back, inside constraining classes, what was the concrete leaf class of the hierarchy in order to check whether its implementation satisfied the constraints or not. For instance, let us assume that we have a concrete class `image2d` inheriting from a constraining class `Image`. Let us now see how one would use the SCOOP paradigm to implement it. First we need to implement the satellite constraints around the image type which are related to the underlying point and domain of the image.

```
template <class P>
struct Point { // concept checking class
  int (P::*m_ptr1) const = & P::dim;
};

struct point2d : Point<point2d> { // CRTP
  int x, y;
  int dim() const { return 2; }
};

template <class D>
struct Domain { // concept checking class
  using value_type = typename D::value_type;
  // ...
};
```

```cpp
struct box2d : Domain<box2d> { // CRTP
  using value_type = point2d;
  // ...
};
```

For the sake of brevity, we are omitting the implementation of domain functions related to iterating over all the points of the domain. With those concepts defined, we can now dig into how we would implement an image class with the SCOOP paradigm.

```cpp
template <class I>
struct Image { // concept checking class
  using value_type = typename I::value_type;
  using point_type = Point<typename I::point_type>;
  using domain_type = Domain<typename I::domain_type>;

  const domain_type& (I::*m_ptr1)() const = & I::domain;
  // ...
};

struct image2d : Image<image2d> { // CRTP
  using value_type = uint8_t;
  using point_type = point2d;
  using domain_type = box2d;

  // ... ctors ...

  const domain_type& domain() const {
    return m_dom;
  }
private:
  domain_type m_dom;
};
```

Each time a leaf class in the hierarchy inherit from a base class, the concepts are checked. The syntax is not really intuitive, especially writing the function pointers to check that their prototype is confirming a certain behavior, but it was all what was available at that time. To be completely exhaustive, the library implementing this paradigm, Milena [68, 65] featured another powerful tool which allows one to require only concept class into input data for algorithms. Inside the algorithm it was then possible to get back the concrete class and use it as if it was originally passed as an argument. To do so, it was needed to add a member field to the concept class named **exact_t** that kept track of the leaf class into the concept checking class.

```cpp
template <class I>
class Image {
  // ...
  using exact_t = I;
};
```

Then a simple cast routine would do the trick inside the algorithm:

```cpp
#define EXACT(Ima) \
  typename Ima::exact_t

template <class Ima> // mutable routine
EXACT(Ima)& exact(Ima& ima) { return static_cast<EXACT(Ima)&>(ima); }

template <class Ima> // const routine
const EXACT(Ima)& exact(const Ima& ima) { return static_cast<const EXACT(Ima)&>(ima); }
```

This way, in image processing algorithms, the implementer would only need to write minimal code for it to work out of the box:

```cpp
template <class I>
void fill(Image<I>& ima, typename image_value_type<Ima>::type val) {
  EXACT(Ima)& ima_ = exact(ima);

  // use the concrete underlying image ima_ and val
}
```

This constrained genericity would be totally transparent as shown with the following code that just works out of the box, thanks to the SCOOP paradigm and its inheritance strategy to constrain classes.

```
image2d ima = /* ... */;
uint8_t val = 0;
fill(ima, val);
```

By extension, this work on Milena was integrated in the image processing platform Olena [161, 77]. This platform centralizes the work that was done around this field of research for a long time [25, 24, 37, 46]. More details can be found about SCOOP, and notably how it enables property based programming (augmentation of types via properties) in the work of Levillain [74, 65, 68, 69, 75, 74, 82, 97].

Those approaches have the advantage of being really flexible and to be able to perform the concept checking work that we wanted to have, be it to constrain implementation or to dispatch to the correct overloaded algorithms depending on specific properties. The disadvantages come in the form of an increased complexity of the design hierarchy of implemented types as they must inherit concepts via CRTP and conform to specific constrains. Also, all the implementation is visible as all the code is generic (template) and all the implementation details are leaked to the user code. For an image processing library which can use several dependencies (for instance, a library that read images from disc from multiple image formats), this is a huge drawback.

With the release of new C++ standards in 2011 [73], then 2014 [95], 2017 [113] where template metaprogramming facilities were greatly improved, it was necessary to review once more this design to improve the design in order to achieve genericity, performance and ease of use. This is the birth of a new library, Pylene [117]. In the end, it was C++20 [73] that marked the shift wanted by Stepanov [7, 12, 23, 66] and Stroustrup [13, 20, 158, 61] for years with the coming of Concepts, and all the new possibilities it brings to the programmer.

2.2.2   Genericity in post C++11 (C++20 and Concepts)

```
1   template <Image Ima, ValueType Val>
2     requires same<Ima::value_type, V>
3   void fill(Ima ima, Val val) {
4     for(Val v : ima)
5       v = val;
6   }
```

Figure 2.7: Fill algorithm, generic implementation.

Most of the algorithms are *generic* by nature. What limits their genericity is the way they are implemented. This statement is justified by the work achieved in the Standard Template Library (STL) [23] in C++ whose algorithms are implemented and designed in a way where they work with all the built-in collections (linked list, vector, etc.). Let us take the example of the algorithm `fill(Collection c, Value v)` which set the same value for all the element of a collection (see fig. 2.7). There are three main requirements here that are not related to the underlying type of *Collection*. First, we check (l.2 2.7) that we are actually filling the collection with the correct type of value. Indeed, it would not make sense, for instance, to assign an RGB triplet color into a pixel from a grayscale image. Secondly, we need to be able to iterate over all the element of the collection (l.4 2.7). Finally, we need to be able to write a value into the collection (l.5 2.7). This requires the collection not to be read-only, or the collection's values not to be yielded on-the-fly. This allows us to deduce what is called a *concept*: a breakdown of all the requirement about the behavior of our collection. When writing down what a *concept* should require, one should always respect this rule: "*It is not the types that define the concepts: it is the algorithms*". Concepts in C++ are not new and there have been a long work to introduce them

that goes back from 2003 [159, 158, 160] to finally appear in the 2020 standard [162] (referred as C++20 [73]). This allows us, as of today, to write code leveraging this facility.

Conceptification

C++ is a multi-paradigm language that enables the developer to write code that can be *object oriented*, *procedural*, *functional* and *generic*. However, there were limitations that were mostly due to the backward compatibility constraint as well as the zero-cost abstraction principle. In particular the *generic programming* paradigm is provided by the *template metaprogramming* machinery which can be rather obscure and error-prone. Furthermore, when the code is incorrect, due to the nature of templates (and the way they are specified) it is extremely difficult for a compiler to provide a clear and useful error message. To solve this issue, a new facility named *concepts* was brought to the language. It enables the developer to constraint types: we say that the type *models* the *concept(s)*. For instance, to compare two images, a function *compare* would restrict its input image types to the ones whose value type provides the *comparison operator* ==. In spite of the history behind the *concept checking* facilities being very turbulent [159, 158, 160], it will finally appear in the next standard [162] (C++20).

The C++ *Standard Template Library* (STL) is a collection of algorithms and data structures that allow the developer to code with generic facilities. For instance, there is a standard way to *reduce* a collection of elements: `std::accumulate` that is agnostic to the underlying collection type. The collection just needs to provide a facility so that it can work. This facility is called *iterator*. All STL algorithms behave this way: the type is a template parameter so it can be anything. What is important is how this type behaves. Some collection requires you to define a `hash` functions (`std::map`), some requires you to set an *order* on your elements (`std::set`) etc. This emphasis the power of genericity. The most important point to remember here (and explained very early in 1988 [7]) is the answer to: "*What is a generic algorithm?*". The answer is: "*An algorithm is generic when it is expressed in the most abstract way possible*". Later, in his book [66], Stepanov explained the design decision behind those algorithms as well as an important notion born in the early 2000s: the concepts. The most important point about concepts is that it constrains the behavior. Henceforth: "*It is not the types that define the concepts: it is the algorithms*". The *Image Processing* and *Computer Vision* fields are facing this issue because there are a lot of algorithms, a lot of different kinds of images and a lot of different kinds of requirements/properties for those algorithms to work. In fact, when analyzing the algorithms, you can always extract those requirements in the form of one or several *concepts*. This section is a preface to the image taxonomy which will be seen more in-depth in chapter 3.

Image processing algorithms, similarly, are *generic* by nature [9, 25, 37, 68, 97]. When writing an image processing algorithm, there is always a way to express it with a high level of genericity. For instance, if is possible to write a morphological dilation in a way that does not care about the underlying value type, the domain nor the structuring element specificities. The most abstract way to write a dilation is shown in 2.8.

```
1   template <Image I, WritableImage O,
2               StructuringElement SE>
3   void dilation(I input, O output, SE se) {
4     assert(input.domain() == output.domain());
5     for(auto pnt : input.points()) {
6       output(p) = input(p)
7       for (nx : se(p))
8         output(p) = max(input(nx), output(p))
9     }
10  }
```

Figure 2.8: Dilation algorithm, generic implementation.

This implementation introduces three concepts at line 1: *Image*, *WritableImage* and *Struc-*

turingElement. Following the behavior of each one of them into the algorithm, we can deduce a list of requirements for each one of them.

Image   is the most basic representation of what an image should be. An image should (a) provide a way to access its domain (l.3 2.8) and (b) a way to iterate over its points (l.4 2.8). This then allows us later to (c) access to the value returned by the image at this point (l.5 2.8). To this point the value is only accessed in read-only. We can then write the following two concepts:

```cpp
template <typename I>
concept Image = requires {
  typename I::point_range;        // needed for a
  typename I::point_type;         // needed for b
  typename I::value_type;         // needed for b
} && ForwardRange<I::point_range>  // needed for a
&& requires (I ima, I::point_type pnt) {
  { ima.points() } -> I::point_range; // a
  { ima(pnt) }     -> I::value_type;  // b
};
```

In reality, more boilerplate code is needed to ensure, for instance that there is no type mismatch between the image's `point_type` and the `point_range`'s value type. For the sake of brevity this boilerplate code is omitted here.

WritableImage   is a more specific concept based on the previous *Image* concept. It requires that the image's value can be (d) accessed to be modified: the user should be able to write into the image's value accessed by a specific point (l.6 2.8). We can then write the following two concepts:

```cpp
template <typename WI>
concept WritableImage = Image<WI>
&& requires (WI wima, I::point_type pnt,
             I::value_type val) {
  { wima(pnt) = val };            // d
};
```

StructuringElement   is an additional input to the image defining the window around each point that will be considered during the dilation (also called the neighborhood). A structuring element should just provide a list of point when input with one (e). From this behavior we can deduce the following concept:

```cpp
template <typename SE, typename I>
concept StructuringElement = Image<I>
&& requires (SE se,I::point_type pnt) {
  { se(pnt) } -> I::point_range;     // e
}
```

This new notion of concept is very important because it decorrelate the requirements on behavior required inside algorithms from the way the data structures are designed. One wan always wrap a specific data structure so that it can behave properly into an algorithm, without needing to rewrite that algorithm.

Simplifying code

The main advantage brought by using modern C++ as the implementation language for an image processing library is to be able to leverage what is called metaprogramming. Metaprogramming is a way to tell the compiler to make decision about which type, which code to generate. These decisions, made at compile time, and then absent from the resulting binary: only the fast and optimized code remains. This brings a new distinction between the static world (what is decided at compile time) and the dynamic world (what is decided at runtime). The more is decided at compile time the smaller, faster the binary will because there is work less to do at runtime.

By following this principle, one can think of some properties that are known ahead of time (at compilation) when writing one's image processing algorithm. For instance, when considering the example of the dilation whose code is shown in 2.8, we can see that the property about the decomposability of the structuring element is linked to the type. This means that when the structuring element's type is of a disc, or a square, the compiler will know at compile time that it is decomposable. To tell the compiler to take advantage of a property at compile time, C++ has a language construct named if-constexpr. The resulting code then becomes:

```cpp
template <Image Img, StructuringElement SE>
auto dilate(Img img, SE se) {
  if constexpr (se.is_decomposable()) {
    lst_small_se = se.decompose();
    for (auto small_se : lst_small_se)
      img = dilate(img, small_se) // Recursive call
    return img;
  } else if (is_pediodic_line(se))
    return fast_dilate1d(img, se) // Van Herk's algorithm;
  else
    return dilate_normal(img, se) // Classic algorithm;
}
```

There are other ways to achieve the same result with different language constructs in C++. There are two "legacy" language constructs which are tag dispatching (or overload) and SFINAE. With the release of C++17 came a new language construct presented above: if-constexpr. Finally, with C++20, it will be possible to use concepts to achieve the same result. To achieve the same result as above with tag dispatching, one would need to write the following code:

```cpp
struct SE_decomp {};
struct SE_no_decomp {};

template <Image Img, StructuringElement SE>
auto dilate(Img img, SE se) {
  // either SE_decompo or SE_no_decomp
  return dilate_(img, se, typename SE::decomposable());
}

auto dilate_(Img img, SE se, SE_decomp) {
  lst_small_se = se.decompose();
  for (auto small_se : lst_small_se)
    // Recursive call
    img = dilate(img, small_se, SE_no_decomp)
  return img;
}
auto dilate_(Img img, SE se, SE_no_decomp) {
  if (is_pediodic_line(se))
    return fast_dilate1d(img, se) // Van Herk's algorithm;
  else
    return dilate_normal(img, se) // Classic algorithm;
}
```

To achieve the same result with SFINAE, one would need to write the following code:

```cpp
// SFINAE helper
template <typename SE, typename = void>
struct is_decomposable : std::false_type {};
template <typename SE>
struct is_decomposable<SE,
  // Check wether the type provides the decompose() method
  std::void_t<decltype(std::declval<SE>().decompose())>
> : std::true_type {};
template <typename SE>
constexpr bool is_decomposable_v =
              is_decomposable<SE>::value;

template <Image Img, StructuringElement SE,
  typename = std::enable_if_t<is_decomposable_v<SE>>>
auto dilate(Img img, SE se) {
```

```cpp
  lst_small_se = se.decompose();
  for (auto small_se : lst_small_se)
    img = dilate(img, small_se) // Recursive call
  return img;
}


template <Image Img, StructuringElement SE,
   typename = std::enable_if_t<not is_decomposable_v<SE>>>
auto dilate(Img img, SE se) {
  if (is_pediodic_line(se))
    return fast_dilate1d(img, se) // Van Herk's algorithm;
  else
    return dilate_normal(img, se) // Classic algorithm;
}
```

Comparing those two last ways of writing static code to the first one comes to an obvious conclusion: the if-constexpr facility is much more readable and maintainable than the two legacy ways of doing it. Finally, there is still another way to handle the issue and it is with C++20's concepts. The following code demonstrates how to leverage this language construct:

```cpp
template <typename SE>
concept SE_decomposable = requires (SE se) {
  se.decompose(); // this method must exist
};


template <typename Img, typename SE>
auto dilate(Img img, SE se) {
 if (is_pediodic_line(se))
    return fast_dilate1d(img, se) // Van Herk's algorithm;
  else
    return dilate_normal(img, se) // Classic algorithm;
}


template <typename Img, typename SE>
  requires SE_decomposable<SE>
auto dilate(Img img, SE se) {
  lst_small_se = se.decompose();
  for (auto small_se : lst_small_se)
    img = dilate(img, small_se) // Recursive call
  return img;
}
```

A best-match mechanic [53] operates under the hood to select the function overload whose concept is the most specialized when possible. The best-match is very interesting to us as it removes completely the need to have mutually exclusive conditions which were required with the SFINAE technique and could result in explosive complexity with the growing number of *and/or* clauses. The best-match mechanic, works on top of another widely used machinery in C++: the function overload set. For a specific function call, all the functions overloads will be considered in the overload set. Then, to select the correct one, they will be sorted according to the best-match procedure. Finally, if there is still two overloads that have the same priority, the compiler will raise a hard error stating that there is an ambiguity it cannot solve on its own. As a matter of fact, adding one more function to the overload set is enough to have it selected when it matches without modifying the already existing functions. One could have added, in our example, an overload for structuring element whose shape is a periodic line, assuming that this is a property that we can detect at compile-time.

2.3   C++ templates in a dynamic world

There are two main categories of programming languages. First are the *compiled* programming languages which requires to feed the source code to a program (a compiler) that will output a binary. This binary will then produce the desired output once the user execute it. Some well known languages of this category are C, C++, Ada, Fortran. Secondly there are the interpreted

programming language which requires to feed the source code to a program (an interpreter) that will directly produce the output as is a binary was executed. Some well known languages of this category are Javascript, Python, Matlab, Common Lisp. There is a third category that tries to combine the best of both world by compiling into bytecode which is an optimized intermediate language that will then be interpreted into a virtual machine. The most famous are indubitably Java and C#. Both categories have advantages as well as drawbacks.

Compiled languages are still widely spread and used as of today. They present a working pipeline which is very classic. First the programmer will write code, then the compiler will build a binary optimized for the target machine and finally the programmer can execute his binary to produce a result. Usually the compilation step is slow whereas executing the binary is fast. There is no additional step when it comes to the binary execution. This, however, has the effect of having a poor portability. Indeed, my binary optimized to use fast and recent SIMD AVX-512 instructions will not work on an old x86 machine that does not support those instructions. When distributing our program, multiple binaries must be produced for each supported CPU architectures. Furthermore, usually compiled languages have very poor support of dynamic language features such as reflection, code evaluation or dynamic typing. It tends to improve with time but solutions are limited to compile-time informations or need to ship a JIT-compiler into the final binary (such as cling [86] for C++) to generate new binary on-the-fly to be executed right after. This has two drawbacks: slowness when the case of needing compilation is presented and increase in the binary size.

Interpreted languages are also widely used, especially in the research area where a fast feedback loop between prototyping and getting results is needed. The compilation time is very fast and allows a program to be almost instantly executed. In fact, the compilation can be done just ahead of program execution not to compile unnecessary code. However the execution time will generally be slow. To the user it is invisible because both compilation step and execution step are blurred together. Furthermore, most of the time a same interpreted program is executed once. Then the programmer will modify it and continue its prototyping process. The real advantage of an interpreted language is the portability. As it is the responsibility of the client to install the correct language interpreter (for the correct version) before running the program from the source code, as long as an interpreter can be installed on a machine, the program can then be run. This also drastically slow distributed package for programs as only source code must be distributed instead of compiled binary. However, the source code is leaked with all the security implication this can have. Finally, interpreted language usually have better memory management (builtin garbage collectors), are easier to debug, have very rich support of dynamic typing, dynamic scoping, reflections facilities, on-the-fly evaluation from the source code or even more like modifying the Abstract Syntax Tree (AST) resulting from the first compilation pass on source code by the interpreter. This last one is implemented in Common Lisp in the name of macros. There are more to say about interpreted languages, especially about those that are compiled into bytecode and tend to get the best of both worlds without the drawbacks but this thesis will not discuss this matter any further.

The main point to understand here is that our main interest is set on C++, a compiled language with slow compilation time and very fast execution time. In C++ there are template metaprogramming to achieve genericity but *templates do not generate any binary code.* Why? Because when the compiler meet a templated type or a templated routine, it does not know which type it will be instantiated with when it is used. Therefore, it can not calculate stuff like type size, alignement, can not select which assembly instruction to select to do an addition or a division (fixed vs. floating point arithmetic). This is why the compiler does not generate any binary code when it first meet templates. The code is generated only it is used with a concrete and known type. This is a huge problem. Now, if a library implementer wants to distribute his

generic library, he must distribute source code and have the user compile it. For a language like C++, with no standard dependency management, it can be a massive turn off. Furthermore, it may not be reasonable for the user to have C++ compiling facilities when the target client is embedded devices with limited storage space. Indeed, C++ intermediary compilation artefacts tend to use a lot of disk space before it is linked into a smaller binary. What solution do we have then?

SWILENA [17, 161]   is a Python bindings wrapper using Swig for the Olena C++ generic library. This wrapper will enumerate all the common use cases and implement a binding for them. The compiler will then generate binary code from the templated generic code for each use case enumerated in the wrapper. This way, we have given access into the dynamic world (Python) to generic code (C++ template). But it is still limited to the supported types. Each type a new combinaison of type needs to be supported from Python, it needs to be explicitly declared and compiled in the wrapper. Other image processing libraries, such as VIGRA [27], chose this solution.

VCSN [88]   is a novel solution that essentially take the same base as SWILENA but goes beyond the boundary to implement a handmade facility that do system compiler calls to compile and link needed code on-the-fly when the binding does not exist. It then leverage the code hotloading feature to plug new dynamic libraries (.dll on windows and .so on linux) into the wrapper to provide the user its needed bindings.

Cython [71]   attempt to solve the issue of the Python inherent language slowness due to its interpreted nature by providing a facility able to transpile a Python program into a C program so that a genuine C compiler (with extensions) is able to compile it and to link it against the Python/C API in order to achieve a huge performance gain at the cost of near zero knowledge of the complex Python/C API for the user. This novel solution essentially bypass the work of a JIT compiler (that would be used by a programming language using bytecode such as Java or C#) and just offload it onto well known/proven solution: the machine's C compiler.

Autowig [118], Cppyy [107] and Xeus-cling [147]   are all solutions aiming to generate automatically Python bindings on-the-fly using different solutions. Autowig has in-house code based on LLVM/Clang to parse C++ code in order to generate and compile a Swig Python binding using the Mako templating engine. Cppyy will generate Python bindings but can also directly interpret C++ code from Python code thanks to being base on LLVM/Cling, a Clang-base C++ interpreter. Finally, Xeus-cling is a Jupyter [106] kernel allowing to directly interpret C++ into a Jupyter notebook. Like Cppyy, it is based on LLVM/Cling. Those three projects are very promising and improve greatly the scope of possibilities for the future.

2.4   Summary

In this chapter we have presented the origin of generic programming, which goes as far as 1988, year. We then proceeds to demonstrate how genericity can be achieved within libraries and what programming languages features enables it. Second, we discuss how true generic facilities are builtin programming languages thanks to template metaprogramming techniques in a first time. In a second time, we discuss how new builtin programming language features such as concepts are finally reaching a state of generic programming as describe in [23]. Finally we discuss how the current state of generic programming via template metaprogramming, which is bound to the static world, interacts with the dynamic world.

Rework summary

Part II

Contribution

Chapter 3

Taxonomy of Images and Algorithms

In this thesis, we have pursued research into how to apply all those new generic facilities from the C++ language into the Image processing area. This allows us to test them in a practical way on our predilection area while remembering our past work, both success and failures in this matter. However, as we saw in the previous chapter 2, birthing concepts from code is something that is done in an emerging way. Henceforth, the first work will be to do an inventory of all existing image algorithms as well as an inventory of all image processing algorithms (both basic and more complex) we can think of. This way, we will notice behavior patterns emerging from similar image types or similar algorithms. We will then be able to extract behavioral patterns from this inventory in order to produce a full taxonomy in the form of a framework of concepts related to image processing. This chapter is structured as followed. First we will study how to extract behavioral pattern from a simple algorithm in order to refine it into one or multiple concepts. Second we will study the theory set behind images types, their conjunctions, disjunctions. We will also produce an inventory of image processing algorithms limited to mathematical morphologies that we can leverage for the final step. Third, we will study the intrinsic genericity of algorithms to produce canvas taking advantage of properties. Finally, we will then study behavioral pattern related to the inventory in the form of a taxonomy indexing a framework of concepts about image processing.

3.1 Rewriting an algorithm to extract a concept

3.1.1 Gamma correction

Let us take the gamma correction algorithm as an example. The naive way to write this algorithm can be:

```cpp
template <class Image>
void gamma_correction(Image& ima, double gamma)
{
  const auto gamma_corr = 1.f / gamma;

  for (int x = 0; x < ima.width(); ++x)
    for (int y = 0; y < ima.height(); ++y)
    {
      ima(x, y).r = 256.f * std::pow(ima(x, y).r / 256.f, gamma_corr);
      ima(x, y).g = 256.f * std::pow(ima(x, y).g / 256.f, gamma_corr);
      ima(x, y).b = 256.f * std::pow(ima(x, y).b / 256.f, gamma_corr);
    }
}
```

This algorithm here performs the transformation correctly but it also makes a lot of hypothesis. Firstly, we suppose that we can write in the image via the = operator (l.9-11): it may not be true if the image is sourced from a generator function. Secondly, we suppose that we have a 2D image via the double loop (l.6-7). Finally, we suppose we are operating on 8bits range (0-255)

RGB via '.r', '.g', '.b' (l.9-11). Those hypothesis are unjustified. Intrinsically, all we want to say is "*For each value of ima, apply a gamma correction on it.*". Let us proceed to make this algorithm the most generic possible by lifting those unjustified constraints one by one.

Lifting RGB constraint: First, we get rid of the 8bits color range (0-255) RGB format requirement. The loops become:

```cpp
using value_t = typename Image::value_type;

const auto gamma_corr = 1.f / gamma;
const auto max_val = std::numeric_limits<value_t>::max();

for(int x = 0; x < ima.width(); ++x)
  for(int y = 0; y < ima.height(); ++y)
    ima(x, y) = max_val * std::pow(ima(x, y) / max_val, gamma_corr);
```

By lifting this constraint, we now require the type Image to define a nested type `Image::value_type` (returned by `ima(x, y)`) on which `std::numeric_limits` and `std::pow` are defined. This way the compiler will be able to check the types at compile-time and emit warning and/or errors in case it detects incompatibilities. We are also able to detect it beforehand using a `static_assert` for instance.1

Lifting bi-dimensional constraint: Here we need to introduce a new abstraction layer, the *pixel*. A *pixel* is a couple (*point*, *value*). The double loop then becomes:

```cpp
for (auto&& val : ima.values())
  val = max_val * std::pow(val / max_val, gamma_corr);
```

This led to us requiring that the type *Image* requires to provide a method `Image::pixels()` that returns *something* we can iterate on with a range-for loop: this *something* is a *Range* of *Pixel*. This *Range* is required to behave like an *iterable*: it is an abstraction that provides a way to browse all the elements one by one. The *Pixel* is required to provide a method `Pixel::value()` that returns a *Value* which is *Regular* (see section 3.1.3). Here, we use `auto&&` instead of `auto&` to allow the existence of proxy iterator (think of `vector<bool>`). Indeed, we may be iterating over a lazy-computed view chapter 4.

Lifting writability constraint: Finally, the most subtle one is the requirement about the *writability* of the image. This requirement can be expressed directly via the new C++20 syntax for *concepts*. All we need to do is changing the template declaration by:

```cpp
template <WritableImage Image>
```

In practice the C++ keyword `const` is not enough to express the *constness* or the *mutability* of an image. Indeed, we can have an image whose pixel values are returned by computing $cos(x+y)$ (for a 2D point). Such an image type can be instantiated as *non-const* in C++ but the values will not be *mutable*: this type will not model the *WritableImage* concept.

Final version

```cpp
template <WritableImage Image>
void gamma_correction(Image& ima, double gamma)
{
  using value_t = typename Image::value_type;

  const auto gamma_corr = 1 / gamma;
  const auto max_val = numeric_limits<value_t>::max();

  for (auto&& pix : ima.pixels())
    pix.value() = std::pow((max_val * pix.value()) / max_val, gamma_corr);
}
```

When re-writing a lot of algorithms this way: lifting constraints by requiring behavior instead, we are able to deduce what our *concepts* needs to be. The real question for a *concept* is: "*what behavior should be required?*"

3.1.2   Dilation algorithm

To show the versatility of this approach, we will now attempt to deduces the requirements necessary to write a classical *dilate* algorithm. First let us start with a naive implementation:

```
1   template <class InputImage, class OutputImage>
2   void dilate(const InputImage& input_ima, OutputImage& output_ima)
3   {
4     assert(input_ima.height() == output_ima.height()
5       && input_ima.width() == output_ima.width());
6
7     for (int x = 2; x < input_ima.width() - 2; ++x)
8       for (int y = 2; y < input_ima.height() - 2; ++y)
9       {
10        output_ima(x, y) = input_ima(x, y)
11        for (int i = x - 2; i <= x + 2; ++i)
12          for (int j = y - 2; j <= y + 2; ++j)
13            output_ima(x, y) = std::max(output_ima(x, y), input_ima(i, j));
14      }
15  }
```

Here we are falling into the same pitfall as for the *gamma correction* example: there are a lot of unjustified hypothesis. We suppose that we have a 2D image (l.7-8), that we can write in the `output_image` (l.10, 13). We also require that the input image does not handle borders, (cf. loop index arithmetic l.7-8, 11-12). Additionally, the *structuring element* is restricted to a 5×5 window (l.11-12) whereas we may need to dilate via, for instance, a 11×15 window, or a sphere. Finally, the algorithm does not exploit any potential properties such as the *decomposability* (l.11-12) to improve its efficiency. Those hypothesis are, once again, unjustified. Intrinsically, all we want to say is "For each value of `input_ima`, take the maximum of the $X \times X$ window around and then write it in `output_ima`".

To lift those constraints, we need a way to know which kind of *structuring element* matches a specific algorithm. Thus, we will pass it as a parameter. Additionally, we are going to lift the first two constraints the same way we did for *gamma correction*:

```
template <Image InputImage, WritableImage OutputImage, StructuringElement SE>
void dilate(const InputImage& input_ima, OutputImage& output_ima, const SE& se)
{
  assert(input_ima.size() == output_ima.size());

  for(auto&& [ipix, opix] : zip(input_ima.pixels(), output_ima.pixels()))
  {
    opix.value() = ipix.value();
    for (const auto& nx : se(ipix))
      opix.value() = std::max(nx.value(), opix.value());
  }
}
```

We now do not require anything except that the *structuring element* returns the neighbors of a pixel. The returned value must be an *iterable*. In addition, this code uses the `zip` utility which allows us to iterate over two ranges at the same time. Finally, this way of writing the algorithm allows us to delegate the issue about the border handling to the neighborhood machinery. Henceforth, we will not address this specific point deeper in this paper.

3.1.3   Concept definition

The more algorithms we analyze to extract their requirements, the clearer the *concepts* become. They are slowly appearing. Let us now attempt to formalize them. The formalization of the

Let `Ima` be a type that models the concept `Image`. Let `WIma` be a type that models the concept `WritableImage`. Then `WIma` inherits all types defined for `Image`. Let `SE` be a type that models the concept `StructuringElement` . Let `DSE` be a type that models the concept `Decomposable`. Then `DSE` inherits all types defined for `StructuringElement`. Let `Pix` be a type that models the concept `Pixel`. Then we can define:

	Definition	Description	Requirement
Image	`Ima::const_pixel_range`	type of the range to iterate over all the constant pixels	models the concept *ForwardRange*
	`Ima::pixel_type`	type of a pixel	models the concept *Pixel*
	`Ima::value_type`	type of a value	models the concept *Regular*
Writable Image	`WIma::pixel_range`	type of the range to iterate over all the non-constant pixels	models the concept *ForwardRange*

Table 3.1: Concepts formalization: definitions

Let `cima` be an instance of `const Ima`. Let `wima` be an instance of `WIma`. Then all the valid expressions defined for `Image` are valid for `WIma`. Let `cse` be an instance of `const SE`. Let `cdse` be an instance of `const DSE`. Then all the valid expressions defined for `StructuringElement` are valid for `const DSE` Let `cpix` be an instance of `const Pix`. Then we have the following valid expressions:

	Expression	Return Type	Description
Image	`cima.pixels()`	`Ima::const_pixel_range`	returns a range of constant pixels to iterate over it
Writable Image	`wima.pixels()`	`WIma::pixel_range`	returns a range of pixels to iterate over it
Structuring Element	`cse(cpix)`	`WIma::pixel_range`	returns a range of the neighboring pixels to iterate over it
Decomposable	`cdse.decompose()`	`implementation defined`	returns a range of structuring elements to iterate over it

Table 3.2: Concepts formalization: expressions

concept Image from the information and requirements we have now is shown in table 3.1 for the required type definitions and in table 3.2 for the required valid expressions.

The *concept Image* does not provide a facility to write inside it. To do so, we have refined a second *concept* named *WritableImage* that provides the necessary facilities to write inside it. We say "*WritableImage* refines *Image*".

The *sub-concept ForwardRange* can be seen as a requirement on the underlying type. We need to be able to browse all the pixels in a forward way. Its *concept* will not be detailed here as it is very similar to *concept* of the same name [163, 165] (soon in the STL). Also, in practice, the *concepts* described here are incomplete. We would need to analyze several other algorithms to deduce all the requirements so that our *concepts* are the most complete possible. One thing important to note here is that to define a simple *Image concept*, there are already a large amount of prerequisites: *Regular*, *Pixel* and *ForwardRange*. Those *concepts* are basic but are also tightly linked to the *concept* in the STL [164]. We refer to the STL *concepts* as *fundamental concepts*. *Fundamentals concepts* are the basic building blocks on which we work to build our own *concepts*. We show the C++20 code implementing those *concepts* in 3.1.

```cpp
template <class Ima>
concept Image = requires {
    typename Ima::value_type;
    typename Ima::pixel_type;
    typename Ima::const_pixel_range;
  } && Regular<Ima::value_type>
  && ForwardRange<Ima::const_pixel_range>
  && requires(const Ima& cima) {
    { cima.pixels() }
      -> Ima::const_pixel_range;
  };


template <class I>
using pixel_t = typename I::pixel_type;
template <class SE, class Ima>
concept StructuringElement = Image<Ima>
  && requires(const SE& cse,
       const pixel_t<Ima> cpix){
    { se(cpix) } -> Ima::const_pixel_range;
  };
```

```cpp
template <class WIma>
concept WritableImage = requires Image<WIma>
  && requires {
    typename WIma::pixel_range;
  } && ForwardRange<WIma::pixel_range>
  && ForwardRange<WIma::pixel_range,
       WIma::pixel_type>
  && requires(WIma& wima) {
    { wima.pixels() } -> WIma::pixel_range;
  };


template <class DSE, class Ima>
concept Decomposable =
  StructuringElement<DSE, Ima>
  && requires(const DSE& cdse) {
    { cdse.decompose() }
      -> /*impl. defined*/;
  };
```

Figure 3.1: Concepts in C++20 codes

3.2 Images types viewed as Sets: version, specialization & inventory

Achieving true genericity in a satisfactory way is a complex problem that has components of different levels. The first goal is to natively support as many sets of image type as possible. Natively means that there is no need for a conversion from one type to a supertype under the hood. The second step is to support an abstraction layer above the underlying data type for each pixel. Indeed, the structure of an image is decorrelated from the underlying data type. The third step is to write image processing algorithms for each set of image type. Fourthly, the performance trade-off shall be negligible if not null. Finally, the final step is to provide a high degree of friendliness for the end user. Ease of use is always to be considered.

Figure 3.2: Set of supported image type.

Considering the available options to achieve our goal, the parametric polymorphism approach is the way to go. This allows the implementer to design image types and algorithms with behavior in mind. To illustrate this remark, let us consider the set of supported set of image types shown in figure 3.2. We usually refer to different image families as sub-type. Indeed, an image with a LUT is sub-type of the (global) Image type. This notion of sub-typing is important because it may be abstracted behind an interface. A user can manipulate an image type without knowing its specific sub-type. This induces that the sub-typing facility may be handled internally in the library with dynamic dispatching code (with run-time overhead). Each image processing library has its own way of handling sub-typing. More generally, the model used to handle it is previously described in section 2.1.1. Indeed, sub-types are handled the same way as their super-type but the fact that they have additional properties the algorithms can take advantage of.

To implement a basic image algorithm such as `fill` there really are two distinct ways of writing it. For the set of images type whose data type is encoded into each pixel, one must traverse the image and set each pixel's color to the new one. However, for the set of images type whose data type is encoded in a look-up table, one only has to traverse the look-up table to set

each color to the new one. This translates into two distinct algorithms shown in fig. 3.3. We can represent the diagram outlining that those two algorithms are two distinct version in fig. 3.4.

$$fill(I, v)\colon \forall p \in \mathcal{D}, I(p) = v \qquad\qquad fill(I, v)\colon \forall i \in I.LUT, i = v$$

(a) Writable image fill algorithm   (b) Image LUT fill algorithm

Figure 3.3: Comparison of implementation of the `fill` algorithm for two families of image type.

More generally, we consider that the set of image type is formed of several subset of image types. In the example there are two subsets: images whose pixel are writable and images whose data type are ordered in a look-up table. *For each one of these subsets, if there is a way to implement an algorithm then we have a* version *of this algorithm.*

Sometimes, it is possible to take advantage of a property on a particular image set, that may be correlated to an external data, to write the algorithm in a more efficient way. When those properties are linked to the types, this is call *specialization*. For instance, when considering a dilation algorithm, if the structuring element (typically the disc) is decomposable then we can branch on an algorithm taking advantage of this opportunity: decompose the dilation disc into small vectors and apply each one of them on the image through multiple passes. The speed-up comparing to a single pass with a large dilation disc is really significant (illustrated in 2.6). The code in 3.5 illustrate how an algorithm can be written to take advantage of the structuring element's decomposability property. The algorithm will first decompose the structuring element into smaller 1D periodic lines. It will then recursively call itself with those lines to do the multi-pass and thanks to known optimizations on periodic lines [10], it will be much faster. The dispatch diagram outlining the different specialization of the dilation algorithm used is show alongside in fig. 3.5.

Figure 3.4: Diagram of the two versions of the fill algorithm.

```
template <Image Img, StructuringElement SE>
auto dilate(Img img, SE se) {
  if (se.is_decomposable()) {
    lst_small_se = se.decompose();
    for (auto small_se : lst_small_se)
      img = dilate(img, small_se) // Recursive call
    return img;
  } else if (is_pediodic_line(se))
    return fast_dilate1d(img, se) // Van Herk's algorithm;
  else
    return dilate_normal(img, se) // Classic algorithm;
}
```
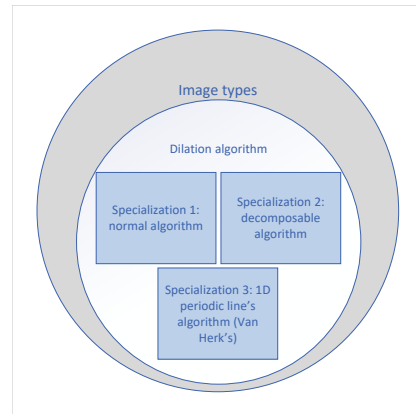

Figure 3.5: Dilate algorithm (left) with decomposable structuring element and its specialization diagram (right).

The figure 3.6 shows how an algorithm specialization may exists in a set of algorithms version. In this figure there exists a specialization of algorithms when it is known that the data buffer has the following property: its memory is contiguous. This implies that, for example, an algorithm

like `fill` can be implemented using low level and fast primitives such as `memset` to increase its efficiency.

Figure 3.6: Algorithm specialization within a set.

Making a full inventory of images types is not possible as there are many families of image types, each family may intersect with each other, images may have some particular properties at some points, those properties may appear in several family of image types. We can nonetheless cite a few to illustrate our points. There exists image-type whose values are cubical complexes [41] or layered as hexagonal or triangular grid instead of square grids (such as meshes), that are represented as graphes, etc. All of those interleaves with different research areas (computer vision for meshes, topological classification for complex cells [138, 31], graph theory and hierarchy for image graphes [92, 100, 129]). Finally, this image type inventory does not help when it comes to designing an image processing library. Instead, what helps is making an inventory of image processing algorithms as well as what properties those algorithms can leverage to speed up their execution.

By making the inventory of image processing algorithms (limited to mathematical morphology), we distinguish three main types of image processing algorithms. The first one are the pixel-wise algorithms. Those algorithms are the basic pixel-wise transformation. They are the base of any image processing toolkit. Indeed, being able to extract the green or red color channel of an image is mandatory. Thresholding as well as the previously seen gamma correction are such algorithm. The second one are the local algorithms. Those algorithms algorithms perform transformations by considering all the neighboring pixels around a given pixel. In order to operate, those algorithms need additional data in the form of image's extension (defining border's behavior) and structuring element (disc, rectangle around the considered pixel). Such algorithms are widely used in mathematical morphology. Dilation, erosion, closing, opening, gradient, rank filter are all local algorithms, as well as stencil-type algorithms such as union-find, max-tree or skeletonize. Finally, the third set of image processing algorithms is the global one. This set contains algorithms where computing the current pixel requires to know what was previously computed on the previous pixels as a prerequisite. Champfer distance transform, labeling, watershed, hierarchy structures related algorithms are in this category.

3.3 Generic aspect of algorithm: canvas

Genericity is always referred to with this sentence "write once, work for every types, run everywhere". However, very quickly we learned that the *run* aspect can be a combination of:

- as fast as possible on a single CPU unit;

- as fast as possible on thanks to using many CPU units;

- as fast as possible on thanks to using many GPU units;

- as fast as possible on thanks to using many computers (cloud) and their CPU/GPU units.

How do we decide what is the most efficient way to do? There is no simple answer to this question but we can start by studying the morphology of algorithms with two goals in mind: (i) find what can be parallelized/distributed and (ii) find one or several algorithms abstraction. Indeed, in image processing there are a lot of common patterns when one looks at algorithms implementations, the most famous being *for all pixels of an image, do something to each pixel*. But there

are other more high-level similarities that we can leverage to have more generic algorithm. Let us first study the different programmatic model there are commonly used to process images.

Rework these 3 paragraphs

The pipeline   is the classic way of doing the work. It consists of an imbrication of different operators (algorithms) taking as input one or several image, maybe additional data as well (such as labels, adjacency map, etc.) in order to process the data "from left to right". The result will show at the end of the pipeline and the optimization opportunities are located inside the smaller operators and in the form of correctly managing data (no useless copies, locality, etc.)

Kernels and tiling   is the trendy way of the last decade. It consists in breaking the original image into small tiles in order to feed those tiles into a massively multicore GPU (via CUDA, Halide [91]). Processing will happen concurrently on those core, but it is costly to swap concents inside those GPU from the RAM memory. It's then preferred to design pipeline for those tiles so that the work directly happens on each core to minimize the number of back and forth copies from the RAM.

Cloud computing   is a huge deal in image processing these last years as well, especially used in Deep learning applications. Deep neural networks are usually executed on cluster of computer through network (cloud) and thus are leveraging combinations of successive MapReduce [62] under the hood. They can also be offloaded onto or local available computers, or directly locally available GPU hardware thanks to frameworks that are dynamically dispatching the workload (such as Tensorflow/Keras [102, 101, 111], PyTorch [127], etc.).

In every case there is a notion of pipeline where the user pipe algorithms into each others in order to achieve a result. Those algorithms can leverage all the heterogenous ressources (cloud, GPU, CPU) they can to *map* the input data. Algorithms will then finally aggregate the results (*reduce*) to output them into another algorithm, or save them, or display them. It is important to dissociate the route the data will go through and the processing pipeline logic. Both have their own specificities. In this thesis, we make a parallel, at small scale, between processing pipeline logic and image processing algorithms. First let us study two basic algorithm: dilation and erosion. The python code of such algorithm is naively be given in figure 3.7.

```python
def dilate(img, se, out):
  for pnt in img.points():
    out(pnt) = img(pnt)
    for nx in se(pnt):
      out(pnt) = max(out(pnt), :
```
(a) Dilation

```python
def erode(img, se, out):
  for pnt in img.points():
    out(pnt) = img(pnt)
    for nx in se(pnt):
      out(pnt) = min(out(pnt), :
```
(b) Erosion

Figure 3.7: Dilate vs. Erode algorithms.

The algorithms are almost written the same way. The only change is the operation *min* and *max* when selecting the value to keep. As such, we can easily see a way to factorize code by passing the operator as an argument. The algorithms can then be rewritten as shown in figure 3.8. In this last example, we have one piece of code which is in charge of abstracting the way an image is traversed: it is the canvas. We also have other pieces of code which carry the logic of the operations, calling the canvas an providing the logic to apply from within the canvas. This way of decomposing the code enables the possibility of writing more specific heterogeneous logic for the traversing code so that the other parts of the code that carry the operator logic remains unaware and unburdened of possible implementation details.

```python
def local_op(img, se, op, out):
  for pnt in zip(img.points()):
    out(pnt) = img(pnt)
    for nx in se(pnt):
      out(pnt) = op(out(pnt), img(nx))
```

(a) Local algorithm with custom operator

```python
def dilate(img, se, out):
  local_op(img, se, max, out)
```

```python
def erode(img, se, out):
  local_op(img, se, min, out)
```

(b) Dilation (delegated)

(c) Erosion (delegated)

Figure 3.8: New Dilate vs. Erode algorithms.

3.3.1 Taxonomy and canvas

This approach is very much compatible with the inventory of the algorithms we did in the previous section 3.2. Indeed, for the point-wise family of algorithms, it is hardly an issue as they can all be written in term of views chapter 4. For the second family, which consists in all the local algorithm, they may be abstracted behind an algorithm canvas where the user provides the work to do at each point of the algorithm. For instance, a single pass local algorithm will always have shape given in 3.9. Finally for third family which consists in all the algorithm that propagate their computation while traversing the image. It is way less easy to provide an algorithmic canvas as each algorithm has its own way to propagate changes curing its computation.

```python
1   def local_canvas(img, out, se):
2     # do something before outer loop
3     for pnt in img.points():
4       # do something before inner loop
5       for nx in se(pnt):
6         # do something inner loop
7       # do something after inner loop
8     # do something after outer loop
```

Figure 3.9: Local algorithm canvas.

This canvas can be customized to do a specific job, especially at the lines 2, 4, 6, 7 and 8. The user would then provide callbacks and the canvas would do the job. This is especially useful when knowing that the canvas would handle the border management (the user would provide a handling strategy like mirroring the image or filling it with a value). The canvas would also take advantage of optimization opportunities (such as the decomposability of a structuring element) that the user would probably forget, or not know, when first writing his local algorithm. Another advantage is the opportunity to do more complex optimization such as parallelizing the execution or offloading part of the calculation on a GPU. More generally, all optimization done through heterogeneous computing would be available by default even if the user is not an area expert.

Despite all these advantages, one big disadvantage is the readability of the algorithm user-side. For instance, the dilation algorithm is rewritten in figure 3.10.

This way of thinking algorithms is far less readable than the classic way. The user does not see the loops happening and it can becomes very messy when several passes are happening (closing, opening, hit or miss, etc.)

3.3.2 Heterogeneous computing: a partial solution, canvas

One of the key aspect driving genericity is performance. We still have the following mantra: "write once, work for every types, run everywhere". However when considering the *run* aspect, one has a lot to do. Indeed, nowadays, leveraging the available resources to their maximum is long standing issue. There are many ongoing work on the subject, such as SyCL [123, 122] which is a standard for heterogeneous computing model edited by Kronos. This standard currently has

```python
def dilate(img, out, se):
  do_nothing = lambda *args, **kwargs: None

  def before_inner_loop(img, out, pnt):
    out(pnt) = img(pnt)

  def inner_loop(ipix, opix, nx):
    out(pnt) = max(out(pnt), img(nx))

  local_canvas(img, out, se,
    before_outer_loop = do_nothing,
    before_inner_loop = before_inner_loop,
    inner_loop        = inner_loop,
    after_inner_loop  = do_nothing,
    after_outer_loop  = do_nothing
  )
```

Figure 3.10: Local algorithm canvas.

four implementations: Codeplay's ComputeCpp [153], Intel's LLVM/clang implementation [154], triSYCL [156] led by Xilinx and hipSYCL [133] led by Heidelberg University. There also exists smaller libraries such as Boost.SIMD [93] or even VCL [89] for easing how to write SIMD code. After taking some distance to study the subject, we can infer that there are three main aspects to consider when optimizing performance.

The first one, the most important one is the algorithm to use in function of certain set of data. This aspect is covered by the C++ language and its builtin genericity tool: template metaprogramming. Indeed, we select the most optimized algorithm for a particular set of data.

The second one is the ability for the code to be understood by the compiler so that is is further optimized during the generation of the binary. Indeed, when compiling for the native architecture of a recent processor, one can use the most recent assembly instructions to use wide vectorized registers (AVX512). The use of a recent compiler also brings the help much needed.

Finally, the third aspect is not as trivial as the first two ones. It consists in studying the structure of an algorithm to allow distributed computation. It also consists in studying the different architectures to select the most efficient algorithm for a given architecture. Sometimes algorithm are friendly to be distributed on several processing units that compute a part of the result concurrently. This is what we call parallelism. There exists several way to take advantage of parallelism. First there is the use of several CPU units on the host computer. Then there is the use of GPU units working in combination with the CPU units to take advantage of the massive amount of core a graphic card can provide. Finally there is the use of cloud computing which consists in using several "virtual" computers, each of them offering of CPU and GPU units in order to compute a result. One should be aware that each time we introduce a new layer of abstraction, there is a cost to orchestrate the computation, send the input data and retrieve the results. It is thus very important to study case by case what is needed. Some solutions exist that abstract away completely the hardware through a DSL [1] such as Halide [91]: the DSL compiler's job will be to try very hard to make the most out of both the available (or targeted) hardware and the code. Those solutions are not satisfactory for us as we want to avoid DSL and remain at code level. We are not developing a compiler: we are working with it.

There is one true issue when studying parallel algorithm: it is whether they can be parallelized or not. Not all algorithms can be parallelized. Some just intrinsically cannot, typically, algorithms that immediately need the result at the previous iteration to compute the next iteration. There may exist specific ways to re-arrange a specific algorithm, for instance, taking advantage of some algebraic property, in order to rewrite it in a parallelisable way, but it is not trivial (on a per-algorithm basis), and not generic. As such, it is out of the scope of this thesis. What interest us are the algorithms whose structure is an accumulation over a data type that can

---

[1]Domain Specific Language

be defined as a monoid. We assert that every algorithm that can be rewritten as an accumulation over a monoid can be parallelized and/or distributed. This model that consists in distributing computation like an accumulation over a monoid data structure is also call the map-reduce. This model has two steps: the distribution (map) and then the accumulation (reduce).

The map step will dispatch computation on sub-units with small set of data. The reduce step will retrieve and accumulate all those resulting data, as soon as they are ready.

The accumulation algorithm has this form:

```
1   template <class In, class T, class Op>
2   auto accumulate(In input, T init, Op op)
3   {
4     for(auto e : input)
5     {
6       init = op(init, e);
7     }
8     return init;
9   }
```

The loop line 4 can be split into several calculation units which are going to be distributed, and then be accumulated later once the units have finished their computation.

The issue left here is the monoid. What exactly is a monoid here? A monoid is a data structure which operates over a set of values, finite or infinite. This data structure must provide a binary operation which is closed and associative. Finally, this data structure must also provide a neutral element (aka the identity). Some trivial monoids comes to mind:

- boolean. For binary operation "and", identity is "true" whereas for binary operation "or", identity is "false".

- integer. For binary operation "-" and "+", identity is "0" whereas for binary operation "*" and "/", identity is "1".

- string. For concatenation, identity is empty string.

- optional value (also known as monadic structure in haskell programming language).

There are many more monoids, less trivial but very handy, such as the unsigned integer/max/0 set and the signed integer/min/global max set.

This theory is extremely beneficial to image processing as the most commonly used algorithms, the local algorithms, can all be written in the form of an accumulation over the pixels of an image. The fact that finding an identity for the operation processed by the algorithm is often quite trivial led us to the idea of canvas. A canvas is a standard way to write an iteration over an image which abstracts the underlying data structure. A canvas is a tool for the user to provide its computation model based on events such as: "entering inner loop" or "exiting inner loop". The user can then provide its operations as if he was writing his algorithm himself (restricted to the accumulation model). As the maintainer of the library provides the canvas of execution, he can know also make change to take advantage of it. For instance, computing a CUDA kernel at one point and dispatching it on GPU units is totally within scope and transparent for the user of the library. Although there is a caveat: rewriting our algorithm in an accumulation form and chunking it in fragments to feed to the canvas is definitely not intuitive. Indeed, we require our user to change his way of thinking from the procedural paradigm to the event-driven paradigm. This approach is not new and is used in other libraries such as Boost.Graph [34] for similar purposes. Dean talks about this recurring monoid pattern more in-depth in his talk [124].

In image processing, we quickly come to identify *local* algorithms, that reason about a group of pixel around a given of coordinate. All those algorithms can be abstracted behind an accumulation of some sort and they all have the same morphology. Thus leading to the following abstraction:

```cpp
template <class In, class Out, class SE, class T, class Op>
auto local_accumulate(In input, Out output, SE se, T init, Op op)
{
    auto zipped_imgs = ranges::view::zip(input.pixels(), output.pixels())
                                                    // (1)
    for(auto&& rows : ranges::rows(zipped_imgs))
    for(auto [px_in, px_out] : rows)
    {
        auto v = op(init, px_in.val());         // (2)
        for(auto nb : se(px_in))
        v = op(v, nb.val());                    // (3)
        px_out.val() = v;                       // (4)
    }
                                                // (5)
  }
```

From this code we can deduce some very useful and easy monoids by the following triplets:

- $(type = boolean, operator = and, neutral = true)$ is a binary erosion

- $(type = boolean, operator = or, neutral = false)$ is a binary dilation

- $(type = unsigned integer, operator = max, neutral = 0)$ is a dilation

- $(type = signed integer, operator = min, neutral = globalmax)$ is an erosion

Now if we want to rewrite the `local_accumulate` in an event driven paradigme, we need to identify the different callbacks to expose our user on the call site. Especially, what will be of the callback parameters. There are five callback event we have identified:

1. before entering outer loop (no work is done)

2. before entering inner loop (iteration over the pixel's neighbor)

3. inner loop (actual operation to perform, result is accumulated)

4. after exiting inner loop (iteration over the neighbor is over, what to do with the accumulated result?)

5. after exiting outer loop (iteration over the image is over)

```cpp
template <class In, class Out, class SE, class T, class BeforeOuterLoopCB,
          class BeforeInnerLoopCB, class InnerLoopCB class AfterInnerLoopCB,
          class AfterOuterLoopCB>
auto local_accumulate(In input, Out output, SE se, T init,
                      BeforeOuterLoopCB bolCB, BeforeInnerLoopCB bilCB,
                      InnerLoopCB ilCB, AfterInnerLoopCB ailCB,
                      AfterOuterLoopCB aolCB)
{
  auto zipped_imgs = ranges::view::zip(input.pixels(), output.pixels()
  bolCB(input, output);                         // (1)
  for(auto&& row : ranges::rows(zipped_imgs))
    for(auto [px_in, px_out] : rows)
    {
      bilCB(px_out.val(), init, px_in.val())    // (2)
      for(auto nb : se(px_in))
        ilCB(px_out.val(), px_out.val(), nb.val())  // (3)
      ailCB(px_out.val(), init, px_in.val())    // (4)
    }
  aolCB(input, output)                          // (5)
}
```

In this code, we can see that all the callbacks do not take the same type and/or number of parameters. Here is what the call site would like if the user wants to perform a dilation:

```
local_accumulation(
  input,                                   // input image
  output,                                  // output image
  se,                                      // structuring element
  0,                                       // monoid's neutral element
  [](auto I, auto& O) { /* do nothing */ }, // (1) entering outer loop callback
  [](auto& o, auto init, auto in){         // (2) entering inner loop callback:
    o = std::max(init, in);                // initialize with neutral element
  },
  [](auto& o, auto cur, auto nbh) {        // (3) inner loop callback:
    o = std::max(cur, nbh);                // keep the local maximum
  },
  [](auto& o, auto init, auto in) {        // (4) exiting inner loop callback
    /* do nothing */
  },
  [](auto I, auto& O) { /* do nothing */ } // (5) exiting outer loop callback
);
```

It is very verbose and non-intuitive but hopefully, once the compiler optimize out the empty callbacks, the generated code is as fast as a non-generic handwritten dilation.

3.4 Library concepts: listing and explanation

Let us now delve into concepts related to the Image Processing area. Indeed, this domain has his specificities, and we want to improve generic image processing library design by learning from our past experiments and working with new techniques. The most basic usage of an image is the famous algebraic formula $y = f(x)$ where y is a *value* generated by the *image f* for the *emph x*. Aside from generating a value, an image can also *store* a value, as in $f(x) = y$ where the value is *assigned* to the image for a given point. Those notions are the basis of our work and will drive the entire design.

3.4.1 The fundamentals

First, let us introduce the fundamental concepts deriving from the basis notion. The *Value* concept is refined into three distinct one in appendix B.1.1. There are the basic *Value* but also the *ComparableValue* and the *OrderedValue* which are useful when it comes to comparison or ordering algorithms. The need behind those three concepts derives from the algorithms who need an ordering relation in order to function properly. Most of mathematical morphology requires it. For instance, the ordering relation for a gray-scale image is trivial whereas it is a field of research for colored (rgb-8bits) images.

The second fundamental brick is the concept of *Point*, detailed in appendix B.1.2 which is a bit less open than the concept *Value* as it must be totally ordered. Indeed, when accessing a value stored in an image, whether it is about reading or mutating, it is important that there is only one accessed value.

Now we introduce an abstract way to represent this relation $Value \times Point$: the Pixel. This is a well known notion in image processing, and it represents a couple (*point, value*). This abstraction layer is easy to move around and contains facilities to read and mutate the pixel's value if possible. Indeed, not all pixels are able to mutate their value. If the pixel is yielded by an image that only generates values on the fly then it cannot be mutated. Henceforth, we introduce two new concepts: *Pixel* and *OutputPixel* in appendix B.1.3. Those two concepts have a very similar interface described in appendix B.1.3. They can both access the stored informations: the point and the value. On top of that, the *OutputPixel* can mutate the value. When interacting with pixels, the user will want to be able to write code as followed:

```
auto pix = Pixel();    // Get a pixel
auto val = pix.val();  // yield the pixel value
auto pnt = pix.point(); // yield the pixel point
pix.val() = 42;        // Assign a value
```
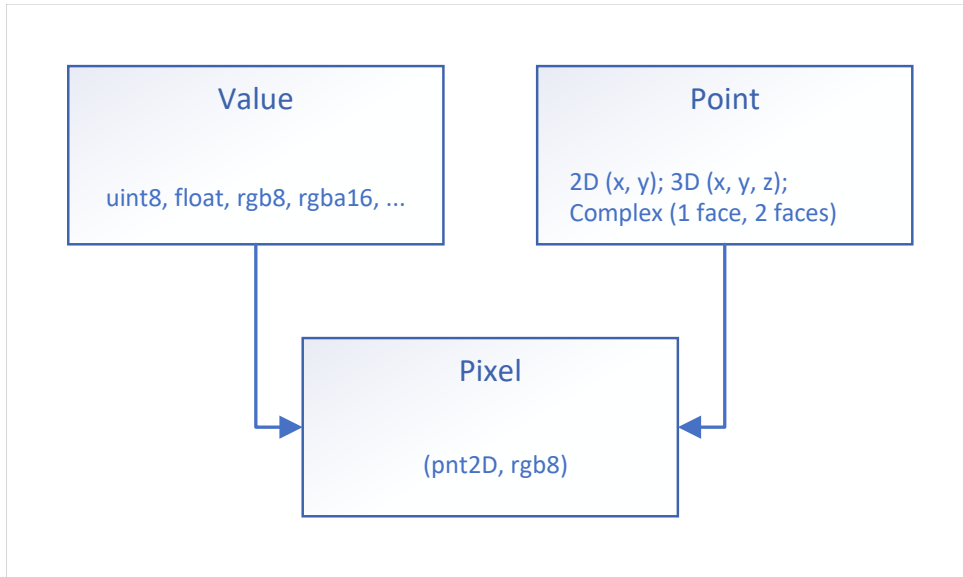
Figure 3.11: Pixel concept.

We show how those three fundamental concepts interact with each other in the diagram fig. 3.11.

Now we need a helper concept: the ranges. Ranges [157, 120, 165, 163] are a set of concepts defined in the C++ standard library shipped with the ISO C++20 norm in 2020 [136]. They allow the user to abstract away iterators to only iterate over one object: the range. This allow the user to migrate his source code from:

```cpp
template <class IteratorBegin, class IteratorEnd>
void my_algorithm(IteratorBegin beg, IteratorEnd end)
{
  for(; beg != end; ++beg)
    // ...
}
```

To:

```cpp
template <class Range>
void my_algorithm(Range rng)
{
  for(auto e : rng)
    // ...
}
```

In image processing, we refine further this concept by introducing multidimensional ranges (*MDRange*). Indeed, in image processing the user is used to write double loop to iterate over a bi-dimensional image. And abstracting away this aspect under standard ranges induces performance loss. That is why we needed this concept to exist. A multidimensional range can be split with a library function, `mln::ranges::rows(mdrng)` to fit the double-loop pattern and keep its performance. This topic is tackle in-depth later in section 4.5.1. For now, let us consider multidimensional ranges as an image processing extension for performance for the image traversing pattern. They are defined in appendix B.1.4 and their interface is the same as standard ranges, as seen in appendix B.1.4. They are designed so that the user code looks like this:

```cpp
auto mdrng = MDRange(); // Get a multi-dimensional range of values
auto rows = mln::ranges::rows(mdrng);
for(auto row : rows) // double loop pattern
  for(auto val : row)
    // use(val)
```

From an algebraic point of view, the definition of an image is not complete without considering a definition domain on which it is defined. In image processing, the same rule applies. We cannot

consider an image without considering the set of points that are valid for this image. Henceforth we must define the concept of *Domain* in *table:concept.domain.definitions*. The *Domain* concept is refined into two sub-concepts which are *SizedDomain* and *ShapedDomain*. This emphasis the chance of existence of possible infinite domain as well as domains that may be defined over non-continuous intervals in space. This enables algorithms to require the domain to have certain shape if needed. The domain behavior is described in appendix B.1.5.

In practice, a domain is used to get information about the points constituting the image. Indeed, we can write code like this:

```
auto dom = Domain(); // Get a domain
auto pnt = Point(..., ...); // Get a random point
bool ret = dom.has(pnt); // Check wether the domain contains the point
bool is_empty = dom.empty(); // Check wether the domain is empty
auto dim = dom.dim(); // Yield the domain's dimension information
```

We show how the concept *Domain* flows from the previous concepts in the diagram shown in fig. 3.12.
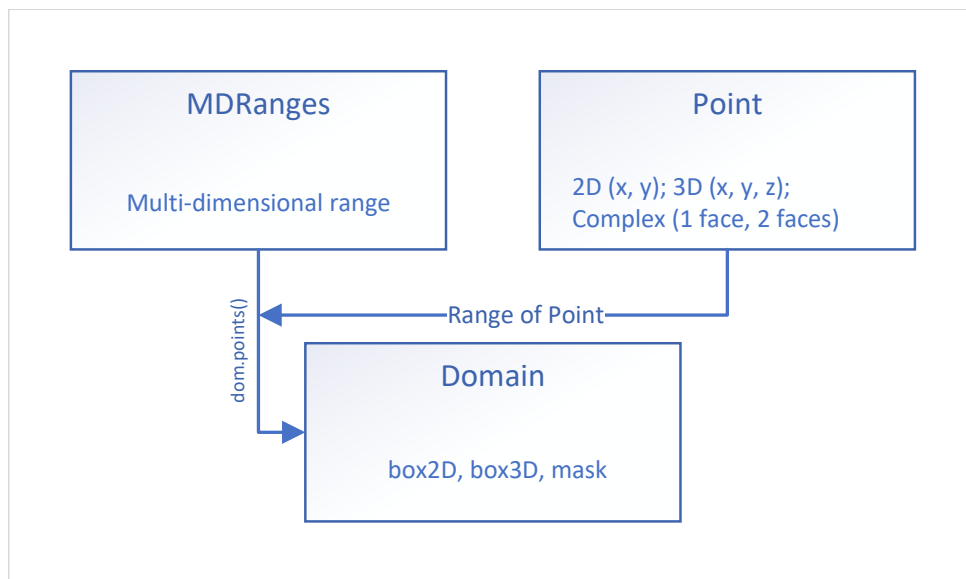
Figure 3.12: Domain concept.

Now we have all the tools to introduce our main concept: *Image*. As for *Pixel*, we have the distinction over image whose value can be mutated in a sub-concept named *WritableImage*. These concepts are defined in table B.9. In addition to this definition we can infer the behavior described in appendix B.1.6. There are complicated requirements written in template metaprogramming code. But in the end it is just to requires the ranges returned by the member functions `pixels()` and `values()` to iterate over element whose type are the same as those declared in the parent image. In addition, we introduce two facilities which are the member function `concretize()` and `ch_value<V>()`. The first is a way to turn a view into a concrete type. This will be seen more in-depth in chapter 4. The last is a way to cast values from one type to another. It forms a new image type whose underlying values will be returned after being casted to a new value type. This last facility is extremely useful when one only wants to mutate the underlying type while keeping all the other details about one type (such as the dimension). For instance, when working with labeling algorithm, we know our algorithm will return an image similar to the input one with the exception of the underlying type which will be the type of the label. The following code shows how it is applied:

```
using label_t = int; // label type

template <class I> // Input image of type I
```

```
auto my_labeling_algorithm(I input_image)
  -> image_ch_value_t<I, label_t> // Output image is Input image (I)
                                  // whose underlying type is label_t
{
  // ...
}
```

We show the diagram building up the image concept from the previous concepts in fig. 3.13.
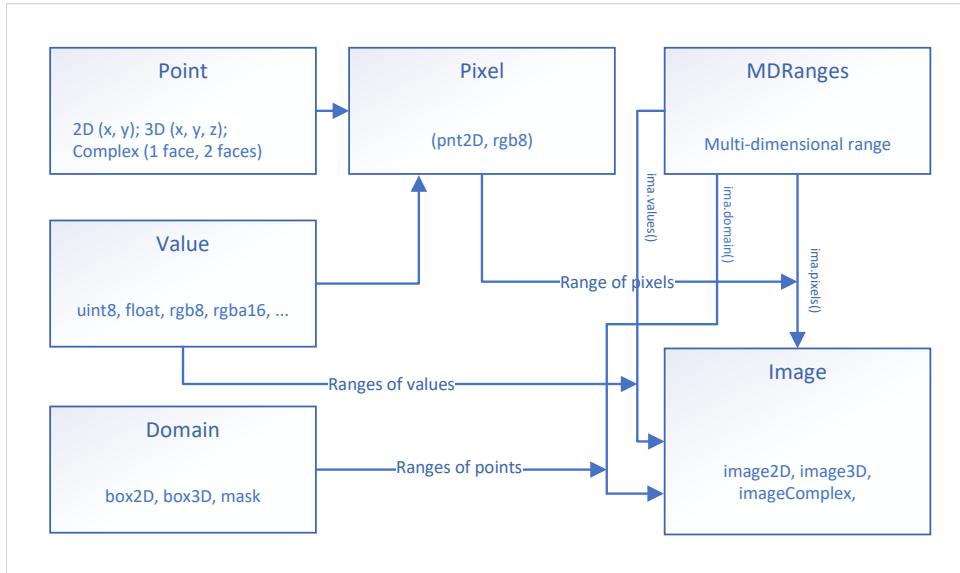
Figure 3.13: Image concept.

3.4.2   Advanced way to access image data

While being able to iterate over ranges of pixels or values is good, we are still lacking fundamental facilities to access an element directly from the image. In order to solve this predicament, first we need to define the concept of *Index* in appendix B.2.1 which we will use afterwards. This is a very simple concept that encapsulate an integral value. This value can be negative as we want to be able to do negative indexing in case our image has an extension, cf. section 3.4.2. The first advanced tool is represented as an *IndexableImage*. An element can be accessed simply by providing its index number. This concept is defined in table B.12. This introduces a simple behavioral pattern described in appendix B.2.2. With this concept, we are able to write code as followed:

```
auto ima = IndexableImage(); // Get an indexable image
int k = 15; // Get an index
auto val = ima[k]; // Get the element at the index
ima[k] = 255; // Mutate the element at the index
```

Being able to traverse an image through indexes is especially useful for algorithms that are aware of the number of elements in the image. We chose to be flexible with our indexing method (i.e. allowing signed indexes) not to fall into the same pitfall as the C++ core language [166]. Indeed, requiring that the type `std::size_t` is unsigned lead to loads of issues, the first one being a conversion issue when writing a simple for-loop. Solving these issues lead to the appearance of new member function `.ssize()` (for signed size) and the new type `std::ptrdiff_t` to store the result of a soustraction between two `std::size_t`. Furthermore, in our specific area (image processing), it may be well defined to access a buffer with negative indexes when, for instance, we are accessing the value of the extension of an image. This is why our indexes are signed.

Additionally, we want to be able to access a value by providing a point, the same way as in the algebraic definition $val = image(point)$. To do so, we introduce the concept of accessibility

through *AccessibleImage*. This concept is defined in table B.14. This introduces new behavior that is described in appendix B.2.3. We can notice facilities specifically including bound checking. Indeed, we suppose, for fast access, that the user is always picking element from the image's domain, but it is possible to bound check elements if needed on access for specific usages. With this concept, we are now able to write code as followed:

```
auto ima = AccessibleImage(); // Get an accessible image
auto p = Point(); // Get a point
auto val = ima(p); // Get a value from a point
auto val = ima.at(p) // Same with no bound checking
auto pix = ima.pixel(p) // Get a pixel from a point
auto pix = ima.pixel_at(p) // Same with no bound checking
ima(p) = 42; // Assign a value from a point
ima.at(p) = 42; // Same with no bound checking
ima.pixel(p).val() = 42; // Assign a pixel value from a point
ima.pixel_at(p).val() = 42; // Same with no bound checking
```

Being able to traverse an image through points is especially useful for algorithms relying on restricting/expanding definition domain that are exclusively yielding points.

Once we know that an image is both *indexable* and *accessible* we can introduce new behaviors (described in appendix B.2.4) that we put behind the concept of *IndexableAndAccessibleImage* defined in *table:concept.image.definitions.4*. This behavior is related to accessing index from points and vise versa. Indeed, it is possible to now write such code:

```
auto pnt = ima.point_at_index(k); // Get the point from an index
auto k = ima.index_of_point(pnt); // Get the index from a point
// Get the index difference for a shift of delta_point
auto delta_idx = ima.delta_index(delta_pnt);
```

Additionally it useful, for propagating algorithms, to be able to traverse images in both a forward way and a backward way. As it may not be possible for all images, this notion needs to be refined into a new concept *BidirectionalImage*. This concept is defined in table B.18 and its behavior is described in appendix B.2.5. Thanks to this concept, we are able to write code as followed:

```
template <class I>
my_algorithm(I input)
{
  // forward pass
  for(auto pix : input.pixels())
    // ...

  auto backward = mln::ranges::views::reverse(input.pixels())
  for(auto pix : backward)
    // ...
}
```

Finally, we need a way, when possible, to iterate over a continuous data buffer for very fast and optimized calculation. That is what the concept of *RawImage* is for: an image whose data buffer can be accessed, as well as its mutable counterpart. It is defined defined in table B.20. Having a raw image whose data buffer can be accessed allow ust to expose two more member function to access the data buffer and its strides for correct pointer arithmetic. They behave as described in appendix B.2.6. This enables writing code as followed:

```
auto ima = Image(); // Image of int
const int* data = ima.data(); // Access the underlying buffer
auto dim = ima.domain().dim(); // Get the dimension of the image
// Retrieve informations about strides
auto strides = std::vector<std::ptrdiff_t>(0, dim)
for (int i = 0; i < dim; ++i)
  strides[i] = ima.stride(i)

// Now use data and strides to traverse the raw buffer
// ...
```

We show how those concepts are defined from one another in the diagram shown i fig. 3.14.
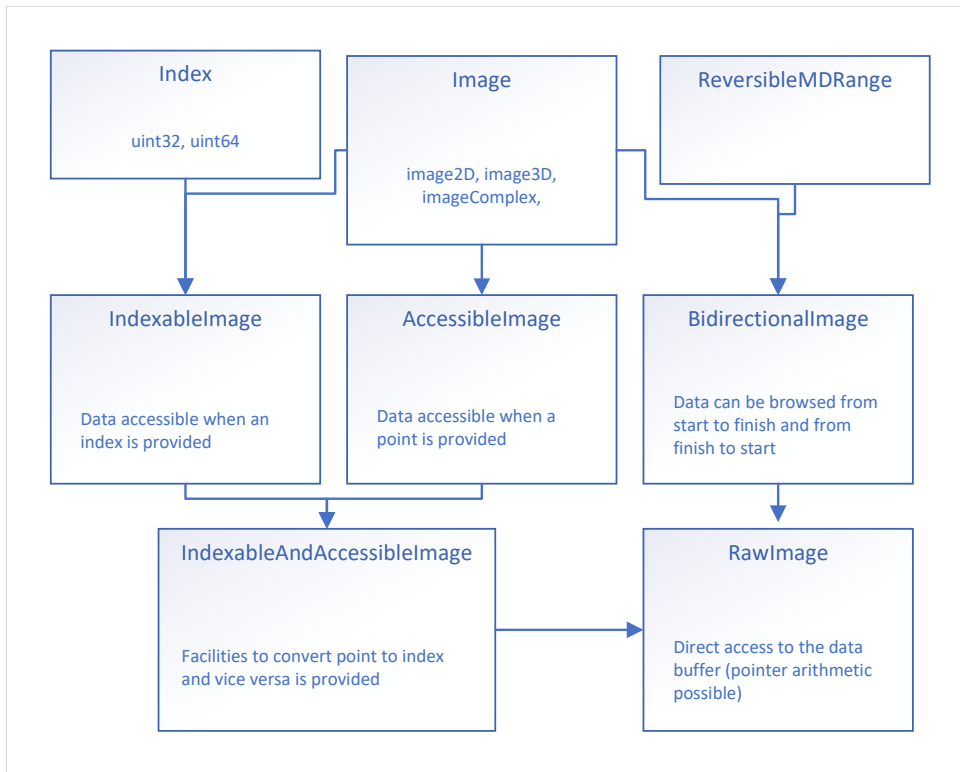
Figure 3.14: All images concepts. Donner du sens aux flèches

3.4.3   Local algorithm concepts: structuring elements and extensions

From the beginning concepts are emerging from behavioral patterns extracted from algorithms. In image processing, there is a family of algorithms called the *local algorithms*. They work by considering a specific pixel as well as all the pixels among a window having a specific shape centered in this first pixel. The window is called the *structuring element* and the pixels considered by this window are called the *neighborhood*. This leads us to introduce the concept of *StructuringElement* which is defined in table B.22.

This concept is refined into three sub-concepts that are related to properties the structuring element can offer. Those properties are:

- decomposability: ability to split a complex structuring element into several smaller and simpler structuring element. There is an equivalence in behavior when the algorithm is recursively run for each smaller structuring element one after another, in a multi-pass way.

- separability: ability to split a complex structuring element into several smaller and simpler structuring element. There is an equivalence in behavior when the convolution is recursively run for each smaller structuring element one after another, in a precise order, in a multi-pass way.

- incremental: ability to tell the points that are added to or removed from the range when the structuring element is shifted by a basic displacement (e.g. for a *2D point*, the basic displacement is $(0, 1)$). Usually used to compute attributes over a sliding structuring element in linear time.

Give exemples

The behavioral for those concepts is described in appendix B.3.1. Being able to manipulate structuring element allows us to write the following code:

```
auto se = se::disc(.radius=3); // get a structuring element
for(auto pix : ima.pixels())   // traverse image
  for(auto nb : se(pix))       // traverse neighboring pixels
    // ...
```

Additionally, we introduce the concept of *Neighborhood* in table B.24. This concept has facilities to know what points/pixels are placed before or after another point/pixel inside the window of a specific structuring element. It behaves as described in appendix B.3.2. This concept is useful when one wants to only consider a certain part of the neighboring pixels within a structuring element. This allow to write the following code:

```
auto se = se::disc(.radius=3); // get a structuring element
for(auto pix : ima.pixels())   // traverse image
  for(auto nb : se.before(pix))// traverse neighboring pixels located before pix
    // ...
```

And the last concept we need to introduce is the extension. Indeed, extension management is very important when dealing with local algorithm as pixels on the border need to be processed too, and the behavior near the border of the image must be defined and well-specified. There are several strategies when it comes to borders and extension. We refine concept for each strategy we identified:

- fillable: fill the border with a specific value.

- mirrorable: mirror the image as if there was an axial symmetry, with the border being the axis.

- periodizeable: repeat the image, as if a modulo size was applied to the coordinates.

- clampable: extend the value at the image's border into the extension.

- extent with: used when tiling. it consider the current image as a sub-image of another bigger image and pick the extension values there.

Those concepts are defined in table B.26 and their behavior is described in appendix B.3.3. All those concepts allow us to introduce the final refined image concepts: *WithExtensionImage* as well as *ConcreteImage* and *ViewImage*. Those two last will be seen in detail in the next chapter 4. Those concepts are defined in table B.28. Their behavior is described in appendix B.3.4. It is now possible to write the following code:

```
template <class I, class SE>
my_local_algorithm(I input, SE se) {
  // if the extension is large enough to function with the passed structuring element
  if(input.extension().fit(se)) {
    // ...
  }
}
```

We show how those three concepts (structuring elements, neighborhood and extension) interact with each other in the diagram shown in fig. 3.15.

Finally, we introduce a helper concept to centralize the detection of the "writability" of an image. Indeed, we do not want the user to have to use the writable counterpart of each concept for each and every case. That is why we introduce this final concept, *OutputImage* in appendix B.3.5, that will tell the user if an image is well-specified.

The correct way to use it is:

```
template <class Img>
requires RawImage<Img> && OutputImage<Img>
void my_algorithm(Img img) {
  // write data in img ...
}
```
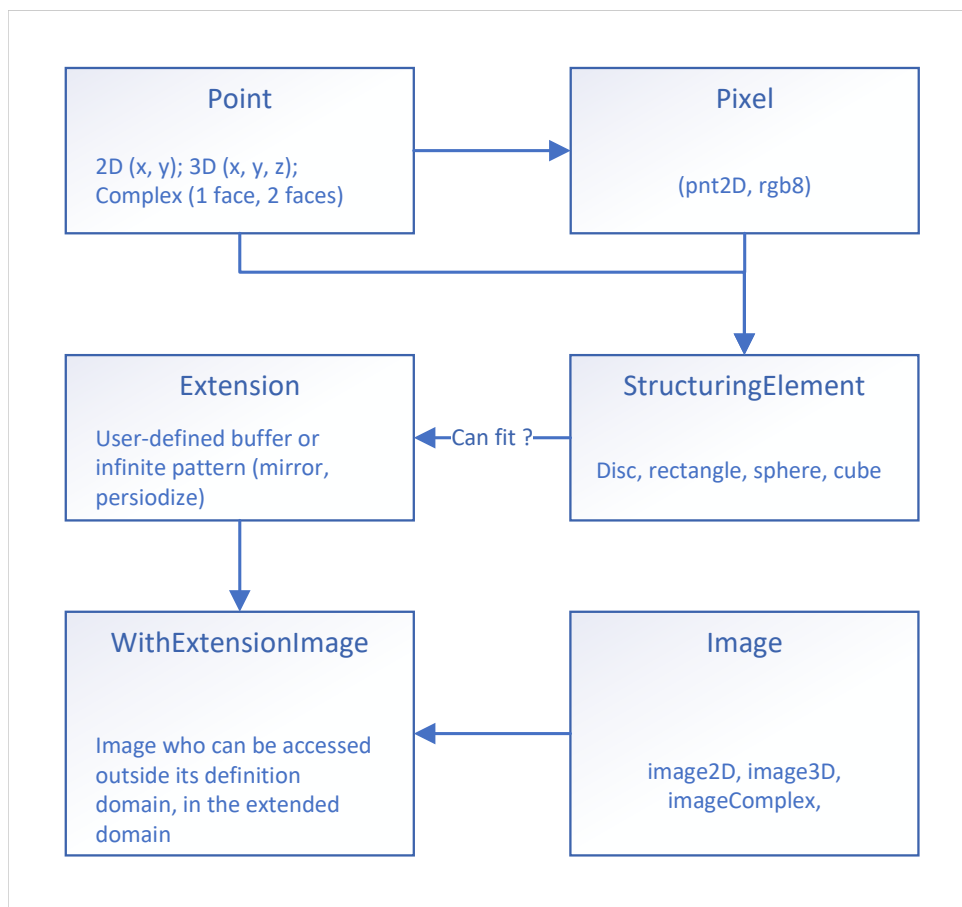
Figure 3.15: Structuring element and Extension concepts.

3.5 Summary

In this chapter, we have demonstrated how it is possible to extract concepts from algorithms. We have then offered a taxonomy of image processing algorithms in order to extract those concepts. We then discuss and dig further the generic aspects of algorithms by introducing the notion of canvas which enables taking advantage of opportunities to do heterogeneous computation while remaining user friendly. Finally, we offer and justify a detailed taxonomy of concepts related to image processing area.

rework summary

Chapter 4

Image views

This concept of views is not new and naturally appeared in Image processing with Milena under the name of morphers [65, 77]. It was always useful to be able to project an image through a prism that could extract specific information about it without the need to copy the underlying data buffer. In modern days, the language C++ (20) also introduces this mechanism with the ranges [157] facilities for *non-owning collections*. It is named *views* and allows the user to access the content of a container (vector, map) through a prism. In Pylene, we decided to align the naming system after what was decided in C++20 in order not to confuse the user. This way, a `transform` view in image processing will do the same thing on an image that the transform view in the standard range library does on a container. *Views* feature the following properties: *cheap to copy*, *non-owner* (does not *own* any data buffer), *lazy evaluation* (accessing the value of a pixel may require computations) and *composition*. When chained, the compiler builds a *tree of expressions* (or *expression template* as used in many scientific computing libraries such as Eigen [67]), thus it knows at compile-time the type of the composition and ensures a 0-overhead at evaluation. We will first motivate the usage of *views* in image processing. We will then present the main views used in image processing. Then will be discussed how image views differ from the one used in C++'s ranges and their main properties (especially how they keep/discard the properties from the parent image) through a concrete example: the management of border and extension policies. Finally, we will discuss the limitations of such a design.

4.1 The Genesis of a new abstraction layers: Views

In image processing an algorithm is naively written by taking one or several inputs' data (among which is the input image(s)), by performing work on this input data and then by returning the resulting data (or an error). Let us take for example the alpha blending example which can be implemented in naive C++ code as followed:

```cpp
void blend_inplace(const uint8_t* ima1, uint8_t* ima2, float alpha,
int width, int height, int stride1, int stride2) {
  for (int y = 0; y < height; ++y) {
    const uint8_t* iptr = ima1 + y * stride1;
    uint8_t* optr = ima2 + y * stride2;
    for (int x = 0; x < width; ++x)
      optr[x] = iptr[x] * alpha + optr[x] * (1-alpha);
  }
}
```

This code has several flaws. It makes strong hypothesis about the input images: its data buffer contiguity and its shape (2D). Let us suppose that our user now wants to restrict the algorithm to a specific region inside the image. The maintainer would have then to provide a version of the algorithm with one additional input argument corresponding to the region of interest. Let us suppose that the user now wants to support manipulate 3D images. The maintainer would now have to provide two additional with an additional stride argument (one for the base algorithm,

one for the region of interest-restricted algorithm). Let us now suppose that the user only wants to manipulate the red color channel. Now the maintainer must support and additional versions of the algorithm for each channel and/or color type. The complexity increases manyfold for each kind of customization points the maintainer wants to offer to the user. Of course, it is possible to prevent code duplication through clever usage of computer engineering techniques (code factorization etc.) but the complexity would still leak through the API anyway. That is way the other solution is to make the user able to perform those restriction upstream from the algorithm transparently so that the downstream algorithm is easy to write, understand and maintain. In order to achieve this, we need to raise the abstraction level around images by one layer so that we can work at the image level. The alpha blending algorithm would then be written as shown in fig. 4.1.

Figure 4.1: Alpha-blending algorithm written at image level.

This way to express an algorithm is achieved by introducing *views* to image processing. An image now is a view and can be restricted/projected/manipulated however the user need before feeding it to an algorithm. Even the whole alpha blending algorithm can be rewritten in terms of views entirely, as shown in fig. 4.2.


```cpp
auto alphablend =
  [](auto ima1, auto ima2, float alpha) {
    return alpha * ima1 + (1 - alpha) * ima2;
  };
```

Figure 4.2: Alpha-blending, generic implementation with *views*, expression tree.

Being able to perform powerful manipulation on images before feeding them to algorithms completely nullify the initial problem of having several versions of the same algorithm while maintaining and documenting all the associated optional arguments. Indeed, in order to perform the alpha blending transformation on the base input image, all that the user must do is:

```cpp
auto ima1, ima2 = /* ... */;
auto ima_blended = alphablend(ima1, ima2, 0.2);
```

If the user wants to restrict the region to be blended, or the color channel to work on, he just has to write the following modification:

```cpp
auto roi = /* ... */;
auto blended_roi = alphablend(view::clip(ima1, roi), view::clip(ima2, roi), 0.2);
auto blended_red = alphablend(view::red(ima1), view::red(ima2), 0.2);
```

The restriction is done upstream from the algorithm and propagated downstream without increasing the code complexity. This way, view greatly increase what the user can do by writting less code.

4.2 Views for image processing

There are four fundamental kinds of views, inspired by functional programming paradigm: `transform(input, f)` applies the transformation f on each pixel of the image *input*, `filter(input, pred)` keeps the pixels of *input* that satisfy the predicate *pred*, `clip(input, domain)` keeps the pixels of *input* that are in the *domain*, `zip(input_1, input_2, ..., input_n)` allows to pack several pixels of several images to iterate on them all at the same time. From those four fundamentals come out more very useful views such as `cast<T>(input)` or `mask(input, msk)` that are more specific to the image processing area.

In Pylena, the practitioner can use a large array of views. Those views come into different form and allow the practitioner to seamlessly use arithmetic or logic operators on images like he would when using expression template. We separate the available views in two main families: the views that perform a restriction of the domain (clip, filter) and the views that transform the values (transform, zip).

4.2.1 Domain-restricting views

The filter view is also a fundamental view which consists in keeping only the values that satisfy a predicate. This is very useful when working with thresholds as shown in the following code:

```
auto my_threshold = 145;
auto inferior_to [my_threshold](uint8_t val) { return val <= my_threshold; };
auto superiorstrict_to = [](uint8_t val) { return not inferior_to(val); };
mln::image2d<uint8_t> ima_grayscale = /* ... */;
auto ima_inferior = mln::view::filter(ima_grayscale, inferior_to);
auto ima_superiorstrict = mln::view::filter(ima_grayscale, superiorstrict_to);
mln::fill(ima_inferior, 0u8);
mln::fill(ima_superiorstrict, 255u8);
```

This code shows a way to binarise `ima_grayscale` with a custom threshold using the filter view. It is important to note that the resulting filtered image has its domain of definition changed. And the new domain of definition will most likely not be in a regular usual shape (such as a 2D rectangle). This implies that the usage of this view inside certain algorithms may be limited.

The clip view is a convenient way to extract a sub-image from a base image. This view essentially redefine the domain of definition to restrict it into a smaller one. It does not change anything else which means it proxies every access to the image. For instance, we make use of this view to easily subdivide a 2D-image into 4 tiles as shown in the code below:

```
mln::image2d<mln::rgb8> large_image = /* ... */;
point2d shape = large_image.domain().shape();
auto middle_pnt = point2d{shape.x() / 2, shape.y() / 2};
auto tl = large_image.domain().tl(); // top-left point
auto br = large_image.domain().br(); // bottom-right point
auto four_tiles = std::tuple{
  mln::view::clip(ima, mln::box2d{tl, middle_pnt}),   // top-left tile
  mln::view::clip(ima, mln::box2d{                    // top-right tile
    point2d{middle_pnt.x(), tl.y()},
    point2d{br.x(), middle_pnt.y()}
    }),
  mln::view::clip(ima, mln::box2d{                    // bottom-left tile
    point2d{tl.x(), middle_point.y()},
    point2d{middle_pnt.x(), br.y()}
  }),
  mln::view::clip(ima, mln::box2d{middle_pnt, br})    // bottom-right tile
};
```

The mask view is very image-processing oriented as it allows the practitioner to provide a boolean image the same size as the original image to select only the pixels whose corresponding value in the mask is true. Its usage is shown in the following code:

```
mln::image2d<mln::rgb8> ima = /* ... */;
auto mask = ima > 127;
mln::fill(mln::view::mask(ima, mask), 255);
```

This code set all the values that are superior to 127 to the max value 255. It shows that it can both be used with read and write access.

4.2.2   Value-transforming views

The transform view   is the most important view of all. It consists in applying a function to each image's pixel. For instance, writing the grayscale algorithm with a transform view is as simple as the following code:

```
auto grayscale_transform = [](mln::rgb8 val) -> uint8_t {
  return 0.2126 * v[0]   // red
       + 0.7152 * v[0]  // green
       + 0.0722 * v[0]; // blue
};
mln::image2d<mln::rgb8> ima_rgb = /* ... */;
mln::image2d<uint8_t>   ima_grayscale = mln::view::transform(ima_rgb, grayscale_transform);
```

There is no loop in this code, just the pixel-wise transformation function. Furthermore, the code will not compute the resulting image. The computation will happen on-the-fly each time a value from `ima_grayscale` is yielded. This view allows the practitioner to quickly write and adapt any pixel-wise algorithm he needs for his more complex calculation in an efficient way.

The zip view   is one of the most useful view and allow the practitioner to iterate over a set of image at the same time. The basic use-case consists in iterating over a set of input image and the output image to be able to consistently assign output values to a resulting computation from input values. Its usage is shown in the following code:

```
mln::image2d<uint8_t> input = /* ... */;
mln::image2d<uint8_t> output{input.domain()};
auto zipped_ima = mln::view::zip(input, output);
for (auto&& [v_in, v_out] : zipped_ima.values())
  v_out = v_in < 145 ? 0 : 255; // binarisation
```

This code is another example of how to compute a binary threshold image.

The channel/RGB views   is a projector to access a specific color channel of an image. There exists image with many more channels than just the standard red/green/blue ones, from the astrophysics or medical area for instances. This view is a tool to restrict an image and only access a specific channel. Its usage is shown in the following code:

```
mln::image2d<mln::rgb8> ima = /* ... */;
mln::copy(mln::view::red(ima), mln::view::green(ima));
```

This code copies the red component into the green component. It shows that the view can be used in both read and write access. Another more generic view exists; `mln::view::channel(ima, k)`, that access the k-th channel in `ima`.

The cast views   is a way to convert an image's underlying type to another type, by performing a cast. As this does not modify the underlying value in itself, the access cannot be a write access. This view can be used as shown in the following code:

```
mln::image2d<double> ima = /* ... */;
mln::image2d<uint8_t> ima_8bits = mln::view::cast<uint8_t>(ima);
```

The arithmetical operators $+, -, *, /, \%$ are implemented in the form of transformation views that operate point-wise between two images whose size is identical. For instance, writing the following code:

```
mln::image2d<uint8_t> ima1 = /* ... */;
mln::image2d<uint8_t> ima2 = /* ... */;
auto ret = ima1 + ima2;
```

Is equivalent to writing the following code:

```
auto ret = mln::view:transform(ima1, ima2, [](auto v1, auto v2){ return v1 + v2; });
```

It is important to note that the - unary operator is also supported: `-ima1`.

The logical operators $<, <=, ==, != =, >, >=$ are implemented in the same way that arithmetical operators are. Both unary and binary operators are expressed as transform views such as writing the following code:

```
auto ret = !ima1 && ima2;
```

is equivalent to writing the following code:

```
auto tmp = mln::view::transform(ima1, [](auto v){ return !v; });
auto ret = mln::view::transform(tmp, ima2, [](auto v1, auto v2){ return v1 && v2; });
```

It is far more expressive and more comprehensible by the practitioner. Also, a new facility is introduced to express the logic behind a ternary expression (if C then A else B): the operator $ifelse(C, A, B)$. The rationale is to be able to swap between values depending on a boolean mask. This way, a mathematical morphology algorithm such as *hit or miss* can be implemented in the following simple manner:

```
mln::image2d<uint8_t> ima = /* ... */;
auto ero = erode(ima);
auto dil = dilate(ima);
uint8_t zero = 0;
auto ret = mln::view::ifelse(dil < ero, ero - dil, zero);
```

Everything is taken care of and the practitioner just has to write down his algorithm to get it done.

The mathematical operators are implemented in the form of views that operates point-wise. The supported mathematical operators are the following: abs, pow, sqr, cbrt, sqrt, sum, prod, min, max, dot, cross, l0norm, l1norm, l2norm, l2norm_sqr, linfnorm, lpnorm, l0dist, l1dist, l2dist, l2dist_sqr, linfdist, lpdist. Calling an operator onto an image is equivalent to calling a transform view on each value of this image:

```
auto ima = /* ... */;
auto ret = view::maths::abs(ima);
```

Is equivalent to calling:

```
auto ima = /* ... */;
auto ret = view::transform(ima, [](auto v){ return std::abs(v); });
```

4.3   View properties

Views feature interesting properties, especially how they keep/discard the properties of the concrete image they are based on. However, before talking about those properties, it is important to draw the line and point the main differences between the C++20 ranges views and our image views.

4.3.1   Differences between C++20 ranges views and image views

C++20 ranges views are a new abstraction layer introduced on top of the already existing iterators. This means that a view is created from an existing container from its iterators. There are also special views such as `std::views::iota` that are able to generate an infinite sequence of number. Those last are the generator views and are designed to be used the same way as a container, except that they do not own any data. In C++20 the views are mainly constructed from a container such as `std::vector`, `std::map` or `std::list`. For instance, the way to create a view featuring all the elements of a container is shown in the following code:

```
auto vec = std::vector { /* ... */ };
auto vec_vw = std::views::all(vec);
```

This induces issues regarding dangling references when passing temporary views or when the container owning the data expires. To summarize the model of ranges in C++20, they are a new abstraction layer much more friendly and powerful than iterators and can construct non-owning, cheap-to-copy views from an owning container.

4.3.2   Data ownership

The concept of *View* brought to us a fundamental issue when dealing with images: "*What is an image?*". More precisely: should an image always be the owner of its data buffer? Should we have a shared ownership of the data buffer between all the images using it? Then what happens when the data changes? The issue about the semantic of an image is crucial but also very similar to the issue there is to differentiate a *container* (such as `std::vector`, that is to say the data buffer) and a *view*, as seen in section 4.3.1.

From here we have considered two approaches. The first one is to have *shared ownership* of the data buffer for the image and its derived views. However, this does not allow the differentiation between an already computed image and a lazy image. To be able to make this differentiation is crucial in an *Image Processing library* as we want to make the most out of the data we already have, and we do not want to compute data we do not need. Also, we cannot distinguish when the *copyability* property is required. This is the main reason why we did not adopt this approach.

The second one is to make the differentiation between a *concrete image* which owns the data (like the standard containers) and the *views* that are lightweight cheap-to-copy objects. However this would imply we have to distinguish both image families when writing algorithms and we do not want that.

This is why we have chosen to take a path where we mix both approaches at the same time. We assert that all image are cheap-to-copy (including views), even the concrete images. The concrete image will have a shared ownership semantic related to its data buffer and will remain cheap-to-copy. It will also behave the same way a view behaves. That is why, with our semantic, all images are views. Ultimately, there is still a way to distinguish a view from a concrete image, if needed, and we introduce two new concepts `ViewImage` and `ConcreteImage` for this end:

```
template <typename I>
concept ConcreteImage =
  Image<I> &&
  std::concepts::semiregular<I> &&  // A concrete image is default constructible
  not image_view_v<I>;

template <typename I>
concept ViewImage =
  Image<I> &&
  image_view_v<I>;
```

Having images as views is a very important property as it simplify greatly the reasoning when performance is needed. This allows the user to pass his images everywhere without worrying about dangling data buffer expiring around the corner. It also enables us to have a library design similar to the C++'s standard library which the user is familiar with and, why not, have

standard algorithm and standard view work on our images types. All of these are the main reason why we decided to adopt this design. However, unlike the standard library, as we are not required to work with iterator due to backward compatibility.

```
auto h = [](int x) { 0    if x < 150
                      255  if x ≥ 150

auto u =

auto v = transform( u , h ) →
```
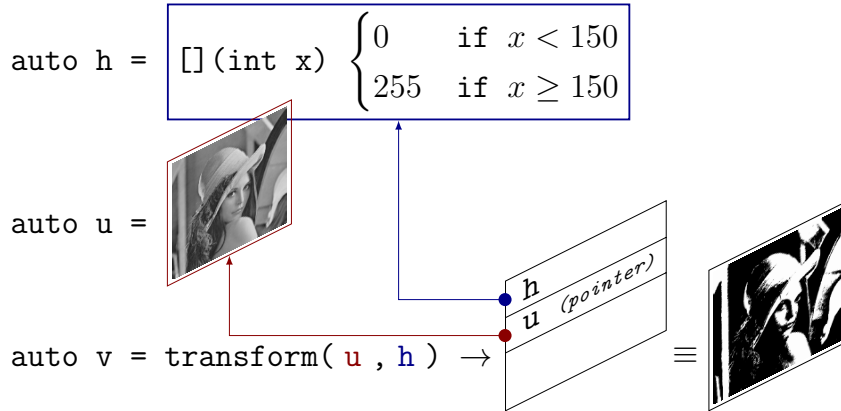
Figure 4.3: An image *view* performing a thresholding.

In our design, all images are lightweight (cheap-to-copy) objects with shared ownership over the data. A view image is a non-owning image that only stores pointers, as shown in fig. 4.3. A concrete image stores the data. The only difference between a view and a concrete image is given by a trait `image_view` which will check the `view` property of the given image to tell whether it is a concrete owning image type or not. Compared to the C++20 ranges views model, mechanism prevent errors resulting from dangling references and confusing ownership of data. It is especially adapted to image processing as the user generally wants to avoid deep copy of its data. When the user wants to deep-copy his image (clone) he wants to do it explicitly.

This design induces one major property which is the *lazy-evaluation* of the views.

4.3.3 Lazy evaluation, composability and chaining

The key point of views is the lazy evaluation. When a concrete image is piped through a view, no computation is done. The computation happens when the practitioner requests a value by doing $val = V(p)$. The implications are multiples: an image can be piped into several computation-heavy views, some of which can be discarded later on, and it will not impact the performance. Also, when processing large images, applying a transformation on a part of the image (such as clipping or filtering) is as simple as restricting the domain with a view and applying the transformation to this resulting sub-image.

Lazy-evaluation combined with the view *chaining* allows the user to write clear and very efficient code whose evaluation is delayed till very last moment as shown in fig. 4.4 (see [119] for additional examples). Neither memory allocation nor computation are performed; the image i has just recorded all the operations required to compute its values.

The tree of type resulting from this view chaining is illustrated by fig. 4.5. It illustrates how chaining views with each other result in the formation of an abstract tree that records the operations to perform. This model enables building complex computational trees via views while keeping efficient performances at run-time. However, those trees are complex for the compiler to process and can induce heavy compilation time overhead.

4.3.4 Preserving image properties

Views will also try to preserve properties of the original image when they can. That means that views can preserve the ability of the practitioner to, for instance, write into this image. This may be a trivial property to preserve when considering a view that restrict a domain, but when

```
image2d<rgb8>  ima1 = /* ... */;
image2d<uint8_t> ima2 = /* ... */;

// Projection: project the red channel value
auto f = view::transform(ima1, [](auto v) {
  return v.r;
});

// Lazy-evaluation of the element-wise
// minimum
auto g = view::transform(view::zip(f, ima2),
  [](auto value) {
    return std::min(std::get<0>(value),
            std::get<1>(value));
});
```

```
// Lazy-Filtering: keep pixels whose value
// is below < 128
auto h = view::filter(g, [] (auto value) {
  return value < 128;
}));

// Lazy-evaluation of a gamma correction
using value_t = typename Image::value_type;
constexpr float gamma = 2.2f;
constexpr auto max_val =
  std::numeric_limits<value_t>::max();
auto i = view::transform(h,
  [gamma_corr = 1 / gamma, max_val] (auto value) {
    return std::pow(value / max_val,
            gamma_corr) * max_val;
});
```

Figure 4.4: Lazy-evaluation and *view* chaining.
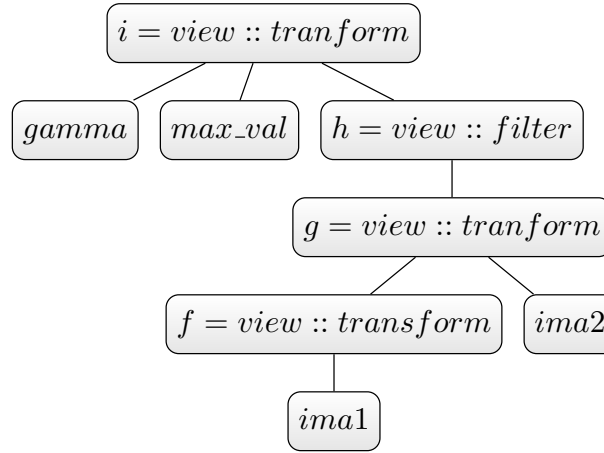
Figure 4.5: Abstract Syntax Tree of the types chained by the code in fig. 4.4

considering a view that transforms the resulting values, it is not. Let us consider the projection $h : (r, g, b) \mapsto g$ that selects the green component of an RGB triplet. When piping the resulting view into, for instance, a blurring algorithm, the computation will take be performed in place thanks to the fact we still have the ability to write into the image. A legacy way of obtaining the same result would have been to create a temporary single-channel image to extract the green channel of the original RGB image so that the temporary image could then be blurred. The final step would have then been to copy the values of the temporary blurred image back into the green channel of the original image. The comparison, including the memory used, between the legacy way and the in-place way of doing this computation is shown in fig. 4.6.

On the other hand, when considering the view $g : (r, g, b) \mapsto 0.2126 * r + 0.7152 * g + 0.0722 * b$ that compute the gray level of a color triplet (as shown in fig. 4.7), the ability to write a value into the image cannot be preserved. Indeed, one would need an inverse function that is able to deduce the original color triplet from the gray level to be able to write back into the original image. This operation alone is a whole field of research on its own [110, 48, 40]

Similarly, a view can apply a restriction on an image domain. In fig. 4.8, we show the adaptor `clip(input, roi)` that restricts the image to a non-regular `roi` and `filter(input, predicate)` that restricts the domain based on a predicate. All subsequent operations on those images will only affect the selected pixels. In this case of restriction, the ability to write data back into the original image is preserved through the view.

Views feature many interesting properties that change the way we program an image processing application. To illustrate those features, let us consider the following image processing pipeline: (Start) Load an input RGB-8 2D image (a classical HDR photography) (A) Convert

(a) Legacy pipeline with copy

(b) Modern pipeline with view
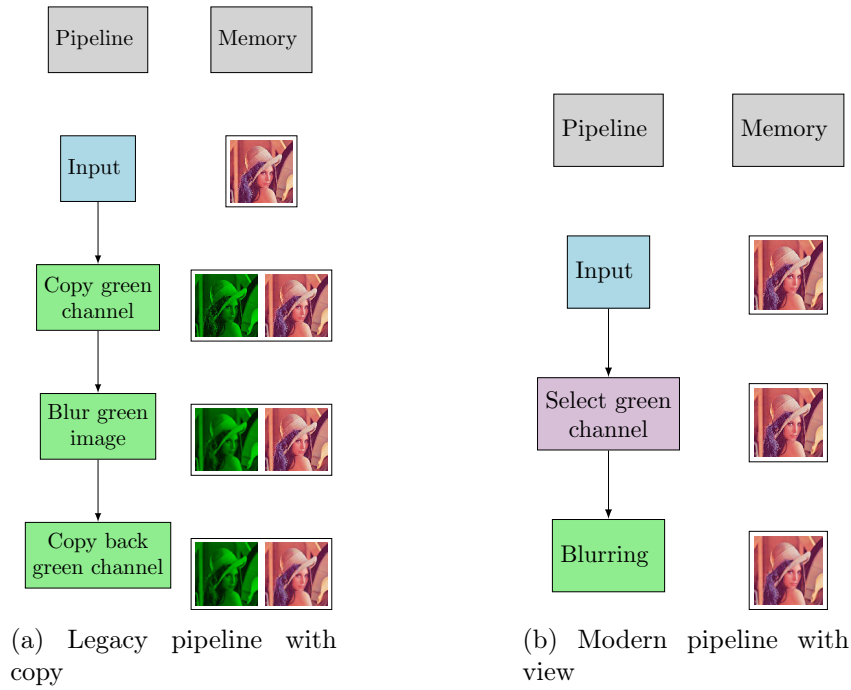
Figure 4.6: Comparison of a legacy and a modern pipeline using *algorithms* (green) and *views* (purple).

Figure 4.7: Usage of transform view: grayscale.

(a) Clip view

(b) Filter view

Figure 4.8: Clip and filter image adaptors that restrict the image domain by a non-regular ROI and by a predicate that selects only even pixels.

Figure 4.9: Example of a simple image processing pipeline.

it in grayscale (B) Sub-quantize to 8-bits (C) Perform the grayscale dilation of the image (End) Save the resulting 2D 8-bit grayscale image; as described in fig. 4.9. This pipeline is expressed with two notions. The first notion is composition of algorithms (A -> B -> C) in order to achieve the desired result. The second notion is the composition of views (Input -> A -> B) which overlaps partially with the algorithm part. This express that part of the algorithm is performed lazily only when we perform the last part (C). Those two notions are illustrated by the fig. 4.10.

Figure 4.10: Example of a simple image processing pipeline illustrating the difference between the composition of algorithms and image views.

There are six properties one want to keep track when working with views: *forward, writable, accessible, indexable, bidirectional* and *raw*. Those properties echo to the concepts seen in section 3.4. An image is *forward* when it can be traversed in a forward way. It is *writable* when the values are mutable. It is *accessible* whenever it allows to access the value associated to a point (i.e. it allows to write the expression $v = ima(p)$). It is *indexable* whenever its values can be accessed through an index localizer (i.e. it allows to write the expression $v = ima[idx]$). Usually, accessing through an index is faster than accessing by a point. It is *bidirectional* when it can be traversed in both a forward and a backward way. Finally, an image is *raw* when its data buffer can is contiguous and can directly be accessed with information about strides. The table 4.1 presents all the views and how they preserve the base properties of a concrete image.

Table 4.1: Views: property conservation

View type / Expression	Property	Forward	Bidirectional	Raw	Writable	Accessible	Indexable
Image	`ima1, ima2`	✓	✓	✓	✓	✓	✓
Cast	`cast<T>(ima)`	✓	✓	✗	✗	✓	✓
Transform	`transform(ima, func)`	✓	✓	✗	✓[1]	✓	✓
Filter	`filter(ima, pred)`	✓	✓	✗	✗	✓	✓
Clip	`clip(ima, dom)`	✓	✓	✗	✓	✓	✓
mask	`mask(ima, mask)`	✓	✓	✗	✓	✓	✓
Zip	`zip(ima1, ima2)`	✓	✓	✗	✓	✓	✓
Channel	`red(ima)`	✓	✓	✗	✓	✓	✓
Arithmetic	`ima1 + ima2`	✓	✓	✗	✗	✓	✓
Logical	`ima > 125`	✓	✓	✗	✗[2]	✓	✓
Mathematical	`abs(ima)`	✓	✓	✗	✗	✓	✓

[1]: writability is preserved only if `func` is a projection.
[2]: writability not preserved except for the expression `ifelse(ima, ima1, ima2)`.

Also, we may want to extend the property preservation discussion to other concepts we saw in the previous chapter 3, especially the concepts of structuring element and extension.

4.4 Decorating images to ease border management

When looking at local algorithms, we notice that a long recurring issue is about the behavior on the border of the image. There are many ways of dealing with this problem. One is to allocate additional memory for the border and paste values in it. Another is to check the bounds when looping over the neighbors inside the computational window. We can also decorate the image to return a correct lazily computed value when accessing out-of-image-bound value still inside the extension. The point is: all these methods have advantages as well as disadvantages.

Memory allocated border The border width is fixed at the image's creation and cannot be augmented without doing a reallocation. There is also a cost when computing border's values (to fill it) which is proportional to the border's width and to the image's size. On the other hand, the access time of a border value during the algorithm unrolling is as fast as a native access time within the image itself. The last issue remaining would be that the border is not infinite. We cannot process a local algorithm with a structuring element that does not fit in the extension. This method is especially adapted when there is medium structuring elements with a known size which will yield a lot of out-of-image's bound accesses. When speed is required, this method is a de facto standard.

Bound checking Assuming there is no border, and we are not allowed to access out-of-image-bound values, a bound check is required when accessing each values. Another way to do would be to decorate the facility that yields the neighbors of a pixel: do not yield out-of-image-bound pixels. This removes the need to bound check for each pixel's value which is relatively faster. The caveats of this method are that it induces a slight slow down when yielding the pixel's neighbors from the structuring element, and that it is not always viable: some algorithms do need to access values in an extension to produce proper results.

Image decoration The border is infinite, and we make a view of our image to decorate it with the required extension. This is achieved using *views*: the original image is chained into a view that will add the required feature to the image. For instance, let us consider the following image:

```
struct borderless_image {
  // ...
  // NO extension_type subtype
  // NO extension() method
};
```

Attempting to use this image in a local algorithm that works with a structuring element as the structuring element does not fit inside the image when considering the behavior on the borders. However, instead of narrowing the region of interest, it is possible to make a view that will return an image from which the behavior at the border is well-defined. Referring to the taxonomy from the previous chapter 3 we remember that we can construct a custom extension type for the sake of an example (as described in section 3.4.3). This example will decorate the image so that the border is always filled with a specific value. The following code shows how we can write such an extension:

```
template <class ValueType>
struct FillExt {
  using support_fill = std::true_type; // Support the fill policy
  bool is_fill_supported() const { return true; }

  using value_type = ValueType;        // Underlying value_type

  // Always fit structuring element of any size
  template <class StructuringElement>
  bool fit(const StructuringElement& se) const { return true; }
```

```cpp
  void fill(ValueType v) { v_ = v; }  // Assign the filled value
  // ...
  // Yield the value for a given point (always return filled value)
  template <class PointType>
  ValueType val(PointType) const { return v_; }
  // ...
  private:
    ValueType v_;
};
```

Now that our extension type is written, we introduce a new image type which will adapt our base `borderless_image` into a view which features out `FillExt` extension type.

```cpp
template <class BaseImage>
struct filled_border_image : image_view_adaptor<BaseImage>{
  // ...
  using extension_type = FillExt;
  extension_type& extention() const { ext_; }
  // ...

  value_type at(point_type pnt) {
    if(!domain().has(pnt)) {
      return extension().val()
    }
  }
  // ... adapt all the methods that can make out-of-bound access and fallback on
  // the extension's value ...
private:
  extension_type ext_;
};
```

Finally, all that is left to do is to write the function that will construct the view from the base image:

```cpp
template <class BaseImage>
auto fill_extension_view(BaseImage ima) {
  // call to the image_view_adaptor ctor
  auto with_fill_ext_ima = filled_border_image<BaseImage>(ima);
  return with_fill_ext_ima; // <-- this is a view
}
```

This simple function enables very powerful usage as illustrated in the code below:

```cpp
auto ima = image2d<uint8_t>{
  {0, 1, 2},
  {3, 4, 5}
  {6, 7, 8}
};
ima.at({1, 1}); // OK, 4
ima.at({5, 5}); // ERROR, out-of-bound
// Get a view
auto ima_with_filled_border = fill_extension_view(ima);
// Fill border with value 255
ima_with_filled_border.extension().fill(255);
ima_with_filled_border.at({5, 5}); // OK, 255
```

For the sake of brevity we have simplified the implementation in our example. In practice the implementation of such a pattern is more complex as there are many strategies to support, the interfaces of the extension may be different, the decoration of the image type may not be enough, notably for the *none* strategy where it is required to decorate the structuring element.

At the end, this method has the advantage to *always work*. Given any structuring element of any size, any algorithm will work. The disadvantage is that we need to check for out-of-bound access at the image level, and lazily compute the value in case of out-of-image-bound access. The slowness induced is not negligible and should be weighted carefully.

It is important to note the very close relation between an image's domain (to perform out-of-bound checks), the structuring element (notably its size) and the extension (its width). A user may require, for a specific set of those three elements, to decorate the image, and/or the

structuring element and/or to perform computation and/or reallocation. To resolve this issue, we decided to provide the user with a new facility: the *border manager* whose job is to prepare a suitable pair (image and structuring element) given a set of configuration wanted by the user.

We designed the configuration to be constructed from a given set of a policy and a method. We currently offer two policies: native and auto.

- Native: if the border is large enough: forward the image as-is to the algorithm to allow the fastest access possible. Otherwise, the border manager fails and halt the program.

- Auto: if the border is large enough: forward the image as-is to the algorithm to allow the fastest access possible. Otherwise, decorate the image with a view whose extension will emulate what is required by the algorithm with the given structuring element.

We also provide seven different methods to fill up our extension with the wanted values. It is important to note that not all the methods are available for both policies. The policies are: none, fill, mirror, periodize, clamp, image and user.

The *none* policy enforces a policy where there is no border to use. The method cannot fail as it enforces the border to vanish. To enforce this method, the border manager decorates the structuring element in a view that checks the domain inclusion of each neighboring point. The *fill* policy enforces that the border is filled with a specific value. The *mirror* policy enforces that the border is filled with a mirrored value from an axial symmetry relative to the image's edges. The *periodize* policy enforces that the border replicate the image, like a mosaic. The *clamp* policy enforces that the border is filled with values expanded from the values at the image's edge. The *image* policy enforces all points out of the current image's domain are to be picked inside another image. A basic use-case is preparing tiles from a large image. The position of our image can be offset in the image acting as an extension which ease the usage when, for instance, clipping a sub-image. The fig. 4.11 shows how a sub-image (tile) can consider the base image as its border. Finally, the *user* policy assumes the user knows what he is doing and do not touch nor decorate the given image in any way. Both policies lead to the same behavior: check whether the structuring element fit and then forward the image as-is if it fits. An exception is raised if it does not. The fig. 4.11 illustrates all the other border policy mentioned.

As a consequence the usage of a local algorithm becomes very simple:

```cpp
// default border width is 3
image2d<int> ima = {{0, 1, 0}, {0, 1, 1}, {0, 1, 0}};
auto disc_se = se::disc{1}; // radius is 1
auto bm = extension::bm::fill(0); // fill border with 0 with policy auto

local_algorithm(ima, disc_se, bm); // will handle the border for you
```

The border manager `bm` is set with the method fill (with value 0) and policy `auto` (which is the default policy). To use the policy `native`, one would write `extension::bm::native::fill(0)` instead.

In the implementation of the local algorithm, a dispatch is made with the pattern *visitor*, relying on the standard facilities `std::visit` and `std::variant` so that the performance overhead as well as the complexity of use remain minimal. Let us assume we have a local algorithm implemented this way:

```cpp
template <class Ima, class SE>
local_algorithm(Ima ima, SE se)
{
  // assume ima has a large enough border for the given se
  // use ima & se in loop
}
```

We can rewrite it leveraging the border manager facility this way :

(a) Border *None*   (b) Border *fill* w/ 0   (c) Border *mirror*

(d) Border *periodize*   (e) Border *clamp*   (f) Border *image*
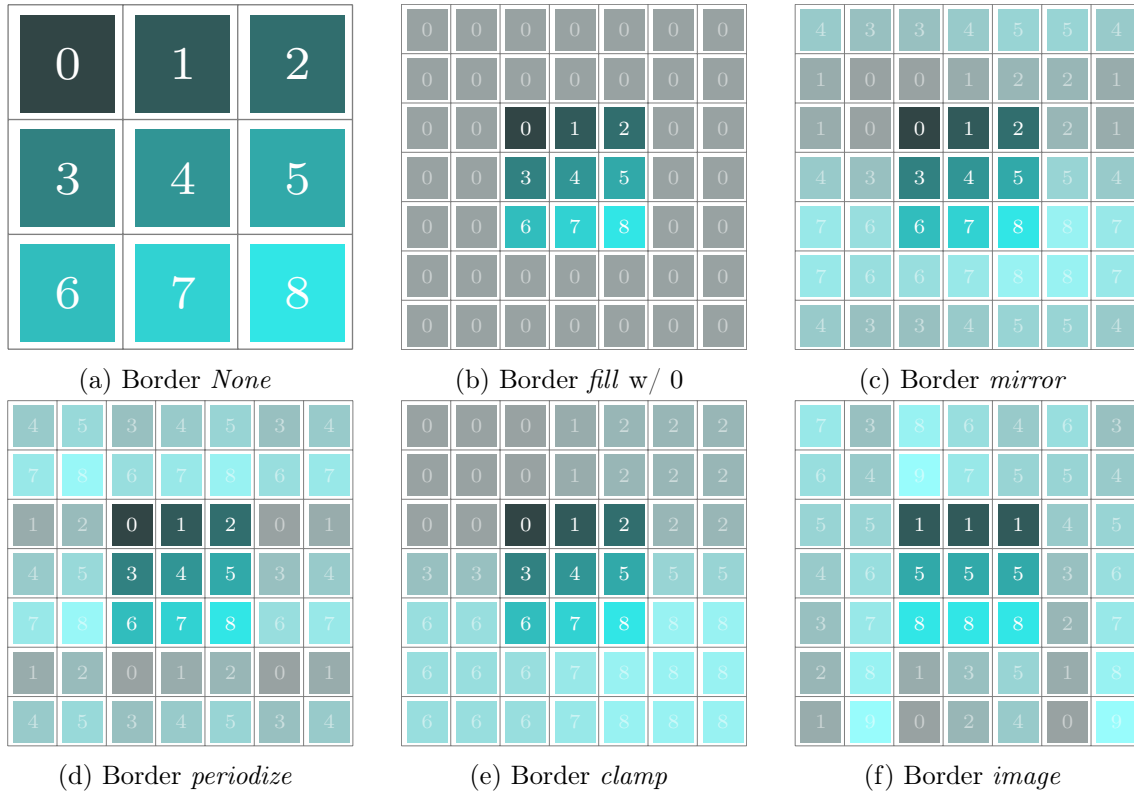
Figure 4.11: Border methods' breakdown.

```cpp
template <class Ima, class SE, class BM>
local_algorithm(Ima ima, SE se, BM bm)
{
  auto [managed_ima, managed_se] = bm.manage(ima, se);
  std::visit([&](auto&& ima_, auto&& se_) {
    // use ima_ & se_ in loop
    }, managed_ima, managed_se);
}
```

The overhead is kept minimal thanks to using `std::variant` and `std::visit` and the algorithm implementer delegate the border management to the border manager. This is made possible thanks to the views. Indeed, under the hood the border manager may pipe the original image into a view that will behave accordingly to the policy chosen by the user. This will be transparent from both practitioner and maintainer points of views.

4.5   Views limitations

Views can be of tremendous use in our area however it relies on metaprogramming techniques which are infamous for greatly increasing the compilation time of source code. Also, when one starts to chain views a lot, combining different image type (via *zip* for instance), combined with the overhead induced via the border manager using `std::variant`, the compilation time can really become an issue. Indeed, C++ developers tend to minimize the cost of compilation time because once the program is compiled, the binaries can be distributed and are really fast to execute. However, we are not exactly in that case as our library is generic. That means we distribute source code to our user and our user compiles it when prototyping their program. This is an issue every library developer faces: distributing heavily templated source code as a library can be a deal-breaker. In the industry, it was even to the point that people refused to use boost in their code line. The boost maintainers had to modularize their library, so that user

is able to cherry-pick the parts he needs without pulling half of the library which was a disaster for the compile time of a project.

The ranges for C++20's standard library and its views face the same issue. It was not rare for someone to need 90sec to compile the calendar toy example of the library which just contains code that displays a given month in the classic printed format (day number-of-month correctly displayed in column corresponding to the day of week label). This huge compile time is due to early compiler implementation needing massive RAM usage for template type and having to swap when the computer running the compilation was out of RAM. Now compilers have optimized the whole process, but the combination behind the types can still be an issue. Introducing complex view code in a program that is compiled often may not be a good idea. However, a program that is rarely compiled but is run a great number of time may take advantages of all the optimizations the compiler do to be efficient.

4.5.1   Image traversing with ranges

Lastly, views usage should be measured when used at critical points. We learned from experience that one simple change can make the compiler miss optimization opportunities which can greatly impact the resulting performances. Let us illustrate our remark with a concrete example: image traversing. In previous version of our library, we used macro for image traversing. `mln_concrete`, `mln_piter`, `mln_qiter`, `for_all` and `mln_value` are all macros aiming at hiding the underlying complexity. Our goal were to replace those macros with existing C++ core language code to improve the user experience as well as ease the maintenance, contribution and further improvement of the library. To do so, we based our image traversing on `std::ranges`. Let us take the old implementation we had of our dilation algorithm as an example:

```cpp
template<class I, class SE>
mln_concrete(I) dilate(const I& f, const SE& se)
{
  mln_concrete(I) g;
  initialize(g, f);
  mln_piter(I) p(f.domain());
  mln_qiter(SE) q(se, p);
  for_all(p) // for all p in f domain
  {
    mln_value(I) v = f(p);
    for_all(q) // for all q in se(p)
      if(f.has(q) and f(q) > v)
        v = f(q);
    g(p) = v;
  }
  return g;
}
```

This code features the macro mentioned above and, while being explicit, may be quite obscure with regard to its internals for a non-initiated user. However, it uses a double loop to traverse the image under the hood. Rewriting the algorithm using ranges results in the following code:

```cpp
template<class I, , class SE>
auto dilate(I input, const SE& se)
{
  auto output = input.concretize(); // clone image
  for(auto [in_px, out_px] : view::zip(f.pixels(), g.pixels()))
  {
    out_px.val() = out_px.val();
    for(auto nhx : se(in_px))
      out_pix.val() = std::max(nhx.val(), out_px.val());
  }
  return output;
}
```

This code use the *zip* view to iterate over two images (the input image and the resulting output image) simultaneously. This is native code, and it should, in theory, be efficient than the old

version of the code (with macros) as it enables compiler optimizations such as vectorization or inner loops unrolling. But through benchmarking, we have learned that this solution does not mix well [21] with the multidimensional nature of images. The issue originates from the fact that we have no way to explicitly say in the code that the multidimensional range is made of chunk of contiguous rows of memory. Indeed, for each element we have to compute an index originating from potentially N dimensions. This disables critical optimizations such as vectorization. We solved this problem by augmenting range-v3's ranges with our own multidimensional ranges. Indeed, we only need to have contiguity on the last dimension to provide the compiler code it can optimize. Which means that each for-loop that traverses the whole n-dimensional image can be transformed into a double for-loop whose inner loop is guaranteed to be a contiguous row. This way we have now an outer range as well as an inner range, as illustrated in fig. 4.12.

(a)   (b)

Figure 4.12: Range-v3's ranges (a) vs. multidimensional ranges (b).

Thanks to this new design we can now rewrite our algorithm with a double for-loop for the image traversing. Hopefully it stays really similar to what one would be used to when working with the classical two dimension image. As an example, we can rewrite the dilation algorithm this way:

```cpp
template<class I, class SE>
auto dilate(I input, const SE& se)
{
  auto output = input.concretize(); // clone image
  // this line is needed to avoid dangling reference
  auto zipped_pixels = view::zip(input.pixels(), output.pixels());
  for(auto&& row : ranges::rows(zipped_pixels)) // unroll the contiguous segments
    for(auto [in_px, out_px] : row)              // optimized traversing of the segment
    {
      out_px.val() = out_px.val();
      for(auto nhx : se(in_px))
        out_pix.val() = std::max(in_px.val(), out_px.val());
    }
  return output;
}
```

The highlight of this code is the usage a new tool: **ranges::rows** to bring out an inner range (contiguous) from the multidimensional outer range.

4.5.2   Performance discussion

In order to have a relevant discussion on performances, we decided to implement a real world image processing pipeline: the background subtraction. It is used to detect changes in image sequences [115]. It is mainly used when regions of interest are foreground objects. The

Framework	Compute Time	Memory usage	ΔMemory usage
Pylene (w/o views)	2.11s (± 144ms)	106 MB	*+0%*
OpenCV	2.41s (± 134ms)	59 MB	*-44%*
Pylene (views)	2.13s (± 164ms)	51 MB	*-52%*

Table 4.2: Benchmarks of the pipeline fig. 4.13 on a dataset (12 images) of 10MPix images. Average computation time and memory usage of implementations with/without *views* and with OpenCV as a baseline.

pipeline components include: subtraction, Gaussian filtering, threshold, erode and dilate, as shown in 4.13.

Figure 4.13: Background substraction pipeline using algorithms and views .

The first thing that we notice is that the implementation of the pipeline using views is transcribed very explicitly in the code, as shown in fig. 4.14. There is a direct correspondence between the graphic pipeline and the code.

For our benchmark, we have decided to run the algorithm on an original set of image to detect a changing foreground. We have considered 10 data set. We present in fig. 4.15 three of them for the sake of brevity.

We have run benchmarks on this set comparing multiple ways of achieving this result, both using Pylene and OpenCV as well as varying the size and the shape of the structuring element window. The breakdown of these benchmarks are presented in fig. 4.16. In table 4.2, we benchmark the computation time and the memory usage [1]of these implementations (all single-threaded) with an opening of disc of radius 32 on 10 MPix RGB images (the minimum of many runs is kept).

The results should not be misunderstood. They do not say that OpenCV is faster or slower but shows that implementations all have the same order of processing time (the algorithms used in our implementation are not the same as those used in OpenCV for blur and dilation/erosion) so that the comparison makes sense. It allows us to validate experimentally the advantages of views in pipelines. First, we have to be cautious about the real benefit in terms of processing

---

[1]Memory usage is computed with *valgring/massif* as the difference between the memory peak of the run and the memory peak without any computation (just setup and image loading)

```
float kThreshold = 150; float kVSigma = 10;
float kHSigma = 10;  int kOpeningRadius = 32;
auto img_grey = view::transform(img_color, to_gray);
auto bg_grey  = view::transform(bg_color, to_gray);
auto bg_blurred = gaussian2d(bg_grey,  kHSigma, kVSigma);
auto tmp_grey = img_grey - bg_blurred; /
auto thresholdf = [](auto x) { return x < kThreshold; };
auto tmp_bin = view::transform(tmp_grey, thresholdf); /
auto ero = erosion(tmp_bin, disc(kOpeningRadius));
dilation(ero, disc(kOpeningRadius), output);
```

Figure 4.14: Pipeline implementation with *views* . Highlighted code uses *views* by prefixing operators with the namespace `view`.

Background   Candidate   Result
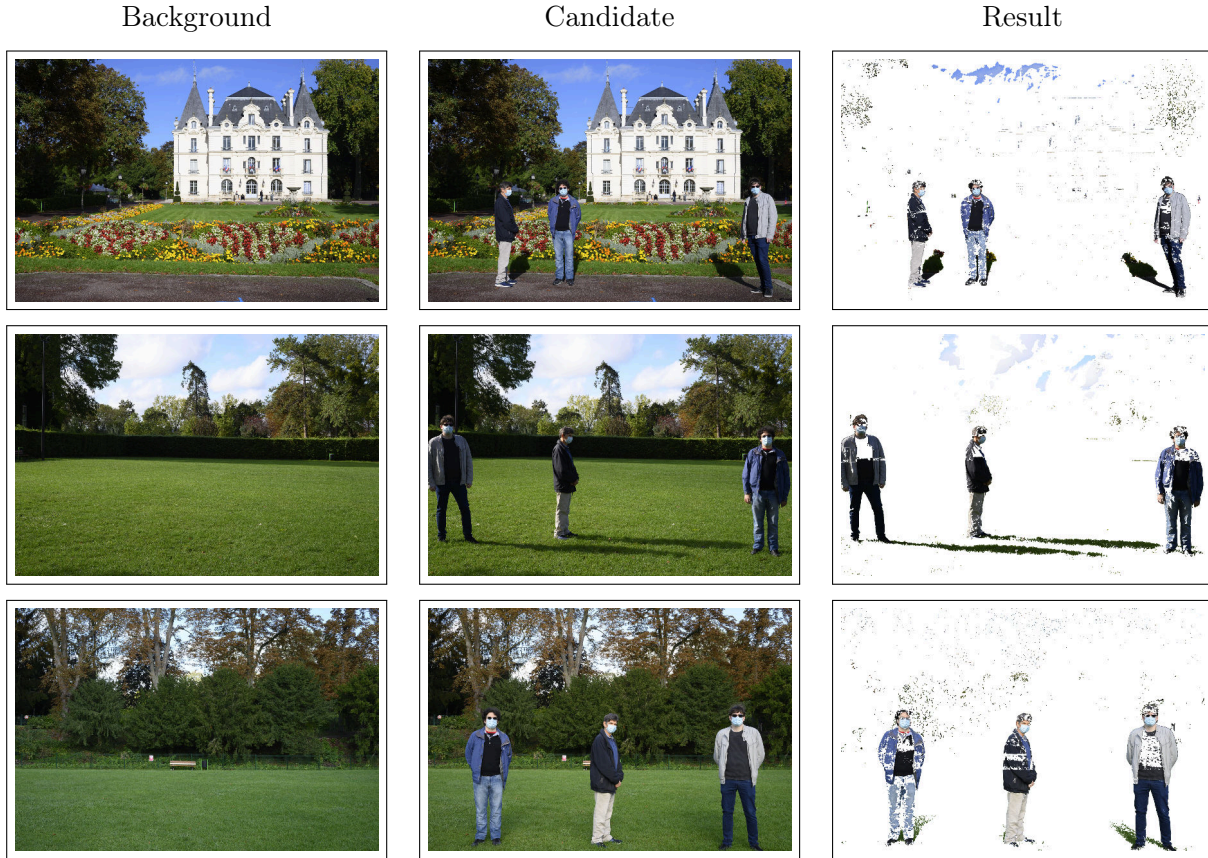
Figure 4.15: Background detection: data set samples.

time. Here, most of the time is spent in algorithms that are not eligible for view transformation. Thus, depending on the operations of the pipeline, views may not improve processing time. Nevertheless, using views does not degrade performance neither (only 1% in this experiment). It seems to show that using views does not introduce performance penalties and may even be beneficial in lightweight pipelines as the one in fig. 4.9. On the memory side, views reduce drastically the memory usage (also seen in fig. 4.6) which is beneficial when developing applications which are memory constrained. From the developer standpoint, it requires only few changes in the code as shown in fig. 4.14 — the implementation of the algorithms remain the same — which is a real advantage for software maintenance.

Result	Benchmark	Benchmark Pln only
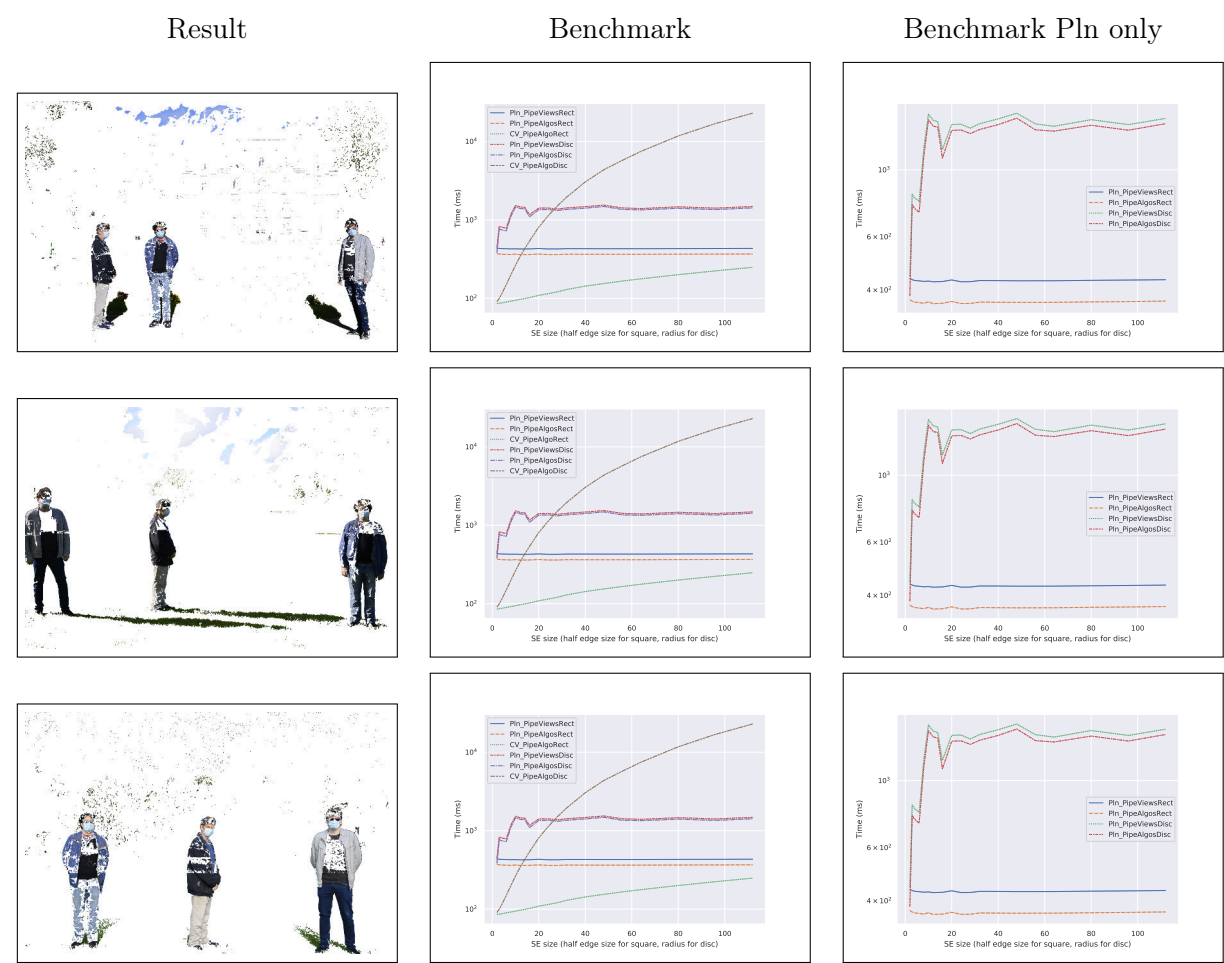
Figure 4.16: Background detection: garden results.

4.6   Summary

Views are composable.   One of the most important feature in a pipeline design (generally, in software engineering) is *object composition*. It enables composing simple blocks into complex ones. Those complex blocks can then be managed as if they were still simple blocks. In fig. 4.9, we have 3 simple image operators *Image → Image* (the grayscale conversion, the sub-quantization, the dilation). As shown in fig. 4.10, algorithm composition would consider these 3 simple operators as a single complex operator *Image → Image* that could then be used in another even more complex processing pipeline. Just like algorithms, image views are composable, e.g. a view of the view of an image is still an image. In fig. 4.10, we compose the input image with a grayscale transform view and a sub-quantization view that then feeds the dilation algorithm.

Views improve usability.   The code to compose images in fig. 4.10 is almost as simple as:
```
auto input = imread(...);
auto A = transform(input, [](rgb16 x) -> float {
    return (x.r + x.g + x.b) / 3.f; }; );
auto MyComplexImage = transform(A, [](float x)
    -> uint8_t { return (x / 256 + .5f); }; );
```

People familiar with functional programming may notice similarities with these languages where *transform* (*map*) and *filter* are sequence operators. Views use the functional paradigm and are created by functions that take a function as argument: the operator or the predicate to apply for each pixel; we do not iterate by hand on the image pixels.

Views improve re-usability.   The code snippets above are simple but not very re-usable. However, following the functional programming paradigm, it is quite easy to define new views, because some image adaptors can be considered as *high-order functions* for which we can bind some parameters. In fig. 4.17, we show how the primitive *transform* can be used to create a view summing two images and a view operator performing the grayscale conversion as well as the sub-quantization which can be reused afterward[2].

Views for lazy computing.   Because the operation is recorded within the image view, this new image type allows fundamental image types to be mixed with algorithms. In fig. 4.17, the creation of views does not involve any computation in itself but rather delays the computation until the expression `v(p)` is invoked. Because views can be composed, the evaluation can be delayed quite far. Image adaptors are *template expressions* [16, 30] as they record the *expression* used to generate the image as a template parameter. A view actually represents an expression tree (fig. 4.2).

------

[2]These functions could have been written in a more generic way for more re-usability, but this is not the purpose here.

```
auto operator+(Image A, Image B) {
  return transform(A, B, std::plus<>());
}
auto togray = [](Image A) { return transform(A, [](auto x)
  { return (x.r + x.g + x.b) / 3.f; };)
};
auto subquantize16to8b = [](Image A) { return transform(A,
  [](float x) { return uint8_t(x / 256 +.5f); });
};

auto input = imread(...);
auto MyComplexImage = subquantize16to8b(togray(A));
```

Figure 4.17: Using high-order primitive views to create custom view operators.

Views for performance. With a classical design, each operation of the pipeline is implemented on "its own". Each operation requires memory to be allocated for the output image and also, each operation requires that the image is fully traversed. This design is simple, flexible, composable, but is not memory efficient nor computation efficient. With the lazy evaluation approach, the image is traversed only once (when the dilation is applied) that has two benefits. First, there are no intermediate images which is very memory effective. Second, traversing the image is faster thanks to a better memory cache usage. Indeed, in our example (fig. 4.9), processing a RGB16 pixel from the dilation algorithm directly converts it in grayscale, then sub-quantize it to 8-bits, and finally makes it available for the dilation algorithm. It acts *as if* we were writing an optimal operator that would combine all these operations. This approach is somewhat related to the kernel-fusing operations available in some HPC specifications [126] but views-fusion is optimized by the C++ compiler only [116].

Views for productivity. All point-wise image processing algorithms can (and should) be rewritten intuitively by using a one-liner view. The *transform* views is the key enabling that point. This implies that there exist a new abstraction level available to the practitioner when prototyping their algorithm. The time spent implementing features is reduced, thus the feedback-loop time is reduced too. This brings the practitioner to a productivity gain.

Chapter 5

Static dynamic bridge

In the programming world, there are three main families of programming language [28]. There are (i) *compiled* programming languages, such as C, C++, Rust or Go, (ii) *interpreted* programming languages, such as Python, PHP, Lisp or Javascript, and (iii) hybrid programming languages, such as Java or C#. The latter have a fast compilation pass that compiles the source code into an intermediate bytecode. Then, this bytecode is interpreted via an interpreter on the host (runner) machine.

5.1 Introducing the static and dynamic bridge

Many studies have been carried out to compare the advantages and disadvantages of each family of programming languages [4]. In this thesis we focus mainly on comparing the burden that are shouldered by the maintainers and the end-user.

5.1.1 Languages types

Rework arguments for compiled and interpreted languages.

Compiled languages   From the maintainer point of view, there is a lot of burden to shoulder. First is the binary generation which can be problematic in itself. Indeed, to generate a binary, there are many steps, as illustrated in fig. 5.1. First the compiler does a pass to generate machine code for each translation unit. There can be as many machine code as there are architecture and/or operating system supported. The maintainer may want to support different operating systems (last two windows and OSX version, a handful linux or unix distributions, maybe mobile phone portages). Each of these OS requires its own bundle. Additionally, the hardware may change, or the maintainer may want to take advantage of some specific hardware when available (like vectorized SIMD instructions such as SSE4, XOP, FMA4, AVX-512, etc.): this also requires the maintainer to multiply the number of binaries he compiles and distributes. Finally, the linker resolves the dependencies of the program and assembles the final binary. At that time the maintainer has to sort out how he wants to bundle the dependencies of his program. Should he statically link them alongside the binary and distribute them, at the risk of having the size of his binary exploding? Or should he state that the user has to install the dependencies on his system (via a package manager for instance) so that the binary can run? The burden of handling the dependencies is then shifted to the end-user when installing the program. Usually the package manager, such as *apt*, *yum* or *pacman*, solves this transparently for the end-user and the programs works "out-of-the-box" once installed, as illustrated in fig. 5.1. The downside is that the maintainer has to publish many bundles of his package where he wants it to be available.

(a) Compile-time   (b) Run-time
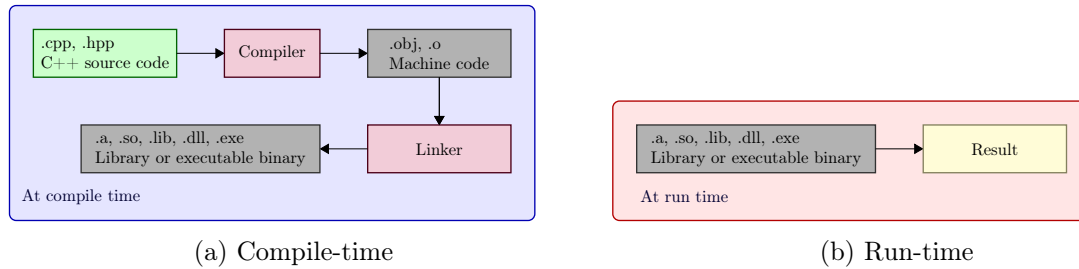
Figure 5.1: Compiled languages: compile-time (a) vs. run-time (b).

Interpreted languages From the maintainer point of view, this is the ideal standpoint. It is easier to distribute software via an interpreted language because only the source code and, a dependency tree and the assets are released in the bundle. All the burden about resolving the dependencies and installing the framework to run the script is shouldered by the end-user that is using the program. Indeed, as shown in fig. 5.2, everything happens at runtime. The main advantage of this approach is to build a very rich ecosystem as distributing, maintaining and using programs is very easy once integrated in a package manager (often delivered alongside the language SDK natively, e.g. pip for python). However, the most notable disadvantage is the performance which is explained by the fact that the source code is not compiled into optimized assembly code ready to be executed by the computer. Instead, the interpreter must do all the work in one go and very often this is slow (at least the first pass). Nowadays, interpreters differentiate two use cases. One is opening the console interpreter from the command line and typing commands to get the immediate interpreted results. This use-case usually does not provide heavy optimization because the user is likely prototyping his script and thus does not need it in the first place. The second use case is when the interpreter parses files and/or libraries as a whole. In this use-case, it is likely that the file will not change, but they will be used a lot. It is then relevant to pay a pass to generate intermediate bytecode that can be interpreted faster for the future passes onward. As an example, the Python programming language has several implementations: CPython, PyPy, Jython or IronPython. CPython generate intermediate bytecode in `*.pyc` files while Jython IronPython and PyPy embed a Just-In-Time (JIT) compiler to generate resp. JVM bytecode, CLR (.NET) bytecode, a large variety of bytecode format.
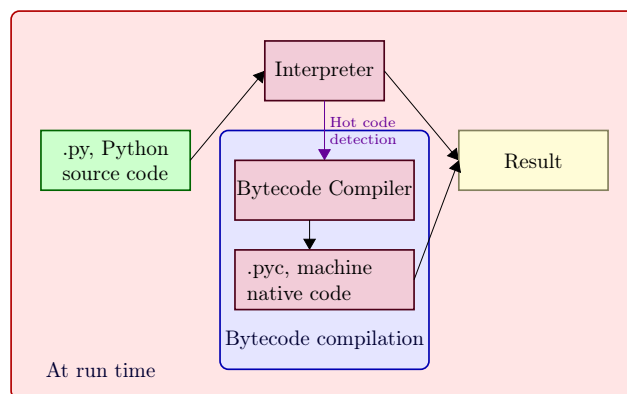
Figure 5.2: Interpreted languages: run-time

Hybrid languages The burden is shared evenly between the maintainer and the user, while remaining minimal. Indeed, languages such as Java or C# are in this category. Those programming languages need a compiling pass which is designed to be fast, so that the feedback loop while prototyping remains fast. The result of the compilation is bytecode which is then executed on an hosting Virtual Machine (VM) that the user must install on his computer. The main

advantages of this solution are the portability and the small distributed binary size. Indeed, in theory, any machine supporting the VM may also support the program. Also, the as the VM execute the bytecode and resolve system dependencies, the distributed binary does not need to embed any system dependency. Finally, the user has the advantage of running a compiled program which provides fast user experience. The goal of hybrid languages is to ally advantages of both compiled and interpreted languages; no dependency management for the user, one small binary to distribute for the maintainer, good execution performance, and fast feedback loop when prototyping (fast incremental compilation), while minimizing the downsides; usually a garbage collector is working inside the VM to handle memory allocations and de-allocations. In this regard, both Java and C# manage to achieve this feat quite elegantly. In theory, VM can further increase performances by implementing hot code detection which would compile the bytecode into native optimized machine code. This area is still a field of research to this day (c.f. Java HotSpot [155, 64, 104]).

To summarize Interpreted and hybrid programming languages produces more portable artefacts and therefore are easier to deploy in a dynamic environment. We summarize in fig. 5.3 the different approaches in order to be able to execute a binary from source code.
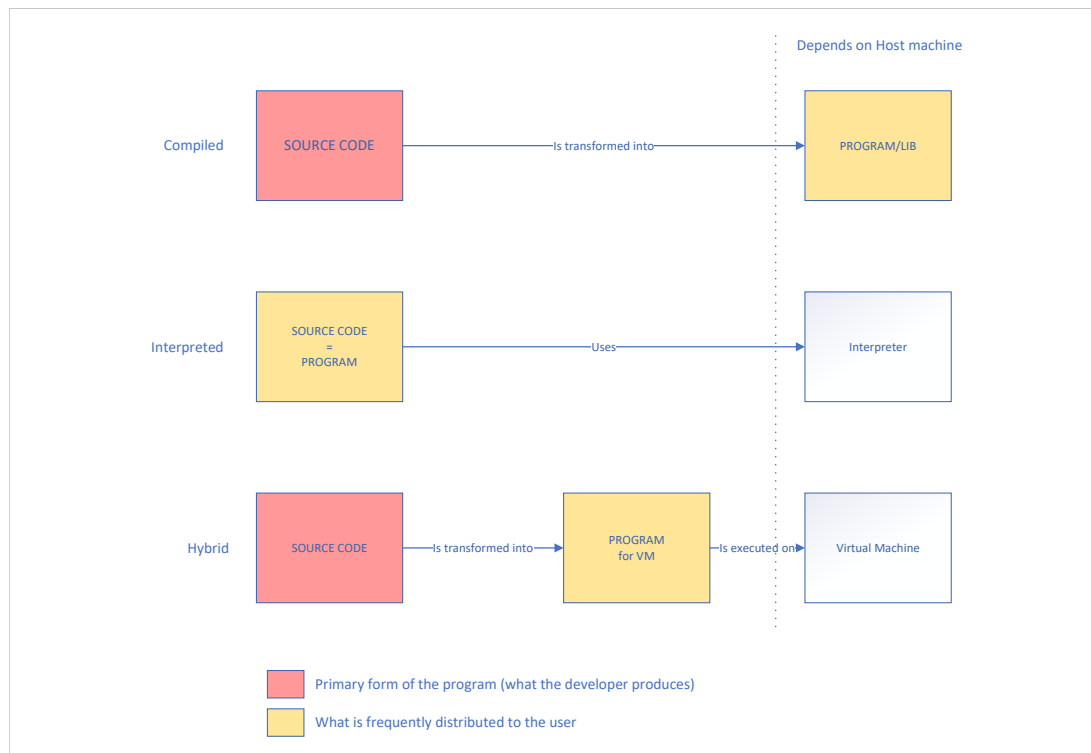
Figure 5.3: Languages types: summary diagram

5.1.2 Static and dynamic informations

Compiles programming languages usually have poor support of introspection facilities. At best, static reflection is available at compile-time, but dynamic reflection is not an option. The structure of the program does not change at runtime. Some flexibility exists when delving into the area of hot-swapping dynamic libraries at run-time, however this techniques drags alongside security-related issues that injecting possible foreign machine code into one's program may generate. Interpreted programming languages usually have very developed introspection facilities. Dynamic reflection at runtime is possible and some language, such as Common Lisp, even go further by allowing the developer to mutate the program Abstract Syntax Tree (AST) at runtime

(macro). This allows very powerful integrations such as defining one's own DSL as if it was part of the core language itself. Hybrid programming languages usually offers very good static and dynamic introspection facility at both compile-time and runtime, even if it means that runtime facilities will hurt performances. Also, those languages are usually designed to be able to hot-swap code at runtime. It is then possible to have the application running, recompile part of the binary of the application, replace the old running binary by the new compiled one, all the while the application is running.

The next important step is to classify what information is known at *compile-time* (machine/bytecode code generation): we call it *static* information; and what information is known at *runtime* (program execution): we call it *dynamic* information. In image processing, we have, on the one hand, knowledge about the following static informations:

- Image's value type (unit8, rgb8, complex, etc.),

- Image's dimension size (1D, 2D, 3D, etc.),

- Architecture of the hardware hosting the program (x86, ARM, PowerPC, GPU, etc.).

This means that, while the information may not be known at compile-time, we are able to write (or, more accurately, generate) code dedicated for those common types that we know constitutes a large portion of the use cases. Furthermore, we can write optimized portion of code dedicated to handle some particular known types that the program will use when it encounters them at run-time. On the other hand, the following informations are always dynamic:

- Image's actual values,

- Image's actual size,

- Architecture of the hardware hosting the program (x86, ARM, PowerPC, GPU, etc.).

The library needs both information type (static and dynamic) however, even if some informations are missing, it is not a fatality and the library can still recover and work efficiently at runtime.

On another note, we notice that the architecture hosting the program is both an information which is static and dynamic. This translates the complexity of this information. Indeed, the maintainer needs to guess the array of architectures he wants to support and generate binaries for them (static). Also, the program needs to detect at runtime (dynamically) on which hardware he is running to possibly take advantage of it to increase performances. This is an area of research on its own called heterogeneous computing [122, 123].

5.1.3   Introducing our hybrid solution

Image processing communities like to have bridges with interpretable language such as Python or Matlab, to interface with their favorite tools, algorithms and/or facilities. As an example, with Python, the module NumPy [57] is a community standard which is heavily used. Henceforth, to broaden the usage of our library, we should be able to provide a way to communicate between our library and NumPy. There is always a need for genecity in both C++ and Python. Indeed, in C++ genericity is achieved via template programming, which is static, whereas in Python genericity is achieved via duck typing, which is dynamic, as shown in fig. 5.4.

On the one hand, static polymorphism induces no indirection in the generated code as the type is known at compile-time. It is then possible to generate optimized code for specific types. It is not possible to add a new supported type at runtime as the code has already been compiled. On the other hand, the dynamic polymorphism implies that there will be indirections when executing the code. Indeed, the code first needs to dispatch onto the appropriate function handling the input types to perform the operation properly. Nevertheless, it is possible to add new supported type at runtime without recompiling the library binary.

```
template <typename T>
T add(T a, T b) {
  return a + b;
}
```

```
def add(a, b):
    return a + b
```

(a) C++ static genericity   (b) Python dynamic genericity

Figure 5.4: C++ Static (a) vs. Python Dynamic (b) genericity.

From the maintainer point of view, however, only distributing the C++ templated source code is a showstopper to the usability of his library by a Python user, because he does not hand over binaries. Indeed, one caveat of using C++ template programming is that the C++ compiler cannot generate a binary until it knows which type (of image, of value) will be used. But the maintainer does not know this information and the user (on Python's end) does not want to recompile the generic library code each time he has another set of types to exercise against. From here, there are still multiple ways to achieve our goal.

The first option is to embed and distribute, alongside the library, a JIT compiler whose job would be to generate the binaries and bindings just as they are used. This solution brings speed (excluding the first run that includes the compilation time) and unrestrained genericity. However it bounds both user and maintainer to the specificities of a compiler vendor and loose platform portability.

Another option is to type-erase (dynamic polymorphism) our types to enable the use of various concrete types through a single generic interface. This would translate into a class hierarchy whose concrete classes are the leaves (thus, whose value-types and dimensions are known). This induces a non-negligible performance overhead but enables us to keep the genericity and portability at the cost of maintaining the class hierarchy.

Type generalization can also be considered. It is possible to cast everything into a super-type that is suitable for the vast majority of cases. For instance, we could say that we have a super-type `image4D<double>` into which we can easily promote sub-types such as `image2D<int>` or `image3D<float>`. Of course we would loose the generic aspect and induce non negligible speed cost. Although portability is kept.

And finally there is the dynamic dispatch. It consists in embedding dynamic information at runtime about types, and dispatch (think of switch/case) to the correct facility which can handle those types. The obvious caveat is the cost of maintenance induced by the genericity as we would have a number of possible dispatches that grow in a multiplicative way with the number of handled types. Which is not very generic. On the other hand there is almost no speed loss and the portability is guaranteed. Theoretical models exist that could bring solutions to lower the number of dispatcher to write, such as multi-method [70]. Unfortunately they are currently not part of C++.

5.2   Designing the hybrid solution

In Pylene we have chosen a hybrid solution between type-erasure and dynamic dispatch. The aim is to have a set of known types for which we have no speed cost as well as continuing to handle other types to remain generic. In [125] we provide a facility to expose our generic code to Python. As seen in the previous chapter, it is not possible to bind C++ source code to Python. We need to have a compiled binary implementing Python binding (we chose Pybind11 [114]) in order to be able to call C++ code from Python. In order to achieve the binding without sacrificing the genericity and the performances, we have designed a solution in two steps. We do not want to provide an abstract interface that will resolve the calls to access data on the call-site via virtual call because it would be very slow when the C++ code is executed. This would defeat the purpose of having to rely on C++ in a first place. However, it is possible to

convert an abstract class into an instantiated concrete generic class whose template parameter are known. This requires, however, to enumerate all the possible cases. With modern C++, it has become possible to design $n * n$ dispatch without gigantic switch-case clauses.

5.2.1   First step: converting back and forth

The first step of our solution consists in designing a buffer class that holds all the informations about an image: dimension, underlying type, strides and pointer to data buffer. This class is named `ndimage_buffer`. When interfacing with Python, it is necessary to convert the Python image which is a `NumPy.array` into our image type and vise versa, converting our C++ image type back into a python image. The purpose of this buffer image is to holds all the information from the `NumPy.array` to then instantiate a concrete C++ type. This process is illustrated in fig. 5.5. The first pitfall here is due to a limitation from the abstraction interface used in Python. Indeed, when using, for instance *Scikit-Image*, it is not possible to differentiate a 2D multichannel image from a 3D grayscale image. Indeed, the image is always broken down to its most simple value and a 2D multichannel image is turned into a 3-dimensional `NumPy.array` containing a single 8-bits channel, the last dimension contains only 3 elements at max but can theoretically contain more. Indeed, there is no limitation embed in the used abstraction which does not prevent that. To prevent this confusion, the C++ wrapper code may chose between two strategies; first is to consider all 3D/1-channel image as 2D/RGB images by default, second is to let the user give the information. For the sake of simplicity, we have chosen the first strategy.

Figure 5.5: Bridge from Python to C++ via Pybind11 and a type-erased C++ class.

From the point of view of a practitioner, the code on the call-site (python side) should be as follow:

```python
from skimage import data
import numpy as np
import Pylena as pln # our python binding
img = data.astronaut() # 2D-rgb8 image -> Numpy.array(ndim=3, dtype='uint8')
# pln.<any_algorithm>(img)
```

The C++ code contains lots of glue code necessary to expose the module to Python. In this thesis we have chosen to work with Pybin11 [114] which provides a modern API and is actively being maintained and improved. The glue code exposing the Python module from C++ is given in appendix C.4. In order to have a seamless interaction between Python's `numpy.array` and C++, we need to define a proper strategy to convert the Python type into the C++ type without copying all data around. Pybind offers us two possibilities to achieve this. The first one is to use the buffer protocol to pass around numpy's buffer information to C++ in way so that C++ can properly bind the data into a proper C++ class. The second one is to use a custom type-caster to implicitly convert the python type into a C++ type each time it is needed.

With the first method, one would need to write the following code on python's side:

```python
np_img = data.astronaut() # 2D-rgb8 image -> Numpy.array(ndim=3, dtype='uint8')
pln_img = pln.ndimage(np_img) # conversion into the C++ image type
pln_img_ret = pln.<any_algorithm>(pln_img) # call to algorithm
np_img_ret = pln_img.to_numpy(); # convert back into numpy.array
# use np.<...>(np_img_ret) # use resulting image with numpy
```

Whereas the second method would require the user to write the following code on python's side:

```python
np_img = data.astronaut() # 2D-rgb8 image -> Numpy.array(ndim=3, dtype='uint8')
np_img_ret = pln.<any_algorithm>(np_img) # implicit conversion with custom type-caster on C++ side
# use np.<...>(np_img_ret) # use resulting image with numpy
```

Removing this conversion step is the major reason we have chosen the second method: the custom type-caster. The C++ code for this part is given in appendix C.1.

We are now all set and are able to convert back-and-forth a python image into a C++ image and vice versa.

5.2.2 Second step: multi-dispatcher (a.k.a. $n * n$ dispatch)

The second step of our hybrid solution is to dispatch the abstract buffer type coming from Python to an efficient generic code. The naive way of doing so would be to include a gigantic switch-case clause in each algorithm implementation and dispatch to the correct instantiated generic algorithm from there. Aside from being a nightmare to maintain, the size of those clauses would grow several fold depending on the cardinality of the generic implementation. For instance, for a generic dilation, there are 3 axis of cardinality: the underlying type, the dimension and the structuring element shape. In the case where the library support 5 different structuring element shape, 10 underlying types and 6 dimension for the image, the switch-case statement would need to dispatch over 300 clauses. Also, each supported algorithm would need to have dispatcher. This solution defeat the purpose of genericity which is to write less code in the first place. We needed to design a solution to implement those dispatchers while keeping our code short and efficient. The idea we took to solve this problem comes from the design of a C++ feature, the variant, and especially the visitor, applied to image processing, as in [72] for instance. We need to have a way to write the implementation of the algorithm once while enumerating all the possible cases. Also, if possible, the list of supported types should be written once at one place for maintenance purpose.

Simple dispatcher We then had the idea of writing a dispatcher. This dispatcher lists all the supported types and call the given callbacks forwarding the given arguments by instantiating a specific type. Let us try to expose to python, for instance, the generic existing algorithm for thresholding a binary image. The Python call-site code will look like this:

```python
img_grayscale = skimage.data.grass()
pln.operators.binary_threshold(img_grayscale)
```

On the C++ side, we want to avoid writing code that looks like this:

```cpp
mln::ndbuffer_image binary_threshold(mln::ndbuffer_image input) {
  auto dim = input.dim();
  auto tid = input.tid();
  switch(dim) {
    case 1: // 1D image
      switch(tid) {
        case UINT8:
          if(auto* image_ptr = input.template cast_to<uint8_t, 1>(); image_ptr)
            return mln::binary_threshold(*image_ptr);
        case RGB8:
          // error support only RGB8 images
      }
      break;
    case 2: // 2D image
      // ...
      break;
    // ... 3D, 4D, ...
  }
}
```

Instead, it is possible to separate the dispatching code and the logical code entirely by using a templated operator, the same way we use lambdas in the pattern `std::variant`/`std::visit`. For our binary threshold example, the operator be implemented just by writing the following code:

```
// Operator templated by the dimension
template <auto Dim>
struct binary_threshold_op_t {
  // Function templated by the image type
  template <typename Img>
  mln::ndbuffer_image operator()(Img&& img) const {
    // Cast to a grayscale (information known) of the correct dimension
    if (auto* image_ptr = std::forward<Img>(img).template cast_to<std::uint8_t, Dim>(); image_ptr)
      // ACTUAL call to the generic algorithm
      return mln::binary_threshold(*image_ptr);
    else {
      std::runtime_error("Unable to convert the image to the required type.");
      return {};
    }
  }
};
```

This code allows us to dispatch over any number of dimensions. We are required to pass a grayscale image for the algorithm so here the example is limited to dispatching over just one cardinality: the dimension. Let us now take a look at how we can implement the dispatcher for our example to work. The dispatcher must take the dimension as first parameter and any number of arguments to forward to the instantiated operator. The dispatcher then looks like the following code:

```
template <template <auto> class Op, typename... Args>
auto dispatch_v(std::size_t dim, Args&&... args) {
  switch (dim) {
    case (1):
      return Op<1>{}(std::forward<Args>(args)...);
    case (2):
      return Op<2>{}(std::forward<Args>(args)...);
    case (3):
      return Op<3>{}(std::forward<Args>(args)...);
    /* ... */
  }
}
```

The operator `Op` is instantiate with the correct dimension number and the the `operator()` (parenthesis) is called while being forwarded the correct arguments. In our case, it will instantiate the type `binary_threshold_op_t<2>` and then call the function `binary_threshold_op_t<2>.operator()(input)`, forwarding the input image to the underlying algorithm. Indeed, using the dispatcher is as simple as writing `dispatch_v<binary_threshold_op_t>(input.dim(), input)`.

The main advantage of this approach is that we respect all the requirements. First the logical code is bounded in the operator, second, the supported types are all listed in one place, once. Also, while our example is limited to one cardinality, any number of dispatcher can be piped to one after another to achieve the cardinality wanted.

Double dispatcher   Let us push our example to implement the mathematical morphology operator dilation. We now have two more generic axis to cover: the structuring element shape and the underlying datatype. First, let us take a look at what the Python code may look like:

```
img_grayscale = skimage.data.grass()
rect = pln.se.rect2d(width=3, height=3)
img_dil = pln.operators.dilate(img_grayscale, se)
```

The first thing to notice is the need to add additional bindings to expose our C++ structuring elements. The glue code to achieve this is given in appendix C.2. Let us take a look at our dilation operator:

```
template <auto Dim, typename T>
struct dilate_operator_t {
  template <typename Img, typename SE>
  mln::ndbuffer_image operator()(Img&& img, SE se) const {
    if (auto* image_ptr = std::forward<Img>(img).template cast_to<T, Dim>(); image_ptr)
      // ACTUAL call to the generic algorithm
```

```
      return mln::dilation(*image_ptr, se);
    else {
      std::runtime_error("Unable to convert the image to the required type.");
      return {};
    }
  }
};
```

This operator needs double dispatch over two cardinalities: the dimension `Dim` and the value type `T`. We can skip the dispatch of the structuring element's shape as we have made a `std::variant` of all the supported structuring element supported for the sake of simplicity. Dispatching over the supported structuring elements can then be offloaded upstream from the call to the double dispatch, just by calling `std::visit`. We can immediately notice that there is an issue with our dilation operator. Indeed, there are two template parameters and our dispatcher `dispatch_v` does only handle one. We solve this issue by writing another intermediate operator dispatcher `dilate_operator_intermediate_t` serving as trampoline operator that will partially instantiate the final operator `dilate_operator_t` along the dimension template parameter to feed it to the last dispatcher, `dispatch_t`:

```
template <auto Dim>
struct dilate_operator_intermediate_t {
  template <typename Img, typename SE>
  mln::ndbuffer_image operator()(Img&& img, SE&& se) const {
    // Partial instantiation
    return double_dispatch_t<dilate_operator_t, Dim>(
            input.sample_type(), std::forward<Img>(input), std::forward<SE>(se));
  }
};
```

Dispatching the operator alongside two cardinalities (even three including the structuring element handled by `std::variant`) would then become as simple as calling:

```
// dispatch the structuring elements through using std::visit for std::variant
return std::visit(
[&input](const auto& se_) { // dispatch over the trampoline
    return dispatch_v<dilate_operator_intermediate_t>(input.dim(), input, se_);
  }, se);
```

In order for this to work, we need to piece together the final part of our puzzle, which is the double dispatch function that will handle the last dispatch along the underlying data while forwarding the first dispatch along the dimension. This dispatcher works the same as the simple one (`dispatch_v`) but take an additional template parameter (here `Dim`) that will be forwarded as-is to the given operator `Op`. The implementation then looks like this:

```
template <template <auto, typename> class Op, auto Dim, typename... Args>
auto double_dispatch_t(type_id tid, Args&&... args) {
  switch (tid) {
    case (type_id::INT8):
      return Op<Dim, std::int8_t>{}(std::forward<Args>(args)...);
    case (type_id::UINT8):
      return Op<Dim, std::uint8_t>{}(std::forward<Args>(args)...);
    case (type_id::DOUBLE):
      return Op<Dim, double>{}(std::forward<Args>(args)...);
      /* ... */
  }
}
```

We have now presented all the techniques and design required to build operators that are agnostic from the supported data-types, dimensions and/or additional data such as structuring elements. Indeed, the maintainer has gathered all the logic about listing the supported data types and dimension in one place: the custom dispatcher. He just needs to maintain those to enable full support for all exposed algorithm, by default. This hybrid solution mixes type-erasure and modern C++ facilities to allow maximum performance. Indeed, the dispatch is done before entering algorithms and the custom type-caster facility allows us to plug directly into the

Python image without having any unnecessary copies. The only caveat would be the code bloat incurred by all the explicit instanciation leading to compiling a larger and larger binary the more algorithms are being exposed. This can lead to performances issues due to pre-fetching memory optimization missed and code locality issues [32]. Another point not covered right now would be a way to support arbitrary data types, possibly injected from Python, into C++. Indeed, our hybrid solution only support the types provided by the library and listed in the dispatchers. It will instantiate all the code relative to them and support all of the combinations but the user may be tempted to plug a user-defined type from Python as an underlying image data-type. To allow this use-case, we introduce a new concept: the *value-set*. The value-set is a standard way manipulate the underlying values. Through type-erasure, we can either manipulate known underlying value type with native facilities (near-zero overhead), or fallback to a virtual call that may report an error, or callback user-provided Python routine to manipulate unknown user values.

5.2.3   Third and final step: type-erasure & the value-set

rework motivation of this part: what do we want to write and why.

As common thread in this section, we will work on the *stretch* algorithm which is naively defined in fig. 5.6.

```cpp
template <class T>
mln::image2d<float> naive_stretch(const mln::image2d<T>& src) {
  mln::image2d<float> res  = mln::transform(src, [](auto val) -> float {
    auto max = std::numeric_limits<T>::max();
    return static_cast<float>(val) / static_cast<float>(max);
  });
  return res;
}
```

Figure 5.6: Stretch algorithm, naive C++ version.

Figure 5.7: Python to C++ pipeline algorithm through the $n * n$ dispatcher.

We can represent the pipeline calling this algorithm from Python in fig. 5.7.

Introducing the value-set   The *value-set* is an abstraction layer around common operations needed when implementing an image processing algorithm such as an addition, a multiplication, a type conversion, getting the maximum etc. It can be defined in C++ as a class template whose parameter is the manipulated type. The following code shows how to define a value-set:

```cpp
template <class T = void>
struct value_set {
  template <class U>
  U cast(T v) const { return static_cast<U>(v); }

  T max() const noexcept { return std::numeric_limits<T>::max(); }
  T min() const noexcept { return std::numeric_limits<T>::min(); }
  /* inf, sup, ... */

  T add(T l, T r) const noexcept { return l + r; }
```

```cpp
  T sub(T l, T r) const noexcept { return l - r; }
  /* mod, pow, min, max, ... */
};
```

We can see that the default parameter of the class template is `void`. Indeed, we are inspired by what was implemented in the standard library for `std:less` and providing a default (void) specialization in order to improve the usability. The following code shows how to implement this specialization:

```cpp
template <>
struct value_set<void> {
  template <class U, class T>
  U cast(T&& t) const { return static_cast<U>(std::forward<T>); }

  template <class T, class U>
  auto add(T&& l, U&& r) const noexcept { return std::forward<T>(l) + std::forward<U>(r); }
  template <class T, class U>
  auto sub(T&& l, U&& r) const noexcept { return std::forward<T>(l) - std::forward<U>(r); }
  /* ... */
};
```

The full code of the value-set is given in appendix C.3.2. The template parameter is shifted from the class to the member functions. It is also important to note that the member function are not static, which requires to instantiate the `value-set` before using it. It may sound like a disadvantage at first glance but it can be turned into an advantage later on. Indeed, this design allows a subclass to hold member variables which will be crucial for injecting user-types from python.

Now that we have designed how our value-set is intended to work, we can deduce that an image is able to provide its own value-set. Indeed, an image knows what values it holds and thus is able to instantiate the proper value-set corresponding to this type. The member function returning the value-set in the class template `ndimage<T, D>` is then implemented as follow:

```cpp
template <class T, std::size_t D>
class ndimage {
  /* ... */
  auto get_value_set() const noexcept {
    return value_set<T>{};
  }
};
```
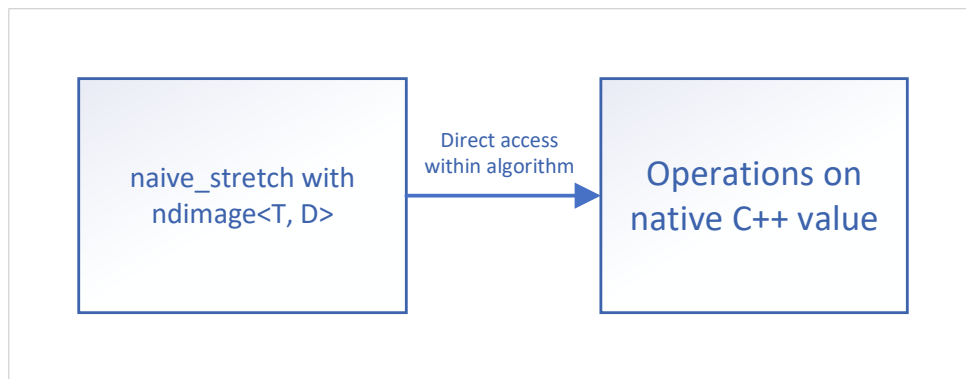

Figure 5.8: Naive stretch algorithm, pipeline to perform operations on values.

Let us recall our naive stretch algorithm presented earlier. The pipeline representing the operations on values inside the algorithm is presented in fig. 5.8. Typically, this algorithm performs three operations that are the responsibility of a value-set: getting the max, performing a cast, and performing a division. The first step is then to use the value set shown above to abstract away those operations. The new algorithm is shown in fig. 5.9.

```cpp
 1  template <class T>
 2  mln::image2d<float> fast_stretch(const mln::image2d<T>& src) {
 3    auto              vs  = src.get_value_set();      // value-set for T
 4    auto              vs_f = mln::value_set<float>{}; // fast value-set for float
 5    mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
 6      auto max  = vs.max();                          // returns T
 7      auto fval = vs.template cast<float>(val); // returns float
 8      auto fmax = vs.template cast<float>(max); // returns float
 9      return vs_f.div(fval, fmax);              // div directly returns float
10    });
11    return res;
12  }
```

Figure 5.9: Stretch algorithm, fast C++ version.
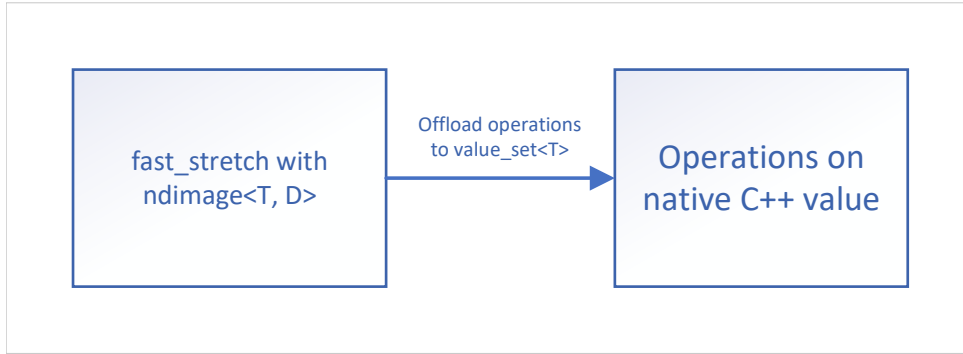
Figure 5.10: Fast stretch algorithm, pipeline to perform operations on values.

We instantiates the value-set we will need (for `T` values and `float` values) on lines 3 and 4. Then the maximum value is obtained via the value-set of `T` on line 6. Then a cast is performed via the value-set to get floating point value on lines 7 and 8. Finally a floating point division is performed on line 9 via the formerly instantiates floating point value-set defined on line 4. Indeed, using the value-set instantiated for `T` would have performed an Euclidean division and which would have resulted in a bug. Calling this algorithm from python is as simple as `ret = pln.fast_stretch(input)` once exposed. The pipeline representing the actual operations in the algorithm to access values is shown in fig. 5.10. The glue code to expose this algorithm is given in appendix C.3.3.

Type-erasure interface for value-set Moving one step further, we ultimately want to be able to dynamically inject a value-set into our algorithm. In order to achieve this feat, we need to design an abstract value-set as an interface so that a user can sub-class it and provide his own value set. This abstract value-set needs to work with `std::any` as input value-type in order to provide a generic interface. Here, the genericity is achieved via type-erasure (the input value is type-erased into a `std::any`). The interface would look like (full code is given in appendix C.3.2):

```cpp
struct abstract_value_set {
  virtual ~abstract_value_set() {}

  virtual std::any max() const = 0;
  /* min, sup, inf, ... */

  virtual std::any add(const std::any& l, const std::any& r) const  = 0;
  virtual std::any div(const std::any& l, const std::any& r) const  = 0;
  /* sub, mult, ... */
};
```

The important part here is to notice that all values (returned and passed as argument) are now type-erased behind a `std::any`. From this interface, it is trivial to define the canonical sub-classes for trivial types. We do this by defining the class template `concrete_value_set` which

is able to generate a concrete interface for every given template type. The implementation will look like this:

```cpp
template <typename T>
struct concrete_value_set : abstract_value_set {
  ~concrete_value_set() override {}

  template <class U>
  std::any cast(std::any v) const { return {static_cast<U>(std::any_cast<T>(v))}; }

  std::any max() const override { return {std::numeric_limits<T>::max()}; };
  /* min, sup, inf, ... */

  std::any add(const std::any& l, const std::any& r) const override {
    return {std::any_cast<T>(l) + std::any_cast<T>(r)};
  }
  std::any div(const std::any& l, const std::any& r) const override {
    return {std::any_cast<T>(l) / std::any_cast<T>(r)};
  }
  /* sub, mult, ... */
};
```

This concrete value-set implement simple dispatch via casting the type-erased `std::any` into the wanted value type to properly perform the operation. Thanks to this implementation, we are able to rewrite our stretch algorithm using this value-set to perform its operations. The new algorithm is shown in fig. 5.11.

```cpp
1   template <class T>
2   mln::image2d<float> virtual_dispatch_stretch(const mln::image2d<T>& src) {
3     auto               vs   = mln::concrete_value_set<T>{};      // value-set for T
4     auto               vs_f = mln::concrete_value_set<float>{}; // value-set for float
5     mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
6       auto anymax  = vs.max();                         // returns std::any
7       auto fanyval = vs.template cast<float>(val);      // cast to float in std::any
8       auto fanymax = vs.template cast<float>(anymax);   // cast to float in std::any
9       return std::any_cast<float>(vs_f.div(fanyval, fanymax)); // div returns float
10    });
11    return res;
12  }
```

Figure 5.11: Stretch algorithm, virtual dispatch version.

Figure 5.12: Virtual dispatch stretch algorithm, pipeline to perform operations on values.

This algorithm uses two value sets (l.3 and l.4), one to get the maximum value with regard of the original underlying value type in the image `T` (l.6) and to properly cast the values into float values (l.7-8). The other one is used to perform the floating point division (l.9). This algorithm is almost identical to the one given for `fast_stretch` earlier. Only the lines instantiating the value-set changes (l. 3-4) as well as the line handing over the result (l. 9) where we added a downcast from `std::any`. The pipeline presenting the operations needed to access values is presented in fig. 5.12.

Going forward: type-erasing the value entirely   The need may arise where the user wants to handle images where the input value is abstracted away behind a type-erased type. This

happens when handling image with heterogeneous value-type (different channel number from one value to another, optional information embed with the value, etc.). In order to work, this value-type needs to be aware of its own value-set to perform operations. To achieve this, the value must embed the value set (as a pointer for instance) directly alongside the value. This embedding would allow writing such code possible:

```cpp
class type_erased_value {
  type_erased_value add(const type_erased_value& rhs);
  /* ... */
};

void my_algo(ndimage<type_erased_value, Dim> img1,
             ndimage<type_erased_value, Dim> img2) {
  ndimage<type_erased_value, Dim> img_out = /* ... */;
  for (auto [v1, v2, out] : zip(img1.values(), img2.values(), img_out.values())) {
    out = v1.add(v2); // the value knows how to perform the addition
  }
}
```

This code enables fallback when all the supported values have been exhausted. For instance, when considering the previously defined `abstract_value_set`, when attempting to dispatch over all the supporting native C++ type unwrapped from the given `std::any`, if no type is matched, we can attempt one final unwrap to this type-erased value which is aware itself of how to perform the operation. In pseudo-code it breaks down to the following logic:

```cpp
1   std::any value_set<type_erased_value>::add(const std::any& lhs, const std::any& rhs){
2     abstract_value_set* vs = lhs.get_embedded_vs();
3     if (lhs.type() == typeid(int) && rhs.type() == typeid(int)) {
4       auto ret = std::any_cast<int>(lhs) + std::any_cast<int>; // unwrap, do the work
5       return std::any{ret}; // rewrap with the vs
6     } else if (lhs.type() == typeid(float) && rhs.type() == typeid(float)) {
7       /* ... same ... */
8     } else {
9       auto te_lhs = std::any_cast<type_erased_value>(lhs); // last attempt
10      return lhs.add(rhs); // fallback on embedded value-set
11    }
12  }
```

First step is to conditionally attempt to cast the type-erased value over the supported native values (lines 3 and 6). When the type is supported then we unwrap it and perform the operation before rewrapping it in a `std::any` and returning it. If we could not unwrap the value into a supported value-type, we make a last attempt on line 9 into our type-erased value-type. If this attempt succeeds, we rely on the fact that this value-type is aware of its own value-set and embed it to perform the required operation on line 10. The full code of the implementation for this value-set aware value-type is given in appendix C.3.2. We have relied on metaprogramming techniques in order to efficiently write the code that will do the work.

Thanks to this technique, we are able to write another version of our stretch algorithm which is shown in fig. 5.13.

This version is verbose because we do not yet support `mln::image2d<type_erased_value>` as a proper image type. Thus we need to first convert the value of type `T` into the value-set aware type-erased value type. This work is done on lines 8 to 12. Then on lines 13 to 20 we perform the actual work on the values. This is where the dispatch mechanic, whose pipeline is shown in detail in fig. 5.14, will happen. Finally, on lines 22 and 23, we unwrap the previously wrapped type-erased value into a float to return for the algorithm.

Everything comes to a full circle: injecting a value-set from Python   We have seen how to write an algorithm independently from ths its underlying types on C++ side thanks to relying on an abstraction layer: the value-set. This enables one fundamental feature which is code injection from Python. Indeed, if we suppose every python value is a `pybind11::object`,

```
 1  template <class T>
 2  mln::image2d<float> stretch_virtual_dispatch_type_erased_value(const mln::image2d<T>& src) {
 3    auto                  vs   = mln::concrete_value_set<T>{};      // value-set for T
 4    auto                  vs_f = mln::concrete_value_set<float>{}; // value-set for float
 5    mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
 6      // simulate having an image<type_erased_value>
 7      auto anyval = std::any{val}; // std::any of T
 8                                   // type_erased_value of std::any of T aware of value-set of T
 9      auto abs_anyval = mln::type_erased_value{anyval, vs};
10      // instantiate a value-set for type_erased_value
11      auto abs_vs = mln::value_set<mln::type_erased_value>{abs_anyval};
12      auto anyabs_anymax =
13          abs_vs.max(); // returns std::any of type_erased_value
14                        // cast underlying std::any of type_erased_value of std::any of T into
15                        // std::any of type_erased_value of std::any of float aware of value-set for float
16      auto anyabs_fanyval = abs_vs.template cast<T, float>(std::any{abs_anyval}, &vs_f);
17      auto anyabs_fanymax = abs_vs.template cast<T, float>(anyabs_anymax, &vs_f);
18      // dispatch on known type, find a type_erased_value, then call
19      // anyabs_fanyval.div(anyabs_fanymax) to perform division which will call
20      // the underlying value-set for float for this operation
21      auto anyabs_fanyret = abs_vs.div(anyabs_fanyval, anyabs_fanymax);
22      // convert result back into float for returning to the image
23      auto anyfret = std::any_cast<mln::type_erased_value>(anyabs_fanyret).val();
24      return std::any_cast<float>(anyfret);
25    });
26    return res;
27  }
```

Figure 5.13: Stretch algorithm, virtual dispatch with a type-erased value version.
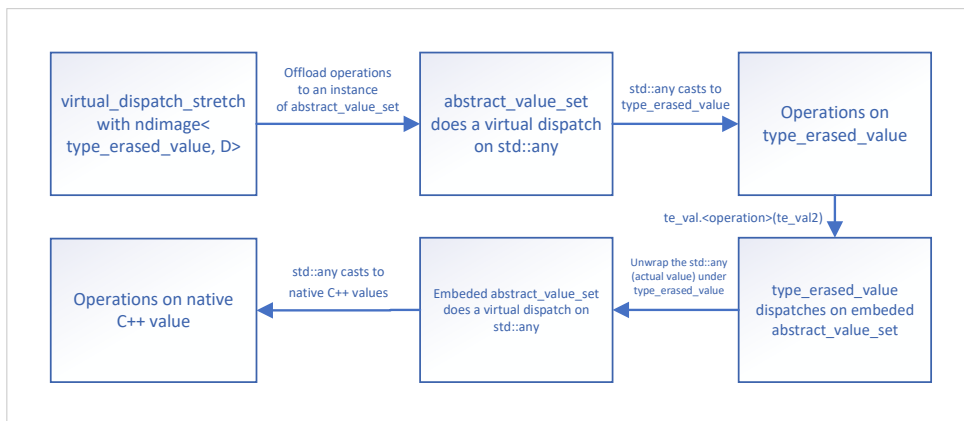
Figure 5.14: Virtual dispatch stretch algorithm with a type-erased value, pipeline to perform operations on values.

which is the generic way to refer to a non-trivially-convertible Python type we can write a value-set handling this value type. The value-set handling python data-type would then provide the following interface:

```cpp
template <>
struct value_set<pybind11::object> : abstract_value_set
{
  value_set(pybind11::object python_vs_instance): vs_instance_(python_vs_instance) {}
  ~value_set() override {}

  template <typename U>
  std::any cast(const std::any& v) const { /* ... */ }

  std::any max() const override { {vs_instance_.attr("max")()}; }
  /* min, sup, inf, ... */

  std::any add(const std::any& l, const std::any& r) const override {
    auto pyl = std::any_cast<pybind11::object>(l);
    auto pyr = std::any_cast<pybind11::object>(r);
    return {vs_instance_.attr("add")(pyl, pyr)};
  }
  std::any div(const std::any& l, const std::any& r) const override {
    auto pyl = std::any_cast<pybind11::object>(l);
    auto pyr = std::any_cast<pybind11::object>(r);
    return {vs_instance_.attr("div")(pyl, pyr)};
  }
  /* max, min, sub, mult, ... */
private:
  pybind11::object vs_instance_;
};
```

In this code we can clearly see at lines 4 and 27 that we are storing our Python's value-set instance into our class. This is possible due to the fact that our value-set abstraction is not providing static class function but member function. Hence, it is possible to offload the work of the value-set to a member variable at lines 10, 15-17 and 21-23 that will call the Python's value-set and get the wanted result. Also, at line 8 we use multiples techniques at once to get the correct resulting cast from a Python type. Indeed, the cast is done on Python side before being casted back into the corresponding C++ type. As such, we have to translate the wanted C++ type information into Python type information to request the cast. Once the cast is done, we need to unwrap the python type into the corresponding C++ type and rewrap it into our, now favorite, `std::any`. The full implementation of this facility is given in appendix C.3.2.

On this particular matter, the user will find a Python abstract class to implement in order for his value-set to be usable by the library. This abstract class is defined by the following Python code:

```python
class AbstractValueSet(ABC):
  @abstractmethod
  def cast(self, value: Any, type_): pass
    if type_ in ["int", "float", "bool", "str"]:
      module = importlib.import_module('builtins')
      cls = getattr(module, type_)
      return cls(value)
    else:
      raise ValueError()
  @abstractmethod
  def max(self): return math.inf

  # ... min, sup, inf, ...

  @abstractmethod
  def add(self, lhs: Any, rhs: Any) -> Any: return lhs + rhs
  @abstractmethod
  def div(self, lhs: Any, rhs: Any) -> Any: return lhs / rhs

  # ... sub, mult, ...
```

This abstract class provide a facility to cast a value into a given type from its representation as a string. It also provides default/standard way of computing values. Those methods needs to be overridden by a child class as they are all tagged with the `@abstractmethod` attribute. The full code of this data structure is given in appendix C.5.

Let us assume our user wants to build an image whose value type are his own custom Python data structure. For the sake of this example, we will name this specific value type will be named `MyStruct`. The Python structure will look like this:

```python
class MyStruct:
  v_: Any
  def __init__(self, v: Any): self.v_ = v
  def getV(self) -> Any:      return self.v_
  def setV(self, v: Any):     self.v_ = v
```

It is just a strong wrapper around a value with setters and getters. If we want to use this data structure as an image value-type and downstream it to our C++ library, we need to provide a value-set which is aware of how to handle this value type properly. If the user wants to write such a code:

```python
img = np.array([MyStruct(1), MyStruct(2), MyStruct(6.5), MyStruct(3.14)] * 10, ndmin=2)
pln_img = pln.stretch(img) # Error, C++ does not know how to handle a value of type MyStruct
```

Then the user must provide his own value-set defined after the Python `AbstractValueSet` seen earlier. Such a value-set would look like this:

```python
class MyValueSet(AbstractValueSet):
  def get_MyStruct__(self, v: Any):
    return v.getV() if isinstance(v, MyStruct) else v
  def cast(self, value: Any, type_):
    v = self.get_MyStruct__(value)
    return super().cast(v, type_)
  def max(self):
    return 255

  # ... min, sup, inf, ...

  def add(self, lhs: Any, rhs: Any) -> Any:
    l = self.get_MyStruct__(lhs)
    r = self.get_MyStruct__(rhs)
    return MyStruct(super().add(l, r))
  def div(self, lhs: Any, rhs: Any) -> Any:
    l = self.get_MyStruct__(lhs)
    r = self.get_MyStruct__(rhs)
    return MyStruct(super().div(l, r))

  # ... sub, mult, ...
```

The full code of this implementation is given in appendix C.5. With this information, we can now write the following Python code:

```python
img = np.array([MyStruct(1), MyStruct(2), MyStruct(6.5), MyStruct(3.14)] * 10, ndmin=2)
pln_img = pln.stretch(img, value_set=MyValueSet()) # This works !
```

And this will work finely.

On the C++ side, the maintainer just needs to write another version of the stretch algorithm supporting the Python value-set which is shown in fig. 5.15.

The `mln::value_set<pybind11::object>` is constructed from the passed `value_set` argument on python side as it is a `pybind11::object`. The dispatcher then forward it to the algorithm as an additional parameter. We then have actual Python call on lines 4, 6 and 7. The line 5 just wrap the image value into a python object so that it can then be forwarded to the python code downstream, as described in the pipeline shown in fig. 5.16. The line 7 cast the python type into an actual C++ type (wrapped into our `std::any` wrapper). Finally, the line 8 unwrap the float for returning it to the algorithm.

```
1   template <class T>
2   mln::image2d<float> slow_stretch(const mln::image2d<T>& src, const mln::value_set<pybind11::object>& py_vs) {
3     mln::image2d<float> res = mln::transform(src, [&py_vs](auto val) -> float {
4       auto anymax  = py_vs.max();                      // returns std::any of pybind11:object
5       auto anyval  = std::any{pybind11::cast(val)};    // converts to std::any of pybind11::object
6       auto anyret  = py_vs.div(anyval, anymax);        // returns std::any of pybind11::object
7       auto anyfret = py_vs.template cast<float>(anyret); // returns std::any of float
8       return std::any_cast<float>(anyfret);            // returns float
9     });
10    return res;
11  }
```

Figure 5.15: Stretch algorithm, injected value-set from Python version.

Figure 5.16: Stretch algorithm with an injected value-set from Python code, pipeline to perform operations on values.

5.3   Summary and continuation

5.3.1   Performances & overhead

In this chapter, we have designed many solutions to solve several kind of issues. However, layering one abstraction layer after another, or even calling Python code does come with performance cost. This is why we have run a benchmark to outline the cost of our solutions. The full code of the benchmark is given in appendix C.6. This benchmark compares the four version of our stretch algorithm. The result is shown in table 5.1.

Dispatch type	Compute Time	ΔCompute Time
Native value-set with native C++ value-type (baseline) 5.9	$0.0093s$	0
Value-set with virtual dispatch with native C++ value-type 5.11	$0.1213s$	$\times 13$
Value-set with virtual dispatch with C++ type-erased values 5.13	$1.0738s$	$\times 115$
Injected Python value-set with native C++ value-type 5.15	$21.5444s$	$\times 2316$

Table 5.1: Benchmarks of all our version of the stretch algorithm.

This benchmark shows that each time an abstraction layer is added on top of the baseline, the user must expect a $10\times$ slowness factor in his code performance. Also, calling Python code is immensely slower ($2300\times$!) than the baseline. This renew the interest there is to recompile the templated C++ library with an additional known type than injecting it from Python long time running code. Being able to inject Python code ease prototyping and increase the speed at which the user can write his code. However the benchmark shows that this is not a viable solution once the prototype needs to scale to a production environment.

5.3.2 Continuation: JIT-based solutions, pros. and cons.

Our hybrid solution certainly has advantages but the huge disadvantage is the slowness of injecting our own types from the Python side. There exists another solution that this thesis did not have the opportunity to study in-depth. This solution is based on a known technology: the Just-In-Time (JIT) compilation which has been previously illustrated in fig. 5.2. Indeed, it is a technology already used by interpreted languages such as Java or PHP to generate on-the-fly native and optimized machine code for the section of the source code that is considered "hot" by the interpreter. A source code is "hot" when it is executed a lot: the end-user would gain paying the compilation time once to have this code executed faster several times later on. When applying this strategy to our problematic, it would mean that the user must be able to compile native machine code from the templated generic C++ code by injected the requested type when it is used. Such an operation shift heavily the burden on the user and it is well-known that compiling C++ code is notably *complicated* and *slow*. In addition, the library needs to be able to auto-generate python-binding once the code is compiled. There are several solutions to achieve this process.

The first solution is to basically use system call to the compilers to actually *compile* C++ code once the templated types are known and explicitly instanciated in the source code. This solution requires careful code-generation design and that the user actually possess a compiler on his computer. Furtheremore, the user must resolve all the library dependencies, such as *freeimage* for IO etc. This solution was engineered in the library [88]. Indeed, each time the user declared a new automata in his jupyter notebook, corresponding source code is compiled in the foreground and then cached. It is a very perilous solution to implement when the final execution environment (OS, installed software) is not well-known in advance. Nowadays, the issue may be lesser, however, it still requires to maintain both the library and the container solution to use it.

The second solution is to use Cython [71]. It is a transpiling infrastructure which transform a Python source code directly into C-language source code so that it can be compiled by a standard C compiler just by linking against the Python/C API. This remove the burden of writting the careful code-generation routine, system-calls to the C++ compiler and removes the need to resolve all the dependencies. This infrastructure takes care of everything for the user. Also, by transpiling it into C code, it is faster because a C compiler is faster than a C++ compiler. Cython even support C++ template code [150] which is mandatory for our use-case.

The third solution consists in relying on recent projects that are all relying on the LLVM infrastructure. We can notably note Autowig [118], Cppyy [107] and Xeus-cling [147]. Autowig has in-house code based on LLVM/Clang to parse C++ code in order to generate and compile a Swig Python binding using the Mako templating engine. Autowig, coupled with Cython would permit the user to, for instance, generate C code related to a custom Python structure. Then a simple call to Autowig will parse the C code and inject it into the C++ library to generate the appropriate bindings for the user. As for Cppyy, it is based on LLVM/Cling, a C++ interpreter, and can directly interpret C++ code from a python string. This allow for easy injection of custom types, be they in Python code (transpiled with Cython) or C++ code (directly interpreted by Cling). Afterwards, the infrastructure generates the appropriate binding from the templated C++ library for the injected type. Finally, Xeus-cling is a ready-to-use jupyter kernel and allow the usage of C++ code directly from within a notebook. This completely bypass the need of a Python binding in the first place and allow the user to use the library from within the notebook as if he was using a Python library. However all those infrastructure come with a hefty cost in term of binary size. Indeed, a C++ compiler is not small and embarking it alongside the image processing library can easily impact greatly the final binary. Without the LLVM infrastructure the binary may weight around 3MB. With the LLVM infrastructure, the binary weight at the bare minimum 50MB. Also, these solutions may not be immediately faster. Indeed, when prototyping back and forth with a variety of types, the user may not be eager to wait for long compilations times each time he is testing with a an iteration of his work. Despite those facts, those solutions

offers great avenue of research for the future and the author is eager to thread those paths.

Part III

Continuation

Chapter 6

Conclusion

The work presented in this thesis by the author followed a very clear narrative arc. The emphasis was first shown on presenting what is the notion of genericity, its story and how anyone can relate to its day to day usage especially when applied to image processing. Genericity is a 4-decades year old notion that has evolved and found usage in very modern area of our society. In particular, in image processing, it is widely used to build modern applications used around the world. However, it was demonstrated how difficult it can be to implement solutions relying on genericity by stating the rule of three related to genericity, performance and easy of use. The rule states that one can only have two of those items by sacrificing the third one. If one wants to be generic and efficient, then the naive solution will be very complex to use with lot of parameters. If one wants a solution to be generic and easy to use, then it will be not very efficient by default. If one wants a solution to be easy to use and efficient then it will not be very generic. In this thesis, we try to demonstrate how to break this rule in three steps.

The first step was to realize an inventory of image types and families as well as different image processing algorithms. The aim was to produce a comprehensive taxonomy of images and algorithms related to image processing in order to be able to write concepts (in the sense of C++ concepts). This first step is to make the perimeter of what the author means by genericity very clear. From this starting point, it becomes easier to write image processing algorithms by default, just by relying on those concepts. Furthermore, different concepts exists to enable algorithm implementers to leverage properties (structuring elements' decomposability, image's buffer contiguous, ...) in order to achieve maximum performances.

A this point, we are still reasoning at low level and need to design an abstraction layer in order to enable fast prototyping for simple operations while guarantying very small memory footprint and near-zero performance impact. We expand the concept of *views* from the C++ standard to images and design what is an image view. We also make the design choice to have cheap-to-copy (because of shared data buffer) image by default in order to merge concrete image and views from the user point of view. The lazy-evaluation, that systematically happens when using views allows performance gain when clipping larges images. In the case where the whole image is processed, we were able to still retain very satisfactory performances that remain stable. Also, we show through concrete use-case, such as pixel-wise algorithm and border management how the usage of views simplify greatly how to write more complex image processing algorithms that are efficient by default.

Finally, this thesis focused its attention on how it is possible to distribute this software to the image processing community which is mainly working with Python. The last work concentrate its effort to how best design a static (templated C++) - dynamic (Python notebook) bridge to bring those concepts and performances to the practitioner. This last work explores the dilemmas and offers to address them with one hybrid solution whose design is explained in-depth. This hybrid solution rely on a type-erased type which offers compatibility with a *NumPy.array* that then is able to cast itself inside $n * n$ dispatcher (dimension and underlying type) into an optimized

templated C++ type. This solution also explain how to write very simply the glue code to enable already-existing algorithms (in C++) to be exposed in Python thanks to a dispatch mechanic heavily inspired from the C++ standard (`std::visit`, `std::variant`). The aim of this solution was to regroup at a single place in the code all the supported types into the dispatchers for maintenance purpose as well as demanding minimal work from algorithm implementer to expose their algorithm, all this while keeping the native performances. Indeed, no-copy is performed thanks to *pybind11*'s *buffer protocol* facility, and one cast is done from the type-erased type to the native one. All the work that is done in the algorithm is performed on native optimized type. Finally this solution offers a way to inject custom Python types into the library for prototyping purpose at the cost of heavy performance thanks to a new abstraction layer, the *value-set* explained in-depth in the chapter. The downside of this solution is obviously the code bloat with the resulting binary size exploding exponentially with the number of supported types multiplied by the number of algorithms multiplied by the number of additional supported data (structuring elements, label map, etc.)

We conclude this thesis by offering new avenue of research, the JIT-compilation. The author think that this avenue is worth exploring, especially with the already promising existing tools (xeus-cling, cppyy, Cython, autowig) in order to solve the code bloat issue. We would only compile what the user needs. But the entry price may be to statically link a C++ interpreter (LLVM/cling ?) into the binary which in itself would greatly bloat it. It may be possible to rely on user's system-wide infrastructure however so that the maintenance does not distribute a whole C++ interpreter/compiler alongside his image processing library binary. This is still a new area of research and the author would very much want to delve into it to study what is possible to achieve as of today with those tools for the image processing community.

Part IV

Appendices

Appendix A

Bibliography

References

[1] Russell R. Atkinson, Barbara H. Liskov, and Robert W. Scheifler. "Aspects Of Implementing CLU". In: *Proceedings of the 1978 Annual Conference*. ACM '78. Washington, D.C., USA: Association for Computing Machinery, 1978, pp. 123–129. ISBN: 0897910001. DOI: 10.1145/800127.804079. URL: https://doi.org/10.1145/800127.804079 (cit. on p. 29).

[2] John Backus. "Can Programming Be Liberated from the von Neumann Style? A Functional Style and Its Algebra of Programs". In: *Commun. ACM* 21.8 (Aug. 1978), pp. 613–641. ISSN: 0001-0782. DOI: 10.1145/359576.359579. URL: https://doi.org/10.1145/359576.359579 (cit. on p. 19).

[3] ANSI. *ANSI/MIL-STD-1815A: Programming languages — Ada*. American National Standards Institute, June 1983, p. 333. URL: https://www.iso.org/standard/16028.html (cit. on p. 29).

[4] Barry W. Boehm. "Software Engineering Economics". In: *IEEE Transactions on Software Engineering* SE-10.1 (Jan. 1984), pp. 4–21. ISSN: 1939-3520. DOI: 10.1109/TSE.1984.5010193 (cit. on p. 89).

[5] ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada*. Geneva, Switzerland: International Organization for Standardization, June 1987, p. 333. URL: https://www.iso.org/standard/16028.html (cit. on p. 30).

[6] Alexander A Stepanov, Aaron Kershenbaum, and David R Musser. *Higher order programming*. March, 1987 (cit. on p. 20).

[7] David R. Musser and Alexander A. Stepanov. "Generic programming". In: *Intl. Symp. on Symbolic and Algebraic Computation*. Springer. 1988, pp. 13–25 (cit. on pp. 20, 23, 30, 36, 37, 42).

[8] David R Musser and Alexander A Stepanov. *The Ada Generic Library linear list processing packages*. Springer-Verlag, 1989 (cit. on p. 20).

[9] Gerhard X. Ritter, Joseph N. Wilson, and Jennifer L. Davidson. "Image algebra: An overview". In: *Computer Vision, Graphics, and Image Processing* 49.3 (1990), pp. 297–331 (cit. on p. 37).

[10] Marcel van Herk. "A fast algorithm for local minimum and maximum filters on rectangular and octagonal kernels". In: *Pattern Recognition Letters* 13.7 (1992), pp. 517–521. ISSN: 0167-8655. DOI: https://doi.org/10.1016/0167-8655(92)90069-C. URL: http://www.sciencedirect.com/science/article/pii/016786559290069C (cit. on p. 50).

[11] Barbara Liskov. "A History of CLU". In: *SIGPLAN Not.* 28.3 (Mar. 1993), pp. 133–147. ISSN: 0362-1340. DOI: 10.1145/155360.155367. URL: https://doi.org/10.1145/155360.155367 (cit. on p. 29).

[12] David R. Musser and Alexander A. Stepanov. "Algorithm-oriented generic libraries". In: *Software: Practice and Experience* 24.7 (1994), pp. 623–642. DOI: 10.1002/spe.4380240703. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/spe.4380240703. URL: https://onlinelibrary.wiley.com/doi/abs/10.1002/spe.4380240703 (cit. on pp. 21, 23, 36).

[13]   Bjarne Stroustrup. *The Design and Evolution of C++*. New York, NY, USA: ACM
       Press/Addison-Wesley Publishing Co., 1994. ISBN: 0-201-54330-3 (cit. on p. 36).

[14]   ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada*. Geneva, Switzerland: In-
       ternational Organization for Standardization, Feb. 1995, p. 511. URL: https://www.iso.
       org/standard/22983.html (cit. on p. 30).

[15]   Alexander Stepanov and Meng Lee. *The standard template library*. Vol. 1501. Hewlett
       Packard Laboratories 1501 Page Mill Road, Palo Alto, CA 94304, 1995 (cit. on p. 20).

[16]   Todd Veldhuizen. "Expression templates". In: *C++ Report* 7.5 (1995), pp. 26–31 (cit. on
       pp. 17, 19, 86).

[17]   David M Beazley et al. "SWIG: An Easy to Use Tool for Integrating Scripting Languages
       with C and C++." In: *Tcl/Tk Workshop*. Vol. 43. 1996, p. 74 (cit. on p. 42).

[18]   James O. Coplien. "Curiously Recurring Template Patterns". In: *C++ Gems*. USA: SIGS
       Publications, Inc., 1996, pp. 135–144. ISBN: 1884842372 (cit. on p. 34).

[19]   ISO. *ISO/IEC 14882:2003: Programming languages — C++*. Geneva, Switzerland: Inter-
       national Organization for Standardization, Sept. 1998, p. 732. URL: https://www.iso.org/
       standard/25845.html (cit. on pp. 20, 30, 34).

[20]   Bjarne Stroustrup. "An overview of the C++ programming language". In: *Handbook of
       Object Technology* (1999) (cit. on p. 36).

[21]   Matthew H. Austern. "Segmented Iterators and Hierarchical Algorithms". In: *Generic
       Programming*. Ed. by Mehdi Jazayeri, Rüdiger G. K. Loos, and David R. Musser. Berlin,
       Heidelberg: Springer Berlin Heidelberg, 2000, pp. 80–90. ISBN: 978-3-540-39953-7 (cit. on
       p. 82).

[22]   G. Bradski. "The OpenCV Library". In: *Dr. Dobb's Journal of Software Tools* (2000) (cit.
       on pp. 17, 25).

[23]   James C. Dehnert and Alexander Stepanov. "Fundamentals of Generic Programming". In:
       Generic Programming. Vol. 1766. LNCS. Springer. 2000, pp. 1–11 (cit. on pp. 20, 21, 23,
       36, 42).

[24]   Alexandre Duret-Lutz. "Olena: a component-based platform for image processing, mixing
       generic, generative and OO programming". In: *symposium on Generative and Component-
       Based Software Engineering, Young Researchers Workshop*. Vol. 10. Citeseer. 2000 (cit. on
       p. 36).

[25]   Thierry Géraud et al. "Obtaining genericity for image processing and pattern recognition
       algorithms". In: *Proceedings of the 15th International Conference on Pattern Recognition
       (ICPR)*. Vol. 4. Barcelona, Spain: IEEE Computer Society, Sept. 2000, pp. 816–819 (cit.
       on pp. 36, 37).

[26]   Øyvind Kolås and et al. *Generic Graphic Library*. Available at http://www.gegl.org. 2000
       (cit. on p. 17).

[27]   Ullrich Köthe. "STL-Style Generic Programming with Images". In: *C++ Report Magazine*
       12.1 (2000). https://ukoethe.github.io/vigra, pp. 24–30 (cit. on pp. 17, 26, 42).

[28]   L. Prechelt. "An empirical comparison of seven programming languages". In: *Computer*
       33.10 (Oct. 2000), pp. 23–29. ISSN: 1558-0814. DOI: 10.1109/2.876288 (cit. on p. 89).

[29]   Jeremy Siek and Andrew Lumsdaine. "Concept checking: Binding parametric polymor-
       phism in C++". In: *First Workshop on C++ Template Programming*. Citeseer. 2000 (cit.
       on p. 34).

[30]   Todd L. Veldhuizen. "Blitz++: The Library that Thinks it is a Compiler". In: *Advances in Software Tools for Scientific Computing*. Ed. by Hans Petter Langtangen, Are Magnus Bruaset, and Ewald Quak. Berlin, Heidelberg: Springer Berlin Heidelberg, 2000, pp. 57–87. ISBN: 978-3-642-57172-5 (cit. on pp. 17, 19, 86).

[31]   M. Allili, K. Mischaikow, and A. Tannenbaum. "Cubical homology and the topological classification of 2D and 3D imagery". In: *Proceedings 2001 International Conference on Image Processing (Cat. No.01CH37205)*. Vol. 2. Oct. 2001, 173–176 vol.2. DOI: 10.1109/ICIP.2001.958452 (cit. on p. 51).

[32]   Abdel-Hameed A. Badawy et al. "Evaluating the Impact of Memory System Performance on Software Prefetching and Locality Optimizations". In: *Proceedings of the 15th International Conference on Supercomputing*. ICS '01. Sorrento, Italy: Association for Computing Machinery, 2001, pp. 486–500. ISBN: 158113410X. DOI: 10.1145/377792.377906. URL: https://doi.org/10.1145/377792.377906 (cit. on p. 98).

[33]   ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada — Technical Corrigendum 1*. Geneva, Switzerland: International Organization for Standardization, June 2001, p. 56. URL: https://www.iso.org/standard/35451.html (cit. on p. 30).

[34]   Andrew Lumsdaine Jeremy Siek Lie-Quan Lee. "The Boost Graph library". In: Addison-Wesley, 2001. URL: http://cds.cern.ch/record/1518180/files/0201729148_TOC.pdf (cit. on p. 55).

[35]   Eric Jones, Travis Oliphant, Pearu Peterson, et al. *SciPy: Open source scientific tools for Python*. 2001. URL: http://www.scipy.org/ (cit. on p. 17).

[36]   L Susan Blackford et al. "An updated set of basic linear algebra subprograms (BLAS)". In: *ACM Transactions on Mathematical Software* 28.2 (2002), pp. 135–151 (cit. on p. 19).

[37]   Jérôme Darbon, Thierry Géraud, and Alexandre Duret-Lutz. "Generic implementation of morphological image operators". In: *Mathematical Morphology, Proceedings of the 6th International Symposium (ISMM)*. Sydney, Australia: CSIRO Publishing, Apr. 2002, pp. 175–184 (cit. on pp. 36, 37).

[38]   Delores M Etter, David C Kuncicky, and Douglas W Hull. *Introduction to MATLAB*. Prentice Hall, 2002 (cit. on p. 26).

[39]   David Vandevoorde and Nicolai M Josuttis. *C++ Templates: The Complete Guide, Portable Documents*. Addison-Wesley Professional, 2002 (cit. on p. 32).

[40]   Tomihisa Welsh, Michael Ashikhmin, and Klaus Mueller. "Transferring Color to Grayscale Images". In: *Proceedings of the 29th Annual Conference on Computer Graphics and Interactive Techniques*. SIGGRAPH '02. San Antonio, Texas: Association for Computing Machinery, 2002, pp. 277–280. ISBN: 1581135211. DOI: 10.1145/566570.566576. URL: https://doi.org/10.1145/566570.566576 (cit. on p. 74).

[41]   Djemel Ziou and Madjid Allili. "Generating cubical complexes from image data and computation of the Euler number". In: *Pattern Recognition* 35.12 (2002). Pattern Recognition in Information Systems, pp. 2833–2839. ISSN: 0031-3203. DOI: https://doi.org/10.1016/S0031-3203(01)00238-2. URL: https://www.sciencedirect.com/science/article/pii/S0031320301002382 (cit. on p. 51).

[42]   Nicolas Burrus et al. "A static C++ object-oriented programming (SCOOP) paradigm mixing benefits of traditional OOP and generic programming". In: *Proceedings of the Workshop on Multiple Paradigm with Object-Oriented Languages (MPOOL)*. Anaheim, CA, USA, Oct. 2003 (cit. on p. 34).

[43]   ISO. *ISO/IEC 14882:2003: Programming languages — C++*. Geneva, Switzerland: International Organization for Standardization, Oct. 2003, p. 757. URL: https://www.iso.org/standard/38110.html (cit. on p. 34).

[44] Alexander Stepanov. *Greatest Common Measure: The Last 2500 Years*. Slides available at http://stepanovpapers.com/gcd.pdf. Talk available on youtube at https://youtu.be/fanm5y00joc. 2003. URL: http://stepanovpapers.com/gcd.pdf (cit. on p. 21).

[45] Todd L. Veldhuizen. *C++ Templates are Turing Complete*. Tech. rep. 2003 (cit. on p. 31).

[46] Jérôme Darbon, Thierry Géraud, and Patrick Bellot. "Generic algorithmic blocks dedicated to image processing". In: *Proceedings of the ECOOP Workshop for PhD Students*. Oslo, Norway, June 2004 (cit. on p. 36).

[47] Jacques Froment. *MegaWave2 user's guide*. 2004 (cit. on p. 16).

[48] Anat Levin, Dani Lischinski, and Yair Weiss. "Colorization Using Optimization". In: *ACM SIGGRAPH 2004 Papers*. SIGGRAPH '04. Los Angeles, California: Association for Computing Machinery, 2004, pp. 689–694. ISBN: 9781450378239. DOI: 10.1145/1186562.1015780. URL: https://doi.org/10.1145/1186562.1015780 (cit. on p. 74).

[49] Stewart Taylor. *Intel integrated performance primitives*. Intel Press, 2004 (cit. on p. 17).

[50] Valérie De Witte et al. "Vector Morphological Operators for Colour Images". In: *Image Analysis and Recognition*. Ed. by Mohamed Kamel and Aurélio Campilho. Berlin, Heidelberg: Springer Berlin Heidelberg, 2005, pp. 667–675. ISBN: 978-3-540-31938-2 (cit. on p. 26).

[51] Guntram Berti. "GrAL–the Grid Algorithms Library". In: *Future Generation Computer Systems* 22.1-2 (2006), pp. 110–122 (cit. on pp. 17, 26).

[52] Lubomir Bourdev and Hailin Jin. *Boost Generic Image Library*. Adobe stlab. Available at https://stlab.adobe.com/gil/index.html. 2006 (cit. on pp. 17, 26).

[53] Gabriel Dos Reis and Bjarne Stroustrup. "Specifying C++ Concepts". In: *Conference Record of the 33rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL '06. Charleston, South Carolina, USA: Association for Computing Machinery, 2006, pp. 295–308. ISBN: 1595930272. DOI: 10.1145/1111037.1111064. URL: https://doi.org/10.1145/1111037.1111064 (cit. on p. 40).

[54] Thierry Géraud. *Advanced Static Object-Oriented Programming Features: A Sequel to SCOOP*. http://www.lrde.epita.fr/people/theo/pub/olena/olena-06-jan.pdf. Jan. 2006 (cit. on p. 34).

[55] Douglas Gregor et al. "Concepts: Linguistic Support for Generic Programming in C++". In: *SIGPLAN Not.* 41.10 (Oct. 2006), pp. 291–310. ISSN: 0362-1340. DOI: 10.1145/1167515.1167499. URL: https://doi.org/10.1145/1167515.1167499 (cit. on p. 20).

[56] Douglas Gregor et al. "Concepts: Linguistic Support for Generic Programming in C++". In: *Proceedings of the 21st Annual ACM SIGPLAN Conference on Object-Oriented Programming Systems, Languages, and Applications*. OOPSLA '06. Portland, Oregon, USA: Association for Computing Machinery, 2006, pp. 291–310. ISBN: 1595933484. DOI: 10.1145/1167473.1167499. URL: https://doi.org/10.1145/1167473.1167499 (cit. on p. 20).

[57] Travis Oliphant. *NumPy: A guide to NumPy*. USA: Trelgol Publishing. 2006. URL: http://www.numpy.org/ (cit. on pp. 17, 92).

[58] Jeremy Siek and Dave Abrahams. *Boost Concept Check*. https://www.boost.org/. Available starting with Boost v1.19.0. 2006 (cit. on p. 34).

[59] Jesús Angulo. "Morphological colour operators in totally ordered lattices based on distances: Application to image filtering, enhancement and analysis". In: *Computer Vision and Image Understanding* 107.1 (2007). Special issue on color image processing, pp. 56–73. ISSN: 1077-3142. DOI: https://doi.org/10.1016/j.cviu.2006.11.008. URL: https://www.sciencedirect.com/science/article/pii/S1077314206002165 (cit. on p. 26).

[60] ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada — Amendment 1.* Geneva, Switzerland: International Organization for Standardization, Mar. 2007, p. 317. URL: https://www.iso.org/standard/45001.html (cit. on p. 30).

[61] Bjarne Stroustrup. "Evolving a Language in and for the Real World: C++ 1991-2006". In: *Proc. of the 3rd ACM SIGPLAN Conf. on History of Programming Languages.* Vol. 4. San Diego, California, 2007, pp. 1–59. ISBN: 9781595937667. DOI: 10.1145/1238844.1238848. URL: https://doi.org/10.1145/1238844.1238848 (cit. on p. 36).

[62] Jeffrey Dean and Sanjay Ghemawat. "MapReduce: Simplified Data Processing on Large Clusters". In: *Commun. ACM* 51.1 (Jan. 2008), pp. 107–113. ISSN: 0001-0782. DOI: 10.1145/1327452.1327492. URL: https://doi.org/10.1145/1327452.1327492 (cit. on p. 52).

[63] Thierry Géraud and Roland Levillain. "Semantics-Driven Genericity: A Sequel to the Static C++ Object-Oriented Programming Paradigm (SCOOP 2)". In: *Proceedings of the 6th International Workshop on Multiparadigm Programming with Object-Oriented Languages (MPOOL).* Paphos, Cyprus, July 2008 (cit. on p. 34).

[64] Thomas Kotzmann et al. "Design of the Java HotSpot™ Client Compiler for Java 6". In: *ACM Trans. Archit. Code Optim.* 5.1 (May 2008). ISSN: 1544-3566. DOI: 10.1145/1369396.1370017. URL: https://doi.org/10.1145/1369396.1370017 (cit. on p. 91).

[65] Roland Levillain, Thierry Géraud, and Laurent Najman. "Milena: Write Generic Morphological Algorithms Once, Run on Many Kinds of Images". In: *Mathematical Morphology and Its Application to Signal and Image Processing — Proceedings of the Ninth International Symposium on Mathematical Morphology (ISMM).* Ed. by Michael H. F. Wilkinson and Jos B. T. M. Roerdink. Vol. 5720. Lecture Notes in Computer Science. Groningen, The Netherlands: Springer Berlin / Heidelberg, Aug. 2009, pp. 295–306 (cit. on pp. 17, 26, 32, 35, 36, 67).

[66] Alexander Stepanov and Paul McJones. *Elements of Programming.* Addison-Wesley Professional, June 2009 (cit. on pp. 20, 23, 36, 37).

[67] Gaël Guennebaud, Benoît Jacob, et al. *Eigen v3.* http://eigen.tuxfamily.org. Available at http://eigen.tuxfamily.org. 2010 (cit. on pp. 17, 19, 67).

[68] Roland Levillain, Thierry Géraud, and Laurent Najman. "Why and How to Design a Generic and Efficient Image Processing Framework: The Case of the Milena Library". In: *Proceedings of the IEEE International Conference on Image Processing (ICIP).* Hong Kong, Sept. 2010, pp. 1941–1944 (cit. on pp. 17, 26, 32, 35–37).

[69] Roland Levillain, Thierry Géraud, and Laurent Najman. "Writing Reusable Digital Geometry Algorithms in a Generic Image Processing Framework". In: *Proceedings of the Workshop on Applications of Digital Geometry and Mathematical Morphology (WADGMM).* Istanbul, Turkey, Aug. 2010, pp. 96–100. URL: http://mdigest.jrc.ec.europa.eu/wadgmm2010/ (cit. on p. 36).

[70] Peter Pirkelbauer, Yuriy Solodkyy, and Bjarne Stroustrup. "Design and evaluation of C++ open multi-methods". In: *Science of Computer Programming* 75.7 (2010). Generative Programming and Component Engineering (GPCE 2007), pp. 638–667. ISSN: 0167-6423. DOI: https://doi.org/10.1016/j.scico.2009.06.002. URL: http://www.sciencedirect.com/science/article/pii/S016764230900094X (cit. on p. 93).

[71] S. Behnel et al. "Cython: The Best of Both Worlds". In: *Computing in Science Engineering* 13.2 (Mar. 2011), pp. 31–39. ISSN: 1521-9615. DOI: 10.1109/MCSE.2010.118 (cit. on pp. 42, 107).

[72] Lubomir Bourdev and Jaakko Järvi. "Efficient run-time dispatching in generic programming with minimal code bloat". In: *Science of Computer Programming* 76.4 (2011). Special issue on library-centric software design (LCSD 2006), pp. 243–257. ISSN: 0167-6423. DOI: https://doi.org/10.1016/j.scico.2008.06.003. URL: http://www.sciencedirect.com/science/article/pii/S0167642308000634 (cit. on p. 95).

[73] ISO. *ISO/IEC 14882:2011: Programming languages — C++*. Geneva, Switzerland: International Organization for Standardization, Sept. 2011, p. 1338. URL: https://www.iso.org/standard/50372.html (cit. on pp. 31, 36, 37).

[74] Roland Levillain. "Towards a Software Architecture for Generic Image Processing". PhD thesis. Marne-la-Vallée, France: Université Paris-Est, Nov. 2011 (cit. on pp. 17, 26, 34, 36).

[75] Roland Levillain, Thierry Géraud, and Laurent Najman. "Une approche générique du logiciel pour le traitement d'images préservant les performances". In: *Proceedings of the 23rd Symposium on Signal and Image Processing (GRETSI)*. In French. Bordeaux, France, Sept. 2011 (cit. on p. 36).

[76] Jacques Froment. "MegaWave". In: *IPOL 2012 Meeting on Image Processing Libraries*. France, June 2012. URL: https://hal.archives-ouvertes.fr/hal-00907378 (cit. on p. 16).

[77] Thierry Géraud. "Outil logiciel pour le traitement d'images: Bibliothèque, paradigmes, types et algorithmes". In French. Habilitation Thesis. Université Paris-Est, June 2012 (cit. on pp. 17, 26, 36, 67).

[78] Klaus Iglberger. *Blaze C++ Linear Algebra Library*. https://bitbucket.org/blaze-lib. 2012 (cit. on pp. 17, 19).

[79] Klaus Iglberger et al. "Expression Templates Revisited: A Performance Analysis of Current Methodologies". In: *SIAM Journal on Scientific Computing* 34(2) (2012), pp. C42–C69 (cit. on pp. 17, 19).

[80] Klaus Iglberger et al. "High Performance Smart Expression Template Math Libraries". In: *Proceedings of the 2nd International Workshop on New Algorithms and Programming Models for the Manycore Era (APMM 2012) at HPCS 2012*. 2012 (cit. on p. 19).

[81] ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada*. Geneva, Switzerland: International Organization for Standardization, Dec. 2012, p. 832. URL: https://www.iso.org/standard/61507.html (cit. on p. 30).

[82] Roland Levillain, Thierry Géraud, and Laurent Najman. "Writing Reusable Digital Topology Algorithms in a Generic Image Processing Framework". In: *WADGMM 2010*. Ed. by Ullrich Köthe, Annick Montanvert, and Pierre Soille. Vol. 7346. Lecture Notes in Computer Science. Springer-Verlag Berlin Heidelberg, 2012, pp. 140–153 (cit. on p. 36).

[83] S. Matuska, R. Hudec, and M. Benco. "The comparison of CPU time consumption for image processing algorithm in Matlab and OpenCV". In: *2012 ELEKTRO*. May 2012, pp. 75–78. DOI: 10.1109/ELEKTRO.2012.6225575 (cit. on p. 28).

[84] Andrew Sutton and Bjarne Stroustrup. "Design of Concept Libraries for C++". In: *Software Language Engineering*. Ed. by Anthony Sloane and Uwe Aßmann. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 97–118. ISBN: 978-3-642-28830-2 (cit. on p. 20).

[85] David Tschumperlé. "The CImg Library". In: *IPOL 2012 Meeting on Image Processing Libraries*. Cachan, France, June 2012, 4 pp. URL: https://hal.archives-ouvertes.fr/hal-00927458 (cit. on pp. 17, 25).

[86] V. Vassilev et al. "Cling — The New Interactive Interpreter for ROOT 6". In: vol. 396. 5. IOP Publishing, Dec. 2012, p. 052071. DOI: 10.1088/1742-6596/396/5/052071. URL: https://iopscience.iop.org/article/10.1088/1742-6596/396/5/052071/pdf (cit. on p. 41).

[87] AMD. *ACML-GPU - drop-in BLAS replacement, multi-GPUs accelerated.* AMD. Available at http://developer.amd.com/wordpress/media/2013/02/ACML-GPUreadme.pdf. 2013 (cit. on p. 19).

[88] Akim Demaille et al. "Implementation Concepts in Vaucanson 2". In: *Implementation and Application of Automata.* Ed. by Stavros Konstantinidis. Berlin, Heidelberg: Springer Berlin Heidelberg, 2013, pp. 122–133. ISBN: 978-3-642-39274-0 (cit. on pp. 42, 107).

[89] Agner Fog. *C++ vector class library.* http://www.agner.org/optimize/vectorclass.pdf. 2013 (cit. on p. 54).

[90] Hans J. Johnson et al. *The ITK Software Guide.* Third. *In press.* Kitware, Inc. 2013. URL: http://www.itk.org/ItkSoftwareGuide.pdf (cit. on pp. 17, 26).

[91] Jonathan Ragan-kelley et al. "Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines". In: *PLDI 2013* (2013) (cit. on pp. 17, 19, 52, 54).

[92] Edwin Carlinet and Thierry Géraud. "A Comparative Review of Component Tree Computation Algorithms". In: *IEEE Transactions on Image Processing* 23.9 (2014), pp. 3885–3895 (cit. on p. 51).

[93] Pierre Estérie et al. "Boost.SIMD: Generic Programming for Portable SIMDization". In: *Proceedings of the 2014 Workshop on Programming Models for SIMD/Vector Processing.* WPMVP '14. Orlando, Florida, USA: ACM, 2014, pp. 1–8. ISBN: 978-1-4503-2653-7. DOI: 10.1145/2568058.2568063. URL: http://doi.acm.org/10.1145/2568058.2568063 (cit. on p. 54).

[94] Matthieu Garrigues and Antoine Manzanera. "Video++, a modern image and video processing C++ framework". In: *Conference on Design and Architectures for Signal and Image Processing (DASIP).* IEEE. 2014, pp. 1–6 (cit. on p. 17).

[95] ISO. *ISO/IEC 14882:2014: Programming languages — C++.* Geneva, Switzerland: International Organization for Standardization, Dec. 2014, p. 1358. URL: https://www.iso.org/standard/64029.html (cit. on p. 36).

[96] Donald E Knuth. *Art of computer programming, volume 2: Seminumerical algorithms.* Addison-Wesley Professional, 2014 (cit. on p. 21).

[97] Roland Levillain et al. "Practical Genericity: Writing Image Processing Algorithms Both Reusable and Efficient". In: *Progress in Pattern Recognition, Image Analysis, Computer Vision, and Applications — Proceedings of the 19th Iberoamerican Congress on Pattern Recognition (CIARP).* Ed. by Eduardo Bayro and Edwin Hancock. Vol. 8827. Lecture Notes in Computer Science. Puerto Vallarta, Mexico: Springer-Verlag, Nov. 2014, pp. 70–79 (cit. on pp. 17, 26, 36, 37).

[98] Alexander A Stepanov and Daniel E Rose. *From mathematics to generic programming.* Pearson Education, 2014 (cit. on p. 21).

[99] Stéfan van der Walt et al. "scikit-image: image processing in Python". In: *PeerJ* 2 (June 2014), e453. ISSN: 2167-8359. DOI: 10.7717/peerj.453. URL: https://doi.org/10.7717/peerj.453 (cit. on pp. 17, 25).

[100] Edwin Carlinet and Thierry Géraud. "MToS: A Tree of Shapes for Multivariate Images". In: *IEEE Transactions on Image Processing* 24.12 (2015), pp. 5330–5342 (cit. on p. 51).

[101] Francois Chollet et al. *Keras.* 2015. URL: https://github.com/fchollet/keras (cit. on p. 52).

[102] Martín Abadi et al. *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems.* Software available from tensorflow.org. 2015. URL: https://www.tensorflow.org/ (cit. on p. 52).

[103] D Coeurjolly, JO Lachaud, and B Kerautret. *DGtal: Digital geometry tools and algorithms library.* 2016 (cit. on pp. 17, 26).

[104] Abderrahmane Nassim Halli. "Optimisation de code pour application Java haute-performance". Available at https://tel.archives-ouvertes.fr/tel-01679740/file/HALLI_2016_archivage.pdf. Theses. Université Grenoble Alpes, Oct. 2016. URL: https://tel.archives-ouvertes.fr/tel-01679740 (cit. on p. 91).

[105] ISO. *ISO/IEC JTC 1/SC 22: Programming languages — Ada — Technical Corrigendum 1.* Geneva, Switzerland: International Organization for Standardization, Feb. 2016, p. 75. URL: https://www.iso.org/standard/69798.html (cit. on p. 30).

[106] Thomas Kluyver et al. "Jupyter Notebooks — a publishing format for reproducible computational workflows". In: *Positioning and Power in Academic Publishing: Players, Agents and Agendas.* Ed. by Fernando Loizides and Birgit Scmidt. Netherlands: IOS Press, 2016, pp. 87–90. URL: https://eprints.soton.ac.uk/403913/ (cit. on pp. 17, 42).

[107] W. T. L. P. Lavrijsen and A. Dutta. "High-Performance Python-C++ Bindings with PyPy and Cling". In: *2016 6th Workshop on Python for High-Performance and Scientific Computing (PyHPC).* Nov. 2016, pp. 27–35. DOI: 10.1109/PyHPC.2016.008 (cit. on pp. 42, 107).

[108] Conrad Sanderson and Ryan Curtin. *Armadillo, C++ library for linear algebra & scientific computing.* Sourceforge. Available at https://gitlab.com/conradsnicta/armadillo-code. 2016 (cit. on pp. 17, 19).

[109] Conrad Sanderson and Ryan Curtin. "Armadillo: a template-based C++ library for linear algebra". In: *Journal of Open Source Software* 1.2 (2016), p. 26 (cit. on p. 19).

[110] Richard Zhang, Phillip Isola, and Alexei A. Efros. "Colorful Image Colorization". In: *Computer Vision – ECCV 2016.* Ed. by Bastian Leibe et al. Cham: Springer International Publishing, 2016, pp. 649–666. ISBN: 978-3-319-46487-9 (cit. on p. 74).

[111] Antonio Gulli and Sujit Pal. *Deep learning with Keras.* Packt Publishing Ltd, 2017 (cit. on p. 52).

[112] Andrew G. Howard et al. *MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications.* 2017. arXiv: 1704.04861 [cs.CV] (cit. on p. 13).

[113] ISO. *ISO/IEC 14882:2017: Programming languages — C++.* Geneva, Switzerland: International Organization for Standardization, Dec. 2017, p. 1605. URL: https://www.iso.org/standard/68564.html (cit. on p. 36).

[114] Wenzel Jakob, Jason Rhinelander, and Dean Moldovan. *pybind11—Seamless operability between C++ 11 and Python.* https://github.com/pybind/pybind11. 2017 (cit. on pp. 93, 94).

[115] *Background Subtraction.* https://docs.opencv.org/3.4/d1/dc5/tutorial_background_subtraction.html. Dec. 2018 (cit. on p. 82).

[116] G. Brown et al. "Introducing Parallelism to the Ranges TS". In: *Proceedings of the International Workshop on OpenCL.* 2018, pp. 1–5 (cit. on p. 87).

[117] Edwin Carlinet et al. *Pylena: a Modern C++ Image Processing Generic Library.* EPITA Research and Developement Laboratory. Available at https://gitlab.lrde.epita.fr/olena/pylene. 2018 (cit. on pp. 17, 23, 26, 36).

[118] Pierre Fernique and Christophe Pradal. "AutoWIG: automatic generation of python bindings for C++ libraries". In: *PeerJ Computer Science* 4 (2018), e149. URL: https://peerj.com/articles/cs-149.pdf (cit. on pp. 42, 107).

[119] Thierry Géraud and Edwin Carlinet. *A Modern C++ Library for Generic and Efficient Image Processing*. Journée du Groupe de Travail de Géométrie Discrète et Morphologie Mathématique, Lyon, France. June 2018. URL: https://www.lrde.epita.fr/~theo/talks/geraud.2018.gtgdmm_talk.pdf (cit. on p. 73).

[120] Eric Niebler and Casey Carter. *P1037R0: Deep Integration of the Ranges TS*. https://wg21.link/p1037r0. May 2018 (cit. on pp. 18, 58).

[121] Conrad Sanderson and Ryan Curtin. "A User-Friendly Hybrid Sparse Matrix Class in C++". In: *Mathematical Software – ICMS 2018*. Ed. by James H. Davenport et al. Cham: Springer International Publishing, 2018, pp. 422–430. ISBN: 978-3-319-96418-8 (cit. on p. 19).

[122] Michael Wong and Hal Finkel. "Distributed & Heterogeneous Programming in C++ for HPC at SC17". In: *Proceedings of the International Workshop on OpenCL*. IWOCL '18. Oxford, United Kingdom: ACM, 2018, 20:1–20:7. ISBN: 978-1-4503-6439-3. DOI: 10.1145/3204919.3204939. URL: http://doi.acm.org/10.1145/3204919.3204939 (cit. on pp. 17, 19, 53, 92).

[123] Gordon Brown, Ruyman Reyes, and Michael Wong. "Towards Heterogeneous and Distributed Computing in C++". In: *Proceedings of the International Workshop on OpenCL*. IWOCL'19. Boston, MA, USA: ACM, 2019, 18:1–18:5. ISBN: 978-1-4503-6230-6. DOI: 10.1145/3318170.3318196. URL: http://doi.acm.org/10.1145/3318170.3318196 (cit. on pp. 17, 19, 53, 92).

[124] Ben Dean. *Identifying Monoids: Exploiting Compositional Structure in Code*. Slides available at https://github.com/boostcon/cppnow_presentations_2019/blob/master/05-09-2019_thursday/Identifying_Monoids_Exploiting_Compositional_Structure_in_Code__Ben_Deane_cppnow_05092019.pdf, Talk available on youtube at https://youtu.be/INnattuluiM. June 2019. URL: https://cppnow2019.sched.com/event/b60160aa659270370279c5acb6196fb6 (cit. on p. 55).

[125] Celian Gossec. "Binding a high-performance C++ image processing library to Python". In: *Student LRDE tech reports*. 2019 (cit. on p. 93).

[126] Khronos Group. *OpenVX*. https://www.khronos.org/openvx/. 2019 (cit. on p. 87).

[127] Adam Paszke et al. "PyTorch: An Imperative Style, High-Performance Deep Learning Library". In: *Advances in Neural Information Processing Systems 32*. Ed. by H. Wallach et al. Curran Associates, Inc., 2019, pp. 8024–8035. URL: http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf (cit. on p. 52).

[128] Jules Penuchot. *CUDA extension for Blaze*. Github. Available at https://github.com/STEllAR-GROUP/blaze_cuda. 2019 (cit. on p. 19).

[129] B. Perret et al. "Higra: Hierarchical Graph Analysis". In: *SoftwareX* 10 (2019), p. 100335. ISSN: 2352-7110. DOI: https://doi.org/10.1016/j.softx.2019.100335. URL: http://www.sciencedirect.com/science/article/pii/S235271101930247X (cit. on p. 51).

[130] PTC. *GraphicsMagick*. Version 6.0. Oct. 1, 2019. URL: https://www.mathcad.com (cit. on p. 16).

[131] Michaël Roynard, Edwin Carlinet, and Thierry Géraud. "An Image Processing Library in Modern C++: Getting Simplicity and Efficiency with Generic Programming". In: *Reproducible Research in Pattern Recognition*. Ed. by Bertrand Kerautret et al. Cham: Springer International Publishing, 2019, pp. 121–137. ISBN: 978-3-030-23987-9 (cit. on pp. 23, 25).

[132] The GIMP Development Team. *GIMP*. Version 2.10.12. June 12, 2019. URL: https://www.gimp.org (cit. on pp. 15, 16).

[133]   Aksel Alpay and Vincent Heuveline. "SYCL beyond OpenCL: The Architecture, Current State and Future Direction of HipSYCL". In: *Proceedings of the International Workshop on OpenCL*. IWOCL '20. Munich, Germany: Association for Computing Machinery, 2020. ISBN: 9781450375313. DOI: 10.1145/3388333.3388658. URL: https://doi.org/10.1145/3388333.3388658 (cit. on p. 54).

[134]   Anaconda, Inc. *Anaconda*. Version 2020.11. Nov. 19, 2020. URL: https://anaconda.com (cit. on p. 17).

[135]   GraphicsMagick Group. *GraphicsMagick*. Version 1.3.36. Dec. 26, 2020. URL: https://http://www.graphicsmagick.org (cit. on p. 16).

[136]   ISO. *ISO/IEC 14882:2020: Programming languages — C++*. Geneva, Switzerland: International Organization for Standardization, Dec. 2020, p. 1853. URL: https://www.iso.org/standard/79358.html (cit. on p. 58).

[137]   MathWorks. *MATLAB*. Version R2020b. Sept. 22, 2020. URL: https://fr.mathworks.com/products/matlab.html (cit. on p. 17).

[138]   Minh Ôn Vũ Ngoc et al. "A New Minimum Barrier Distance for Multivariate Images with Applications to Salient Object Detection, Shortest Path Finding, and Segmentation". In: *Computer Vision and Image Understanding* 197–198 (Aug. 2020). DOI: 10.1016/j.cviu.2020.102993 (cit. on p. 51).

[139]   Scilab Enterprises. *Scilab*. Version 6.1.0. Feb. 25, 2020. URL: https://www.scilab.org (cit. on p. 17).

[140]   wolfram Research. *Mathematica*. Version 12.2. Dec. 16, 2020. URL: https://www.wolfram.com/mathematica/ (cit. on p. 17).

[141]   Adobe. *Adobe Photoshop*. Version 22.2. Feb. 9, 2021. URL: https://photoshop.com/fr (cit. on pp. 15, 16).

[142]   Alex Clark and al. *Pillow*. Version 8.1.2. Mar. 6, 2021. URL: https://python-pillow.org (cit. on p. 17).

[143]   Boost. *Boost C++ Libraries*. Version 1.76.0. Apr. 16, 2021. URL: https://www.boost.org/ (cit. on p. 17).

[144]   GNU Project. *Octave*. Version 6.2.0. Feb. 20, 2021. URL: https://www.gnu.org/software/octave/index (cit. on p. 17).

[145]   Daniel Lemire, Geoff Langdale, and John Keiser. *simdjson*. Version 1.0. Sept. 8, 2021. URL: https://github.com/simdjson/simdjson (cit. on p. 17).

[146]   Niels Lohmann. *Json*. Version 3.10.2. Aug. 26, 2021. URL: https://github.com/nlohmann/json (cit. on p. 17).

[147]   QuantStack. *xeus-cling*. Version 0.12.1. Mar. 16, 2021. URL: https://github.com/QuantStack/xeus-cling (cit. on pp. 42, 107).

[148]   The ImageMagick Development Team. *ImageMagick*. Version 7.0.10. Jan. 4, 2021. URL: https://imagemagick.org (cit. on p. 16).

[149]   The PyPi Developement Team. *PyPi*. Version 21.0.1. Jan. 31, 2021. URL: https://pypi.org (cit. on p. 17).

[150]   S. Behnel et al. *Using C++ in Cython > templates*. https://cython.readthedocs.io/en/latest. Available at https://cython.readthedocs.io/en/latest/src/userguide/wrapping_CPlusPlus.html#templates. 2022 (cit. on p. 107).

[151]   Gaël Guennebaud, Benoît Jacob, et al. *Using Eigen in CUDA kernels*. http://eigen.tuxfamily.org. Available at https://eigen.tuxfamily.org/dox/TopicCUDA.html. 2022 (cit. on p. 19).

[152] Nvidia. *NVBLAS - drop-in BLAS replacement, multi-GPUs accelerated.* Nvidia. Available at https://docs.nvidia.com/cuda/nvblas/. 2022 (cit. on p. 19).

[153] Codeplay. *ComputeCpp.* Version 2.8.0. URL: https://codeplay.com/products/computesuite/computecpp (cit. on p. 54).

[154] Intel Corporation. *SYCL Compiler and Runtimes.* Version master. URL: https://github.com/intel/llvm (cit. on p. 54).

[155] Danke Xie, ChernHsi Jeffrey Wang, and Michael Schimpf. "Improving Java Virtual Machine performance with SIMD technology". In: () (cit. on p. 91).

[156] Xilinx et al. *The triSYCL project.* Version master. URL: https://github.com/trisycl/trisycl (cit. on p. 54).

[157] Andrew Sutton Eric Niebler Sean Parent. *Ranges for the Standard Library: Revision 1.* Oct. 2014. URL: https://ericniebler.github.io/std/wg21/D4128.html (cit. on pp. 58, 67).

[158] Bjarne Stroustrup and Gabriel Dos Reis. *Concepts — Design choices for template argument checking.* WG21, Oct. 2003. URL: https://wg21.link/n1522 (cit. on pp. 36, 37).

[159] Bill Seymour. *LWG Papers to Re-Merge into C++0x After Removing Concepts.* WG21, July 2009. URL: https://wg21.link/n2929 (cit. on p. 37).

[160] Andrew Sutton. *Working Draft, C++ extensions for Concepts.* WG21, June 2017. URL: https://wg21.link/n4674 (cit. on p. 37).

[161] EPITA Research and Developpement Laboratory (LRDE). *The Olena image processing platform.* http://olena.lrde.epita.fr. 2000 (cit. on pp. 17, 26, 36, 42).

[162] Ville Voutilainen. *Merge the Concepts TS Working Draft into the C++20 working draft.* WG21, June 2017. URL: https://wg21.link/p0724r0 (cit. on p. 37).

[163] Eric Niebler and Casey Carter. *Merging the Ranges TS.* WG21, May 2018. URL: https://wg21.link/p0896r1 (cit. on pp. 48, 58).

[164] Casey Carter and Eric Niebler. *Standard Library Concepts.* WG21, June 2018. URL: https://wg21.link/p0898r3 (cit. on p. 48).

[165] Eric Niebler and Casey Carter. *Deep Integration of the Ranges TS.* WG21, May 2018. URL: https://wg21.link/p1037r0 (cit. on pp. 48, 58).

[166] Bjarne Stroustrup. *Don't add to the signed/unsigned mess.* WG21, Feb. 2019. URL: https://wg21.link/p1491r0 (cit. on p. 60).

Further reading

[167] C. A. R. Hoare. "An Axiomatic Basis for Computer Programming". In: *Commun. ACM* 12.10 (Oct. 1969), pp. 576–580. ISSN: 0001-0782. DOI: 10.1145/363235.363259. URL: https://doi.org/10.1145/363235.363259.

[168] Nicolas Bourbaki. *Théorie des ensembles.* Vol. 1. Hermann Paris, 1970.

[169] R. M. Haralick. "Image Access Protocol for Image Processing Software". In: *IEEE Transactions on Software Engineering* SE-3.2 (Mar. 1977), pp. 190–192. ISSN: 2326-3881. DOI: 10.1109/TSE.1977.231124.

[170] Deepak Kapur, David R. Musser, and Alexander A. Stepanov. "Tecton: A language for manipulating generic objects". In: *Program Specification.* Ed. by Jørgen Staunstrup. Berlin, Heidelberg: Springer Berlin Heidelberg, 1982, pp. 402–414. ISBN: 978-3-540-39176-0.

[171] Jon Bentley. "Programming pearls: Writing correct programs". In: *Communications of the ACM* 26.12 (1983), pp. 1040–1045.

[172] Joseph A. Goguen. "Parameterized Programming". In: *IEEE Transactions on Software Engineering* SE-10.5 (Sept. 1984), pp. 528–543. ISSN: 1939-3520. DOI: 10.1109/TSE.1984.5010277.

[173] Jim Piper and Denis Rutovitz. "Data structures for image processing in a C language and Unix environment". In: *Pattern Recognition Letters* 3.2 (1985), pp. 119–129.

[174] David R. Musser and Alexander A. Stepanov. "A Library of Generic Algorithms in Ada". In: *Proceedings of the 1987 Annual ACM SIGAda International Conference on Ada.* SIGAda '87. Boston, Massachusetts, USA: Association for Computing Machinery, 1987, pp. 216–225. ISBN: 0897912438. DOI: 10.1145/317500.317529. URL: https://doi.org/10.1145/317500.317529.

[175] B. Meyer. "Applying 'design by contract'". In: *Computer* 25.10 (Oct. 1992), pp. 40–51. ISSN: 1558-0814. DOI: 10.1109/2.161279.

[176] Delores M Etter, Dolores M Etter, and Delores M Etter. *Engineering problem solving with MATLAB.* Vol. 2. Prentice Hall Englewood cliffs, New Jersey 07632, 1993.

[177] Robert W. Floyd. "Assigning Meanings to Programs". In: *Program Verification: Fundamental Issues in Computer Science.* Ed. by Timothy R. Colburn, James H. Fetzer, and Terry L. Rankin. Dordrecht: Springer Netherlands, 1993, pp. 65–81. ISBN: 978-94-011-1793-7. DOI: 10.1007/978-94-011-1793-7_4. URL: https://doi.org/10.1007/978-94-011-1793-7_4.

[178] Barbara Liskov. "A History of CLU". In: *The Second ACM SIGPLAN Conference on History of Programming Languages.* HOPL-II. Cambridge, Massachusetts, USA: Association for Computing Machinery, 1993, pp. 133–147. ISBN: 0897915704. DOI: 10.1145/154766.155367. URL: https://doi.org/10.1145/154766.155367.

[179] Erich Gamma. *Design patterns: elements of reusable object-oriented software.* Pearson Education India, 1995.

[180] Michael Hanus, Herbert Kuchen, and Juan Jose Moreno-Navarro. "Curry: A Truly Functional Logic Language". In: *Proc. ILPS'95 Workshop on Visions for the Future of Logic Programming*. Vol. 95. 1995, pp. 95–107.

[181] Eva Pärt-Enander et al. *The MATLAB handbook*. Addison-Wesley Harlow, 1996.

[182] Bjarne Stroustrup. "A History of C++: 1979–1991". In: *History of Programming Languages—II*. New York, NY, USA: Association for Computing Machinery, 1996, pp. 699–769. ISBN: 0201895021. URL: https://doi.org/10.1145/234286.1057836.

[183] G. S. Novak. "Software reuse by specialization of generic procedures through views". In: *IEEE Transactions on Software Engineering* 23.7 (July 1997), pp. 401–417. ISSN: 2326-3881. DOI: 10.1109/32.605759.

[184] D. Roman, M. Fisher, and J. Cubillo. "Digital image processing-an object-oriented approach". In: *IEEE Transactions on Education* 41.4 (Nov. 1998), pp. 331–333. ISSN: 1557-9638. DOI: 10.1109/13.728270.

[185] Todd L. Veldhuizen. "Arrays in Blitz++". In: *Computing in Object-Oriented Parallel Environments*. Ed. by Denis Caromel, Rodney R. Oldehoeft, and Marydell Tholburn. Berlin, Heidelberg: Springer Berlin Heidelberg, 1998, pp. 223–230. ISBN: 978-3-540-49372-3.

[186] David Abrahams. "Exception-Safety in Generic Components". In: *Generic Programming*. Ed. by Mehdi Jazayeri, Rüdiger G. K. Loos, and David R. Musser. Berlin, Heidelberg: Springer Berlin Heidelberg, 2000, pp. 69–79. ISBN: 978-3-540-39953-7.

[187] Krzysztof Czarnecki et al. "Generative Programming and Active Libraries". In: *Generic Programming*. Ed. by Mehdi Jazayeri, Rüdiger G. K. Loos, and David R. Musser. Berlin, Heidelberg: Springer Berlin Heidelberg, 2000, pp. 25–39. ISBN: 978-3-540-39953-7. URL: http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.646.8822&rep=rep1&type=pdf.

[188] Thierry Géraud and Alexandre Duret-Lutz. "Generic programming redesign of patterns". In: *Proceedings of the 5th European Conference on Pattern Languages of Programs (EuroPLoP)*. Irsee, Germany: UVK, Univ. Verlag, Konstanz, July 2000, pp. 283–294.

[189] Phillip James Plauger et al. *C++ standard template library*. Prentice Hall PTR, 2000.

[190] Todd Veldhuizen. "Techniques for scientific C++". In: *Computer science technical report* 542 (2000), p. 60.

[191] Terry S Yoo et al. "Engineering and algorithm design for an image processing API: a technical report on ITK-the insight toolkit". In: *Studies in health technology and informatics* (2002), pp. 586–592.

[192] L. Ibanez et al. *The ITK Software Guide*. First. ISBN 1-930934-10-6. Kitware, Inc. 2003. URL: http://www.itk.org/ItkSoftwareGuide.pdf.

[193] Alexander A Stepanov and Matthew A Marcus. *Notes on the Foundations of Programming*. Citeseer, 2004.

[194] L. Ibanez et al. *The ITK Software Guide*. Second. ISBN 1-930934-15-7. Kitware, Inc. 2005. URL: http://www.itk.org/ItkSoftwareGuide.pdf.

[195] Lionel Moisan. "Modeling and image processing". In: *Lecture Notes of ENS Cachan edition* (2005).

[196] Bjarne Stroustrup. "The design of C++ 0x". In: *C/C++ Users Journal* 23.5 (2005), p. 7.

[197] S Wagner and M Affenzeller. "A Generic and Extensible Optimization Environment". In: *Adaptive and Natural Computing Algorithms, Springer Computer Science* (2005), pp. 538–541.

[198] Gabriel Dos Reis and Bjarne Stroustrup. "Specifying C++ Concepts". In: *SIGPLAN Not.*
       41.1 (Jan. 2006), pp. 295–308. ISSN: 0362-1340. DOI: 10.1145/1111320.1111064. URL:
       https://doi.org/10.1145/1111320.1111064.

[199] Jaakko Järvi et al. "Algorithm Specialization in Generic Programming: Challenges of
       Constrained Generics in C++". In: *Proceedings of the 27th ACM SIGPLAN Conference on
       Programming Language Design and Implementation*. PLDI '06. Ottawa, Ontario, Canada:
       ACM, 2006, pp. 272–282. ISBN: 1-59593-320-4. DOI: 10.1145/1133981.1134014. URL: http:
       //doi.acm.org/10.1145/1133981.1134014.

[200] Jaakko Järvi et al. "Algorithm Specialization in Generic Programming: Challenges of
       Constrained Generics in C++". In: *SIGPLAN Not.* 41.6 (June 2006), pp. 272–282. ISSN:
       0362-1340. DOI: 10.1145/1133255.1134014. URL: http://doi.acm.org/10.1145/1133255.
       1134014.

[201] Jeremy Gibbons. "Datatype-Generic Programming". In: *Datatype-Generic Programming:
       International Spring School, SSDGP 2006, Nottingham, UK, April 24-27, 2006, Revised
       Lectures* 4719 (2007), p. 1.

[202] Kenneth E. Iverson. "Notation as a Tool of Thought". In: *ACM Turing Award Lec-
       tures*. New York, NY, USA: Association for Computing Machinery, 2007, p. 1979. ISBN:
       9781450310499. URL: https://doi.org/10.1145/1283920.1283935.

[203] Travis E. Oliphant. "Multidimensional Iterators in NumPy". In: *Beautiful code*. Ed. by
       Andrew Oram and Greg Wilson. O'reilly Sebastopol, CA, 2007. Chap. 19.

[204] Bjarne Stroustrup. "Evolving a Language in and for the Real World: C++ 1991-2006".
       In: *Proceedings of the Third ACM SIGPLAN Conference on History of Programming
       Languages*. HOPL III. San Diego, California: Association for Computing Machinery, 2007,
       pp. 4-1-4–59. ISBN: 9781595937667. DOI: 10.1145/1238844.1238848. URL: https://doi.org/
       10.1145/1238844.1238848.

[205] Fernand Meyer and Jean Stawiaski. "Morphology on graphs and minimum spanning trees".
       In: *Proc. of the Intl. Symp. on Mathematical Morphology (ISMM)*. Vol. 5720. LNCS.
       Springer, 2009, pp. 161–170.

[206] Todd Mytkowicz et al. "Producing Wrong Data without Doing Anything Obviously Wrong!"
       In: *SIGPLAN Not.* 44.3 (Mar. 2009), pp. 265–276. ISSN: 0362-1340. DOI: 10.1145/1508284.
       1508275. URL: https://doi.org/10.1145/1508284.1508275.

[207] Todd Mytkowicz et al. "Producing Wrong Data without Doing Anything Obviously Wrong!"
       In: *Proceedings of the 14th International Conference on Architectural Support for Program-
       ming Languages and Operating Systems*. ASPLOS XIV. Washington, DC, USA: Associ-
       ation for Computing Machinery, 2009, pp. 265–276. ISBN: 9781605584065. DOI: 10.1145/
       1508244.1508275. URL: https://doi.org/10.1145/1508244.1508275.

[208] Pearu Peterson. "F2PY: a tool for connecting Fortran and Python programs". In: *Inter-
       national Journal of Computational Science and Engineering* 4.4 (2009), pp. 296–305.

[209] Thierry Géraud, Hugues Talbot, and Marc Van Droogenbroeck. "Algorithms for Math-
       ematical Morphology". In: *Mathematical Morphology—From Theory to Applications*. Ed.
       by Laurent Najman and Hugues Talbot. ISTE. Wiley-ISTE, July 2010, pp. 323–353.
       ISBN: 978-1-84821-215-2. URL: http://eu.wiley.com/WileyCDA/WileyTitle/productCd-
       1848212151.html.

[210] P Kobalicek. *asmjit-complete x86/x64 JIT assembler for C++ language*. 2011.

[211] Ullrich Köthe. *Generic programming for computer vision: The vigra computer vision li-
       brary*. 2011.

[212]   Thierry Géraud, Roland Levillain, and Guillaume Lazzara. *The Milena Image Processing Library*. IPOL meeting, ENS Cachan, France. June 2012. URL: https://www.lrde.epita.fr/~theo/talks/geraud.2012.ipol_talk.pdf.

[213]   David R Musser and Alexander A Stepanov. *The Ada® generic library: linear list processing packages*. Springer Science & Business Media, 2012.

[214]   Darren Redfern and Colin Campbell. *The MATLAB® 5 Handbook*. Springer Science & Business Media, 2012.

[215]   Bjarne Stroustrup et al. "A Concept Design for the STL". In: *C++ Standards Committee Papers, Technical Report N* 3351 (2012), pp. 12–0041.

[216]   Charlie Curtsinger and Emery D. Berger. "STABILIZER: Statistically Sound Performance Evaluation". In: *SIGARCH Comput. Archit. News* 41.1 (Mar. 2013), pp. 219–228. ISSN: 0163-5964. DOI: 10.1145/2490301.2451141. URL: https://doi.org/10.1145/2490301.2451141.

[217]   Charlie Curtsinger and Emery D. Berger. "STABILIZER: Statistically Sound Performance Evaluation". In: *Proceedings of the Eighteenth International Conference on Architectural Support for Programming Languages and Operating Systems*. ASPLOS '13. Houston, Texas, USA: Association for Computing Machinery, 2013, pp. 219–228. ISBN: 9781450318709. DOI: 10.1145/2451116.2451141. URL: https://doi.org/10.1145/2451116.2451141.

[218]   L. Najman and H. Talbot. *Mathematical Morphology: From Theory to Applications*. ISTE. Wiley, 2013. ISBN: 9781118600856. URL: https://books.google.fr/books?id=9FUlX8YrRvMC.

[219]   Charlie Curtsinger and Emery D. Berger. "Coz: Finding Code That Counts with Causal Profiling". In: *Proceedings of the 25th Symposium on Operating Systems Principles*. SOSP '15. Monterey, California: Association for Computing Machinery, 2015, pp. 184–197. ISBN: 9781450338349. DOI: 10.1145/2815400.2815409. URL: https://doi.org/10.1145/2815400.2815409.

[220]   Serge Guelton et al. "Pythran: Enabling static optimization of scientific python programs". In: *Computational Science & Discovery* 8.1 (2015), p. 014001.

[221]   Hans J Johnson, Matthew M McCormick, and Luis Ibanez. *Template:The ITK Software Guide Book 1: Introduction and Development Guidelines-Volume 1*. Kitware, Inc., 2015.

[222]   Hans J Johnson, Matthew M McCormick, and Luis Ibanez. *Template:The ITK Software Guide Book 2: Design and Functionality-Volume 2*. Kitware, Inc., 2015.

[223]   Sebastian Nanz and Carlo A. Furia. "A Comparative Study of Programming Languages in Rosetta Code". In: *2015 IEEE/ACM 37th IEEE International Conference on Software Engineering*. Vol. 1. May 2015, pp. 778–788. DOI: 10.1109/ICSE.2015.90.

[224]   Yongchao Xu, Thierry Géraud, and Laurent Najman. "Connected filtering on tree-based shape-spaces". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 38.6 (2015), pp. 1126–1140.

[225]   Jon Bentley. *Programming pearls*. Addison-Wesley Professional, 2016.

[226]   Ampomah Ernest Kwame, Ezekiel Mensah Martey, and Abilimi Gilbert Chris. "Qualitative assessment of compiled, interpreted and hybrid programming languages". In: *Communications* 7.7 (2017), pp. 8–13.

[227]   Holly Moore. *MATLAB for Engineers*. Pearson, 2017.

[228]   M. A. Oezkan et al. "A Highly Efficient and Comprehensive Image Processing Library for C++-based High-Level Synthesis". In: *FSP 2017; Fourth International Workshop on FPGAs for Software Programmers*. Sept. 2017, pp. 1–10.

[229] Sean Parent. *2018 : Generic Programming*. Talk available on youtube at https://youtu.be/FtZEU9zv9eM. 2018. URL: https://codedive.pl/2018/generic-programming.

[230] Louis Dionne. *Metabench*. Version head. Sept. 16, 2019. URL: https://github.com/ldionne/metabench/.

[231] Bertrand Kerautret et al., eds. *Reproducible Research in Pattern Recognition — Second International Workshop, RRPR 2018, Beijing, China, August 20, 2018, Revised Selected Papers.* Vol. 11455. Lecture Notes in Computer Science. Springer, 2019.

[232] R. Liu et al. "SunwayImg: A Parallel Image Processing Library for the Sunway Many-Core Processor". In: *IEEE Access* 7 (2019), pp. 128555–128569. ISSN: 2169-3536. DOI: 10.1109/ACCESS.2019.2939940.

[233] Conrad Sanderson and Ryan Curtin. "Practical Sparse Matrices in C++ with Hybrid Storage and Template-Based Expression Optimisation". In: *Mathematical and Computational Applications* 24.3 (2019), p. 70.

[234] Richard Smith. *N4849: Working Draft, Standard for Programming Language C++*. Tech. rep. https://wg21.link/n4849. International Organization for Standardization, Jan. 2020.

[235] Bjarne Stroustrup. "Thriving in a Crowded and Changing World: C++ 2006-2020". In: *Proc. ACM Program. Lang.* 4.HOPL (June 2020). DOI: 10.1145/3386320. URL: https://doi.org/10.1145/3386320.

Appendix B

Concepts & archetypes

B.1 The fundamentals

B.1.1 Value

Concept	Modeling type	Inherit behavior from	Instance of type
Value	`Val`	$\varnothing$	`val`
ComparableValue	`CmpVal`	Value	`cmp_val1, cmp_val2`
OrderedValue	`OrdVal`	ComparableValue	`ord_val1, ord_val2`

Concept	Expression	Return Type	Description
Value	`std::semiregular<Val>`	`std::true_type`	`Val` is a semiregular type. It can be: copied, moved, swapped, and default constructed.
ComparableValue	`std::regular<CmpVal>`	`std::true_type`	`CmpVal` is a regular type. It is a semiregular type that is equality comparable.
	`cmp_val1 == cmp_val2`	`boolean`	Supports equality comparison
OrderedValue	`std::totally_ordered<OrdVal>`	`std::true_type`	`CmpVal` is a totally ordered as well as a regular type. Additionally the expressions must be equality preserving.
	`ord_val1 < ord_val2`	`boolean`	Supports inequality comparisons
	`ord_val1 <= ord_val2, ...`	`boolean`	

Table B.1: Concepts Value: expressions

Concept table

Concept code

```cpp
// Value
template <typename Val>
concept Value = std::semiregular<Val>;

// ComparableValue
template <typename RegVal>
concept ComparableValue =
  std::regular<RegVal>;

// OrderedValue
template <typename STORegVal>
concept OrderedValue =
  std::regular<STORegVal> &&
  std::totally_ordered<STORegVal>;
```

Archetype code

```cpp
struct Value
{
};
```

```cpp
struct ComparableValue
{
};
bool operator==(const ComparableValue&, const ComparableValue&);
bool operator!=(const ComparableValue&, const ComparableValue&);


struct OrderedValue
{
};
bool operator==(const OrderedValue&, const OrderedValue&);
bool operator!=(const OrderedValue&, const OrderedValue&);
bool operator<(const OrderedValue&, const OrderedValue&);
bool operator>(const OrderedValue&, const OrderedValue&);
bool operator<=(const OrderedValue&, const OrderedValue&);
bool operator>=(const OrderedValue&, const OrderedValue&);

static_assert(mln::concepts::Value<Value>, "Value archetype does not model the Value concept!");
static_assert(mln::concepts::ComparableValue<ComparableValue>, "ComparableValue archetype does not model the ComparableValue co
static_assert(mln::concepts::OrderedValue<OrderedValue>, "OrderedValue archetype does not model the OrderedValue concept!");
```

B.1.2 Point

Concept	Modeling type	Inherit behavior from	Instance of type
Point	Pnt	∅	pnt1, pnt2

Concept	Expression	Return Type	Description
Point	std::regular<Pnt>	std::true_type	Pnt is a regular type. It can be: copied, moved, swapped, and default constructed. It also is equality comparable.
	std::totally_ordered<OrdVal>	std::true_type	Pnt is a totally ordered as well as a regular type. Additionally the expressions must be equality preserving.
	pnt1 < pnt2	boolean	supports inequality comparisons
	pnt1 <= pnt2, ...	boolean	

Table B.2: Concepts Point: expressions

Concept table

Concept code

```cpp
// Point
template <typename P>
concept Point =
  std::regular<P> &&
  std::totally_ordered<P>;
```

Archetype code

```cpp
struct Point final
{
};

bool operator==(const Point&, const Point&);
bool operator!=(const Point&, const Point&);
bool operator<(const Point&, const Point&);
bool operator>(const Point&, const Point&);
bool operator<=(const Point&, const Point&);
bool operator>=(const Point&, const Point&);

static_assert(mln::concepts::Point<Point>, "Point archetype does not model the Point concept!");
```

B.1.3   Pixel

Concept	Modeling type	Inherit behavior from	Instance of type
Pixel	`Pix`	∅	`pix`
OutputPixel	`OPix`	Pixel	`opix`

Concept	Definition	Description	Requirement
Pixel	`value_type`	Type of the value contained in the pixel. Cannot be constant or reference.	Models the concept `Value`.
	`reference_type`	Type used to mutate the pixel's value if non-const. Can be a proxy.	Models the concept `std::indirectly_writable` if non-const.
	`point_type`	Type of the pixel's point.	Models the concept `Point`

Table B.3: Concepts Pixel: definitions

Type	Instance of type
`Pix::value_type`	`val`
`Pix::point_type`	`pnt`

Concept	Expression	Return Type	Description
Pixel	`pix.val()`	`Pix::reference_type`	Access the pixel's value for read and/or write purpose.
	`pix.point()`	`Pix::point_type`	Read the pixel's point.
	`pix.shift(pnt)`	`void`	Shift pixel's point coordinate base on `pnt`'s coordinates.
OutputPixel	`opix.val() = val`	`void`	Mutate pixel's value.

Table B.4: Concepts Pixel: expressions

Concept table

Concept code

```cpp
// Pixel
template <class Pix>
concept Pixel =
  std::is_base_of_v<mln::details::Pixel<Pix>, Pix> &&
  std::copy_constructible<Pix> &&
  std::move_constructible<Pix> &&
  requires {
    typename pixel_value_t<Pix>;
    typename pixel_reference_t<Pix>;
    typename pixel_point_t<Pix>;
  } &&
  std::semiregular<pixel_value_t<Pix>> &&
  Point<pixel_point_t<Pix>> &&
  !std::is_const_v<pixel_value_t<Pix>> &&
  !std::is_reference_v<pixel_value_t<Pix>> &&
  requires(const Pix cpix, Pix pix, pixel_point_t<Pix> p) {
    { cpix.point() } -> std::convertible_to<pixel_point_t<Pix>>;
    { cpix.val() }   -> std::convertible_to<pixel_reference_t<Pix>>;
    { pix.shift(p) };
  };

// WritablePixel
template <typename WPix>
concept WritablePixel =
  Pixel<WPix> &&
  requires(const WPix cpix, pixel_value_t<WPix> v) {
    // Not deep-const, view-semantic.
    { cpix.val() = v };
    // Proxy rvalues must not be deep-const on their assignement semantic (unlike tuple...)
    { const_cast<typename WPix::reference const &&>(cpix.val()) = v };
  };

// OutputPixel
template <typename Pix>
concept OutputPixel = detail::WritablePixel<Pix>;
```

Archetype code

```cpp
namespace details
{
  template <class P, class V>
  struct PixelT
  {
    using value_type = V;
    using point_type = P;
    using reference  = const value_type&;

    PixelT()               = delete;
    PixelT(const PixelT&) = default;
    PixelT(PixelT&&)       = default;
    PixelT& operator=(const PixelT&) = delete;
    PixelT& operator=(PixelT&&) = delete;

    point_type point() const;
    reference  val() const;
    void       shift(const P& dp);
  };

  struct OutputPixel : PixelT<Point, Value>
  {
    using reference = Value&;
    reference val() const;
  };


  template <class Pix>
  struct AsPixel : Pix, mln::details::Pixel<AsPixel<Pix>>
  {
  };
} // namespace details

template <class P, class V = Value>
using PixelT      = details::AsPixel<details::PixelT<P, V>>;
using Pixel       = PixelT<Point, Value>;
using OutputPixel = details::AsPixel<details::OutputPixel>;

static_assert(mln::concepts::Pixel<Pixel>, "Pixel archetype does not model the Pixel concept!");
static_assert(mln::concepts::OutputPixel<OutputPixel>, "OutputPixel archetype does not model the OutputPixel concept!");
```

B.1.4   Ranges

Concept	Modeling type	Inherit behavior from	Instance of type
MDRange	`MDRng`	∅	`mdrng`
OutputMDRange	`OMDRng`	MDRange	`omdrng`
ReversibleMDRange	`RMDRng`	MDRange	`rmdrng`

Concept	Definition	Description	Requirement
MDRange	`value_type`	Type of the value contained in the range. Cannot be constant or reference.	Models the concept `Value`.
	`reference_type`	Type used to mutate the pixel's value if non-const. Can be a proxy.	Models the concept `std::indirectly_writable` if non-const.

Table B.5: Concepts Ranges: definitions

Type	Instance of type
`std::ranges::range_value_t<MDRng>`	`val`

Concept	Expression	Return Type	Description
MDRange	`mdrng.begin()`	`unspecified`	Return a forward iterator allowing a traversing of the range.
	`mdrng.end()`	`unspecified`	Return a sentinel allowing to know when the end is reached.
OutputMDRange	`auto it = omdrng.begin()` `*it++ = val`	`void`	Mutate a value inside the range then increment the iterator's position
ReversibleMDRange	`rmdrng.rbegin()`	`unspecified`	Return a forward iterator allowing a traversing of the range starting from the end.
	`rmdrng.rend()`	`unspecified`	Return a sentinel allowing to know when the end is reached.

Table B.6: Concepts Ranges: expressions

Concept table

Concept code

```cpp
template <class C>
concept MDCursor =
  std::ranges::detail::forward_cursor<C> &&
  std::ranges::detail::forward_cursor<std::ranges::detail::begin_cursor_t<C>> &&
  requires (C c)
  {
    { c.read() } -> std::ranges::forward_range;
    c.end_cursor();
  };

template <class C>
concept NDCursor = std::semiregular<C> &&
  requires (C c)
  {
    { C::rank } -> std::same_as<int>;
    c.read();
    c.move_to_next(0);
    c.move_to_end(0);
  };

template <class C>
concept MDBidirectionalCursor = MDCursor<C> &&
  requires (C c)
  {
    c.move_to_prev();
    c.move_to_prev_line();
  };
```

```cpp
template <class R>
concept MDRange =
  requires(R r)
  {
    { r.rows() } -> std::ranges::forward_range;
    { r.begin_cursor() } -> MDCursor;
    { r.end_cursor() } -> std::same_as<std::ranges::default_sentinel_t>;
  };

template <class R>
concept MDBidirectionalRange = MDRange<R> &&
  requires (R r)
  {
    { r.rrows() } -> std::ranges::forward_range;
    { r.rbegin_cursor() } -> MDCursor;
    { r.rend_cursor() } -> std::same_as<std::ranges::default_sentinel_t>;
  };

template <class R>
concept mdrange = MDRange<R> || std::ranges::range<R>;

template <class R, class V>
concept output_mdrange = mdrange<R> && std::ranges::output_range<mdrange_row_t<R>, V>;

template <class R>
concept reversible_mdrange = MDBidirectionalRange<R> || std::ranges::bidirectional_range<R>;
```

Archetype code

```cpp
// TODO
```

B.1.5   Domain

Concept	Modeling type	Inherit behavior from	Instance of type
Domain	`MDRng`	MDRange	`dom`
SizedDomain	`OMDRng`	Domain	`sdom`
ShapedDomain	`RMDRng`	SizedDomain	`shdom`

Table B.7: Concepts Domain: definitions

Type	Instance of type
`Dom::value_type`	`pnt`

Concept	Expression	Return Type	Description
Domain	`Point<Dom::value_type>`	`std::true_type`	Domain's value models the `Point` concept
	`dom.has(pnt)`	`bool`	Check if a points is included in the domain.
	`dom.empty()`	`void`	Read the pixel's point.
	`dom.dim()`	`void`	Returns the domain's dimension.
SizedDomain	`sdom.size()`	`unsigned int`	Returns the number of points inside the domain.
ShapedDomain	`shdom.extends()`	`std::forward_range`	Return a range that yields the number of elements for each dimension.

Table B.8: Concepts Domain: expressions

Concept table

Concept code

```cpp
// Domain
template <typename Dom>
concept Domain =
  mln::ranges::mdrange<Dom> &&
  Point<mln::ranges::mdrange_value_t<Dom>> &&
  requires(const Dom cdom, mln::ranges::mdrange_value_t<Dom> p) {
    { cdom.has(p) }   -> std::same_as<bool>;
    { cdom.empty() }  -> std::same_as<bool>;
    { cdom.dim() }    -> std::same_as<int>;
  };


// SizedDomain
template <typename Dom>
concept SizedDomain =
  Domain<Dom> &&
  requires(const Dom cdom) {
    { cdom.size() } -> std::unsigned_integral;
  };

// ShapedDomain
template <typename Dom>
concept ShapedDomain =
  SizedDomain<Dom> &&
  requires(const Dom cdom) {
    { cdom.extents() }  -> std::ranges::forward_range;
  };
```

Archetype code

```cpp
struct Domain
{
  using value_type = Point;
  using reference  = Point&;

  value_type* begin();
  value_type* end();
```

```cpp
  bool has(value_type) const;
  bool empty() const;
  int  dim() const;
};

static_assert(mln::concepts::Domain<Domain>, "Domain archetype does not model the Domain concept!");

struct SizedDomain : Domain
{
  unsigned size() const;
};

static_assert(mln::concepts::SizedDomain<SizedDomain>,
                         "SizedDomain archetype does not model the SizedDomain concept!");

struct ShapedDomain final : SizedDomain
{
  static constexpr std::size_t  ndim = 1;
  value_type                    shape() const;
  std::array<std::size_t, ndim> extents() const;
};

static_assert(mln::concepts::ShapedDomain<ShapedDomain>,
                         "ShapedDomain archetype does not model the ShapedDomain concept!");
```

B.1.6   Image

Concept	Modeling type	Inherit behavior from	Instance of type
Image (InputImage, ForwardImage)	`Img`	∅	`img`
WritableImage	`WImg`	Image	`wimg`

Concept	Definition	Description	Requirement
	`pixel_type`	Type of the image's pixel.	Models the concept `Pixel`.
	`point_type`	Type of the image's point.	Models the concept `Point`.
	`value_type`	Type of the image's value. Cannot be constant or reference	Models the concept `Value`.
	`domain_type`	Type of the image's domain.	Models the concept `Domain`.
	`reference`	Type used to mutate an image pixel's value if non-const	Models the concept `std::indirectly_writable` if non-const.
	`concrete_type`	Image concrete type (that holds data).	Models the concept `Image`.
	`ch_value_type<V>`	Facility to return a new image type. that casts the underlying `value_type` into `V`	Models the concept `Image`.

Table B.9: Concepts Image: definitions (1)

Concept	Expression	Return Type	Description
Image	`img.concretize()`	`std::convertible_to< concrete_type>`	Return a concrete image that holds data.
	`img.ch_value<V>()`	`std::convertible_to< ch_value_type<V> >`	Return an image whose values are casted to `V`.
	`img.domain()`	`std::convertible_to< domain_type>`	Return the image's domain.
	`img.pixels()`	`MDRange`	Return a range that yields all the image pixels.
	`img.values()`		Return a range that yields all the image values.
	`std::convertible_to< std::ranges::ranges_value_t< decltype(img.pixels())>, pixel_type> std::convertible_to< std::ranges::ranges_value_t< decltype(img.values())>, value_type>`	`std::true_type`	Ranges converts to compatible element types.
WritableImage	`wimg.values()`	`OutputMDRange`	Return a range that yields all the image values (mutable).
	`OutputPixel< std::ranges::ranges_value_t< decltype(wimg.pixels())> >`	`std::true_type`	Ranges whose elements are mutable.

Table B.10: Concepts Image: expressions (1)

Concept table

Concept code

```
template <typename I>
concept Image =
  // Minimum constraint on image object
  // Do not requires DefaultConstructible
  std::is_base_of_v<mln::details::Image<I>, I> &&
  std::copy_constructible<I> &&
  std::move_constructible<I> &&
  std::derived_from<image_category_t<I>, forward_image_tag> &&
  requires {
    typename image_pixel_t<I>;
    typename image_point_t<I>;
    typename image_value_t<I>;
```

```cpp
    typename image_domain_t<I>;
    typename image_reference_t<I>;
    typename image_concrete_t<I>;
    typename image_ch_value_t<I, mln::archetypes::Value>;
    // traits
    typename image_indexable_t<I>;
    typename image_accessible_t<I>;
    typename image_extension_category_t<I>;
    typename image_category_t<I>;
    typename image_view_t<I>;
  } &&
  Pixel<image_pixel_t<I>> &&
  Point<image_point_t<I>> &&
  Value<image_value_t<I>> &&
  Domain<image_domain_t<I>> &&
  std::convertible_to<pixel_point_t<image_pixel_t<I>>, image_point_t<I>> &&
  std::convertible_to<pixel_reference_t<image_pixel_t<I>>, image_reference_t<I>> &&
  // Here we don't want a convertible constraint as value_type is the decayed type and should really be the same
  std::same_as<pixel_value_t<image_pixel_t<I>>, image_value_t<I>> &&
  std::common_reference_with<image_reference_t<I>&&, image_value_t<I>&> &&
  std::common_reference_with<image_reference_t<I>&&, image_value_t<I>&&> &&
  std::common_reference_with<image_value_t<I>&&, const image_value_t<I>&> &&
  requires(I ima, const I cima, image_domain_t<I> d, image_point_t<I> p) {
    { cima.template ch_value<mln::archetypes::Value>() }
        -> std::convertible_to<image_ch_value_t<I, mln::archetypes::Value>>;
    { cima.concretize() } -> std::convertible_to<image_concrete_t<I>>;
    { cima.domain() }     -> std::convertible_to<image_domain_t<I>>;
    { ima.pixels() }  -> mln::ranges::mdrange;
    { ima.values() }  -> mln::ranges::mdrange;
    requires std::convertible_to<mln::ranges::mdrange_value_t<decltype(ima.pixels())>, image_pixel_t<I>>;
    requires std::convertible_to<mln::ranges::mdrange_value_t<decltype(ima.values())>, image_value_t<I>>;
  };

namespace detail
{
  // WritableImage
  template <typename I>
  concept WritableImage =
    Image<I> &&
    OutputPixel<image_pixel_t<I>> &&
    requires(I ima) {
    { ima.values() }  -> mln::ranges::output_mdrange<image_value_t<I>>;
      // Check Writability of each pixel of the range
      requires OutputPixel<
                  std::common_type_t<
                    mln::ranges::mdrange_value_t<decltype(ima.pixels())>,
                    image_pixel_t<I>>>;
    };
} // namespace detail


// InputImage
template <typename I>
concept InputImage = Image<I>;


// ForwardImage
template <typename I>
concept ForwardImage = InputImage<I>;
```

Archetype code

```cpp
namespace details
{
  template <class I>
  struct AsImage : I, mln::details::Image<AsImage<I>>
  {
    using I::I;

    using concrete_type = AsImage<typename I::concrete_type>;
    concrete_type concretize() const;
```

```cpp
  template <typename V>
  using ch_value_type = AsImage<typename I::template ch_value_type<V>>;

  template <typename V>
  ch_value_type<V> ch_value() const;
};


struct Image
{
  using pixel_type = archetypes::Pixel;
  using value_type     = pixel_value_t<mln::archetypes::Pixel>;
  using reference      = pixel_reference_t<mln::archetypes::Pixel>;
  using point_type     =  std::ranges::range_value_t<Domain>;
  using domain_type    = Domain;
  using category_type  = forward_image_tag;
  using concrete_type  = Image;

  template <class V>
  using ch_value_type = Image;

  // additional traits
  using extension_category = mln::extension::none_extension_tag;
  using indexable          = std::false_type;
  using accessible         = std::false_type;
  using view               = std::false_type;

  Image()                    = default;
  Image(const Image&) = default;
  Image(Image&&)      = default;
  Image& operator=(const Image&) = default;
  Image& operator=(Image&&) = default;

  domain_type domain() const;


  struct pixel_range
  {
    const pixel_type* begin();
    const pixel_type* end();
  };
  pixel_range pixels();


  struct value_range
  {
    const value_type* begin();
    const value_type* end();
  };

  value_range values();
};

struct OutputImage : Image
{
  using pixel_type = archetypes::OutputPixel;
  using reference  = pixel_reference_t<mln::archetypes::OutputPixel>;

  struct pixel_range
  {
    const pixel_type* begin();
    const pixel_type* end();
  };

  pixel_range pixels();

  struct value_range
  {
    value_type* begin();
    value_type* end();
  };

  value_range values();
```

```
  };
} // namespace details


using Image        = details::AsImage<details::Image>;
using ForwardImage = Image;
using WritableImage = details::AsImage<details::OutputImage>;
```

B.2   Advanced way to access image data

B.2.1   Index

Concept	Modeling type	Inherit behavior from	Instance of type
Index	`Idx`	$\varnothing$	`idx, idy`

Concept	Expression	Return Type	Description
Index	`std::signed_integral<Idx>`	`std::true_type`	Idx is a signed integral arithmetic type
	`idx + idy, idx - idy, ...`	`Idx`	Supports all trivial arithmetical operations

Table B.11: Concepts Index: expressions

Concept table

Concept code

```cpp
// Index
template <typename Idx>
concept Index = std::signed_integral<Idx>;
```

Archetype code

```cpp
using Index = int;

static_assert(mln::concepts::Index<Index>, "Index archetype does not model the Index concept!");
```

B.2.2 Indexable image

Concept	Modeling type	Inherit behavior from	Instance of type
IndexableImage	`IdxImg`	Image	`idximg`
WritableIndexableImage	`WIdxImg`	IndexableImage, WritableImage	`widximg`

Concept	Definition	Description	Requirement
IndexableImage	`index_type`	Type of the image's buffer index.	Models the concept `Index`.

Table B.12: Concepts Image: definitions (2)

Type	Instance of type
`Img::value_type`	`val`
`IdxImg::index_type`	`k`

Concept	Expression	Return Type	Description
IndexableImage	`idximg[k]`	`std::same_as<reference>`	Access a value at a given index.
Writable IndexableImage	`widximg[k] = val`	`void`	Mutate a value at a given index.

Table B.13: Concepts Image: expressions (2)

Concept table

Concept code

```cpp
// IndexableImage
template <typename I>
concept IndexableImage =
  Image<I> &&
  requires {
    typename image_index_t<I>;
  } &&
  image_indexable_v<I> &&
  requires (I ima, image_index_t<I> k) {
    { ima[k] }  -> std::same_as<image_reference_t<I>>; // For concrete image it returns a const_reference
  };


namespace detail
{
  // WritableIndexableImage
  template <typename I>
  concept WritableIndexableImage =
    WritableImage<I> &&
    IndexableImage<I> &&
    requires(I ima, image_index_t<I> k, image_value_t<I> v) {
      { ima[k] = v }
    };
} // namespace detail
```

Archetype code

```cpp
namespace details
{

  struct IndexableImage : Image
  {
    using index_type = int;
    using indexable  = std::true_type;

    using concrete_type = IndexableImage;

    template <class V>
    using ch_value_type = IndexableImage;
```

```cpp
    reference operator[](index_type);
  };

  struct OutputIndexableImage : OutputImage
  {
    using index_type = int;
    using indexable  = std::true_type;

    using concrete_type = OutputIndexableImage;

    template <class V>
    using ch_value_type = OutputIndexableImage;

    reference operator[](index_type);
  };

} // namespace details

using IndexableImage          = details::AsImage<details::IndexableImage>;
using WritableIndexableImage  = details::AsImage<details::OutputIndexableImage>;
```

B.2.3 Accessible image

Concept	Modeling type	Inherit behavior from	Instance of type
AccessibleImage	`AccImg`	Image	`accimg`
WritableAccessibleImage	`WAccImg`	AccessibleImage, WritableImage	`waccimg`

Table B.14: Concepts Image: definitions (3)

Type	Instance of type
`Img::point_type`	`pnt`

Concept	Expression	Return Type	Description
AccessibleImage	`accimg(pnt)`	`std::same_as<reference>`	Access a value for a given point.
	`accimg.at(pnt)`	`std::same_as<reference>`	Access a value for a given point. No bound checking.
	`accimg.pixel(pnt)`	`std::same_as<pixel_type>`	Access a pixel for a given point.
	`accimg.pixel_at(pnt)`	`std::same_as<pixel_type>`	Access a pixel for a given point. No bound checking.
Writable AccessibleImage	`img(pnt) = val`	`void`	Mutate a value at a given point.
	`waccimg.at(pnt) = val`	`void`	Mutate a value at a given point. No bound checking.
	`OutputPixel<decltype( waccimg.pixel(pnt))>`	`std::true_type`	The returned pixel models `OutputPixel`.
	`OutputPixel<decltype( waccimg.pixel_at(pnt))>`	`std::true_type`	The returned pixel models `OutputPixel`. No bound checking.

Table B.15: Concepts Image: expressions (3)

Concept table

Concept code

```cpp
// AccessibleImage
template <typename I>
concept AccessibleImage =
  Image<I> &&
  image_accessible_v<I> &&
  requires (I ima, image_point_t<I> p) {
    { ima(p) }             -> std::same_as<image_reference_t<I>>; // For concrete image it returns a const_reference
    { ima.at(p) }          -> std::same_as<image_reference_t<I>>; // idem
    { ima.pixel(p) }    -> std::same_as<image_pixel_t<I>>; // For concrete image pixel may propagate constness
    { ima.pixel_at(p) } -> std::same_as<image_pixel_t<I>>; // idem
  };


namespace detail
{
  // WritableAccessibleImage
  template <typename I>
  concept WritableAccessibleImage =
    detail::WritableImage<I> &&
    AccessibleImage<I> &&
    requires(I ima, image_point_t<I> p, image_value_t<I> v) {
      { ima(p) = v };
      { ima.at(p) = v };

      requires OutputPixel<decltype(ima.pixel(p))>;
      requires OutputPixel<decltype(ima.pixel_at(p))>;
    };
} // namespace detail
```

Archetype code

```cpp
namespace details {
  struct OutputAccessibleImage : OutputImage
  {
    using accessible    = std::true_type;
    using concrete_type = OutputAccessibleImage;

    template <class V>
    using ch_value_type = OutputAccessibleImage;


    reference        operator()(point_type);
    reference        at(point_type);
    pixel_type pixel(point_type);
    pixel_type pixel_at(point_type);
  };


  struct AccessibleImage : Image
  {
    using accessible    = std::true_type;
    using concrete_type = AccessibleImage;

    template <class V>
    using ch_value_type = AccessibleImage;

    reference        operator()(point_type);
    reference        at(point_type);
    pixel_type pixel(point_type);
    pixel_type pixel_at(point_type);
  };
} // namespace details

using AccessibleImage           = details::AsImage<details::AccessibleImage>;
using WritableAccessibleImage   = details::AsImage<details::OutputAccessibleImage>;
```

B.2.4 Indexable and accessible image

Concept	Modeling type	Inherit behavior from	Instance of type
IndexableAnd AccessibleImage	IdxAccImg	IndexableImage, AccessibleImage	idxaccimg
WritableIndexable AndAccessibleImage	WIdxAccImg	IndexableAndAccessibleImage, WritableIndexableImage, WritableAccessibleImage	widxaccimg

Table B.16: Concepts Image: definitions (4)

Concept	Expression	Return Type	Description
IndexableAnd AccessibleImage	`img.point_at_index(k)`	`point_type`	Get the point corresponding to the given index.
	`idxaccimg.index_of_point(pnt)`	`index_type`	Get the linear index (offset in the buffer) of multi-dimensional point.
	`idxaccimg.delta_index(pnt)`	`index_type`	Get the linear index offset for the given point.

Table B.17: Concepts Image: expressions (4)

Concept table

Concept code

```cpp
// IndexableAndAccessibleImage
template <typename I>
concept IndexableAndAccessibleImage =
  IndexableImage<I> &&
  AccessibleImage<I> &&
  requires (const I cima, image_index_t<I> k, image_point_t<I> p) {
    { cima.point_at_index(k) }  -> std::same_as<image_point_t<I>>;
    { cima.index_of_point(p) }  -> std::same_as<image_index_t<I>>;
    { cima.delta_index(p) }     -> std::same_as<image_index_t<I>>;
  };


namespace detail
{
  // WritableIndexableAndAccessibleImage
  template <typename I>
  concept WritableIndexableAndAccessibleImage =
    IndexableAndAccessibleImage<I> &&
    detail::WritableImage<I> &&
    detail::WritableIndexableImage<I>;
} // namespace detail
```

Archetype code

```cpp
namespace details {
  struct OutputIndexableAndAccessibleImage : OutputAccessibleImage
  {
    using index_type = int;
    using indexable  = std::true_type;

    using concrete_type = OutputIndexableAndAccessibleImage;

    template <class V>
    using ch_value_type = OutputIndexableAndAccessibleImage;


    reference  operator[](index_type);
    point_type point_at_index(index_type) const;
    index_type index_of_point(point_type) const;
    index_type delta_index(point_type) const;
```

```cpp
  };

  struct IndexableAndAccessibleImage : AccessibleImage
  {
    using index_type = int;
    using indexable  = std::true_type;

    using concrete_type = IndexableAndAccessibleImage;

    template <class V>
    using ch_value_type = IndexableAndAccessibleImage;

    reference   operator[](index_type);
    point_type point_at_index(index_type) const;
    index_type index_of_point(point_type) const;
    index_type delta_index(point_type) const;
  };
} // namespace details

using IndexableAndAccessibleImage         = details::AsImage<details::IndexableAndAccessibleImage>;
using WritableIndexableAndAccessibleImage = details::AsImage<details::OutputIndexableAndAccessibleImage>;
```

B.2.5 Bidirectional image

Concept	Modeling type	Inherit behavior from	Instance of type
BidirectionalImage	BidirImg	Image	bidirimg
WritableBidirectionalImage	WBidirImg	BidirectionalImage, WritableImage	wbidirimg

Table B.18: Concepts Image: definitions (5)

Type	Instance of type
Img::point_type	pnt
Img::value_type	val
IdxImg::index_type	k
int	dim

Concept	Expression	Return Type	Description
BidirectionalImage	bidirimg.pixels()	ReversibleMDRange	Return a reversible range that yields all the image pixels.
	bidirimg.values()		Return a reversible range that yields all the image values.

Table B.19: Concepts Image: expressions (5)

Concept table

Concept code

```cpp
namespace detail
{
  // WritableBidirectionalImage
  template <typename I>
  concept WritableBidirectionalImage =
    WritableImage<I> &&
    BidirectionalImage<I>;
} // namespace detail

// RawImage (not contiguous, stride = padding)
template <typename I>
concept RawImage =
  IndexableAndAccessibleImage<I> &&
  BidirectionalImage<I> &&
  std::derived_from<image_category_t<I>, raw_image_tag> &&
  requires (I ima, const I cima, int dim) {
    { ima.data() }        -> std::convertible_to<const image_value_t<I>*>; // data() may be proxied by a view
    { cima.stride(dim) } -> std::same_as<std::ptrdiff_t>;
  };
```

Archetype code

```cpp
namespace details {
  struct BidirectionalImage : Image
  {
    using category_type = bidirectional_image_tag;

    struct pixel_range
    {
      const pixel_type* begin();
      const pixel_type* end();
      pixel_range        reversed();
    };

    pixel_range pixels();

    struct value_range
    {
      const value_type* begin();
```

```cpp
      const value_type* end();
      value_range       reversed();
    };

    value_range values();
  };

  struct OutputBidirectionalImage : BidirectionalImage
  {
    using pixel_type = archetypes::OutputPixel;
    using reference      = pixel_reference_t<mln::archetypes::OutputPixel>;

    struct value_range
    {
      value_type* begin();
      value_type* end();
      value_range reversed();
    };
    value_range values();


    struct pixel_range
    {
      const pixel_type* begin();
      const pixel_type* end();
      pixel_range           reversed();
    };

    pixel_range pixels();
  };
} // namespace details

using BidirectionalImage = details::AsImage<details::BidirectionalImage>;
using OutputBidirectionalImage = details::AsImage<details::OutputBidirectionalImage>;
```

B.2.6 Raw image

Concept	Modeling type	Inherit behavior from	Instance of type
RawImage	`RawImg`	IndexableAndAccessibleImage, BidirectionalImage	`rawimg`
WritableRawImage	`WRawImg`	RawImage, WritableIndexableImage, WritableBidirectionalImage	`wrawimg`

Table B.20: Concepts Image: definitions (6)

Concept	Expression	Return Type	Description
RawImage	`rawimg.data()`	`std::convertible_to< const value_type*>`	Get a constant pointer to the first element of the domain.
	`rawimg.stride(dim)`	`std::ptrdiff_t`	Get the stride (in number of elements) between two consecutive elements in the given `dim`.
Writable RawImage	`img.data()`	`std::convertible_to< value_type*>`	Get a pointer to the first element of the domain.
	`*(wrawimg.data() + k) = val`	`void`	Access an element from the data buffer at index `k` and mutate it to `val`.
WithExtension Image	`wextimg.extension()`	`std::convertible_to< extension_type>`	Get the extension of the image.

Table B.21: Concepts Image: expressions (6)

Concept table

Concept code

```cpp
// RawImage (not contiguous, stride = padding)
template <typename I>
concept RawImage =
  IndexableAndAccessibleImage<I> &&
  BidirectionalImage<I> &&
  std::derived_from<image_category_t<I>, raw_image_tag> &&
  requires (I ima, const I cima, int dim) {
    { ima.data() }        -> std::convertible_to<const image_value_t<I>*>; // data() may be proxied by a view
    { cima.stride(dim) } -> std::same_as<std::ptrdiff_t>;
  };


namespace detail
{
  // WritableRawImage
  template <typename I>
  concept WritableRawImage =
    WritableIndexableAndAccessibleImage<I> &&
    WritableBidirectionalImage<I> &&
    RawImage<I> &&
    requires(I ima, image_value_t<I> v, image_index_t<I> k) {
      { ima.data() }        -> ::concepts::convertible_to<image_value_t<I>*>;
      { *(ima.data() + k) = v };
    };
} // namespace detail
```

Archetype code

```cpp
namespace details {
  struct RawImage : IndexableAndAccessibleImage
  {
    using category_type = raw_image_tag;
    using pixel_range   = BidirectionalImage::pixel_range;
    using value_range   = BidirectionalImage::value_range;
```

```cpp
    pixel_range pixels();
    value_range values();


    const value_type* data() const;
    std::ptrdiff_t    strides(int) const;
  };

  struct OutputRawImage : OutputIndexableAndAccessibleImage
  {
    using category_type = raw_image_tag;
    using pixel_range   = OutputBidirectionalImage::pixel_range;
    using value_range   = OutputBidirectionalImage::value_range;

    pixel_range pixels();
    value_range values();

    value_type*    data() const;
    std::ptrdiff_t strides(int) const;
  };
} // namespace details

using RawImage          = details::AsImage<details::RawImage>;
using OutputRawImage    = details::AsImage<details::OutputRawImage>;
```

B.3 Local algorithm concepts: structuring elements and extensions

B.3.1 Structuring element

Concept	Modeling type	Inherit behavior from	Instance of type
StructuringElement	`SE, Pnt`	$\varnothing$, Point	`se, pnt`
DecomposableStructuringElement	`DSE, Pnt`	StructuringElement, Point	`dse`
SeparableStructuringElement	`SSE, Pnt`	StructuringElement, Point	`sse`
IncrementalStructuringElement	`ISE, Pnt`	StructuringElement, Point	`ise`

Concept	Definition	Description	Requirement
StructuringElement	`incremental` `decomposable` `separable`	`std::bool_constant`	`std::true_type` if supported `std::false_false` otherwise.

Table B.22: Concepts Structuring Elements: definitions

Type	Instance of type	Requirement
Pix	pix	std::same_as<Pix::point_type, Pnt>
Pnt	pnt	

Concept	Expression	Return Type	Description
StructuringElement	std::regular_invocable<SE, Pnt> std::regular_invocable<SE, Pix>	std::true_type	se can called using std::invoke, is equality preserving and does not modify function object nor arguments.
	se(pnt)	std::forward_range	Return a range that yeilds the neighboring points of pnt.
	se.offsets()	std::forward_range	Return a range that yeilds the neighboring points, in relative coordinates.
	se(pix)	std::forward_range	Return a range that yeilds the neighboring pixels of pix.
	se.radial_extent()	int	Returns the radial extent of the SE, the radius of the disc (square).
	std::convertible_to< std::ranges::range_value_t< decltype(se(pnt))>, Pnt> std::convertible_to< std::ranges::range_value_t< decltype(se.offsets())>, Pnt>	std::true_type	Converts to a compatible point type
	Pixel<std::ranges::range_value_t< decltype(se(pix))> >		Is a range of compatible pixel type
Decomposable StructuringElement	dse.is_decomposable()	std::bool_constant	std::true_type if supported std::false_false otherwise.
	dse.decompose()	std::forward_range	Return a range that yeilds simpler structuring elements.
	StructuringElement< std::ranges::range_value_t< decltype(dse.decompose())> >	std::true_type	Is a range of compatible structuring elements types.
	dse.is_decomposable()	bool	Wether the decompose facility is supported
Separable StructuringElement	dse.is_decomposable()	std::bool_constant	std::true_type if supported std::false_false otherwise.
	sse.separate()	std::forward_range	Return a range that yeilds simpler structuring elements.
	StructuringElement< std::ranges::range_value_t< decltype(sse.separate())> >	std::true_type	Is a range of compatible structuring elements types.
	sse.is_separable()	bool	Wether the separate facility is supported
Incremental StructuringElement	dse.is_decomposable()	std::bool_constant	std::true_type if supported std::false_false otherwise.
	ise.inc()	StructuringElement< SE, Pnt>	Return the next simpler structuring elements.
	ise.dec()	StructuringElement< SE, Pnt>	Return the previous simpler structuring elements.
	ise.is_incremental()	bool	Wether the incremental facility is supported

Table B.23: Concepts Structuring Elements: expressions

Concept table

Concept code

```
namespace details
{
  template <typename SE>
  concept DynamicStructuringElement =
    requires (SE se) {
    { se.radial_extent() }  -> std::same_as<int>;
    };

  constexpr bool implies(bool a, bool b) { return !a || b; }
}
```

```cpp
template <typename SE, typename P>
concept StructuringElement =
  std::convertible_to<SE, mln::details::StructuringElement<SE>> &&
  std::ranges::regular_invocable<SE, P> &&
  std::ranges::regular_invocable<SE, mln::archetypes::PixelT<P>> &&
  requires {
    typename SE::category;
    typename SE::incremental;
    typename SE::decomposable;
    typename SE::separable;
  } &&
  std::convertible_to<typename SE::category, mln::adaptative_neighborhood_tag> &&
  details::implies(std::convertible_to<typename SE::category, mln::dynamic_neighborhood_tag>,
                   details::DynamicStructuringElement<SE>) &&
  requires (SE se, const SE cse, P p, mln::archetypes::PixelT<P> px) {
    { se(p) }           -> std::ranges::forward_range;
    { se(px) }          -> std::ranges::forward_range;
    { cse.offsets() } -> std::ranges::forward_range;

    requires std::convertible_to<std::ranges::range_value_t<decltype(se(p))>, P>;
    requires std::Pixel<std::ranges::range_value_t<decltype(se(px))>>;
    requires std::convertible_to<std::ranges::range_value_t<decltype(cse.offsets())>, P>;
  };

namespace details
{
  template <typename R, typename P>
  concept RangeOfStructuringElement =
    StructuringElement<std::ranges::range_value_t<R>, P>;
}

template <typename SE, typename P>
concept DecomposableStructuringElement =
  StructuringElement<SE, P> &&
  std::convertible_to<typename SE::decomposable, std::true_type> &&
  requires(const SE se) {
    { se.is_decomposable() }  -> std::same_as<bool>;
    { se.decompose() }        -> std::ranges::forward_range;
    requires details::RangeOfStructuringElement<decltype(se.decompose()), P>;
  };

template <typename SE, typename P>
concept SeparableStructuringElement =
  StructuringElement<SE, P> &&
  std::convertible_to<typename SE::separable, std::true_type> &&
  requires(const SE se) {
    { se.is_separable() } -> std::same_as<bool>;
    { se.separate() }     -> std::ranges::forward_range;
    requires details::RangeOfStructuringElement<decltype(se.separate()), P>;
  };

template <typename SE, typename P>
concept IncrementalStructuringElement =
  StructuringElement<SE, P> &&
  std::convertible_to<typename SE::incremental, std::true_type> &&
  requires(const SE se) {
    { se.is_incremental() } -> std::same_as<bool>;
    { se.inc() }  -> StructuringElement<SE, P>;
    { se.dec() }  -> StructuringElement<SE, P>;
  };
```

Achetype code

```cpp
namespace details
{
  template <class P, class Pix>
  requires mln::concepts::Point<P>&& mln::concepts::Pixel<Pix>
  struct StructuringElement
  {
    using category     = adaptative_neighborhood_tag;
    using incremental  = std::false_type;
    using decomposable = std::false_type;
```

```cpp
    using separable    = std::false_type;

      std::ranges::subrange<P*> operator()(P p);

      std::ranges::subrange<Pix*> operator()(Pix px);
      std::ranges::subrange<P*>   offsets() const;
  };


  template <class SE>
  struct AsSE : SE, mln::details::Neighborhood
  helper<AsSE<SE>>
  {
  };
} // namespace details

template <class P = Point, class Pix = PixelT<P>>
using StructuringElement = details::AsSE<details::StructuringElement<P, Pix>>;


namespace details
{
  template <class P, class Pix>
  struct DecomposableStructuringElement : StructuringElement<P, Pix>
  {
    using decomposable = std::true_type;

    bool                                                  is_decomposable() const;
    std::ranges::subrange<mln::archetypes::StructuringElement<P, Pix>*> decompose() const;
  };

  template <class P, class Pix>
  struct SeparableStructuringElement : StructuringElement<P, Pix>
  {
    using separable = std::true_type;

    bool                                                  is_separable() const;
    std::ranges::subrange<mln::archetypes::StructuringElement<P, Pix>*> separate() const;
  };

  template <class P, class Pix>
  struct IncrementalStructuringElement : StructuringElement<P, Pix>
  {
    using incremental = std::true_type;

    bool                              is_incremental() const;
    archetypes::StructuringElement<P, Pix> inc()         const;
    archetypes::StructuringElement<P, Pix> dec()         const;
  };
} // namespace details


template <class P = Point, class Pix = PixelT<P>>
using DecomposableStructuringElement = details::AsSE<details::DecomposableStructuringElement<P, Pix>>;

template <class P = Point, class Pix = PixelT<P>>
using SeparableStructuringElement = details::AsSE<details::SeparableStructuringElement<P, Pix>>;

template <class P = Point, class Pix = PixelT<P>>
using IncrementalStructuringElement = details::AsSE<details::IncrementalStructuringElement<P, Pix>>;
```

B.3.2 Neighborhood

Concept	Modeling type	Inherit behavior from	Instance of type
Neighborhood	`SE, Pnt`	StructuringElement, Point	`se, pnt`

Table B.24: Concepts Neighborhood: definitions

Type	Instance of type	Requirement
Pix	pix	`std::same_as<Pix::point_type, Pnt>`
Pnt	pnt	

Concept	Expression	Return Type	Description
Neighborhood	`se.before(pnt)`	`std::forward_range`	Return a range that yields the points before `pnt`.
	`se.after(pnt)`		Return a range that yields the points after `pnt`.
	`se.before(pix)`		Return a range that yields the pixels before `pix`.
	`se.after(pix)`		Return a range that yields the pixels after `pnt`.
	`std::convertible_to<` `std::ranges::ranges_value_t<` `decltype(se.before(pnt))>, Pnt>` `std::convertible_to<` `std::ranges::ranges_value_t<` `decltype(se.after(pnt))>, Pnt>`	`std::true_type`	Ranges converts to compatible element types.
	`Pixel<` `std::ranges::ranges_value_t<` `decltype(se.before(pix))>, Pix>` `Pixel<` `std::ranges::ranges_value_t<` `decltype(se.after(pix))>, Pix>`		Ranges that are of compatible element types.

Table B.25: Concepts Neighborhood: expressions

Concept table

Concept code

```
template <typename SE, typename P>
concept Neighborhood =
  StructuringElement<SE, P> &&
  requires (SE se, P p, mln::archetypes::PixelT<P> px) {
    { se.before(p) }  -> std::ranges::forward_range;
    { se.after(p) }   -> std::ranges::forward_range;
    { se.before(px) } -> std::ranges::forward_range;
    { se.after(px) }  -> std::ranges::forward_range;

    requires std::convertible_to<std::ranges::range_value_t<decltype(se.before(p))>, P>;
    requires std::convertible_to<std::ranges::range_value_t<decltype(se.after(p))>, P>;
    requires std::Pixel<std::ranges::range_value_t<decltype(se.before(px))>>;
    requires std::Pixel<std::ranges::range_value_t<decltype(se.after(px))>>;
  };
```

Archetype code

```
namespace details
{
  template <class P, class Pix>
  requires mln::concepts::Point<P>&& mln::concepts::Pixel<Pix>
  struct Neighborhood : StructuringElement<P, Pix>
  {
    std::ranges::iterator_range<P*>   before(P p);
    std::ranges::iterator_range<P*>    after(P p);
    std::ranges::iterator_range<Pix*> before(Pix px);
    std::ranges::iterator_range<Pix*> after(Pix px);
  };

  template <class N>
```

```cpp
  struct AsNeighborhood : N, mln::details::Neighborhood<AsNeighborhood<N>>
  {
  };

} // namespace details

template <class P = Point, class Pix = PixelT<P>>
using Neighborhood = details::AsSE<details::Neighborhood<P, Pix>>;
```

B.3.3 Extensions

Concept	Modeling type	Inherit behavior from	Instance of type
Extension	Ext, Pnt	∅, Point	ext, pnt
FillableExtension	FExt, Pnt	Extension, Point	fext
MirrorableExtension	MExt, Pnt	Extension, Point	mext
PeriodizableExtension	PExt, Pnt	Extension, Point	pext
ClampableExtension	CExt, Pnt	Extension, Point	cext
ExtendWithExtension	EwExt, Pnt, U	Extension, Point, Image	ewext, u

Concept	Definition	Description	Requirement
Extension	`value_type`	Type of value contained in the range. Cannot be constant or reference.	Models the concept `Value`.
	`support_fill` `support_mirror` `support_periodize` `support_clamp` `support_extend_with`	`std::bool_constant`	`std::true_type` if supported `std::false_false` otherwise.
ExtendWithExtension	`point_type`	Type of point in the extended image.	Converts to the sub-image points type.

Table B.26: Concepts Extensions: definitions

Type	Instance of type
`SE<Pnt, Pix>`	se
`Ext::value_type`	val
`Ext::point_type`	offset

Concept	Expression	Return Type	Description
Extension	`ext.fit(se)`	`bool`	Wether the extension fit the structuring element.
	`ext.extent()`	`int`	Extension's border width. std::numeric_limits<int>::max() for infinite size.
FillableExtension	`fext.fill(val)`	`void`	Fill the extension with the value `val`.
	`fext.is_fill_supported`	`bool`	Wether the fill facility is supported.
MirrorableExtension	`mext.mirror()`	`void`	Fill the extension with mirrored image's values.
	`mext.is_mirror_supported()`	`bool`	Wether the mirror facility is supported.
PeriodizeableExtension	`mext.periodize()`	`void`	Fill the extension with periodic copies of image's values.
	`mext.is_periodize_supported()`	`bool`	Wether the mirror facility is supported.
ClampableExtension	`mext.clamp()`	`void`	Fill the extension by extending image's values at extremities.
	`mext.is_clamp_supported()`	`bool`	Wether the clamp facility is supported.
ExtendWithExtension	`std::convertible_to<point_type, U::point_type>`	`std::true_type`	Converts to the sub-image points type.
	`mext.extend_with(u, offset)`	`void`	Fill the extension with sub-image `u`'s values starting from `offset`.
	`mext.is_extent_with_supported()`	`bool`	Wether the extend-with-sub-image facility is supported.

Table B.27: Concepts Extensions: expressions

Concept table

Concept code

```cpp
template <typename Ext, typename Pnt>
concept Extension =
  std::is_base_of_v<mln::Extension<Ext>, Ext> &&
  requires {
    typename Ext::support_fill;
    typename Ext::support_mirror;
    typename Ext::support_periodize;
    typename Ext::support_clamp;
    typename Ext::support_extend_with;
  } &&
  Value<typename Ext::value_type> &&
  requires (const Ext cext,
      mln::archetypes::StructuringElement<
        Pnt,
        mln::archetypes::Pixel> se) {
    { cext.fit(se) }       -> std::same_as<bool>;
    { cext.extent() }      -> std::same_as<int>;
  };

template <typename Ext, typename Pnt>
concept FillableExtension =
  Extension<Ext, Pnt> &&
  std::convertible_to<typename Ext::support_fill, std::true_type> &&
  requires {
    typename Ext::value_type;
  } &&
  requires (Ext ext, const Ext cext, const typename Ext::value_type& v) {
    { ext.fill(v) };
    { cext.is_fill_supported() }  -> std::same_as<bool>;
  };

template <typename Ext, typename Pnt>
concept MirrorableExtension =
  Extension<Ext, Pnt> &&
  std::convertible_to<typename Ext::support_mirror, std::true_type> &&
  requires (Ext ext, const Ext cext) {
    { ext.mirror() };
    { cext.is_mirror_supported() }  -> std::same_as<bool>;
  };

template <typename Ext, typename Pnt>
concept PeriodizableExtension =
  Extension<Ext, Pnt> &&
  std::convertible_to<typename Ext::support_periodize, std::true_type> &&
  requires (Ext ext, const Ext cext) {
    { ext.periodize() };
    { cext.is_periodize_supported() }  -> std::same_as<bool>;
  };

template <typename Ext, typename Pnt>
concept ClampableExtension =
  Extension<Ext, Pnt> &&
  std::convertible_to<typename Ext::support_clamp, std::true_type> &&
  requires (Ext ext, const Ext cext) {
    { ext.clamp() };
    { cext.is_clamp_supported() }  -> std::same_as<bool>;
  };

template <typename Ext, typename Pnt, typename U>
concept ExtendWithExtension =
  Extension<Ext, Pnt> &&
  std::convertible_to<typename Ext::support_extend_with, std::true_type> &&
  InputImage<U> &&
  requires {
    typename Ext::point_type;
  } &&
  std::convertible_to<typename U::value_type, typename Ext::value_type> &&
  std::convertible_to<typename Ext::point_type, typename U::point_type> &&
  requires (Ext ext, const Ext cext, U u, typename Ext::point_type offset) {
    { ext.extend_with(u, offset) };
    { cext.is_extend_with_supported() }  -> std::same_as<bool>;
  };
```

Archetype code

```
// TODO
```

B.3.4   Extended image

Concept	Modeling type	Inherit behavior from	Instance of type
WithExtensionImage	`WExtImg`	Image	`wextimg`
ConcreteImage	`CImg`	Image	`cimg`
ViewImage	`VImg`	Image	`vimg`

Concept	Definition	Description	Requirement
WithExtensionImage	`extension_type`	Type of the image's extension.	Models the concept `Extension`.

Table B.28: Concepts Image: definitions (7)

Concept	Expression	Return Type	Description
WithExtension Image	`wextimg.extension()`	`std::convertible_to< extension_type>`	Get the extension of the image.

Table B.29: Concepts Image: expressions (7)

Concept table

Concept code

```
template <typename I>
concept WithExtensionImage =
  Image<I> &&
  requires {
    typename image_extension_t<I>;
  } &&
  Extension<image_extension_t<I>, image_point_t<I>> &&
  not std::same_as<mln::extension::none_extension_tag, image_extension_category_t<I>> &&
  requires (I ima, image_point_t<I> p) {
    { ima.extension() } -> std::convertible_to<image_extension_t<I>>;
  };
```

Archetype code

```
namespace details {
  struct WithExtensionImage : Image
  {
    struct Extension : ::mln::Extension<Extension>
    {
      using support_fill = std::false_type;
      using support_mirror = std::false_type;
      using support_periodize = std::false_type;
      using support_clamp = std::false_type;
      using support_extend_with = std::false_type;
      using value_type = image_value_t<Image>;
      bool fit(mln::archetypes::StructuringElement<image_point_t<Image>, mln::archetypes::Pixel> se) const;
      int extent() const;
    };

    using extension_type = Extension;

    using extension_category = mln::extension::custom_extension_tag;

    extension_type extension() const;
  };
} // namespace details

using WithExtensionImage = details::AsImage<details::WithExtensionImage>;
```

B.3.5 Output image

$$
\begin{aligned}
OutputImage = \; & (Image \implies WritableImage) \wedge \\
& (IndexableImage \implies WritableIndexableImage) \wedge \\
& (AccessibleImage \implies WritableAccessibleImage) \wedge \\
& (IndexableAndAccessibleImage \implies WritableIndexableAndAccessibleImage) \wedge \\
& (BidirectionalImage \implies WritableBidirectionalImage) \wedge \\
& (RawImage \implies WritableRawImage)
\end{aligned}
\tag{B.1}
$$

Figure B.1: Concepts OutputImage: definition

Concept table

Concept code

```
// OutputImage
// Usage: RawImage<I> && OutputImage<I>
template <typename I>
concept OutputImage =
  (not Image<I> || (detail::WritableImage<I>)) &&
  (not IndexableImage<I> || (detail::WritableIndexableImage<I>)) &&
  (not AccessibleImage<I> || (detail::WritableAccessibleImage<I>)) &&
  (not IndexableAndAccessibleImage<I> ||
    (detail::WritableIndexableAndAccessibleImage<I>)) &&
  (not BidirectionalImage<I> || (detail::WritableBidirectionalImage<I>)) &&
  (not RawImage<I> || (detail::WritableRawImage<I>));
```

Appendix C

Static-dynamic bridge

C.1   Ndimage module

Here is the code of the ndimage module, including the helpers for the type conversion.

ndimage.hpp

```cpp
#include <mln/core/image/ndbuffer_image.hpp>

#include <pybind11/numpy.h>
#include <pybind11/pybind11.h>


namespace mln::py
{
  mln::ndbuffer_image from_numpy(pybind11::array arr);

  pybind11::object to_numpy(const mln::ndbuffer_image& img);

  void init_pylena_numpy(pybind11::module& m);
} // namespace mln::py

namespace pybind11::detail
{
  template <>
  struct type_caster<mln::ndbuffer_image>
  {
    PYBIND11_TYPE_CASTER(mln::ndbuffer_image, _("numpy.ndarray"));

    bool load(handle h, bool)
    {
      pybind11::array arr = reinterpret_borrow<pybind11::array>(h);
      value               = mln::py::from_numpy(arr);
      return true;
    }

    static handle cast(const mln::ndbuffer_image& img, return_value_policy, handle)
    {
      return mln::py::to_numpy(img).inc_ref();
    }
  };
} // namespace pybind11::detail
```

ndimage.cpp

```cpp
#include "ndimage.hpp"

#include <mln/core/image/ndimage.hpp>

#include <fmt/core.h>
#include <pybind11/cast.h>

#include <cassert>
```

```cpp
#include <stdexcept>
#include <string>


namespace
{

  namespace details
  {
    template <mln::sample_type_id T>
    static pybind11::dtype dtype_of()
    {
      return pybind11::dtype::of<typename mln::sample_type_id_traits<T>::type>();
    }
  } // namespace details

  pybind11::dtype get_sample_type(mln::sample_type_id type)
  {
    switch (type)
    {
    case mln::sample_type_id::INT8:
      return details::dtype_of<mln::sample_type_id::INT8>();
    case mln::sample_type_id::INT16:
      return details::dtype_of<mln::sample_type_id::INT16>();
    case mln::sample_type_id::INT32:
      return details::dtype_of<mln::sample_type_id::INT32>();
    case mln::sample_type_id::INT64:
      return details::dtype_of<mln::sample_type_id::INT64>();
    case mln::sample_type_id::UINT8:
      return details::dtype_of<mln::sample_type_id::UINT8>();
    case mln::sample_type_id::UINT16:
      return details::dtype_of<mln::sample_type_id::UINT16>();
    case mln::sample_type_id::UINT32:
      return details::dtype_of<mln::sample_type_id::UINT32>();
    case mln::sample_type_id::UINT64:
      return details::dtype_of<mln::sample_type_id::UINT64>();
    case mln::sample_type_id::FLOAT:
      return details::dtype_of<mln::sample_type_id::FLOAT>();
    case mln::sample_type_id::DOUBLE:
      return details::dtype_of<mln::sample_type_id::DOUBLE>();
    case mln::sample_type_id::BOOL:
      return details::dtype_of<mln::sample_type_id::BOOL>();
    case mln::sample_type_id::RGB8:
      return details::dtype_of<mln::sample_type_id::UINT8>();
    default:
      throw std::runtime_error("Invalid sample_type_id");
    }
    return pybind11::none();
  }

  mln::sample_type_id get_sample_type(const std::string& type_format)
  {
    pybind11::dtype type;
    try
    {
      type = pybind11::dtype(type_format);
    }
    catch (const std::exception&)
    {
      return mln::sample_type_id::OTHER;
    }
    if (type.is(details::dtype_of<mln::sample_type_id::INT8>()))
      return mln::sample_type_id::INT8;
    else if (type.is(details::dtype_of<mln::sample_type_id::INT16>()))
      return mln::sample_type_id::INT16;
    else if (type.is(details::dtype_of<mln::sample_type_id::INT32>()))
      return mln::sample_type_id::INT32;
    else if (type.is(details::dtype_of<mln::sample_type_id::INT64>()))
      return mln::sample_type_id::INT64;
    else if (type.is(details::dtype_of<mln::sample_type_id::UINT8>()))
      return mln::sample_type_id::UINT8;
    else if (type.is(details::dtype_of<mln::sample_type_id::UINT16>()))
      return mln::sample_type_id::UINT16;
```

```cpp
      else if (type.is(details::dtype_of<mln::sample_type_id::UINT32>()))
        return mln::sample_type_id::UINT32;
      else if (type.is(details::dtype_of<mln::sample_type_id::UINT64>()))
        return mln::sample_type_id::UINT64;
      else if (type.is(details::dtype_of<mln::sample_type_id::FLOAT>()))
        return mln::sample_type_id::FLOAT;
      else if (type.is(details::dtype_of<mln::sample_type_id::DOUBLE>()))
        return mln::sample_type_id::DOUBLE;
      else if (type.is(details::dtype_of<mln::sample_type_id::BOOL>()))
        return mln::sample_type_id::BOOL;
      return mln::sample_type_id::OTHER;
  }
} // namespace


namespace mln::py
{
  mln::ndbuffer_image from_numpy(pybind11::array arr)
  {
    if (!pybind11::detail::check_flags(arr.ptr(), pybind11::detail::npy_api::NPY_ARRAY_C_CONTIGUOUS_))
      throw std::invalid_argument("Array should be C contiguous");
    auto              base = arr.base();
    const auto        info = arr.request();
    mln::sample_type_id type = get_sample_type(info.format);
    if (type == mln::sample_type_id::OTHER)
      throw std::invalid_argument(fmt::format(
          "Invalid dtype argument (Got dtype format {} expected types: [u]int[8, 16, 32, 64], float, double or bool)",
          info.format));
    const bool is_rgb8 = info.ndim == 3 && info.shape[2] == 3 && type == mln::sample_type_id::UINT8;
    const auto pdim    = info.ndim - (is_rgb8 ? 1 : 0);
    if (pdim > mln::PYLENE_NDBUFFER_DEFAULT_DIM)
      throw std::invalid_argument(
          fmt::format("Invalid number of dimension from numpy array (Got {} but should be less than {})", pdim,
                      mln::PYLENE_NDBUFFER_DEFAULT_DIM));
    int           size[mln::PYLENE_NDBUFFER_DEFAULT_DIM]    = {0};
    std::ptrdiff_t strides[mln::PYLENE_NDBUFFER_DEFAULT_DIM] = {0};
    for (auto d = 0; d < pdim; d++)
    {
      size[d]    = info.shape[pdim - d - 1];
      strides[d] = info.strides[pdim - d - 1];
    }
    const auto sample_type = is_rgb8 ? mln::sample_type_id::RGB8 : type;
    auto       res =
        mln::ndbuffer_image::from_buffer(reinterpret_cast<std::byte*>(info.ptr), sample_type, pdim, size, strides);
    if (base && pybind11::isinstance<mln::internal::ndbuffer_image_data>(base))
      res.__data() = pybind11::cast<std::shared_ptr<mln::internal::ndbuffer_image_data>>(base);
    return res;
  }


  pybind11::object to_numpy(const mln::ndbuffer_image& img)
  {
    const auto&      api   = pybind11::detail::npy_api::get();
    pybind11::object data  = pybind11::none();
    int              flags = pybind11::detail::npy_api::NPY_ARRAY_WRITEABLE_;
    if (img.__data())
    {
      data = pybind11::cast(img.__data());
      assert(data.ref_count() > 0);
    }

    /* For the moment, restrict RGB8 image to 2D image */
    const bool            is_rgb8 = img.pdim() == 2 && img.sample_type() == mln::sample_type_id::RGB8;
    const auto            ndim    = img.pdim() + (is_rgb8 ? 1 : 0);
    std::vector<std::size_t> strides(ndim, 1);
    std::vector<std::size_t> shapes(ndim, 3);
    auto                  descr = get_sample_type(img.sample_type());

    for (auto d = 0; d < img.pdim(); d++)
    {
      strides[d] = img.byte_stride(img.pdim() - d - 1);
      shapes[d]  = img.size(img.pdim() - d - 1);
    }
```

```cpp
    auto res = pybind11::reinterpret_steal<pybind11::object>(api.PyArray_NewFromDescr_(
        api.PyArray_Type_, descr.release().ptr(), ndim, reinterpret_cast<Py_intptr_t*>(shapes.data()),
        reinterpret_cast<Py_intptr_t*>(strides.data()), reinterpret_cast<void*>(img.buffer()), flags, nullptr));

    if (!res)
      throw std::runtime_error("Unable to create the numpy array in ndimage -> array");
    if (data)
      // **Steal** a reference to data (https://numpy.org/devdocs/reference/c-api/array.html#c.PyArray_SetBaseObject)
      api.PyArray_SetBaseObject_(res.ptr(), data.release().ptr());
    return res;
  }

  void init_pylena_numpy(pybind11::module& m)
  {
    if (!pybind11::detail::get_global_type_info(typeid(mln::internal::ndbuffer_image_data)))
    {
      pybind11::class_<mln::internal::ndbuffer_image_data, std::shared_ptr<mln::internal::ndbuffer_image_data>>(
          m, "ndbuffer_image_data");
    }
  }

} // namespace mln::py
```

C.2 Structuring element module

Here is the code of the structuring element module. This module defines what structuring elements available from Python.

se.hpp

```cpp
#include <pybind11/pybind11.h>


namespace mln::py
{

  void init_module_se(pybind11::module& m);

}
```

se.cpp

```cpp
#include "se.hpp"

#include <mln/core/image/ndimage.hpp>
#include <mln/core/se/disc.hpp>
#include <mln/core/se/rect2d.hpp>

#include <pybind11/pybind11.h>


namespace mln::py
{

  void init_module_se(pybind11::module& m)
  {
    pybind11::class_<mln::se::disc_non_decomp>(m, "disc").def(
        pybind11::init([](float radius) { return mln::se::disc_non_decomp{radius}; }));

    pybind11::class_<mln::se::disc>(m, "disc_decomposable").def(pybind11::init([](float radius) {
      return mln::se::disc{radius};
    }));

    pybind11::class_<mln::se::rect2d_non_decomp>(m, "rect2d").def(pybind11::init([](int width, int height) {
      return mln::se::rect2d_non_decomp{width, height};
    }));

    pybind11::class_<mln::se::rect2d>(m, "rect2d_decomposable").def(pybind11::init([](int width, int height) {
      return mln::se::rect2d{width, height};
    }));
  }

} // namespace mln::py
```

C.3   Mathematical morphology module

This module defines the routines that are exposed and available from Python. It is divided in three sub-parts. First is the $n * n$ dispatcher code. Second is the value set mechanic. Third is the code exposing the mathematical morphology algorithms.

C.3.1   The $n * n$ dispatcher

visit.hpp

```cpp
#include <mln/core/image/ndimage.hpp>

#include <stdexcept>

namespace mln::py
{
  template <template <typename> class F, typename... Args>
  auto visit(mln::sample_type_id tid, Args&&... args)
  {
    switch (tid)
    {
    case (mln::sample_type_id::INT8):
      return F<std::int8_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT16):
      return F<std::int16_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT32):
      return F<std::int32_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT64):
      return F<std::int64_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT8):
      return F<std::uint8_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT16):
      return F<std::uint16_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT32):
      return F<std::uint32_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT64):
      return F<std::uint64_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::FLOAT):
      return F<float>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::DOUBLE):
      return F<double>{}(std::forward<Args>(args)...);
    /*case (mln::sample_type_id::RGB8):
      return F<mln::rgb8>{}(std::forward<Args>(args)...);*/
    case (mln::sample_type_id::OTHER):
      [[fallthrough]];
    default:
      throw std::runtime_error("Unhandled data type");
    }
  }

  template <template <typename> class F, typename D, typename... Args>
  auto visit_d(mln::sample_type_id tid, D&& d, Args&&... args)
  {
    switch (tid)
    {
    case (mln::sample_type_id::INT8):
      return F<std::int8_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT16):
      return F<std::int16_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT32):
      return F<std::int32_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::INT64):
      return F<std::int64_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT8):
      return F<std::uint8_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT16):
      return F<std::uint16_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT32):
      return F<std::uint32_t>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::UINT64):
      return F<std::uint64_t>{}(std::forward<Args>(args)...);
```

```cpp
    case (mln::sample_type_id::FLOAT):
      return F<float>{}(std::forward<Args>(args)...);
    case (mln::sample_type_id::DOUBLE):
      return F<double>{}(std::forward<Args>(args)...);
    /*case (mln::sample_type_id::RGB8):
      return F<mln::rgb8>{}(std::forward<Args>(args)...);*/
    case (mln::sample_type_id::OTHER):
      [[fallthrough]];
    default:
      return std::forward<D>(d);
    }
  }
} // namespace mln::py
```

C.3.2   Value-set mechanics

value_set.hpp

```cpp
#include <algorithm>
#include <any>
#include <cassert>
#include <cmath>
#include <limits>
#include <type_traits>
#include <utility>


namespace mln
{

  template <class T = void>
  struct value_set
  {
    template <class U>
    U cast(T v) const
    {
      return static_cast<U>(v);
    }

    T max() const noexcept { return std::numeric_limits<T>::max(); }
    T min() const noexcept { return std::numeric_limits<T>::min(); }
    T sup() const noexcept { return std::numeric_limits<T>::max(); }
    T inf() const noexcept { return std::numeric_limits<T>::min(); }

    T plus(T v) const noexcept { return +v; }
    T minus(T v) const noexcept { return -v; }

    auto max(T l, T r) const noexcept { return std::max(l, r); }
    auto min(T l, T r) const noexcept { return std::min(l, r); }
    auto add(T l, T r) const noexcept { return l + r; }
    auto sub(T l, T r) const noexcept { return l - r; }
    auto mult(T l, T r) const noexcept { return l * r; }
    auto div(T l, T r) const noexcept { return l / r; }
    auto mod(T l, T r) const noexcept { return l % r; }
    auto pow(T l, T r) const noexcept { return std::pow(l, r); }
  };

  template <>
  struct value_set<void>
  {

    template <class U, class T>
    U cast(T&& t) const
    {
      return static_cast<U>(std::forward<T>);
    }

    template <class T>
    T max() const noexcept
    {
      return std::numeric_limits<T>::max();
    }
```

```cpp
  template <class T>
  T min() const noexcept
  {
    return std::numeric_limits<T>::min();
  }
  template <class T>
  T sup() const noexcept
  {
    return std::numeric_limits<T>::max();
  }
  template <class T>
  T inf() const noexcept
  {
    return std::numeric_limits<T>::min();
  }

  template <class T>
  T plus(T&& v) const noexcept
  {
    return +std::forward<T>(v);
  }
  template <class T>
  T minus(T&& v) const noexcept
  {
    return -std::forward<T>(v);
  }

  template <class T, class U>
  auto max(T&& l, T&& r) const noexcept
  {
    return std::max(std::forward<T>(l), std::forward<T>(r));
  }
  template <class T, class U>
  auto min(T&& l, T&& r) const noexcept
  {
    return std::min(std::forward<T>(l), std::forward<T>(r));
  }

  template <class T, class U>
  auto add(T&& l, U&& r) const noexcept
  {
    return std::forward<T>(l) + std::forward<T>(r);
  }
  template <class T, class U>
  auto sub(T&& l, U&& r) const noexcept
  {
    return std::forward<T>(l) - std::forward<T>(r);
  }
  template <class T, class U>
  auto mult(T&& l, U&& r) const noexcept
  {
    return std::forward<T>(l) * std::forward<T>(r);
  }
  template <class T, class U>
  auto div(T&& l, U&& r) const noexcept
  {
    return std::forward<T>(l) / std::forward<T>(r);
  }
  template <class T, class U>
  auto mod(T&& l, U&& r) const noexcept
  {
    return std::forward<T>(l) % std::forward<T>(r);
  }
  template <class T, class U>
  auto pow(T&& l, U&& r) const noexcept
  {
    return std::pow(std::forward<T>(l), std::forward<T>(r));
  }
};

struct abstract_value_set
{
  virtual ~abstract_value_set() {}
```

```cpp
  /*
  template <class U>
  virtual std::any cast(std::any v) const = 0;
  */

  virtual std::any max() const = 0;
  virtual std::any min() const = 0;
  virtual std::any sup() const = 0;
  virtual std::any inf() const = 0;

  virtual std::any plus(const std::any& v) const  = 0;
  virtual std::any minus(const std::any& v) const = 0;

  virtual std::any max(const std::any& l, const std::any& r) const  = 0;
  virtual std::any min(const std::any& l, const std::any& r) const  = 0;
  virtual std::any add(const std::any& l, const std::any& r) const  = 0;
  virtual std::any sub(const std::any& l, const std::any& r) const  = 0;
  virtual std::any mult(const std::any& l, const std::any& r) const = 0;
  virtual std::any div(const std::any& l, const std::any& r) const  = 0;
  virtual std::any mod(const std::any& l, const std::any& r) const  = 0;
  virtual std::any pow(const std::any& l, const std::any& r) const = 0;
};

template <typename T>
struct concrete_value_set : abstract_value_set
{
  ~concrete_value_set() override {}

  template <class U>
  std::any cast(std::any v) const /* override */
  {
    return {static_cast<U>(std::any_cast<T>(v))};
  }

  std::any max() const override { return {std::numeric_limits<T>::max()}; };
  std::any min() const override { return {std::numeric_limits<T>::min()}; };
  std::any sup() const override { return {std::numeric_limits<T>::max()}; };
  std::any inf() const override { return {std::numeric_limits<T>::min()}; };

  std::any plus(const std::any& v) const override { return {+std::any_cast<T>(v)}; };
  std::any minus(const std::any& v) const override { return {-std::any_cast<T>(v)}; };

  std::any max(const std::any& l, const std::any& r) const override
  {
    return {std::max(std::any_cast<T>(l), std::any_cast<T>(r))};
  }
  std::any min(const std::any& l, const std::any& r) const override
  {
    return {std::min(std::any_cast<T>(l), std::any_cast<T>(r))};
  }
  std::any add(const std::any& l, const std::any& r) const override
  {
    return {std::any_cast<T>(l) + std::any_cast<T>(r)};
  }
  std::any sub(const std::any& l, const std::any& r) const override
  {
    return {std::any_cast<T>(l) - std::any_cast<T>(r)};
  }
  std::any mult(const std::any& l, const std::any& r) const override
  {
    return {std::any_cast<T>(l) * std::any_cast<T>(r)};
  }
  std::any div(const std::any& l, const std::any& r) const override
  {
    return {std::any_cast<T>(l) / std::any_cast<T>(r)};
  }
  std::any mod(const std::any& l, const std::any& r) const override
  {
    return {std::any_cast<T>(l) % std::any_cast<T>(r)};
  }
  std::any pow(const std::any& l, const std::any& r) const override
  {
```

```cpp
      return {std::pow(std::any_cast<T>(l), std::any_cast<T>(r))};
  }
};

struct type_erased_value
{
  type_erased_value(const std::any& v, const abstract_value_set& vs);

  const std::any&       val() const;
  const std::type_info& tid() const;

  template <typename From, typename To>
  type_erased_value cast(abstract_value_set* new_vs = nullptr) const
  {
    assert(typeid(From) == tid());
    auto new_val = std::any{static_cast<To>(std::any_cast<From>(v_))};
    return {new_val, new_vs != nullptr ? *new_vs : *vs_};
  }

  type_erased_value max() const;
  type_erased_value min() const;
  type_erased_value sup() const;
  type_erased_value inf() const;

  type_erased_value plus() const;
  type_erased_value minus() const;

  type_erased_value max(const type_erased_value& rhs) const;
  type_erased_value min(const type_erased_value& rhs) const;
  type_erased_value add(const type_erased_value& rhs) const;
  type_erased_value sub(const type_erased_value& rhs) const;
  type_erased_value mult(const type_erased_value& rhs) const;
  type_erased_value div(const type_erased_value& rhs) const;
  type_erased_value mod(const type_erased_value& rhs) const;
  type_erased_value pow(const type_erased_value& rhs) const;

private:
  std::any                  v_;
  const abstract_value_set* vs_;
};

template <>
struct value_set<type_erased_value> : abstract_value_set
{
  value_set(const type_erased_value& v);
  ~value_set() override;

  template <typename From, typename To>
  std::any cast(const std::any& v, abstract_value_set* new_vs = nullptr) const
  {
    auto absv = std::any_cast<type_erased_value>(v);
    return {absv.template cast<From, To>(new_vs)};
  }

  std::any max() const override;
  std::any min() const override;
  std::any sup() const override;
  std::any inf() const override;

  std::any plus(const std::any& v) const override;

  std::any minus(const std::any& v) const override;

  std::any max(const std::any& l, const std::any& r) const override;
  std::any min(const std::any& l, const std::any& r) const override;
  std::any add(const std::any& l, const std::any& r) const override;
  std::any sub(const std::any& l, const std::any& r) const override;
  std::any mult(const std::any& l, const std::any& r) const override;
  std::any div(const std::any& l, const std::any& r) const override;
  std::any mod(const std::any& l, const std::any& r) const override;
  std::any pow(const std::any& l, const std::any& r) const override;

private:
```

```cpp
    const type_erased_value* v_;
  };
} // namespace mln
```

value_set.cpp

```cpp
#include <mln/core/value/value_set.hpp>

#include <algorithm>
#include <any>
#include <cinttypes>
#include <cmath>
#include <limits>
#include <optional>
#include <type_traits>
#include <typeinfo>
#include <variant>


namespace details
{


  template <class... Args>
  auto any_to_variant_cast(std::any a) -> std::optional<std::variant<Args...>>
  {
    if (!a.has_value())
      return std::nullopt;

    std::optional<std::variant<Args...>> v = std::nullopt;

    bool found = ((a.type() == typeid(Args) && (v = std::any_cast<Args>(std::move(a)), true)) || ...);

    if (!found)
      return std::nullopt;

    return {std::move(*v)};
  }

  template <template <typename> class Op, typename... Args>
  auto call_template_op_from_type_info(const std::type_info& tid) -> std::optional<std::variant<Op<Args>...>>
  {
    std::optional<std::variant<Op<Args>...>> v = std::nullopt;

    bool found = ((tid == typeid(Args) && (v = Op<Args>{}, true)) || ...);

    if (!found)
      return std::nullopt;

    return {std::move(*v)};
  }

  template <class... Ts>
  struct type_list
  {
  };

  template <class... Ts>
  auto any_to_variant_cast_type_list(std::any a, type_list<Ts...>) -> std::optional<std::variant<Ts...>>
  {
    return any_to_variant_cast<Ts...>(a);
  }

  template <template <typename> class Op, typename... Ts>
  auto call_template_op_from_type_info_type_list(const std::type_info& tid, type_list<Ts...>)
      -> std::optional<std::variant<Op<Ts>...>>
  {
    return call_template_op_from_type_info<Op, Ts...>(tid);
  }

  using supported_types = type_list<std::uint8_t, std::uint16_t, std::uint32_t, std::uint64_t, std::int8_t,
                                    std::int16_t, std::int32_t, std::int64_t, float, double>;
```

```cpp
  template <class T, class U, class C = std::common_type_t<T, U>>
  C max(const T& a, const U& b)
  {
    return std::max(static_cast<C>(a), static_cast<C>(b));
  }

  template <class T, class U, class C = std::common_type_t<T, U>>
  C min(const T& a, const U& b)
  {
    return std::min(static_cast<C>(a), static_cast<C>(b));
  }

  template <class T, class U>
  long double pow(const T& a, const U& b)
  {
    return std::pow(static_cast<long double>(a), static_cast<long double>(b));
  }

} // namespace details


namespace mln
{

  type_erased_value::type_erased_value(const std::any& v, const abstract_value_set& vs)
    : v_(v)
    , vs_(&vs)
  {
  }

  const std::any&      type_erased_value::val() const { return v_; }
  const std::type_info& type_erased_value::tid() const { return v_.type(); }

  template <typename T>
  struct limit_max
  {
    auto operator()() { return std::numeric_limits<T>::max(); }
  };

  type_erased_value type_erased_value::max() const
  {
    auto op = details::call_template_op_from_type_info_type_list<limit_max>(tid(), details::supported_types{});

    auto ret = std::visit([](auto op_) -> std::any { return {op_()}; }, *op);
    return {ret, *vs_};
  }

  template <typename T>
  struct limit_min
  {
    auto operator()() { return std::numeric_limits<T>::min(); }
  };

  type_erased_value type_erased_value::min() const
  {
    auto op = details::call_template_op_from_type_info_type_list<limit_min>(tid(), details::supported_types{});

    auto ret = std::visit([](auto op_) -> std::any { return {op_()}; }, *op);
    return {ret, *vs_};
  }

  type_erased_value type_erased_value::sup() const
  {
    auto op = details::call_template_op_from_type_info_type_list<limit_max>(tid(), details::supported_types{});

    auto ret = std::visit([](auto op_) -> std::any { return {op_()}; }, *op);
    return {ret, *vs_};
  }

  type_erased_value type_erased_value::inf() const
  {
    auto op = details::call_template_op_from_type_info_type_list<limit_min>(tid(), details::supported_types{});
```

```cpp
  auto ret = std::visit([](auto op_) -> std::any { return {op_()}; }, *op);
  return {ret, *vs_};
}

type_erased_value type_erased_value::plus() const { return {vs_->plus(v_), *vs_}; }
type_erased_value type_erased_value::minus() const { return {vs_->plus(v_), *vs_}; }

type_erased_value type_erased_value::max(const type_erased_value& rhs) const
{
  return {vs_->max(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::min(const type_erased_value& rhs) const
{
  return {vs_->min(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::add(const type_erased_value& rhs) const
{
  return {vs_->add(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::sub(const type_erased_value& rhs) const
{
  return {vs_->sub(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::mult(const type_erased_value& rhs) const
{
  return {vs_->mult(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::div(const type_erased_value& rhs) const
{
  return {vs_->div(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::mod(const type_erased_value& rhs) const
{
  return {vs_->mod(v_, rhs.val()), *vs_};
}
type_erased_value type_erased_value::pow(const type_erased_value& rhs) const
{
  return {vs_->pow(v_, rhs.val()), *vs_};
}


value_set<type_erased_value>::value_set(const type_erased_value& v)
  : v_(&v)
{
}

value_set<type_erased_value>::~value_set() {}

std::any value_set<type_erased_value>::max() const { return {v_->max()}; }
std::any value_set<type_erased_value>::min() const { return {v_->min()}; }
std::any value_set<type_erased_value>::sup() const { return {v_->sup()}; }
std::any value_set<type_erased_value>::inf() const { return {v_->inf()}; }

std::any value_set<type_erased_value>::plus(const std::any& v) const
{

  auto va = details::any_to_variant_cast_type_list(v, details::supported_types{});
  if (va)
    return {std::visit([](auto _v) -> std::any { return {+_v}; }, *va)};
  else
  {
    auto absv = std::any_cast<type_erased_value>(v);
    return {absv.plus()};
  }
}

std::any value_set<type_erased_value>::minus(const std::any& v) const
{
  auto va = details::any_to_variant_cast_type_list(v, details::supported_types{});
  if (va)
    return {std::visit([](auto _v) -> std::any { return {-_v}; }, *va)};
  else
  {
```

```cpp
    auto absv = std::any_cast<type_erased_value>(v);
    return {absv.minus()};
  }
}


std::any value_set<type_erased_value>::max(const std::any& l, const std::any& r) const
{
  auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
  auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

  if (lhs && rhs)
    return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {details::max(lhs_, rhs_)}; }, *lhs, *rhs)};
  else
  {
    auto absl = std::any_cast<type_erased_value>(l);
    auto absr = std::any_cast<type_erased_value>(r);
    return {absl.max(absr)};
  }
}


std::any value_set<type_erased_value>::min(const std::any& l, const std::any& r) const
{
  auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
  auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

  if (lhs && rhs)
    return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {details::min(lhs_, rhs_)}; }, *lhs, *rhs)};
  else
  {
    auto absl = std::any_cast<type_erased_value>(l);
    auto absr = std::any_cast<type_erased_value>(r);
    return {absl.min(absr)};
  }
}


std::any value_set<type_erased_value>::add(const std::any& l, const std::any& r) const
{
  auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
  auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

  if (lhs && rhs)
    return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {lhs_ + rhs_}; }, *lhs, *rhs)};
  else
  {
    auto absl = std::any_cast<type_erased_value>(l);
    auto absr = std::any_cast<type_erased_value>(r);
    return {absl.add(absr)};
  }
}


std::any value_set<type_erased_value>::sub(const std::any& l, const std::any& r) const
{
  auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
  auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

  if (lhs && rhs)
    return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {lhs_ - rhs_}; }, *lhs, *rhs)};
  else
  {
    auto absl = std::any_cast<type_erased_value>(l);
    auto absr = std::any_cast<type_erased_value>(r);
    return {absl.sub(absr)};
  }
}


std::any value_set<type_erased_value>::mult(const std::any& l, const std::any& r) const
{
  auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
  auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

  if (lhs && rhs)
    return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {lhs_ * rhs_}; }, *lhs, *rhs)};
  else
```

```cpp
    {
      auto absl = std::any_cast<type_erased_value>(l);
      auto absr = std::any_cast<type_erased_value>(r);
      return {absl.mult(absr)};
    }
  }

  std::any value_set<type_erased_value>::div(const std::any& l, const std::any& r) const
  {
    auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
    auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

    if (lhs && rhs)
      return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {lhs_ / rhs_}; }, *lhs, *rhs)};
    else
    {
      auto absl = std::any_cast<type_erased_value>(l);
      auto absr = std::any_cast<type_erased_value>(r);
      return {absl.div(absr)};
    }
  }

  std::any value_set<type_erased_value>::mod(const std::any& l, const std::any& r) const
  {
    auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
    auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

    if (lhs && rhs)
      return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {lhs_ % rhs_}; }, *lhs, *rhs)};
    else
    {
      auto absl = std::any_cast<type_erased_value>(l);
      auto absr = std::any_cast<type_erased_value>(r);
      return {absl.mod(absr)};
    }
  }

  std::any value_set<type_erased_value>::pow(const std::any& l, const std::any& r) const
  {
    auto lhs = details::any_to_variant_cast_type_list(l, details::supported_types{});
    auto rhs = details::any_to_variant_cast_type_list(r, details::supported_types{});

    if (lhs && rhs)
      return {std::visit([](auto lhs_, auto rhs_) -> std::any { return {details::pow(lhs_, rhs_)}; }, *lhs, *rhs)};
    else
    {
      auto absl = std::any_cast<type_erased_value>(l);
      auto absr = std::any_cast<type_erased_value>(r);
      return {absl.pow(absr)};
    }
  }

} // namespace mln
```

py_value_set.hpp

```cpp
#include <mln/core/value/value_set.hpp>
#include <pybind11/pybind11.h>

#include <any>
#include <cinttypes>
#include <memory>
#include <string>
#include <string_view>
#include <typeinfo>
#include <variant>


namespace mln
{
  template <>
  struct value_set<pybind11::object> : abstract_value_set
  {
```

```cpp
    value_set(pybind11::object python_vs_instance);

    ~value_set() override;

    template <typename U>
    std::any cast(const std::any& v) const;

    std::any max() const override;
    std::any min() const override;
    std::any sup() const override;
    std::any inf() const override;

    std::any plus(const std::any& v) const override;
    std::any minus(const std::any& v) const override;

    std::any max(const std::any& l, const std::any& r) const override;
    std::any min(const std::any& l, const std::any& r) const override;
    std::any add(const std::any& l, const std::any& r) const override;
    std::any sub(const std::any& l, const std::any& r) const override;
    std::any mult(const std::any& l, const std::any& r) const override;
    std::any div(const std::any& l, const std::any& r) const override;
    std::any mod(const std::any& l, const std::any& r) const override;
    std::any pow(const std::any& l, const std::any& r) const override;

  private:
    pybind11::object vs_instance_;
  };

} // namespace mln


namespace details
{

  template <class U>
  auto get_python_type()
  {
    // C++ type -> Python type
    static std::unordered_map<std::type_index, std::string> type_names{
        {std::type_index(typeid(bool{})), "bool"},          //
        {std::type_index(typeid(int8_t{})), "int"},         //
        {std::type_index(typeid(int16_t{})), "int"},        //
        {std::type_index(typeid(int32_t{})), "int"},        //
        {std::type_index(typeid(int64_t{})), "int"},        //
        {std::type_index(typeid(uint8_t{})), "int"},        //
        {std::type_index(typeid(uint16_t{})), "int"},       //
        {std::type_index(typeid(uint32_t{})), "int"},       //
        {std::type_index(typeid(uint64_t{})), "int"},       //
        {std::type_index(typeid(float{})), "float"},        //
        {std::type_index(typeid(double{})), "float"},       //
        {std::type_index(typeid((char*){})), "str"},        //
        {std::type_index(typeid((const char*){})), "str"}, //
        {std::type_index(typeid(std::string{})), "str"}};
    return type_names[std::type_index(typeid(U{}))];
  }

} // namespace details


namespace mln
{
  template <typename U>
  std::any value_set<pybind11::object>::cast(const std::any& v) const
  {
    pybind11::object v_ = std::any_cast<pybind11::object>(v);
    // if constexpr (std::is_arithmetic_v<U>) // bool, signed/unsigned int, floating point
    if constexpr (std::is_same_v<U, bool>) // bool
    {
      pybind11::bool_ pyv = vs_instance_.attr("cast")(v_, ::details::get_python_type<U>());
      return {static_cast<U>(pyv)};
    }
    else if constexpr (std::is_integral_v<U>) // signed/unsigned int
    {
```

```cpp
      pybind11::int_ pyv = vs_instance_.attr("cast")(v_, ::details::get_python_type<U>());
      return {static_cast<U>(pyv)};
    }
    else if constexpr (std::is_floating_point_v<U>) // floating point
    {
      pybind11::float_ pyv = vs_instance_.attr("cast")(v_, ::details::get_python_type<U>());
      return {static_cast<U>(pyv)};
    }
    else
    {
      pybind11::object ret = vs_instance_.attr("cast")(v, ::details::get_python_type<U>());
      return {*(ret.cast<U*>())};
    }
  }

} // namespace mln
```

py_value_set.cpp

```cpp
#include "py_value_set.hpp"

#include <any>

namespace mln
{

  value_set<pybind11::object>::value_set(pybind11::object python_vs_instance)
    : vs_instance_(python_vs_instance)
  {
  }

  value_set<pybind11::object>::~value_set()
  {
  }

  std::any value_set<pybind11::object>::max() const { return {vs_instance_.attr("max")()}; }

  std::any value_set<pybind11::object>::min() const { return {vs_instance_.attr("min")()}; }

  std::any value_set<pybind11::object>::sup() const { return {vs_instance_.attr("sup")()}; }

  std::any value_set<pybind11::object>::inf() const { return {vs_instance_.attr("inf")()}; }


  std::any value_set<pybind11::object>::plus(const std::any& v) const
  {
    auto pyv = std::any_cast<pybind11::object>(v);
    return {vs_instance_.attr("plus")(pyv)};
  }

  std::any value_set<pybind11::object>::minus(const std::any& v) const
  {
    auto pyv = std::any_cast<pybind11::object>(v);
    return {vs_instance_.attr("minus")(pyv)};
  }


  std::any value_set<pybind11::object>::max(const std::any& l, const std::any& r) const
  {
    auto pyl = std::any_cast<pybind11::object>(l);
    auto pyr = std::any_cast<pybind11::object>(r);
    return {vs_instance_.attr("max")(pyl, pyr)};
  }

  std::any value_set<pybind11::object>::min(const std::any& l, const std::any& r) const
  {
    auto pyl = std::any_cast<pybind11::object>(l);
    auto pyr = std::any_cast<pybind11::object>(r);
    return {vs_instance_.attr("min")(pyl, pyr)};
  }

  std::any value_set<pybind11::object>::add(const std::any& l, const std::any& r) const
  {
```

```cpp
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("add")(pyl, pyr)};
}

std::any value_set<pybind11::object>::sub(const std::any& l, const std::any& r) const
{
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("sub")(pyl, pyr)};
}

std::any value_set<pybind11::object>::mult(const std::any& l, const std::any& r) const
{
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("mult")(pyl, pyr)};
}

std::any value_set<pybind11::object>::div(const std::any& l, const std::any& r) const
{
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("div")(pyl, pyr)};
}

std::any value_set<pybind11::object>::mod(const std::any& l, const std::any& r) const
{
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("mod")(pyl, pyr)};
}

std::any value_set<pybind11::object>::pow(const std::any& l, const std::any& r) const
{
  auto pyl = std::any_cast<pybind11::object>(l);
  auto pyr = std::any_cast<pybind11::object>(r);
  return {vs_instance_.attr("pow")(pyl, pyr)};
}

} // namespace mln
```

C.3.3   Mathematical morphology routines

morpho.hpp

```cpp
#include <pybind11/pybind11.h>


namespace mln::py
{

  void init_module_morpho(pybind11::module& m);

}
```

morpho.cpp

```cpp
#include "morpho.hpp"

#include "ndimage.hpp"
#include "py_value_set.hpp"
#include "visit.hpp"

#include <mln/core/algorithm/transform.hpp>
#include <mln/core/image/ndimage.hpp>
#include <mln/core/se/disc.hpp>
#include <mln/core/se/rect2d.hpp>
#include <mln/morpho/experimental/dilation.hpp>

#include <fmt/core.h>
```

```cpp
#include <pybind11/stl.h>

#include <algorithm>
#include <string_view>
#include <variant>


namespace my
{
  template <class T>
  mln::image2d<float> slow_stretch(const mln::image2d<T>& src, const mln::value_set<pybind11::object>& py_vs)
  {
    mln::image2d<float> res = mln::transform(src, [&py_vs](auto val) -> float {
      auto anymax  = py_vs.max();                      // returns std::any of pybind11:object
      auto anyval  = std::any{pybind11::cast(val)};    // converts to std::any of pybind11::object
      auto anyret  = py_vs.div(anyval, anymax);        // returns std::any of pybind11::object
      auto anyfret = py_vs.template cast<float>(anyret); // returns std::any of float
      return std::any_cast<float>(anyfret);            // returns float
    });
    return res;
  }

  template <class T>
  mln::image2d<float> virtual_dispatch_stretch(const mln::image2d<T>& src)
  {
    auto              vs   = mln::concrete_value_set<T>{};     // value-set for T
    auto              vs_f = mln::concrete_value_set<float>{}; // value-set for float
    mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
      auto anymax  = vs.max();                            // returns std::any
      auto fanyval = vs.template cast<float>(val);        // cast to float in std::any
      auto fanymax = vs.template cast<float>(anymax);     // cast to float in std::any
      return std::any_cast<float>(vs_f.div(fanyval, fanymax)); // div returns float
    });
    return res;
  }

  template <class T>
  mln::image2d<float> stretch_virtual_dispatch_type_erased_value(const mln::image2d<T>& src)
  {
    auto              vs   = mln::concrete_value_set<T>{};     // value-set for T
    auto              vs_f = mln::concrete_value_set<float>{}; // value-set for float
    mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
      // simulate having an image<type_erased_value>
      auto anyval = std::any{val}; // std::any of T
                                   // type_erased_value of std::any of T aware of value-set of T
      auto abs_anyval = mln::type_erased_value{anyval, vs};
      // instantiate a value-set for type_erased_value
      auto abs_vs = mln::value_set<mln::type_erased_value>{abs_anyval};
      auto anyabs_anymax =
          abs_vs.max(); // returns std::any of type_erased_value
                        // cast underlying std::any of type_erased_value of std::any of T into
                        // std::any of type_erased_value of std::any of float aware of value-set for float
      auto anyabs_fanyval = abs_vs.template cast<T, float>(std::any{abs_anyval}, &vs_f);
      auto anyabs_fanymax = abs_vs.template cast<T, float>(anyabs_anymax, &vs_f);
      // dispatch on known type, find a type_erased_value, then call
      // anyabs_fanyval.div(anyabs_fanymax) to perform division which will call
      // the underlying value-set for float for this operation
      auto anyabs_fanyret = abs_vs.div(anyabs_fanyval, anyabs_fanymax);
      // convert result back into float for returning to the image
      auto anyfret = std::any_cast<mln::type_erased_value>(anyabs_fanyret).val();
      return std::any_cast<float>(anyfret);
    });
    return res;
  }

  template <class T>
  mln::image2d<float> fast_stretch(const mln::image2d<T>& src)
  {
    auto              vs   = src.get_value_set();        // value-set for T
    auto              vs_f = mln::value_set<float>{};    // fast value-set for float
    mln::image2d<float> res  = mln::transform(src, [&vs, &vs_f](auto val) -> float {
      auto max  = vs.max();                        // returns T
```

```cpp
      auto fval = vs.template cast<float>(val); // returns float
      auto fmax = vs.template cast<float>(max); // returns float
      return vs_f.div(fval, fmax);              // div directly returns float
    });
    return res;
  }
} // namespace my

namespace
{

  template <typename T>
  struct stretch_operator_t
  {
    template <typename Img>
    mln::ndbuffer_image operator()(Img&& img, const std::optional<pybind11::object>& py_vs) const
    {
      if (auto* image_ptr = std::forward<Img>(img).template cast_to<T, 2>(); image_ptr)
      {
        if (py_vs.has_value())
        {
          auto vs = mln::value_set<pybind11::object>{py_vs.value()};
          return my::slow_stretch(*image_ptr, vs);
        }
        else
        {
          return my::fast_stretch(*image_ptr);
        }
      }
      else
      {
        std::runtime_error("Unable to convert the image to the required type.");
        return {};
      }
    }
  };

  mln::ndbuffer_image stretch(mln::ndbuffer_image input, std::optional<pybind11::object> py_vs)
  {
    if (input.pdim() == 2)
    {
      return mln::py::visit<stretch_operator_t>(input.sample_type(), input, py_vs);
    }
    else
    {
      std::runtime_error("Unsupported dimension.");
      return {};
    }
  }

  template <typename T>
  struct stretch_virtual_dispatch_operator_t
  {
    template <typename Img>
    mln::ndbuffer_image operator()(Img&& img) const
    {
      if (auto* image_ptr = std::forward<Img>(img).template cast_to<T, 2>(); image_ptr)
      {
        return my::virtual_dispatch_stretch(*image_ptr);
      }
      else
      {
        std::runtime_error("Unable to convert the image to the required type.");
        return {};
      }
    }
  };

  mln::ndbuffer_image stretch_virtual_dispatch(mln::ndbuffer_image input)
  {
    if (input.pdim() == 2)
    {
      return mln::py::visit<stretch_virtual_dispatch_operator_t>(input.sample_type(), input);
```

```cpp
  }
  else
  {
    std::runtime_error("Unsupported dimension.");
    return {};
  }
}

template <typename T>
struct stretch_virtual_dispatch_type_erased_value_operator_t
{
  template <typename Img>
  mln::ndbuffer_image operator()(Img&& img) const
  {
    if (auto* image_ptr = std::forward<Img>(img).template cast_to<T, 2>(); image_ptr)
    {
      return my::stretch_virtual_dispatch_type_erased_value(*image_ptr);
    }
    else
    {
      std::runtime_error("Unable to convert the image to the required type.");
      return {};
    }
  }
};

mln::ndbuffer_image stretch_virtual_dispatch_type_erased_value(mln::ndbuffer_image input)
{
  if (input.pdim() == 2)
  {
    return mln::py::visit<stretch_virtual_dispatch_type_erased_value_operator_t>(input.sample_type(), input);
  }
  else
  {
    std::runtime_error("Unsupported dimension.");
    return {};
  }
}

using se_t = std::variant<mln::se::disc, mln::se::disc_non_decomp, mln::se::rect2d, mln::se::rect2d_non_decomp>;

template <typename T>
struct dilate2d_operator_t
{
  template <typename Img, typename SE>
  mln::ndbuffer_image operator()(Img&& img, SE&& se) const
  {
    if (auto* image_ptr = std::forward<Img>(img).template cast_to<T, 2>(); image_ptr)
    {
      return mln::morpho::experimental::dilation(*image_ptr, std::forward<SE>(se));
    }
    else
    {
      std::runtime_error("Unable to convert the image to the required type.");
      return {};
    }
  }
};

mln::ndbuffer_image dilate(mln::ndbuffer_image input, const se_t& se)
{
  if (input.pdim() == 2)
  {
    return std::visit(
        [&input](const auto& se_) { return mln::py::visit<dilate2d_operator_t>(input.sample_type(), input, se_); },
        se);
  }
  else
  {
    std::runtime_error("Unsupported dimension.");
    return {};
  }
}
```

```cpp
} // namespace


namespace mln::py
{

  using namespace pybind11::literals;

  void init_module_morpho(pybind11::module& m)
  {
    m.def("dilate", dilate,
          "Perform a morphological dilation.\n"
          "\n"
          "structuring element must be valid.",
          "Input"_a, "se"_a)

      .def("stretch", stretch,
           "Perform a morphological stretch.\n"
           "Input"_a,
           "vs"_a = pybind11::none())

      .def("stretch_virtual_dispatch", stretch_virtual_dispatch,
           "Perform a morphological stretch.\n"
           "Input"_a)

      .def("stretch_virtual_dispatch_type_erased_value", stretch_virtual_dispatch_type_erased_value,
           "Perform a morphological stretch.\n"
           "Input"_a);
  }

} // namespace mln::py
```

C.4 Exposed Pylena main module

pylena.cpp

```cpp
#include "morpho.hpp"
#include "ndimage.hpp"
#include "se.hpp"

#include <pybind11/pybind11.h>


PYBIND11_MODULE(pylena, m)
{
  mln::py::init_pylena_numpy(m);

  auto mmorpho = m.def_submodule("morpho", "Mathematical morphology module.");
  mln::py::init_module_morpho(mmorpho);

  auto mse = m.def_submodule("se", "Structuring elements module.");
  mln::py::init_module_se(mse);
}
```

C.5 Python value set module

The module has the following directory structure:

```
ValueSetExample/
| __init__.py
| BaseValueSet.py
| ExampleValueSet.py
```

BaseValueSet.py

```python
from abc import ABC, abstractmethod
from typing import Any


class AbstractValueSet(ABC):

    def __init__(self): pass

    @abstractmethod
    def cast(self, value: Any, type_):
        if type_ in ["int", "float", "bool", "str"]:
            import importlib
            module = importlib.import_module('builtins')
            cls = getattr(module, type_)
            return cls(value)
        raise ValueError()

    @abstractmethod
    def max(self):
        import math
        return math.inf

    @abstractmethod
    def min(self):
        import math
        return -math.inf

    @abstractmethod
    def sup(self):
        import math
        return math.inf

    @abstractmethod
    def inf(self):
        import math
        return -math.inf

    @abstractmethod
    def plus(self, value: Any) -> Any:
        return + value

    @abstractmethod
    def minus(self, value: Any) -> Any:
        return - value

    @abstractmethod
    def add(self, lhs: Any, rhs: Any) -> Any:
        return lhs + rhs

    @abstractmethod
    def sub(self, lhs: Any, rhs: Any) -> Any:
        return lhs - rhs

    @abstractmethod
    def mult(self, lhs: Any, rhs: Any) -> Any:
        return lhs * rhs

    @abstractmethod
    def div(self, lhs: Any, rhs: Any) -> Any:
        return lhs / rhs
```

```python
    @abstractmethod
    def mod(self, lhs: Any, rhs: Any) -> Any:
        return lhs % rhs

    @abstractmethod
    def pow(self, lhs: Any, rhs: Any) -> Any:
        return lhs ** rhs
```

ExampleStruct.py

```python
from typing import Any


class MyStruct:
    v_: Any

    def __init__(self, v: Any):
        self.v_ = v

    def getV(self) -> Any:
        return self.v_

    def setV(self, v: Any):
        self.v_ = v
```

ExampleValueSet.py

```python
from ValueSetExample.ExampleStruct import MyStruct
from ValueSetExample.BaseValueSet import AbstractValueSet

from typing import Any


class MyValueSet(AbstractValueSet):

    def __init__(self): pass

    def get_MyStruct__(self, v: Any):
        return v.getV() if isinstance(v, MyStruct) else v

    def cast(self, value: Any, type_):
        v = self.get_MyStruct__(value)
        return super().cast(v, type_)

    def max(self):
        return 255

    def min(self):
        return 0

    def sup(self):
        return 255

    def inf(self):
        return 0

    def plus(self, value: Any) -> Any:
        v = self.get_MyStruct__(value)
        return MyStruct(super().plus(v))

    def minus(self, value: Any) -> Any:
        v = self.get_MyStruct__(value)
        return MyStruct(super().minus(v))

    def add(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
        return MyStruct(super().add(l, r))

    def sub(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
```

```python
        return MyStruct(super().sub(l, r))

    def mult(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
        return MyStruct(super().mult(l, r))

    def div(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
        return MyStruct(super().div(l, r))

    def mod(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
        return MyStruct(super().mod(l, r))

    def pow(self, lhs: Any, rhs: Any) -> Any:
        l = self.get_MyStruct__(lhs)
        r = self.get_MyStruct__(rhs)
        return MyStruct(super().pow(l, r))
```

__init__.py

```python
from ValueSetExample.BaseValueSet import AbstractValueSet
from ValueSetExample.ExampleStruct import MyStruct
from ValueSetExample.ExampleValueSet import MyValueSet
```

C.6   Benchmark source code

benchmark.py

```python
import numpy as np
import skimage
import timeit
import pylena as pln
from ValueSetExample import MyValueSet, MyStruct

grayscale = skimage.data.camera()

vs = MyValueSet()


def test_fast():
    global grayscale
    pln_ret = pln.morpho.stretch(grayscale)
    return pln_ret


def test_slow():
    global grayscale, vs
    pln_ret = pln.morpho.stretch(grayscale, vs)
    return pln_ret


def test_virtual_dispatch():
    global grayscale
    pln_ret = pln.morpho.stretch_virtual_dispatch(grayscale)
    return pln_ret


def test_virtual_dispatch_type_erased_value():
    global grayscale
    pln_ret = pln.morpho.stretch_virtual_dispatch_type_erased_value(grayscale)
    return pln_ret


if __name__ == "__main__":
    print("Benchmarking static-dynamic bridge...")
    cpp_fast_ret = timeit.timeit("test_fast()", globals=locals(), number=10)
    cpp_virtual_dispatch_ret = timeit.timeit(
        "test_virtual_dispatch()", globals=locals(), number=10)
    cpp_virtual_dispatch_type_erased_value_ret = timeit.timeit(
        "test_virtual_dispatch_type_erased_value()", globals=locals(), number=10)
    cpp_python_vs_ret = timeit.timeit(
        "test_slow()", globals=locals(), number=10)
    print("Benchmarking done...")

    print(
        "Native value-set with native C++ value-type: {0:.4f}sec".format(cpp_fast_ret))
    print(
        "Value-set with virtual dispatch with native C++ value-type: {0:.4f}sec".format(cpp_virtual_dispatch_ret))
    print(
        "Value-set with virtual dispatch with C++ type-erased values: {0:.4f}sec".format(cpp_virtual_dispatch_type_erase
    print(
        "Injected Python value-set with native C++ value-type: {0:.4f}sec".format(cpp_python_vs_ret))
```